

Advanced Image Processing

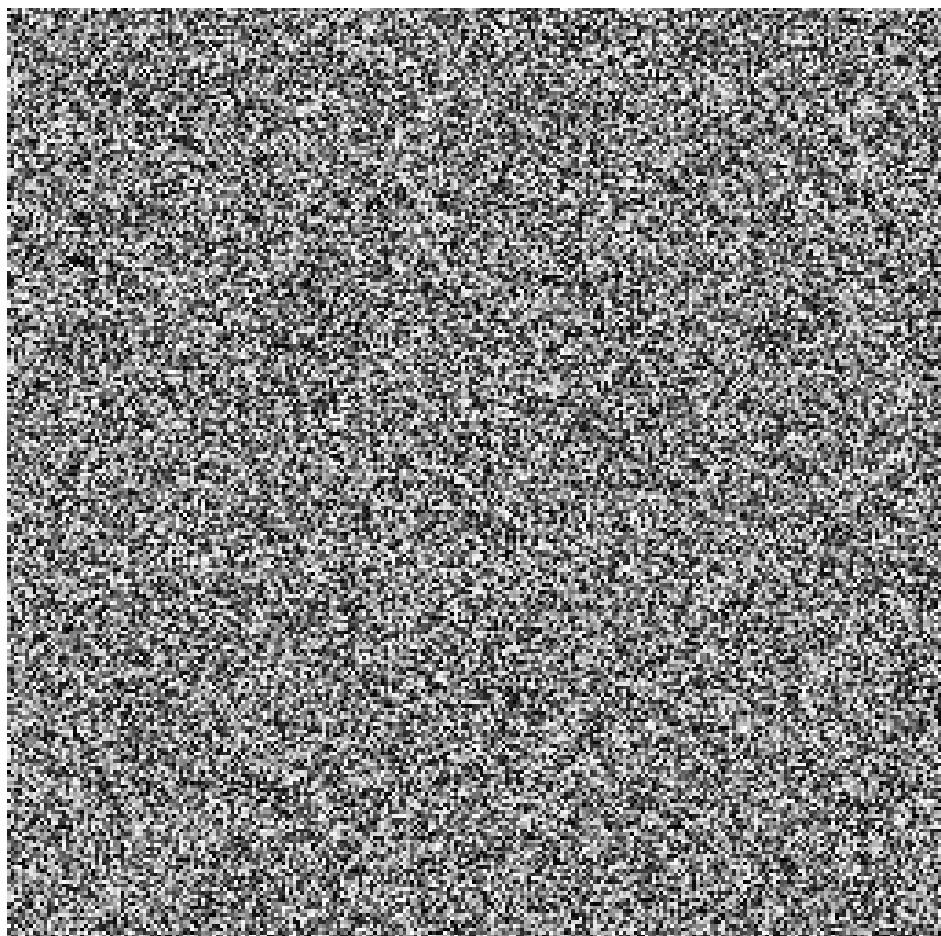
Image Filtering

What Is an Image?

```
>>> from numpy import random as r  
>>> I = r.rand(256,256);
```

- What is this? What does it look like?
- Which values does it take?
- How many values can it take?
- Is it an image?

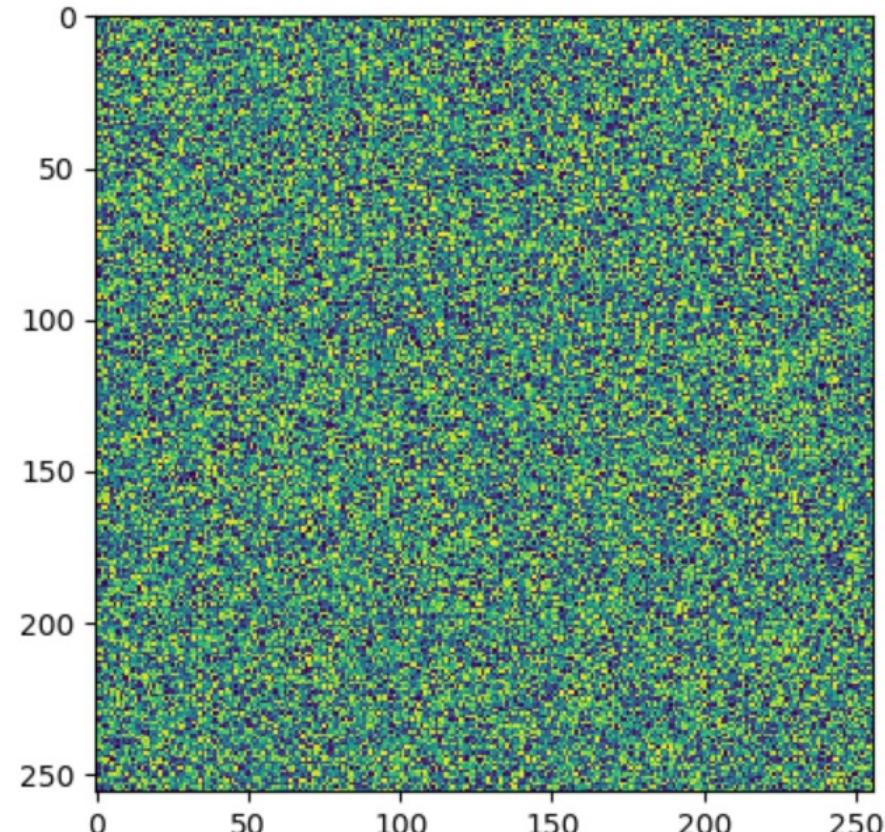
```
>>> from matplotlib import pyplot as p  
>>> I = r.rand(256,256);  
>>> p.imshow(I);  
>>> p.show();
```



Note: matplotlib and colour maps

```
>>> from matplotlib import pyplot as p  
>>> from numpy import random as r  
>>> I = r.rand(256,256);  
>>> p.imshow(I);  
>>> p.show();
```

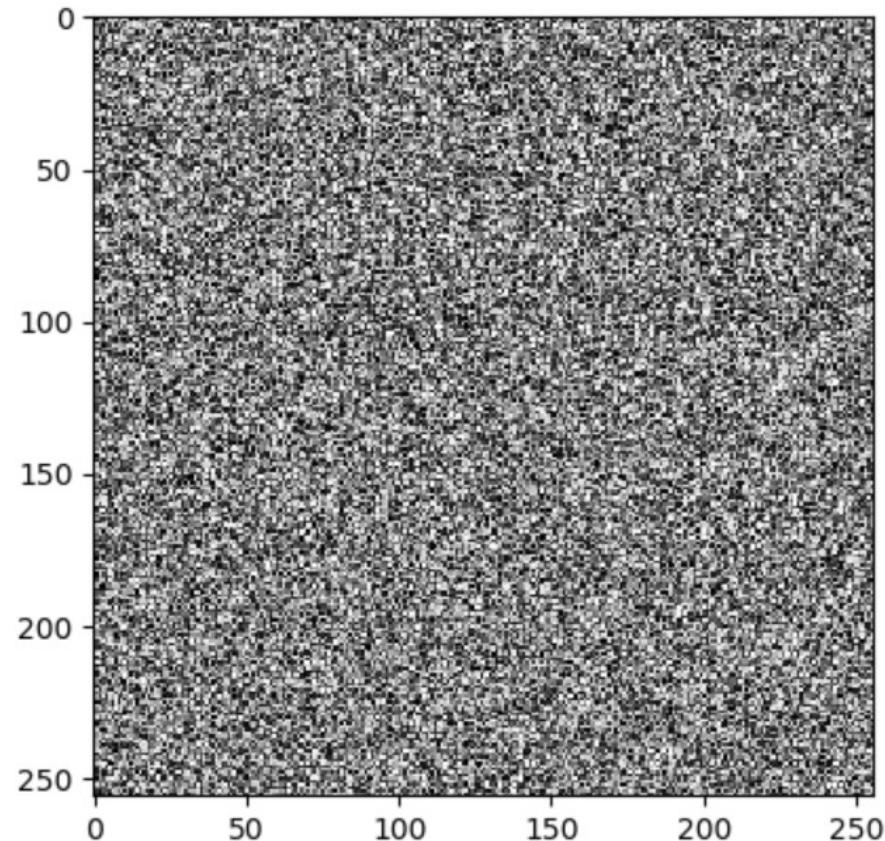
This code will actually produce an image like this:
matplotlib assumes the data is 'scientific' and applies a colourmap.



Note: matplotlib and colour maps

```
>>> from matplotlib import pyplot as p  
>>> from numpy import random as r  
>>> I = r.rand(256,256);  
>>> p.imshow(I,cmap='gray');  
>>> p.show();
```

Better.



Dimensionality of an Image

- @ 8bit = 256 values \wedge 65,536
 - Computer says 'Inf' combinations.
- Some depiction of all possible scenes would fit into this memory.

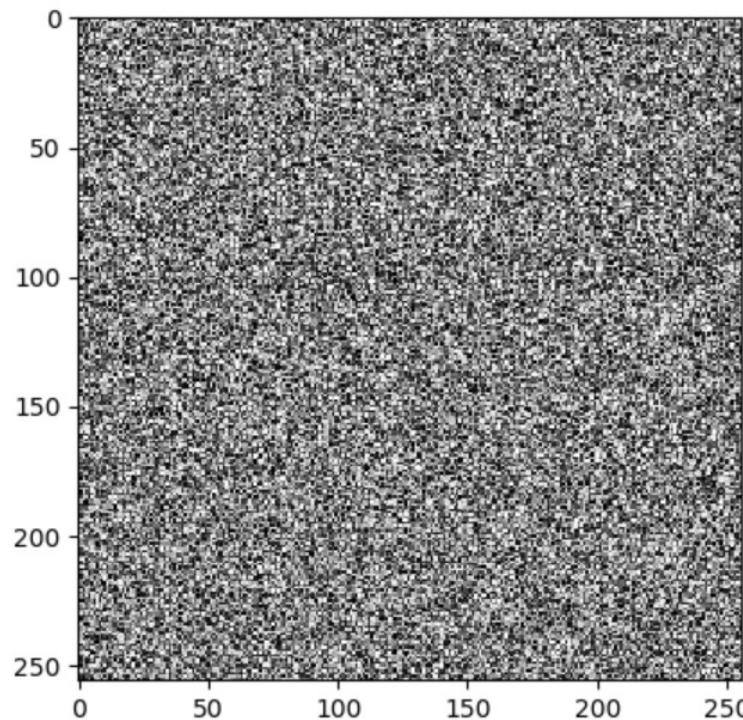
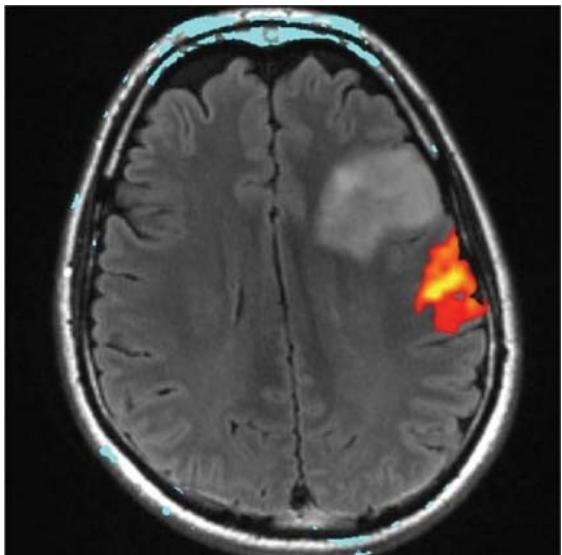
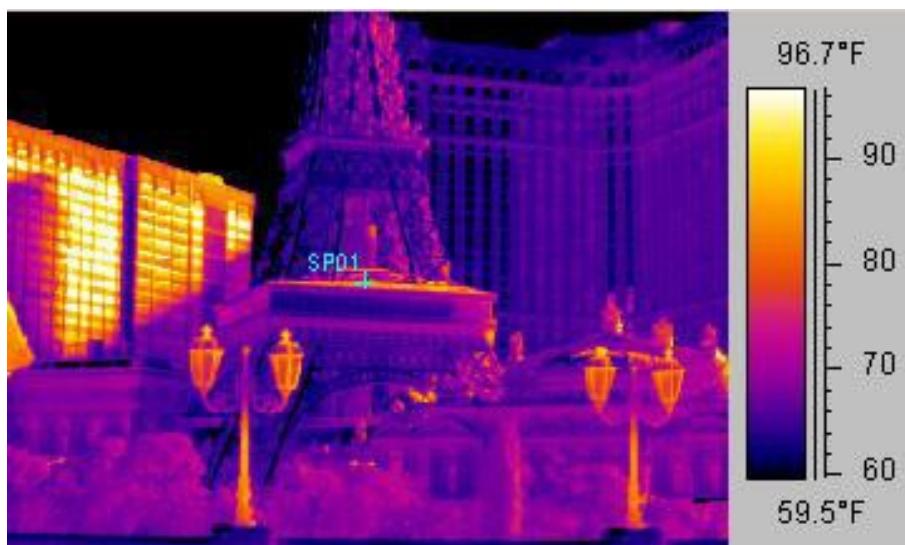


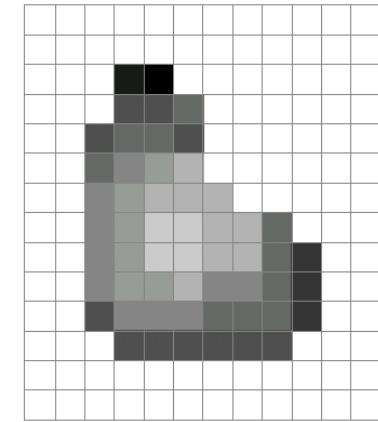
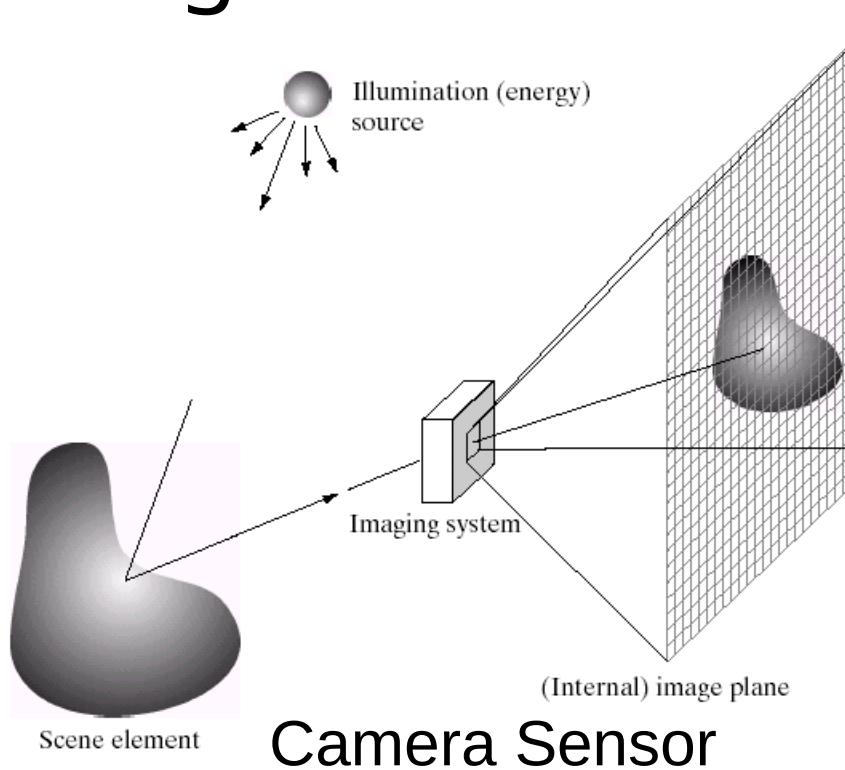
Image as a 2D sampling of signal

- Signal: function depending on some variable with physical meaning.
- Image: sampling of that function.
 - 2 variables: xy coordinates
 - 3 variables: xy + time (video)
 - ‘Brightness’ is the value of the function for visible light
- Can be other physical values too: temperature, pressure, depth ...

Example 2D Images



Integrating light over a range of angles.

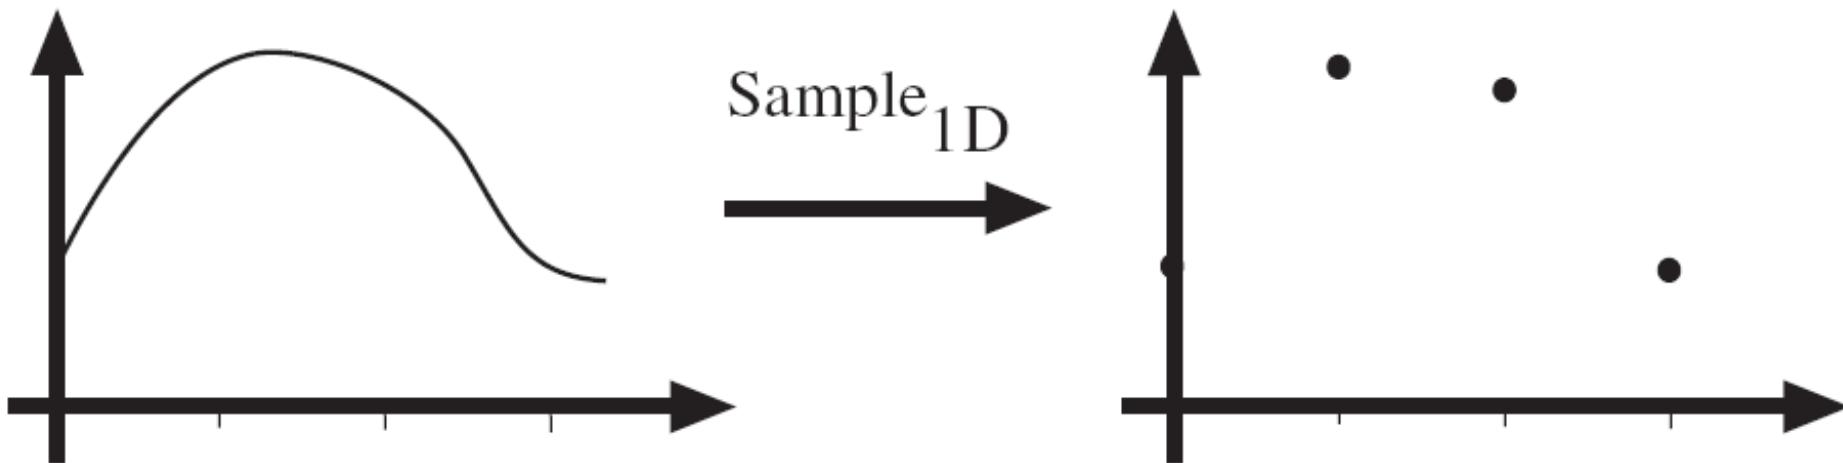


Output Image

Camera Sensor

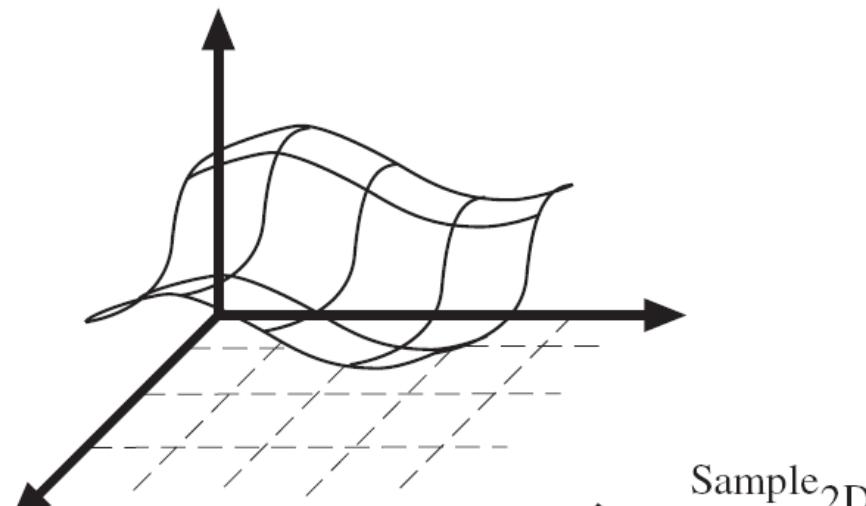
- Two things taking place:
- Sampling
 - Quantization

Sampling in 1D



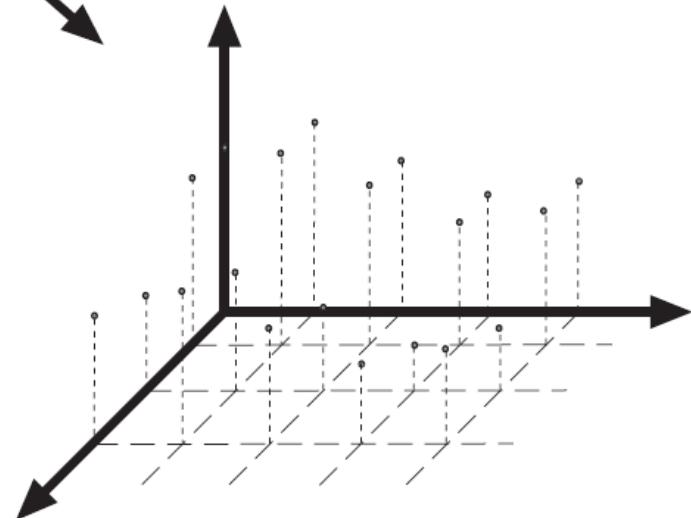
- Sampling in 1D takes a function, and returns a vector whose elements are values of that function at the sample points.

Sampling in 2D

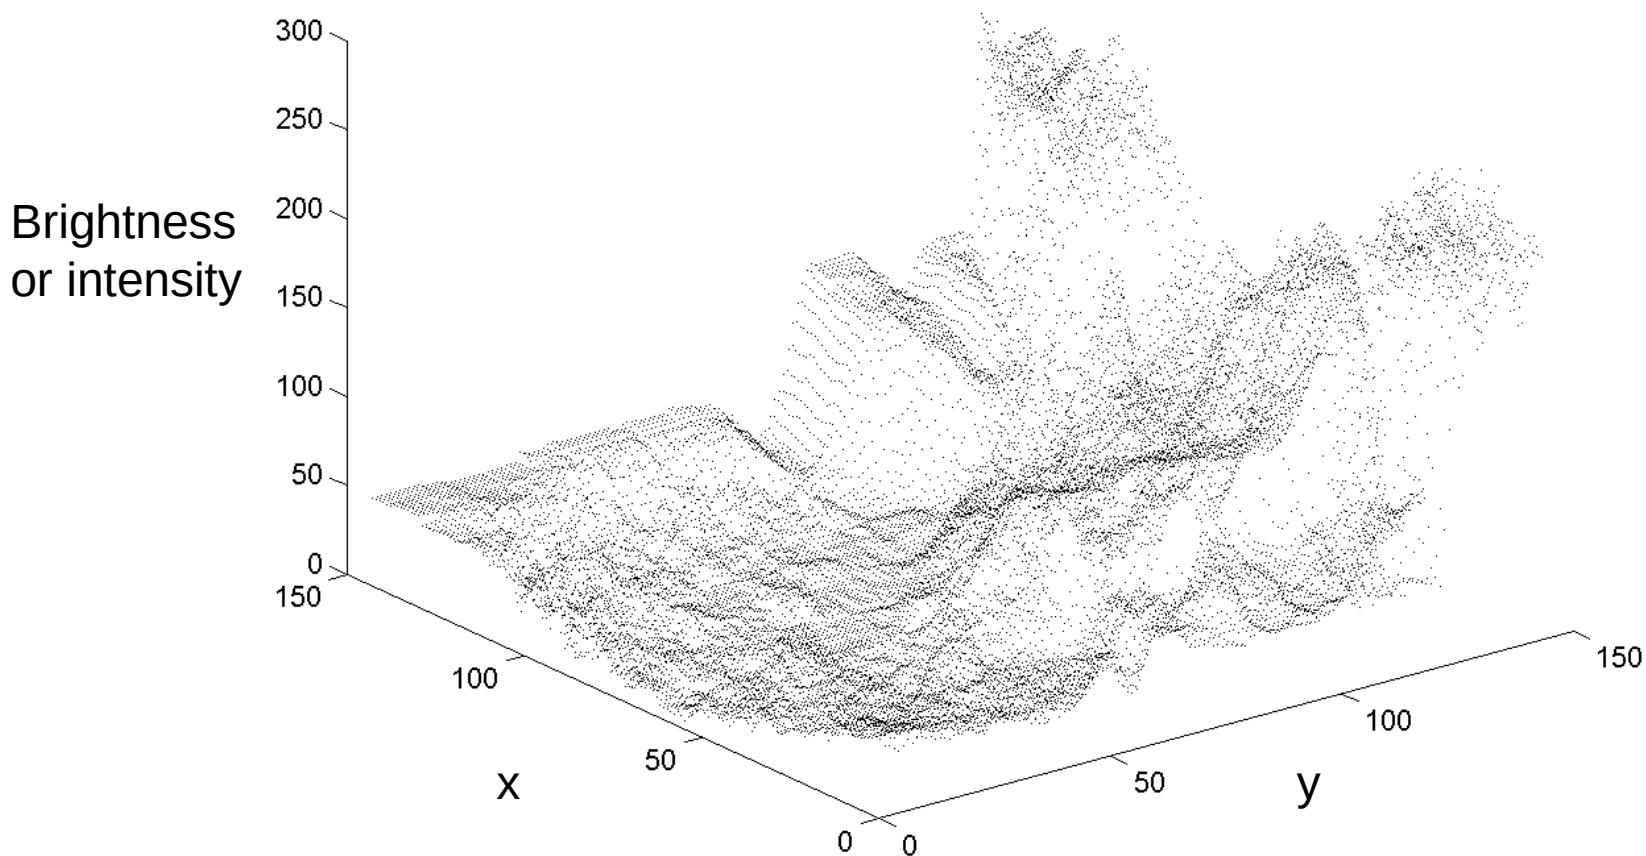


Sample_{2D}

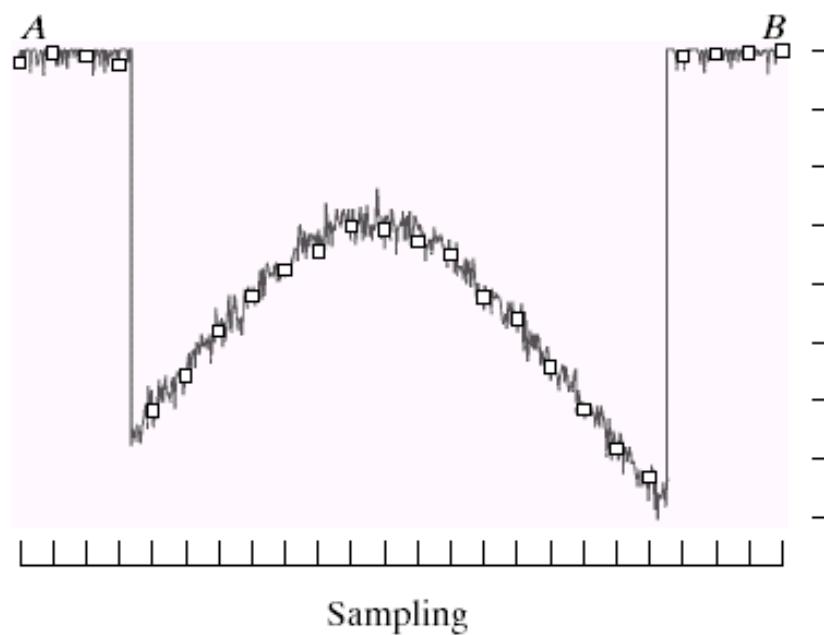
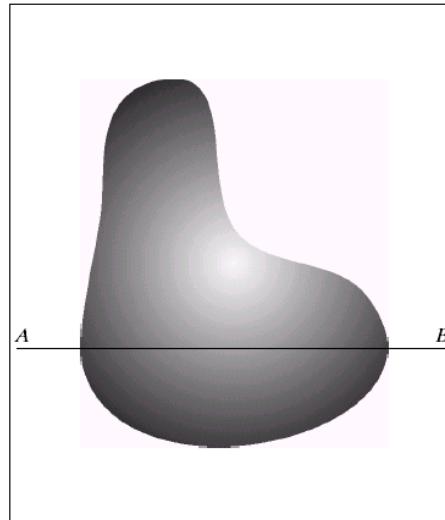
- Sampling in 2D takes a function and returns a matrix.



Grayscale Digital Image



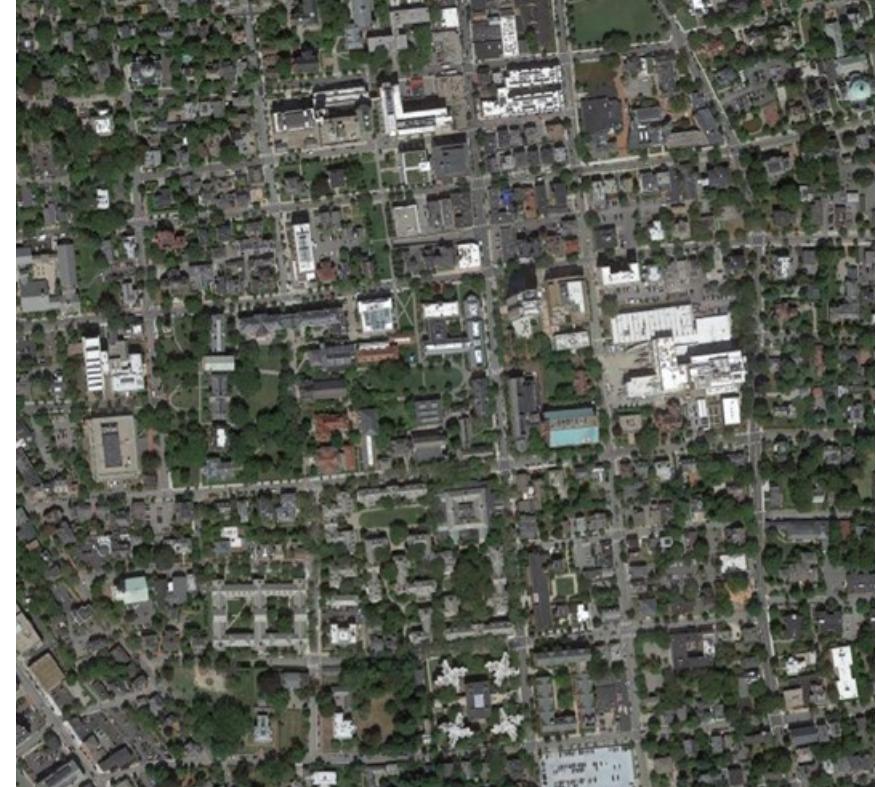
Sampling



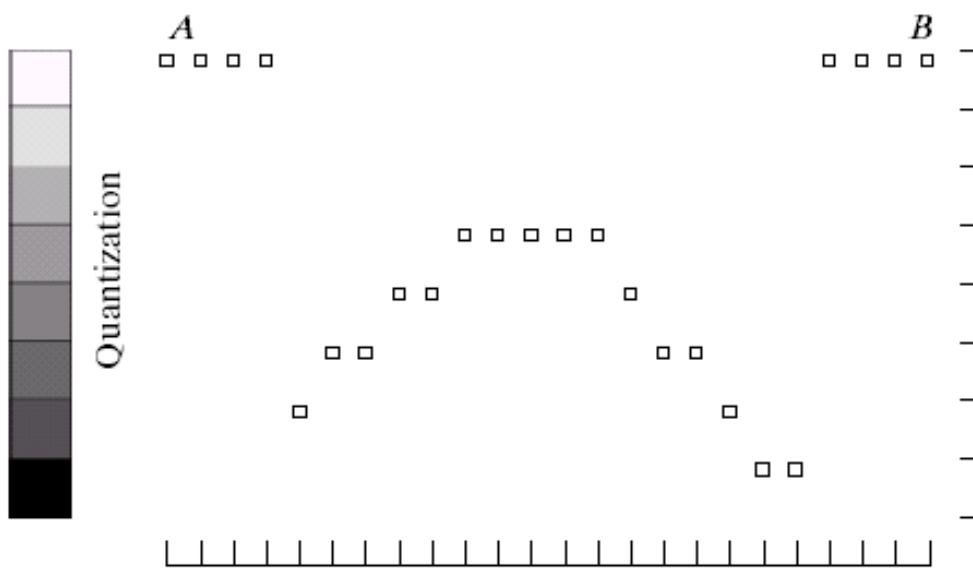
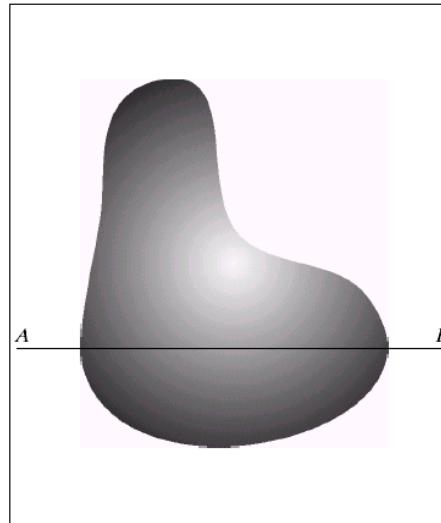
Sampling – spatial or geometric resolution

Both images are ~500x500 pixels

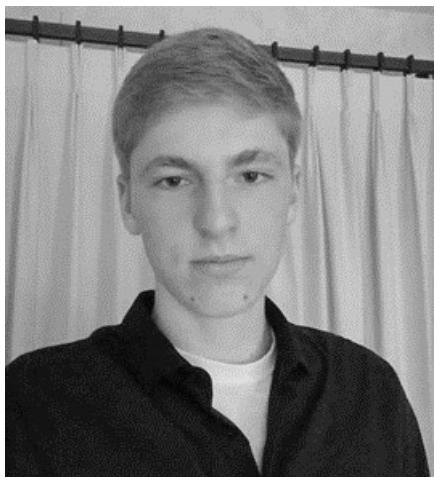
Each pixel area maps to different areas dimensions at the surface of the earth



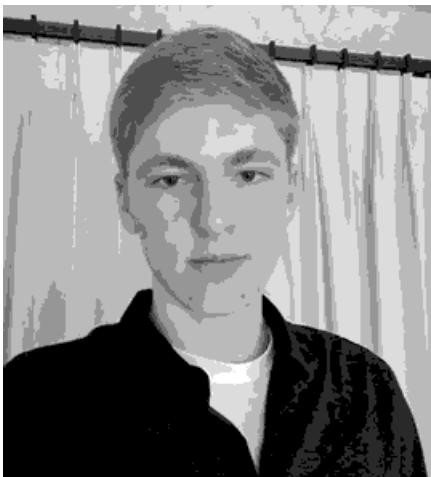
Quantization



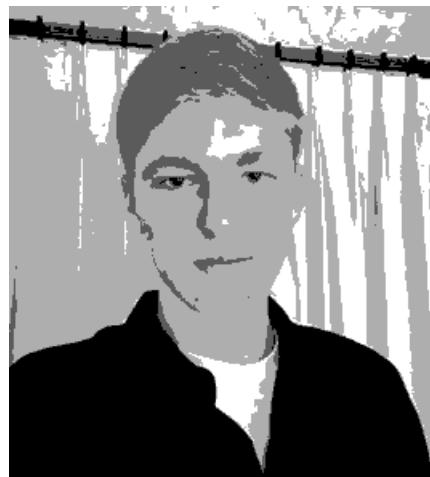
Quantization Effects – Radiometric Resolution



8 bit – 256 levels



4 bit – 16 levels



2 bit – 4 levels

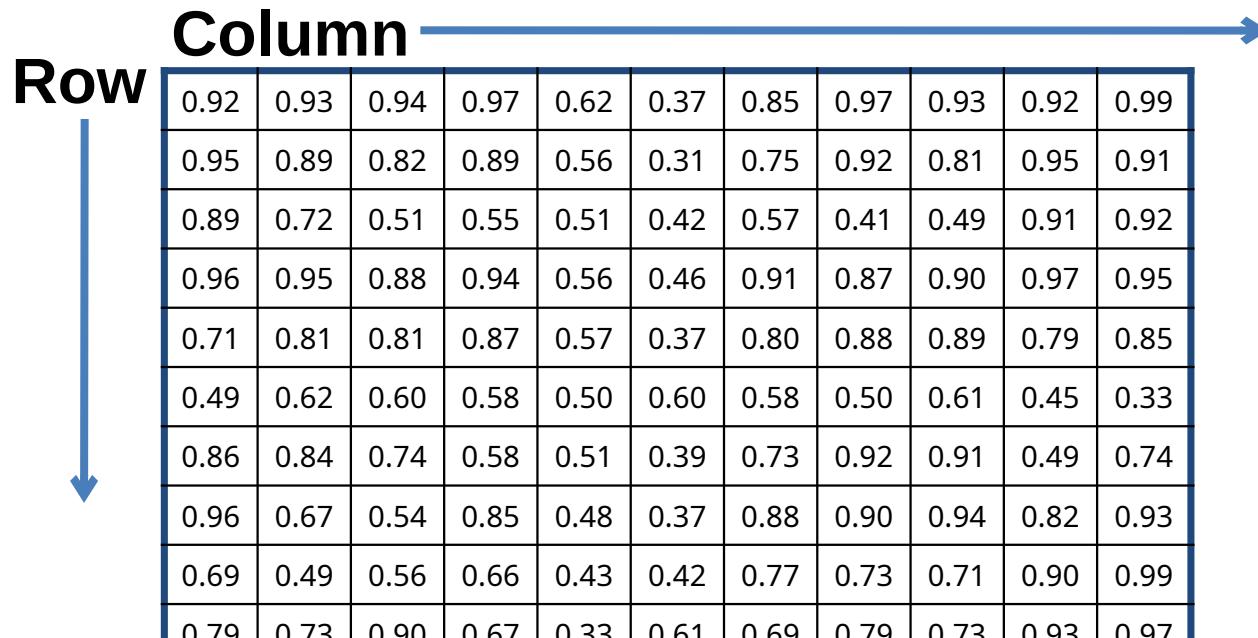


1 bit – 2 levels

Images in Python Numpy

N x M grayscale image “im”

- $\text{im}[0,0]$ = top-left pixel value
- $\text{im}[y, x]$ = y pixels down, x pixels to right
- $\text{im}[N-1, M-1]$ = bottom-right pixel



Row	Column	0.92	0.93	0.94	0.97	0.62	0.37	0.85	0.97	0.93	0.92	0.99
		0.95	0.89	0.82	0.89	0.56	0.31	0.75	0.92	0.81	0.95	0.91
		0.89	0.72	0.51	0.55	0.51	0.42	0.57	0.41	0.49	0.91	0.92
		0.96	0.95	0.88	0.94	0.56	0.46	0.91	0.87	0.90	0.97	0.95
		0.71	0.81	0.81	0.87	0.57	0.37	0.80	0.88	0.89	0.79	0.85
		0.49	0.62	0.60	0.58	0.50	0.60	0.58	0.50	0.61	0.45	0.33
		0.86	0.84	0.74	0.58	0.51	0.39	0.73	0.92	0.91	0.49	0.74
		0.96	0.67	0.54	0.85	0.48	0.37	0.88	0.90	0.94	0.82	0.93
		0.69	0.49	0.56	0.66	0.43	0.42	0.77	0.73	0.71	0.90	0.99
		0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97

Grayscale intensity



Color

R



G



B



copyright 2000 philg@mit.edu

Images in Python Numpy

N x M RGB image “im”

- $\text{im}[0,0,0]$ = top-left pixel value in R-channel
- $\text{Im}[x, y, b]$ = x pixels to right, y pixels down in the b^{th} channel
- $\text{Im}[N-1, M-1, 3]$ = bottom-right pixel in B-channel

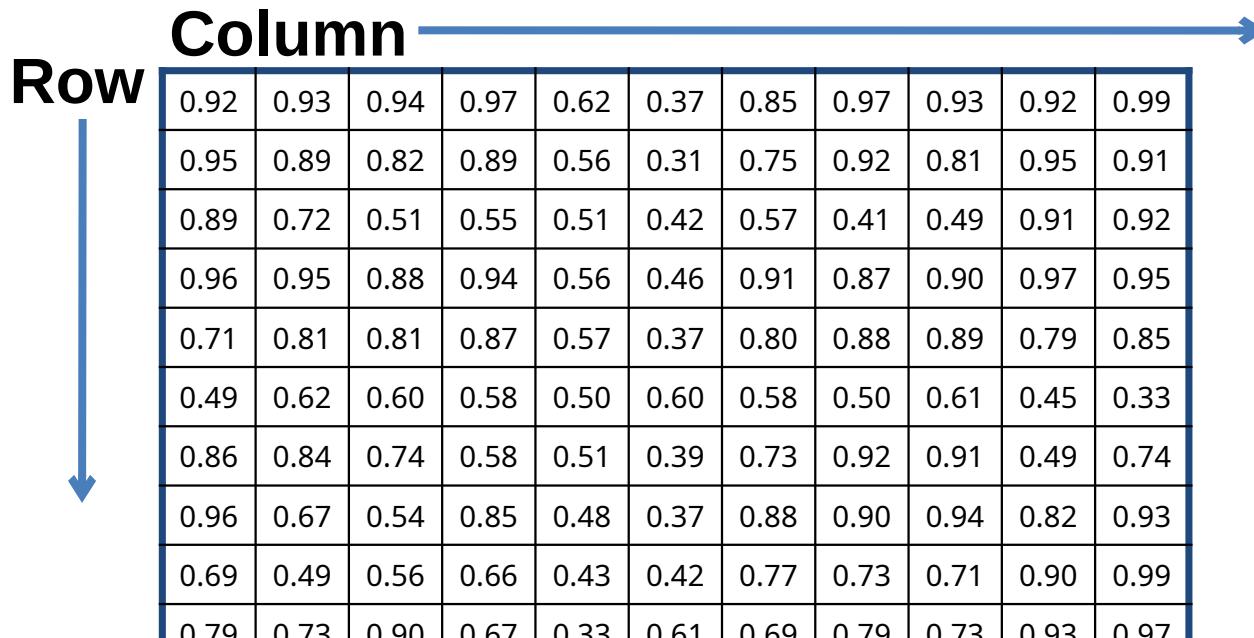
The diagram illustrates a 10x10 grid representing an RGB image. The vertical axis is labeled "Row" with a downward-pointing arrow, and the horizontal axis is labeled "Column" with a rightward-pointing arrow. The grid is divided into three vertical columns, each representing a color channel: R (Red), G (Green), and B (Blue). The R channel contains values from 0.92 to 0.99. The G channel contains values from 0.91 to 0.99. The B channel contains values from 0.79 to 0.97. The grid is composed of 100 cells, representing the total number of pixels in the image.

0.92	0.93	0.94	0.97	0.62	0.37	0.85	0.97	0.93	0.92	0.99
0.95	0.89	0.82	0.89	0.56	0.31	0.75	0.92	0.81	0.95	0.91
0.89	0.72	0.51	0.55	0.51	0.42	0.57	0.41	0.49	0.91	0.92
0.96	0.95	0.88	0.94	0.56	0.46	0.91	0.87	0.90	0.97	0.95
0.71	0.81	0.81	0.87	0.57	0.37	0.80	0.88	0.89	0.79	0.85
0.49	0.62	0.60	0.58	0.50	0.60	0.58	0.50	0.61	0.45	0.33
0.86	0.84	0.74	0.58	0.51	0.39	0.73	0.92	0.91	0.49	0.74
0.96	0.67	0.54	0.85	0.48	0.37	0.88	0.90	0.94	0.82	0.93
0.69	0.49	0.56	0.66	0.43	0.42	0.77	0.73	0.71	0.90	0.99
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93
				0.79	0.73	0.90	0.67	0.33	0.61	0.69
				0.91	0.94	0.89	0.49	0.41	0.78	0.78
					0.79	0.73	0.90	0.67	0.33	0.61
					0.91	0.94	0.89	0.49	0.41	0.78
						0.79	0.73	0.90	0.67	0.33
						0.91	0.94	0.89	0.49	0.41
							0.79	0.73	0.90	0.67
							0.91	0.94	0.89	0.49
								0.79	0.73	0.90
								0.91	0.94	0.89
									0.79	0.73
									0.91	0.94
										0.79

Images in Python Numpy

Take care between types! (io and img_as_float32 are part of skimage library)

- uint8 (values 0 to 255) - io.imread("file.jpg")
- float32 (values 0 to 255) -
io.imread("file.jpg").astype(np.float32)
- float32 (values 0 to 1) - img_as_float32(io.imread("file.jpg"))



0.92	0.93	0.94	0.97	0.62	0.37	0.85	0.97	0.93	0.92	0.99
0.95	0.89	0.82	0.89	0.56	0.31	0.75	0.92	0.81	0.95	0.91
0.89	0.72	0.51	0.55	0.51	0.42	0.57	0.41	0.49	0.91	0.92
0.96	0.95	0.88	0.94	0.56	0.46	0.91	0.87	0.90	0.97	0.95
0.71	0.81	0.81	0.87	0.57	0.37	0.80	0.88	0.89	0.79	0.85
0.49	0.62	0.60	0.58	0.50	0.60	0.58	0.50	0.61	0.45	0.33
0.86	0.84	0.74	0.58	0.51	0.39	0.73	0.92	0.91	0.49	0.74
0.96	0.67	0.54	0.85	0.48	0.37	0.88	0.90	0.94	0.82	0.93
0.69	0.49	0.56	0.66	0.43	0.42	0.77	0.73	0.71	0.90	0.99
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97

Image filtering

- Image filtering:
 - Compute function of local neighborhood at each position

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

Image filtering

- Image filtering:
 - Compute function of local neighborhood at each position

h=output f=filter I=image

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

2d coords=k, l 2d coords=m, n

[]

[]

[]

Example: box filter

$$f[\cdot, \cdot]$$
$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

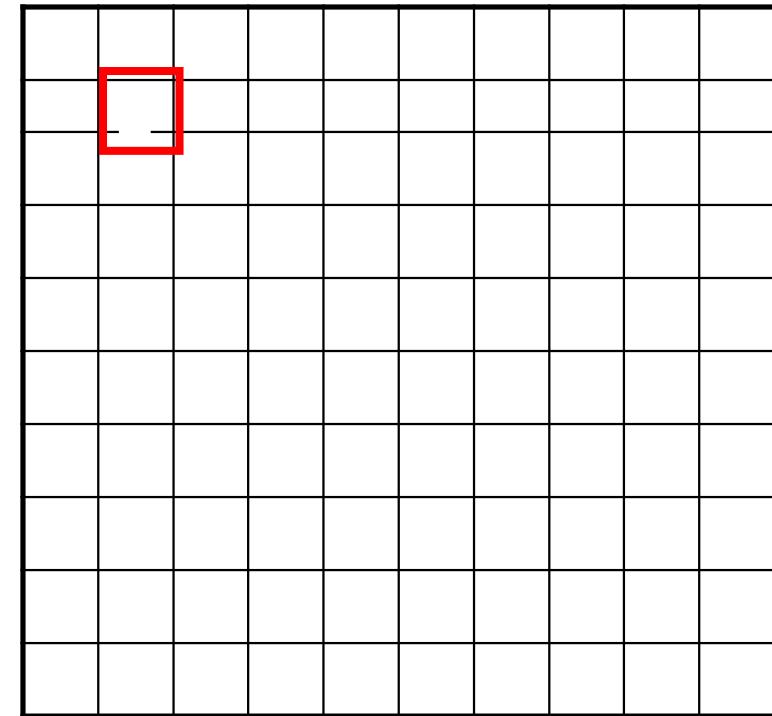
Image filtering

$$f[\cdot, \cdot] \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$I[.,.]$

$h[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0



$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

Credit: S. Seitz

Image filtering

$$f[\cdot, \cdot] \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$$I[.,.]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$h[.,.]$$

0										

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

Credit: S. Seitz

Image filtering

$$f[\cdot, \cdot] \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$$I[.,.]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	0	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$h[.,.]$$

			0	10	20				

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

Credit: S. Seitz

Image filtering

$$f[\cdot, \cdot] \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$I[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$

0 10 20 30

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

Image filtering

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$I[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$

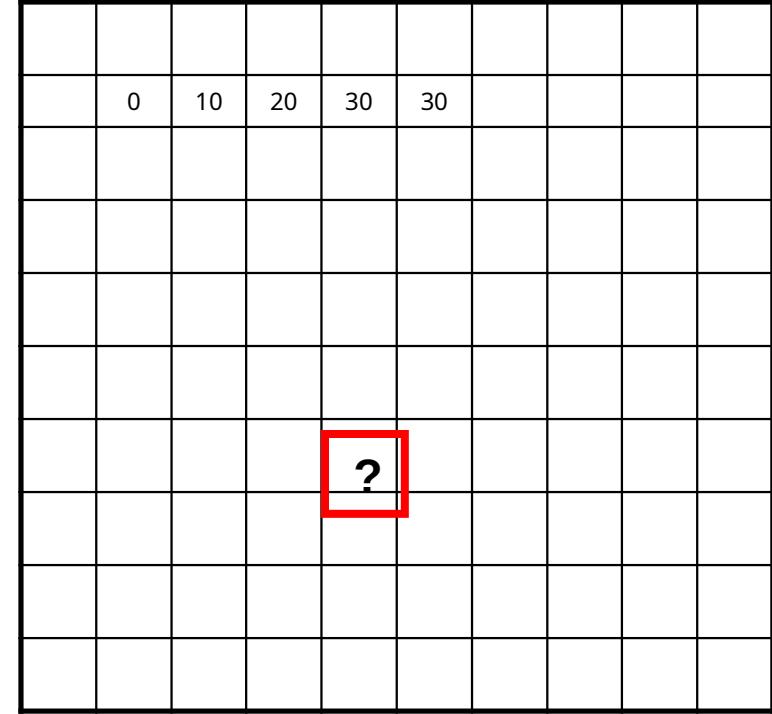
0 10 20 30 30

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

Image filtering

 $I[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

 $h[.,.]$ 

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

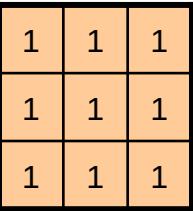


Image filtering

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$$I[.,.]$$

$$h[.,.]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30					

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

Credit: S. Seitz

Image filtering

$$f[\cdot, \cdot] \quad \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$$I[.,.]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$h[.,.]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

Credit: S. Seitz

Box Filter

What does it do?

Replaces each pixel with
an average of its
neighborhood

Achieve smoothing effect
(remove sharp features)

$$\frac{1}{9} \begin{matrix} f[\cdot, \cdot] \\ \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} \end{matrix}$$

Why does it sum to one?

Smoothing with box filter

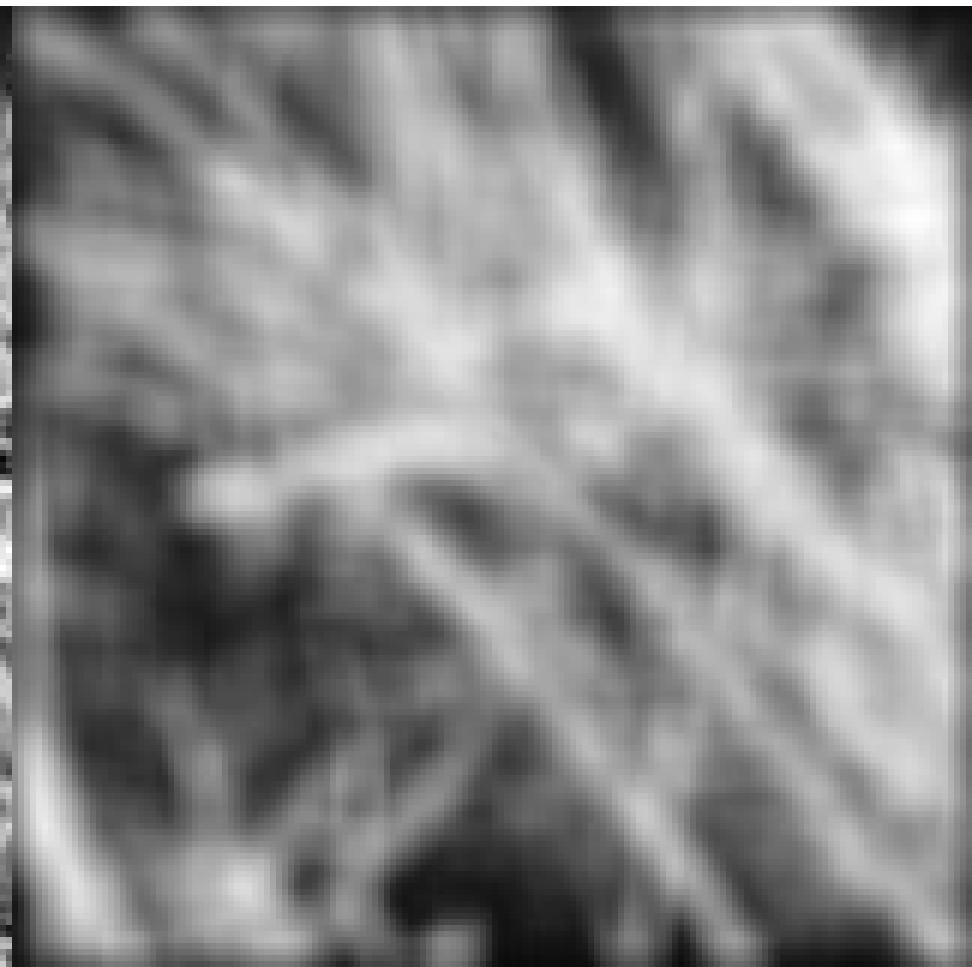
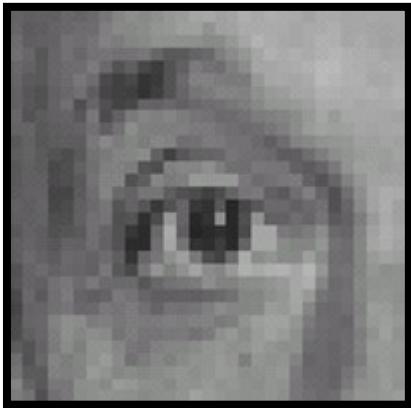
 $f[\cdot, \cdot]$
$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$


Image filtering

- Image filtering:
 - Compute function of local neighborhood at each position

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

- Really important!
 - Enhance images
 - Denoise, resize, increase contrast, etc.
 - Extract information from images
 - Texture, edges, distinctive points, etc.
 - Detect patterns
 - Template matching



1.

0	0	0
0	1	0
0	0	0
2.

0	0	0
0	0	1
0	0	0
3.

1	0	-1
2	0	-2
1	0	-1
4.

0	0	0
0	2	0
0	0	0

 - $\frac{1}{9}$

1	1	1
1	1	1
1	1	1

1. Practice with linear filters



0	0	0
0	1	0
0	0	0

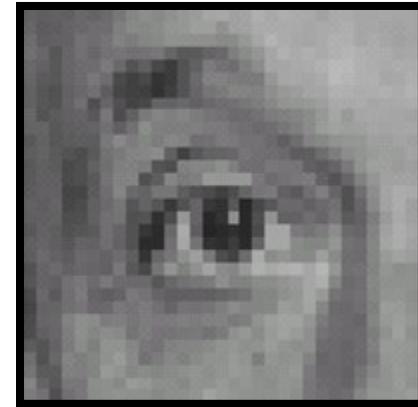
?

1. Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

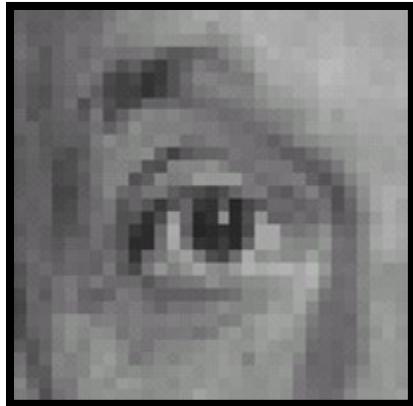
2. Practice with linear filters



0	0	0
0	0	1
0	0	0

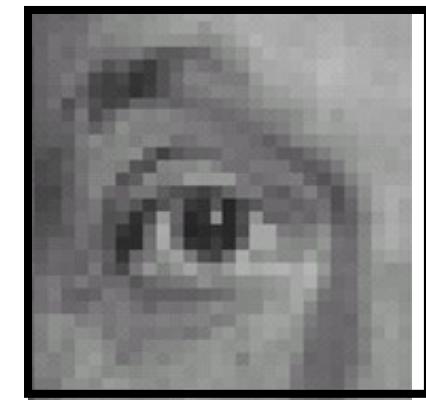
?

2. Practice with linear filters



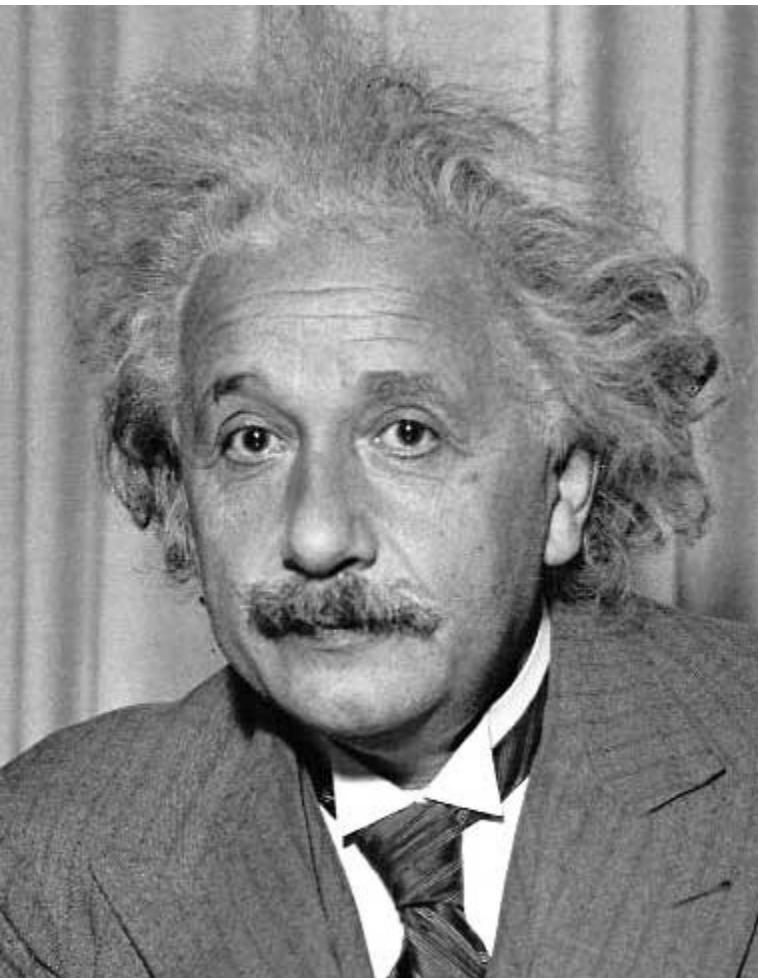
Original

0	0	0
0	0	1
0	0	0



Shifted left
By 1 pixel

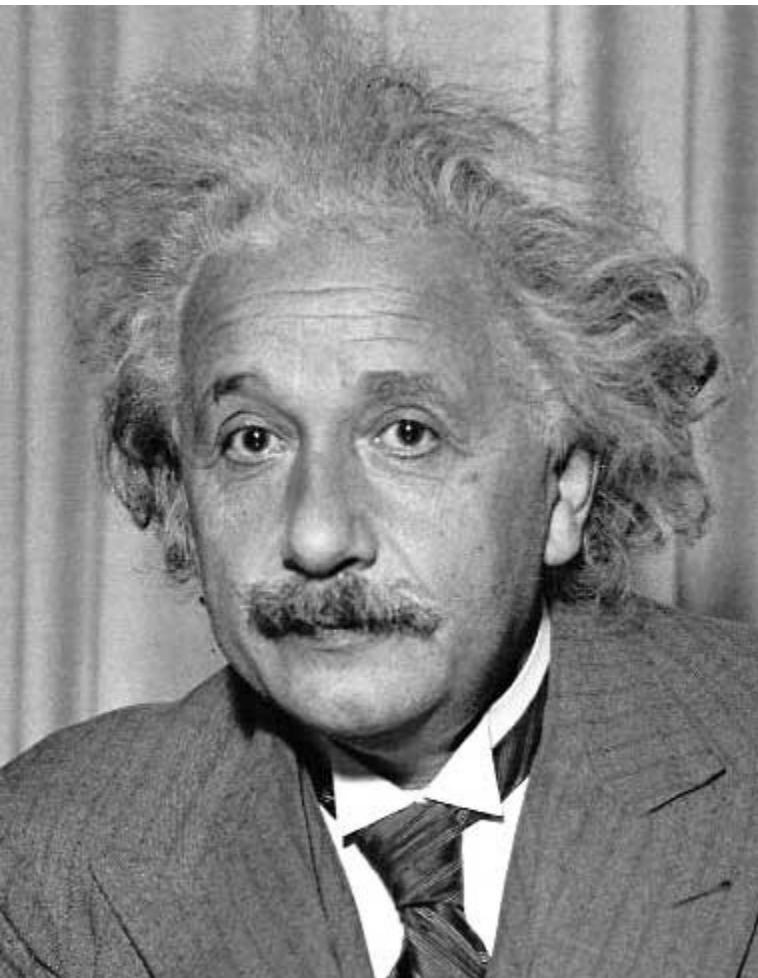
3. Practice with linear filters



1	0	-1
2	0	-2
1	0	-1

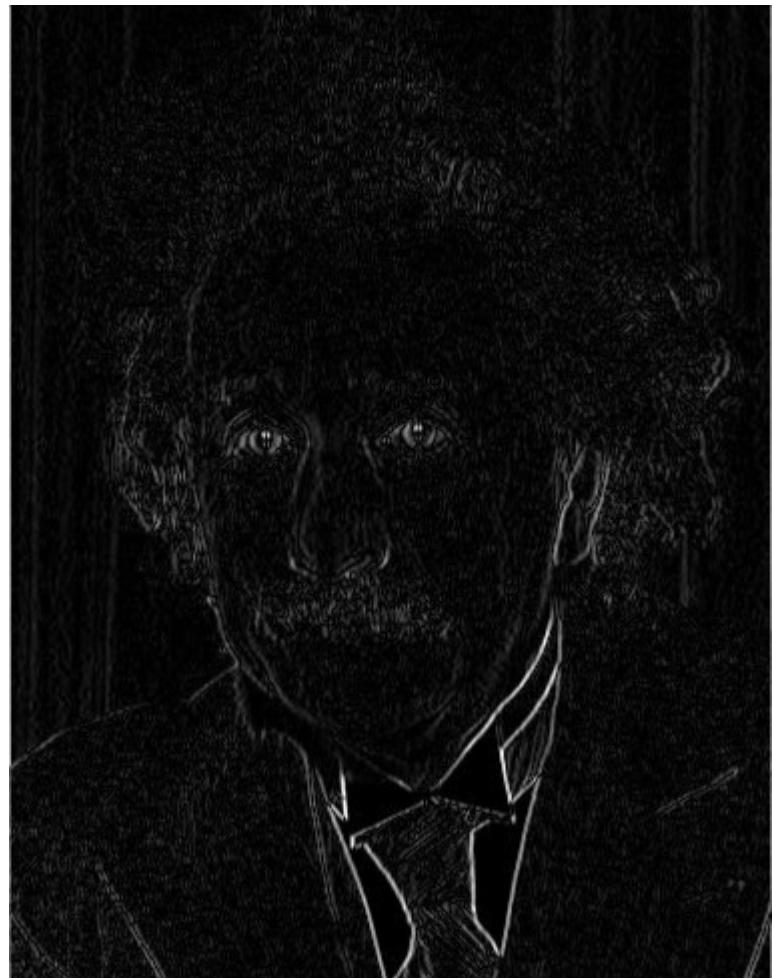
?

3. Practice with linear filters



1	0	-1
2	0	-2
1	0	-1

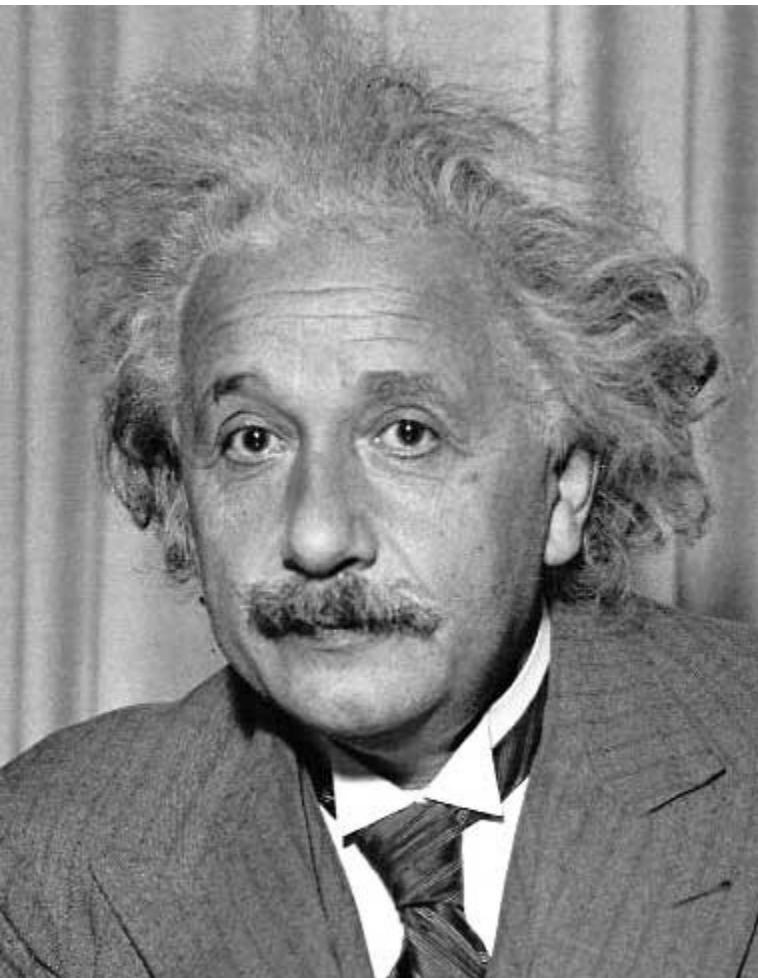
Sobel



Vertical Edge
(absolute value)

David Lowe

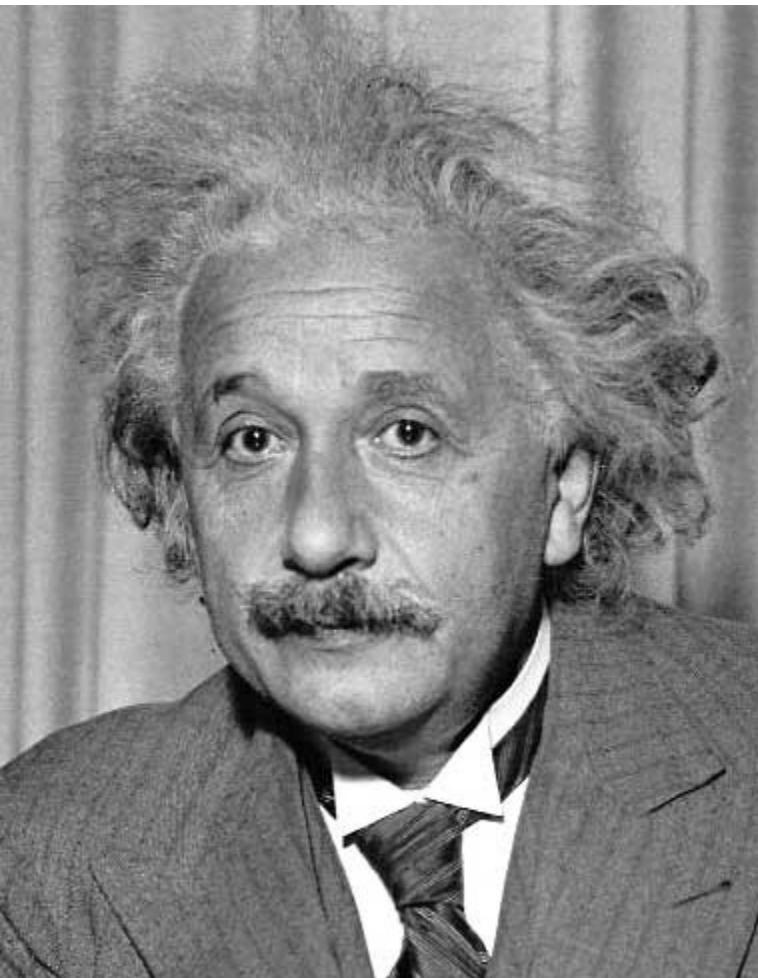
3. Practice with linear filters



1	2	1
0	0	0
-1	-2	-1

?

3. Practice with linear filters



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge
(absolute value)

David Lowe

4. Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$	1	1	1
1	1	1	1
1	1	1	1

?

(Note that filter sums to 1)

4. Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

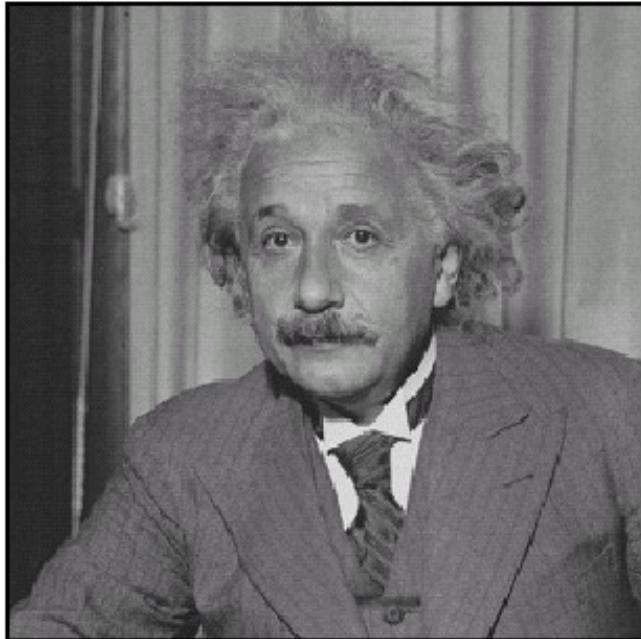
1	1	1
1	1	1
1	1	1



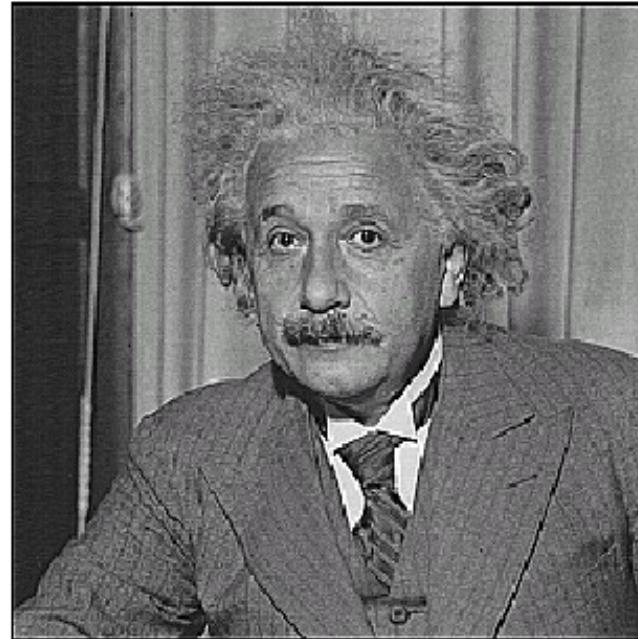
Sharpening filter

- Accentuates differences with local average

4. Practice with linear filters



before



after

Key properties of linear filters

Linearity:

$$\text{imfilter}(I, f_1 + f_2) = \text{imfilter}(I, f_1) + \text{imfilter}(I, f_2)$$

Shift invariance:

Same behavior given intensities regardless of pixel location m, n

$$\text{imfilter}(I, \text{shift}(f)) = \text{shift}(\text{imfilter}(I, f))$$

Any linear, shift-invariant operator can be represented as a convolution.

Correlation and Convolution

- 2d correlation

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

e.g., `h = scipy.signal.correlate2d(f, I)`

Correlation and Convolution

- 2d correlation

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

e.g., `h = scipy.signal.correlate2d(f, I)`

- 2d convolution

$$h[m, n] = \sum_{k,l} f[k, l] I[m - k, n - l]$$

e.g., `h = scipy.signal.convolve2d(f, I)`

Convolution is the same as correlation with a 180° rotated filter kernel.

Correlation and convolution are identical when the filter kernel is symmetric.

Convolution properties

Commutative: $a * b = b * a$

- Conceptually no difference between filter and signal
- But particular filtering implementations might break this equality, e.g., image edges

Associative: $a * (b * c) = (a * b) * c$

- Often apply several filters one after another:
- $((a * b_1) * b_2) * b_3$

Convolution properties

Commutative: $a * b = b * a$

- Conceptually no difference between filter and signal
- But particular filtering implementations might break this equality, e.g., image edges

Associative: $a * (b * c) = (a * b) * c$

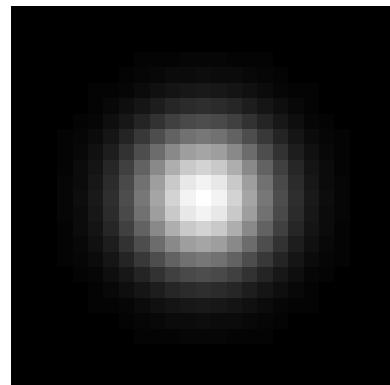
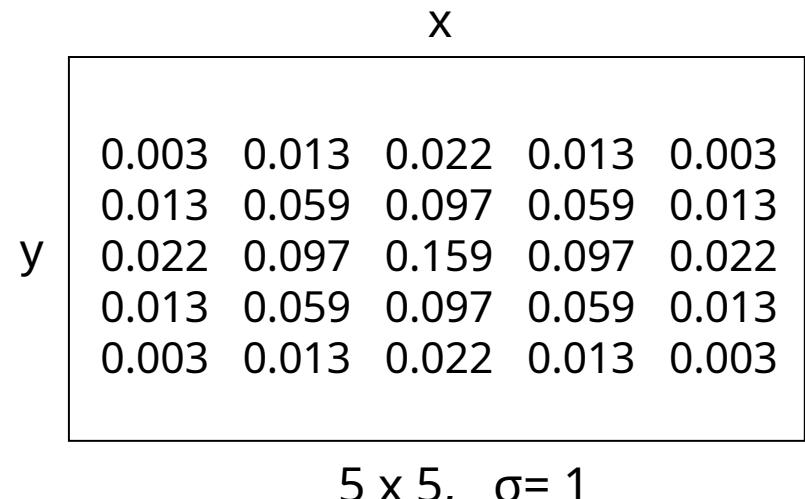
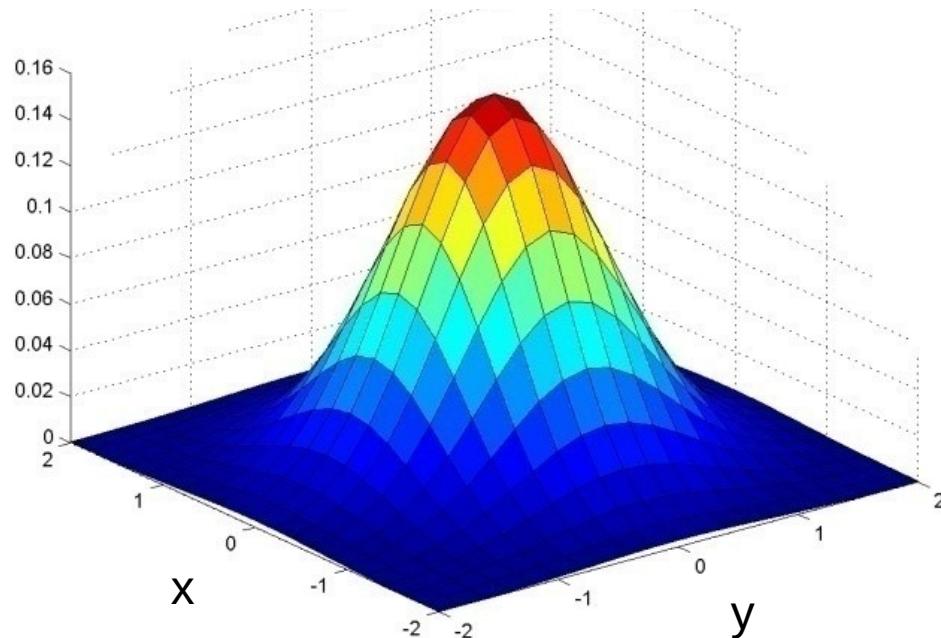
- Often apply several filters one after another:
- $((a * b_1) * b_2) * b_3$
- This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
- Correlation is

Convolution properties

- Commutative: $a * b = b * a$
 - Conceptually no difference between filter and signal
 - But particular filtering implementations might break this equality, e.g., image edges
- Associative: $a * (b * c) = (a * b) * c$
 - Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
 - This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
 - Correlation is not associative (rotation effect)
 - Why important?
- Distributes over addition: $a * (b + c) = (a * b) + (a * c)$
- Scalars factor out: $ka * b = a * kb = k(a * b)$

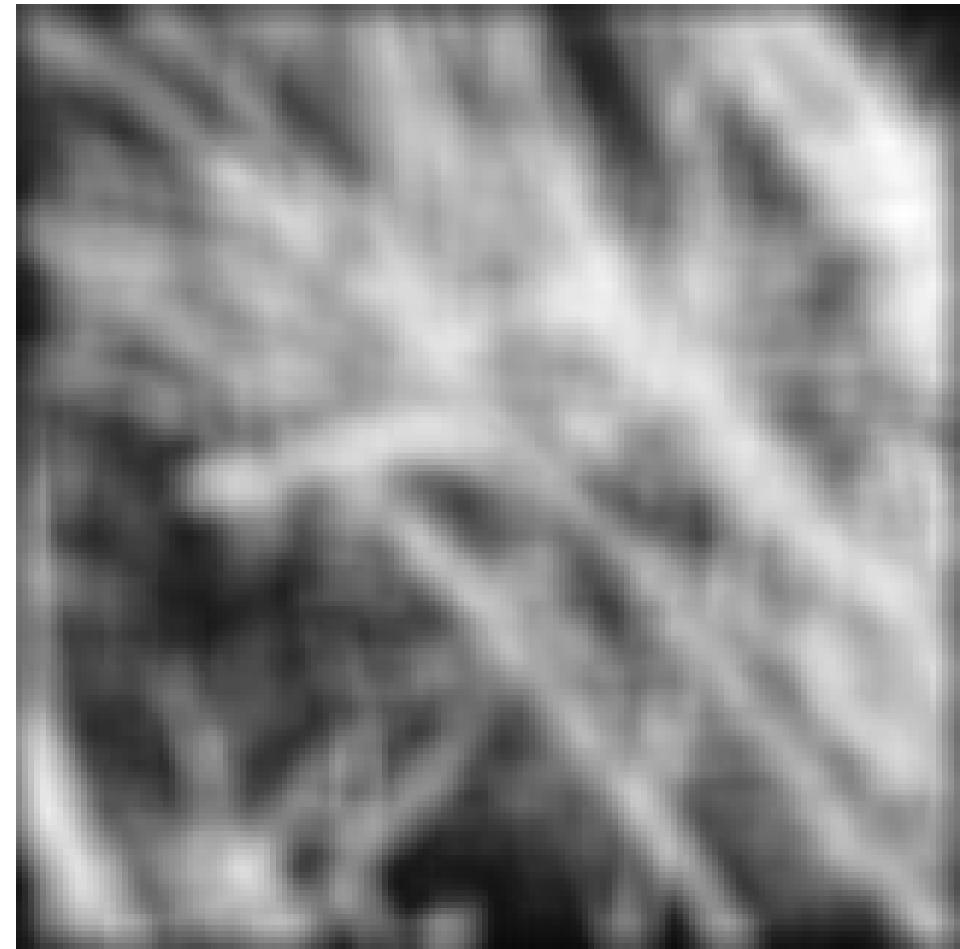
Important filter: Gaussian

Weight contributions of neighboring pixels by nearness

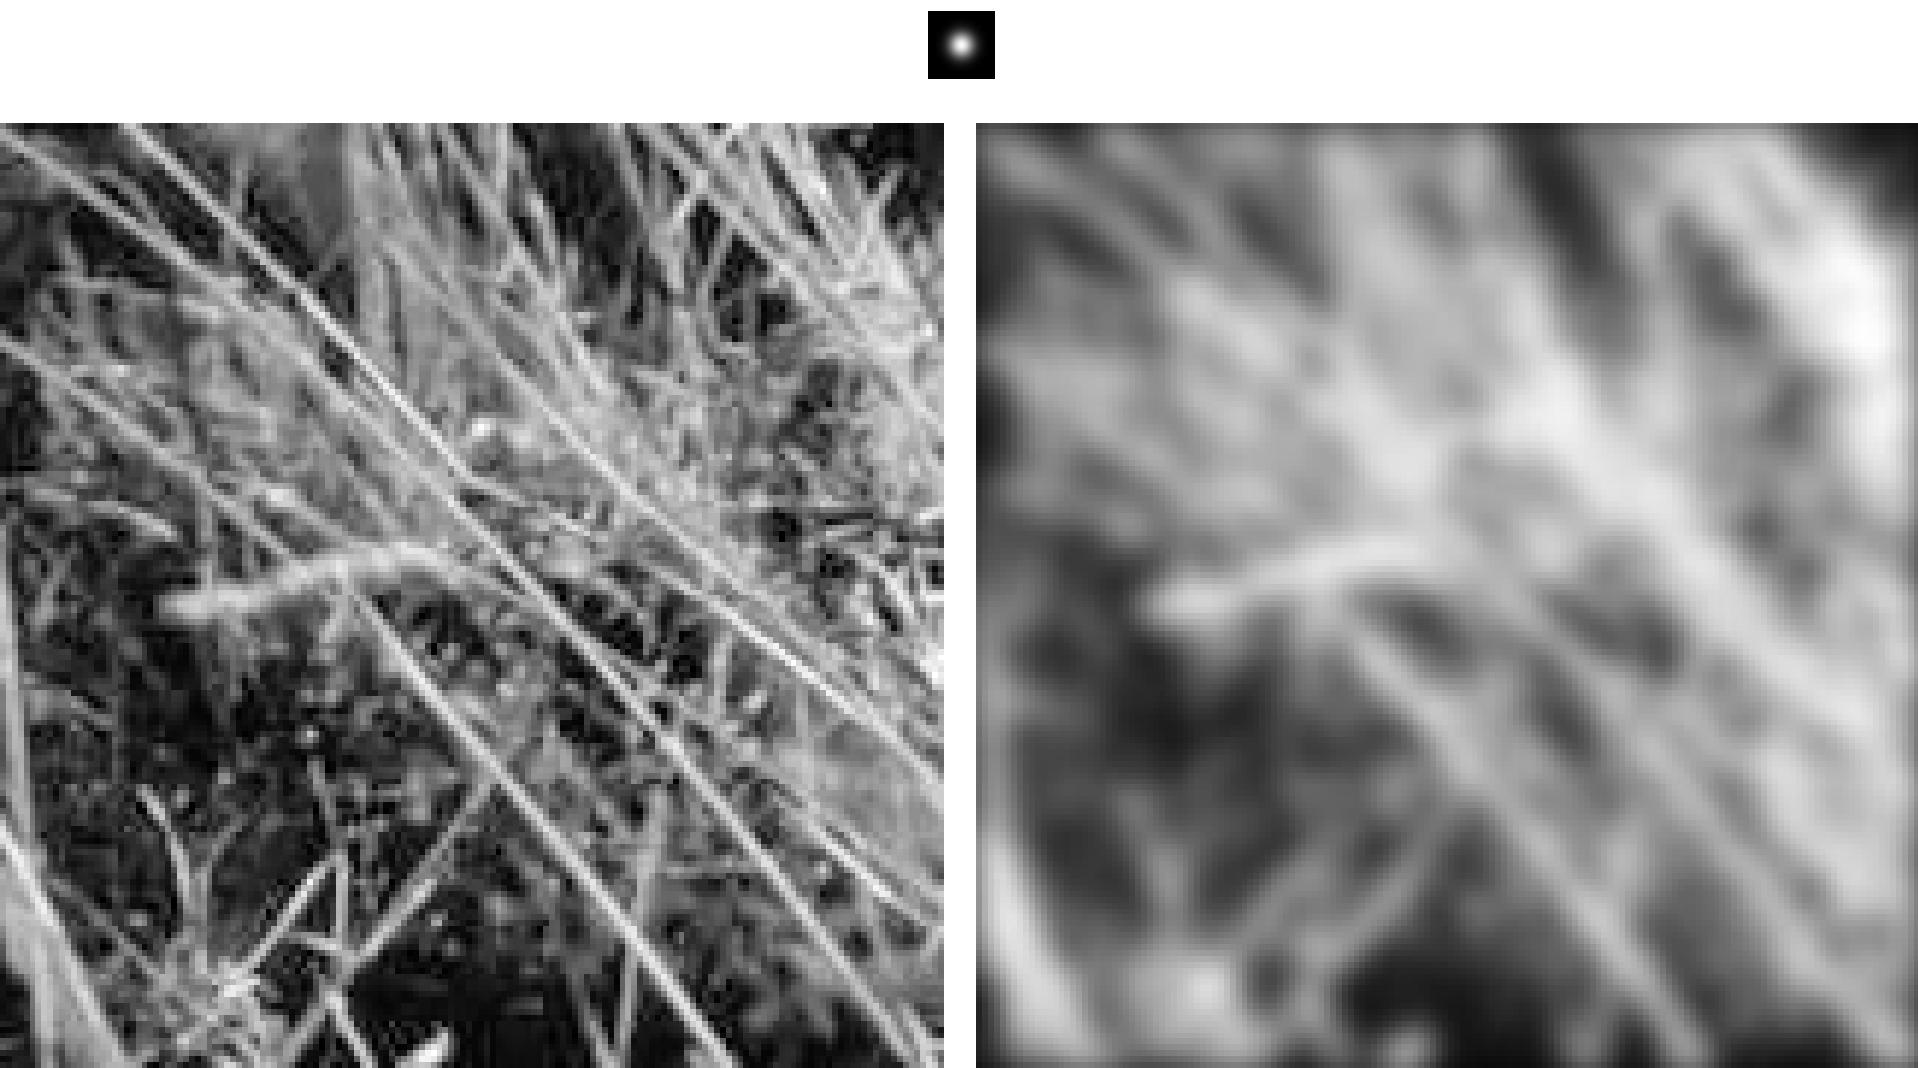


$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

Smoothing with box filter



Smoothing with Gaussian filter



Gaussian filters

- Remove “high-frequency” components from the image (low-pass filter)
 - Images become more smooth
- Gaussian convolved with Gaussian...
 - ...is another Gaussian
 - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
 - Convolving twice with Gaussian kernel of width σ is same as convolving once with kernel of width $\sigma\sqrt{2}$
- *Separable* kernel
 - Factors into product of two 1D Gaussians

Separability of the Gaussian filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

Separability example

2D convolution
(center location only)

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix}$$

The filter factors
into a product of 1D
filters:

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} = \begin{matrix} 1 \\ 2 \\ 1 \end{matrix} \times \begin{matrix} 1 & 2 & 1 \end{matrix}$$

Perform convolution
along rows:

$$\begin{matrix} 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix} = \begin{matrix} 11 & & \\ 18 & & \\ 18 & & \end{matrix}$$

Followed by convolution
along the remaining column:

Separability

Why is separability useful in practice?

MxN image, PxQ filter

- 2D convolution: $\sim MNPQ$ multiply-adds
- Separable 2D: $\sim MN(P+Q)$ multiply-adds

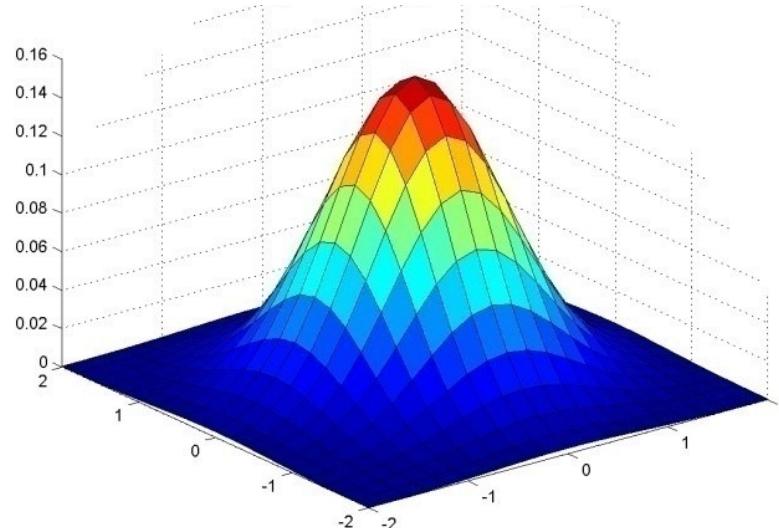
Speed up = $PQ/(P+Q)$

9x9 filter = ~4.5x faster

Practical matters: Parameter Selection

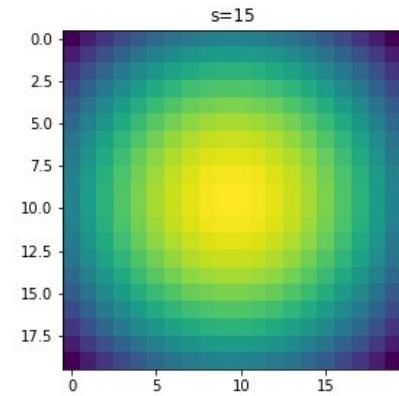
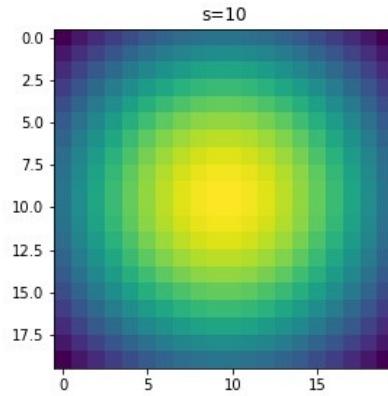
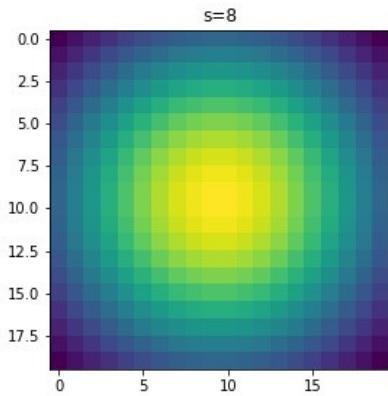
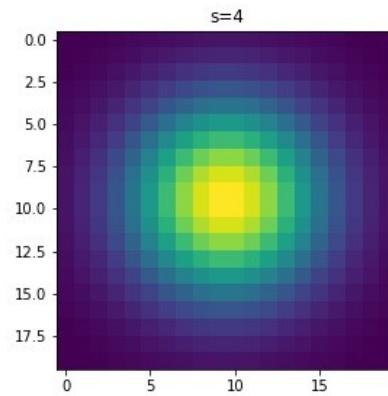
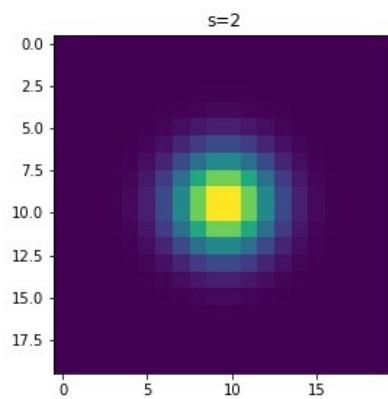
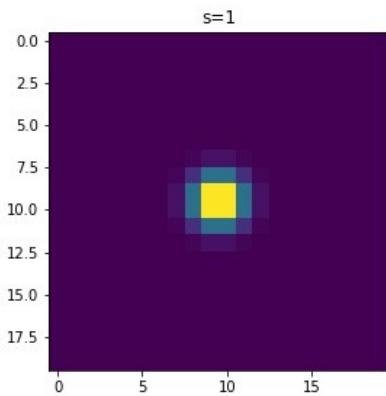
How big should the filter be?

- Values at edges should be near zero
- Gaussians have infinite extent...
- Rule of thumb for Gaussian: set filter half-width to about 3σ



Practical matters: Parameter Selection

Window = 21x21



Practical matters: Image boundaries

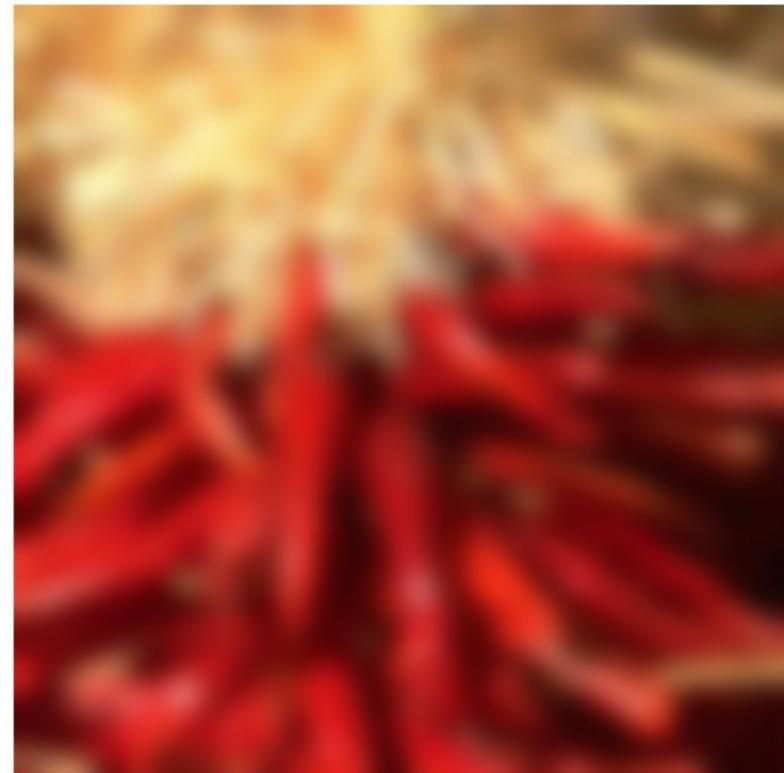
What about near the edge?



Practical matters: Image boundaries

What about near the edge?

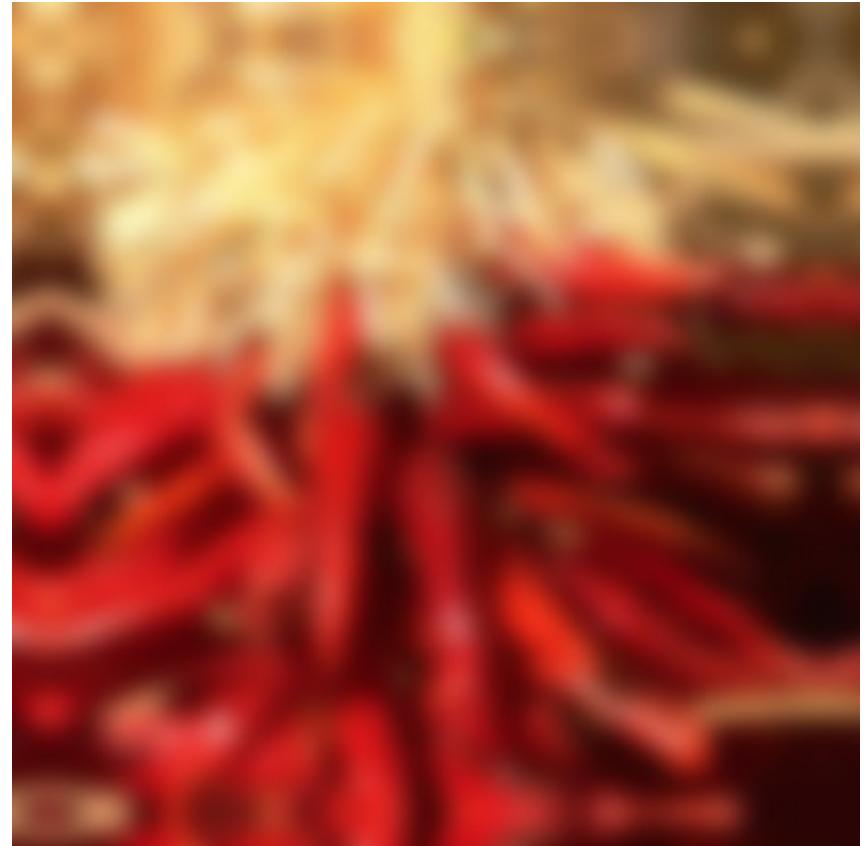
- The filter window falls off the edge of the image



Practical matters: Image boundaries

What about near the edge?

- The filter window falls off the edge of the image
- Need to extrapolate
- methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge



Practical matters: Sobel negative values !

- Throwback: Sobel Filter

```
>> I = img_to_float32( io.imread( 'luke.jpg' ) );  
>> h = convolve2d( I, sobelKernel );
```

1	0	-1
2	0	-2
1	0	-1

Sobel (Vertical edge detector)

```
plt.imshow(h); plt.imshow(h+0.5);
```



Practical matters: Sobel negative values !

$h(:,:,1) < 0$



$h(:,:,1) > 0$



Practical matters: Sobel Negative values !

Solution for display purposes we can:

- Add the mean of the result before displaying it.
 - Shifts the image towards the positive domain, but the contrast would remain and may suffer from saturation.
- Use the absolute values of the result (stretched between 0 and 255).
 - This makes very negative and very positive gradients appear brighter
- Other solution coming when we see edge detection

Exercise

* = Convolution operator

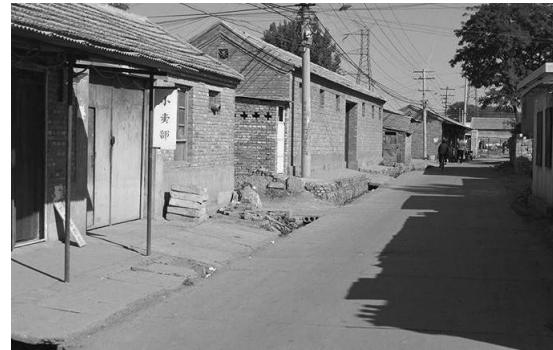
a) $_+ = D * B$

b) $F = D * _+$

c) $_+ = D * D$

d) $A = _+ * _+$

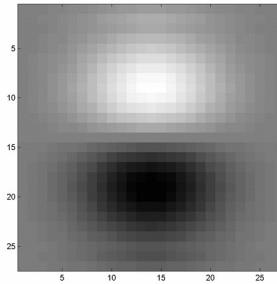
D



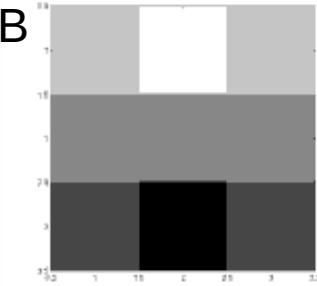
H



A



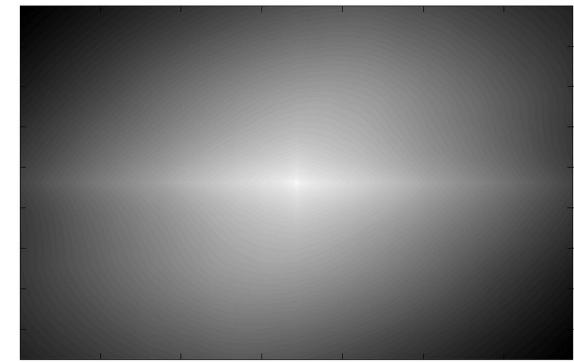
B



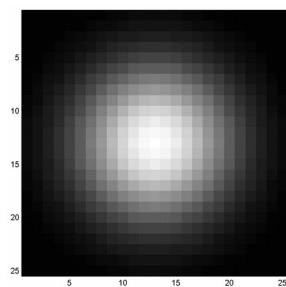
F



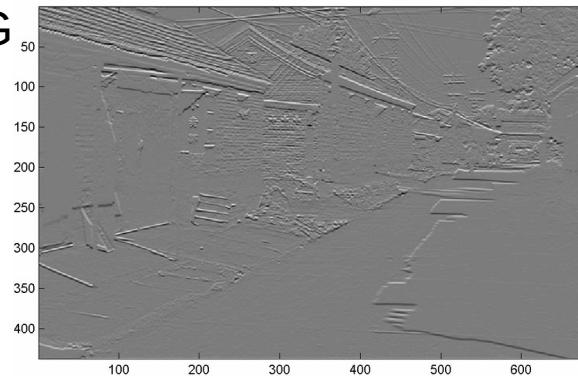
I



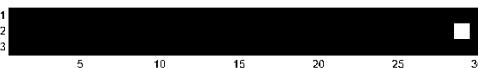
C



G

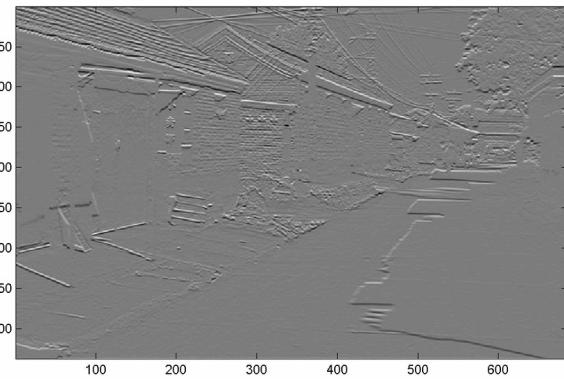


E



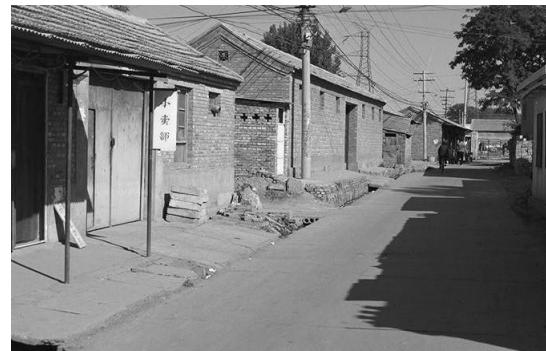
Exercise

- a) $G = D * B$ (some sort of horizontal detection)



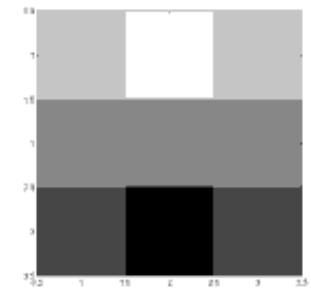
G

=



D

*



B

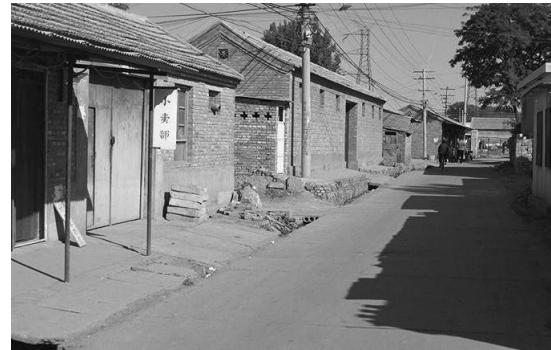
Exercise

b) $F = D * E$ - Shift to the left.



F

=



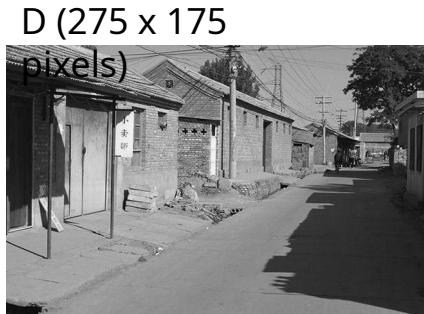
D



E

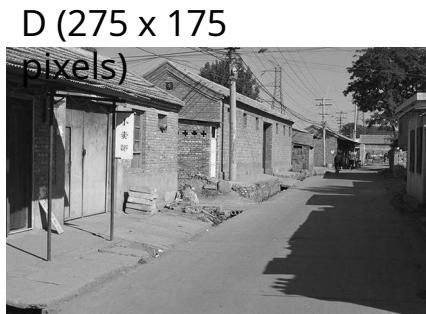
Exercise

$$c) \quad I = D * D$$



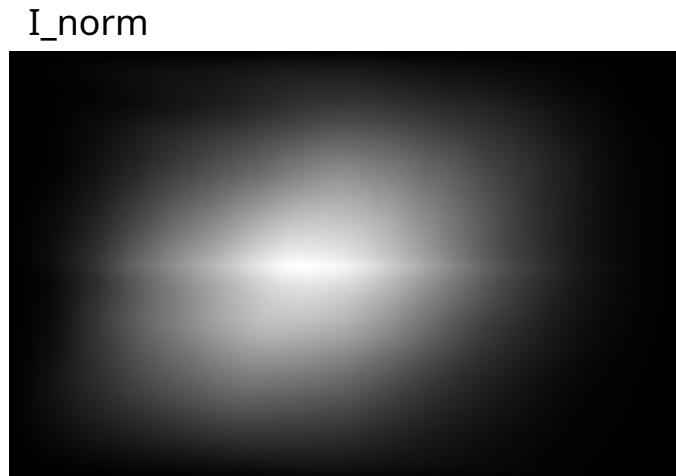
"...something to do with lack of content (black) at edges..."

$$I = D * D$$



```
>> D = img_as_float32( io.imread('convexample.png') )  
>> I = convolve2d( D, D )  
>> np.max(I)  
1.1021e+04
```

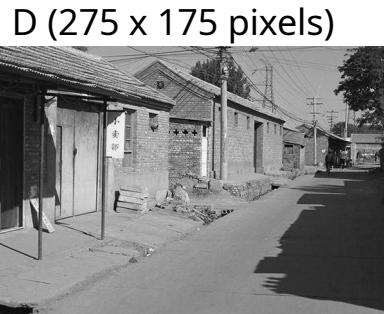
```
# Normalize for visualization  
>> I_norm = (I - np.min(I)) / (np.max(I) - np.min(I))  
>> plt.imshow( I_norm )
```



$$I = D * D$$



For x: $275 + (275-1)/2 + (275-1)/2 = 549$



I_norm (549 x 349 pixels)



Sklearn Convolve2d parameters

mode : str {‘full’, ‘valid’, ‘same’}, optional

A string indicating the size of the output:

full

The output is the full discrete linear convolution of the inputs. (Default)

valid

The output consists only of those elements that do not rely on the zero-padding. In ‘valid’ mode, either *in1* or *in2* must be at least as large as the other in every dimension.

same

The output is the same size as *in1*, centered with respect to the ‘full’ output.

boundary : str {‘fill’, ‘wrap’, ‘symm’}, optional

A flag indicating how to handle boundaries:

fill

pad input arrays with fillvalue. (default)

wrap

circular boundary conditions.

symm

symmetrical boundary conditions.

$$I = D * D$$

D (275 x 175 pixels)

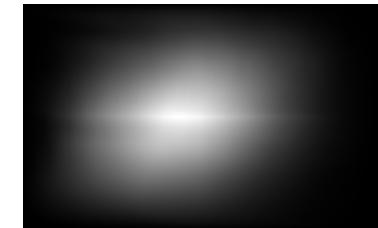


I (from slide – 275 x 175)



>> I = convolve2d(D, D, mode='full')

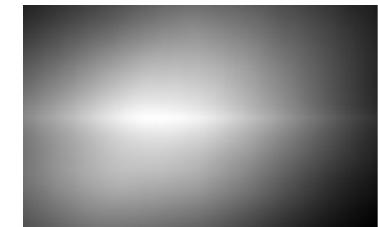
(Default; pad with zeros)



549 x 349

>> I = convolve2d(D, D, mode='same')

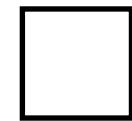
(Return same size as D)



275 x 175

>> I = convolve2d(D, D, mode='valid')

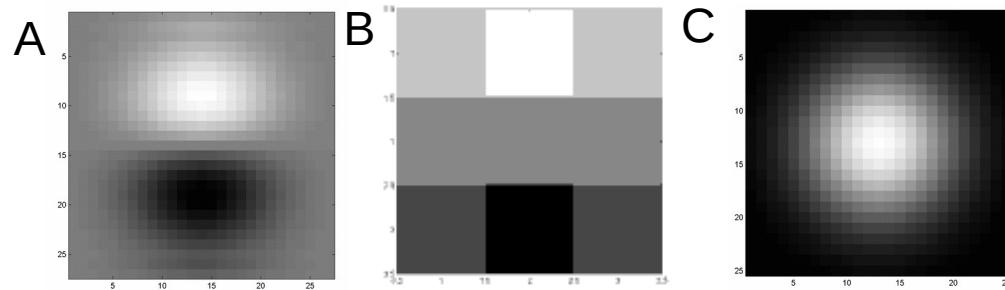
(No padding)



Value = 10528.3
1x1

Exercise

d) $A = B * C$ - “because it kind of looks like it.”



C is a Gaussian filter
(or something close to it),
and we know that it ‘blurs’.

Difference between Convolution and Correlation:

- **Convolution** is the response of one function when subjected to another function.
- **Correlation**: A measure of similarity between two functions. For example: ‘template matching’

For symmetric filters: use either convolution or correlation.

For nonsymmetric filters: correlation is template matching.

Filtering: Correlation and Convolution

- 2d correlation

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

e.g., `h = scipy.signal.correlate2d(f, I)`

- 2d convolution

$$h[m, n] = \sum_{k,l} f[k, l] I[m - k, n - l]$$

e.g., `h = scipy.signal.convolve2d(f, I)`

Convolution is the same as correlation with a 180° rotated filter kernel.

Correlation and convolution are identical when the filter kernel is symmetric.

D (275 x 175 pixels)



OK, so let's test this idea. Let's see if we can use correlation to 'find' the parts of the image that look like a template.

```
>> f = D[ 57:117, 107:167 ]
```

Expect response 'peak' in middle of I



f
61 x 61

```
>> I = correlate2d( D, f, 'same' )
```

OK, so let's test this idea. Let's see if we can use correlation to 'find' the parts of the image that look like a template.

```
>> f = D[ 57:117, 107:167 ]
```

Expect response 'peak' in middle of I

D (275 x 175 pixels)

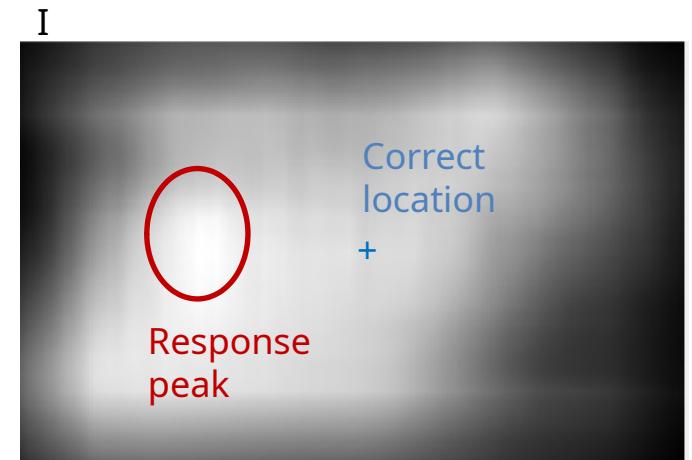


f
61 x 61

```
>> I = correlate2d( D, f, 'same' )
```

Hmm...

That didn't work – why not?



Correlation

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

e.g., `h = scipy.signal.correlate2d(f, I)`

As brightness in I increases, the response in h will increase, as long as f is positive.

Overall brighter regions will give higher correlation response -> not useful!

OK, so let's subtract the mean

```
>> f = D[ 57:117, 107:167 ]  
>> f2 = f - np.mean(f)  
>> D2 = D - np.mean(D)
```

Now zero centered.

*Score is higher only when dark parts
match and when light parts match.*

```
>> I2 = correlate2d( D2, f2, 'same' )
```

D2 (275 x 175)



f2
61 x 61

I2



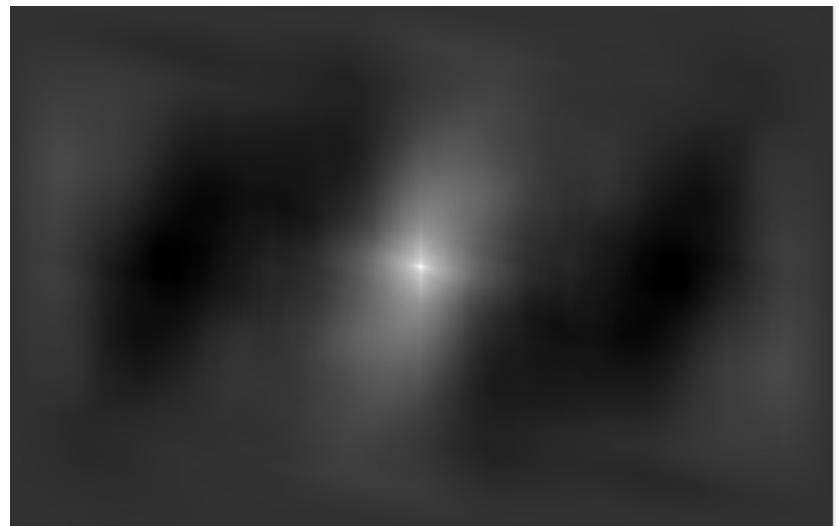
D2 (275 x 175)



Or even

```
>> I3 = correlate2d( D2, D2, 'full' )
```

I3



What happens with convolution?

```
>> f = D[ 57:117, 107:167 ]
```

```
>> f2 = f - np.mean(f)
```

```
>> D2 = D - np.mean(D)
```

```
>> I2 = convolve2d( D2, f2, 'same' )
```

D2 (275 x 175)



f2
61 x 61

What happens with convolution?

```
>> f = D[ 57:117, 107:167 ]
```

```
>> f2 = f - np.mean(f)
```

```
>> D2 = D - np.mean(D)
```

```
>> I2 = convolve2d( D2, f2, 'same' )
```

D2 (275 x 175)



f2
61 x 61

I2



Noise – Salt and Pepper Jack



Mean Jack – 3 x 3 filter



Very Mean Jack - 11 x 11 filter



NON-LINEAR FILTERS

Median filters

- Operates over a window by selecting the median intensity in the window.
- Part of the 'Rank' filter family, based on ordering of gray levels
 - E.G., min, max, range filters

Image filtering - mean

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$$I[.,.]$$

$$h[.,.]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30				

?

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

Credit: S. Seitz

Image filtering - mean

$$f[\cdot, \cdot] \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$$I[.,.]$$

$$h[.,.]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30					

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

Credit: S. Seitz

Median filter?

 $I[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

 $h[.,.]$

								?		

Median filters

- Operates over a window by selecting the median intensity in the window.
- What advantage does a median filter have over a mean filter?

Noise – Salt and Pepper Jack



Mean Jack – 3 x 3 filter



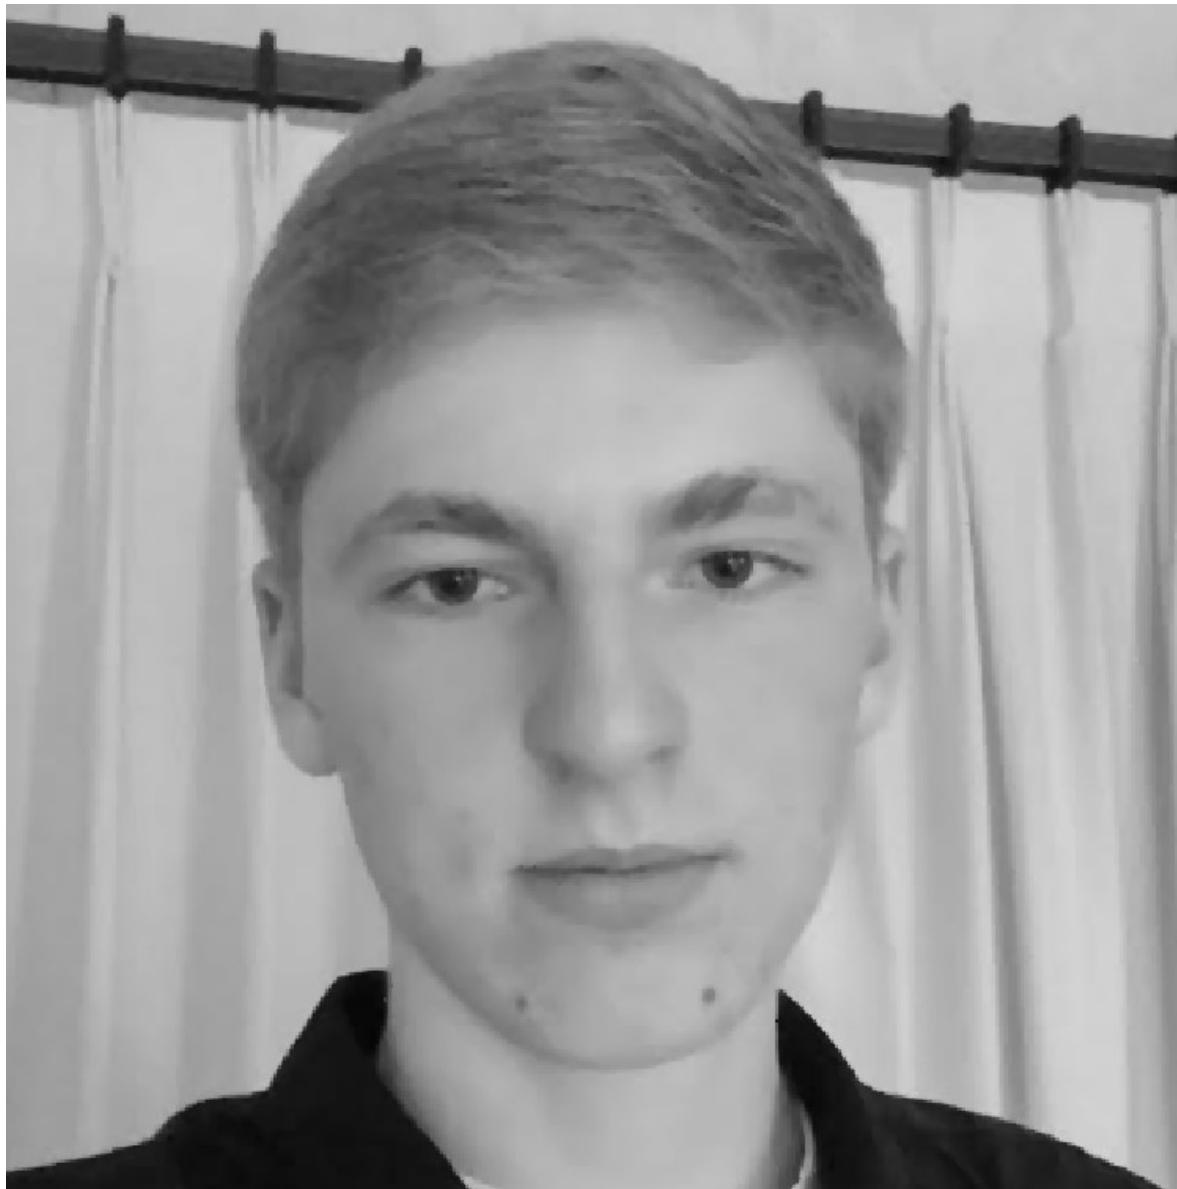
Very Mean Jack - 11 x 11 filter



Noisy Jack – Salt and Pepper



Median Jack – 3 x 3



Very Median Jack – 11 x 11



Median filters

- Operates over a window by selecting the median intensity in the window.
- What advantage does a median filter have over a mean filter?
- Is a median filter a kind of convolution?

Median filters

- Operates over a window by selecting the median intensity in the window.
- What advantage does a median filter have over a mean filter?
- Is a median filter a kind of convolution?

Interpretation: Median filtering is sorting.

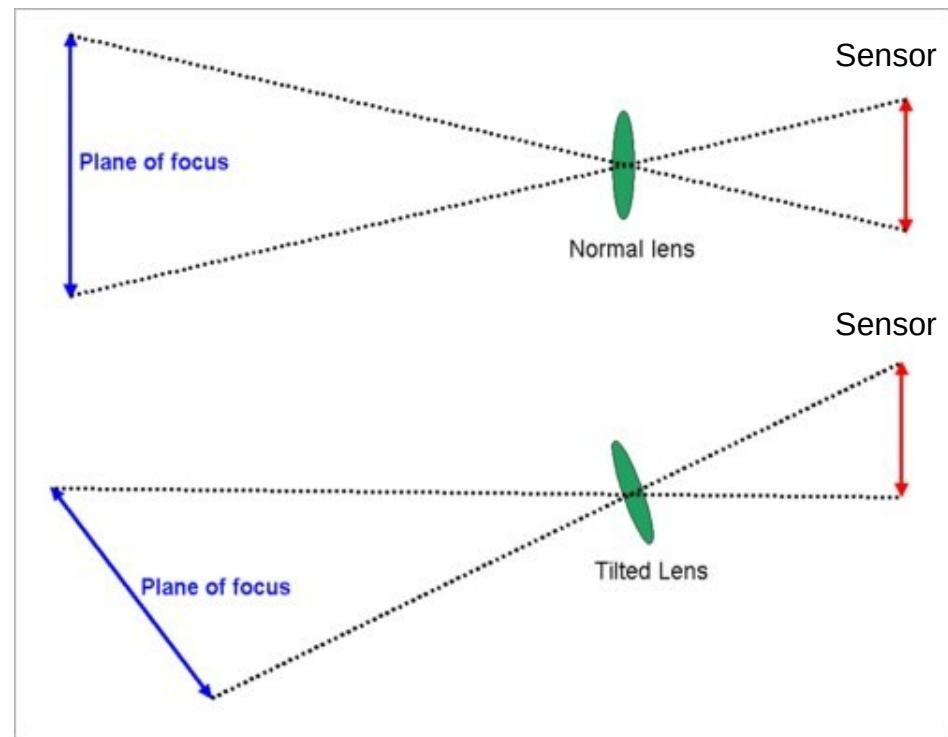
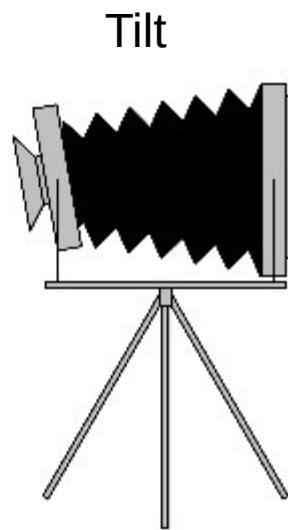
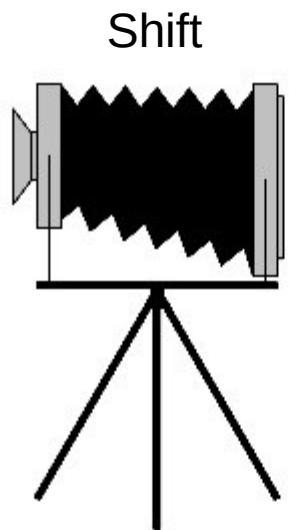


Ben Thomas

Tilt-shift photography



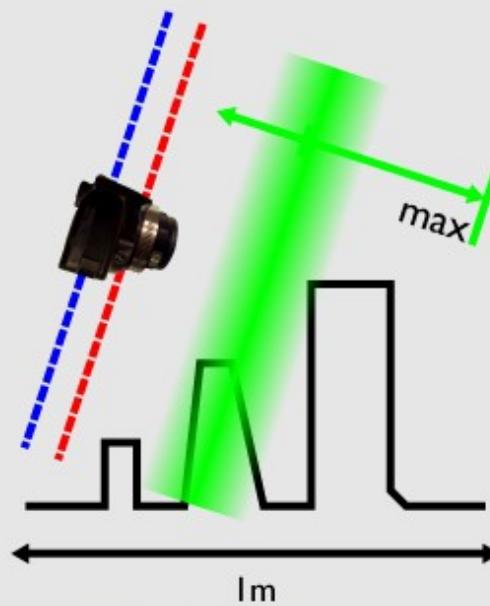
Tilt shift camera



Macro photography



Small scale in-focus volume

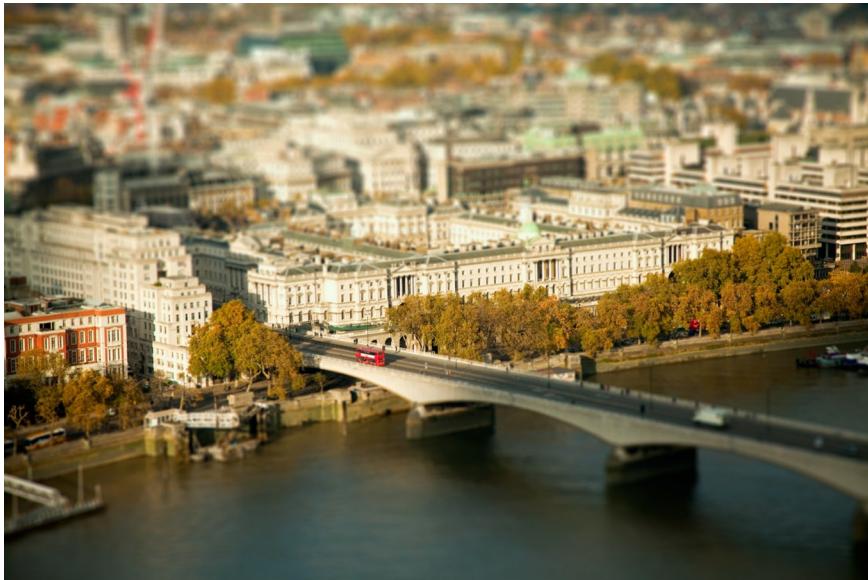


Can we fake tilt shift?

- We need to blur the image
 - OK, now we know how to do that.

Can we fake tilt shift?

- We need to blur the image
 - OK, now we know how to do that.
- We need to blur progressively more away from our ‘fake’ focal point



But can I make it look more like a toy?

- Boost saturation – toys are very colorful
- We'll learn how to do this when we discuss color
- For now: transform to Hue, Saturation, Value instead of RGB



Next class: Thinking in Frequency

