

Advanced Image Processing

Interest points and Corners

Filtering —→ Edges —→ Corners

Feature points


Also called interest points, key points, etc.
Often described as 'local' features.

Szeliski 4.1

Correspondence across views

- Correspondence: matching points, patches, edges, or regions across images.

Feature Matching



Fundamental to Applications

- Feature points are used for:
 - Image alignment
 - 3D reconstruction
 - Motion tracking (robots, drones,
 - Indexing and database retrieval
 - Object recognition
 - ...

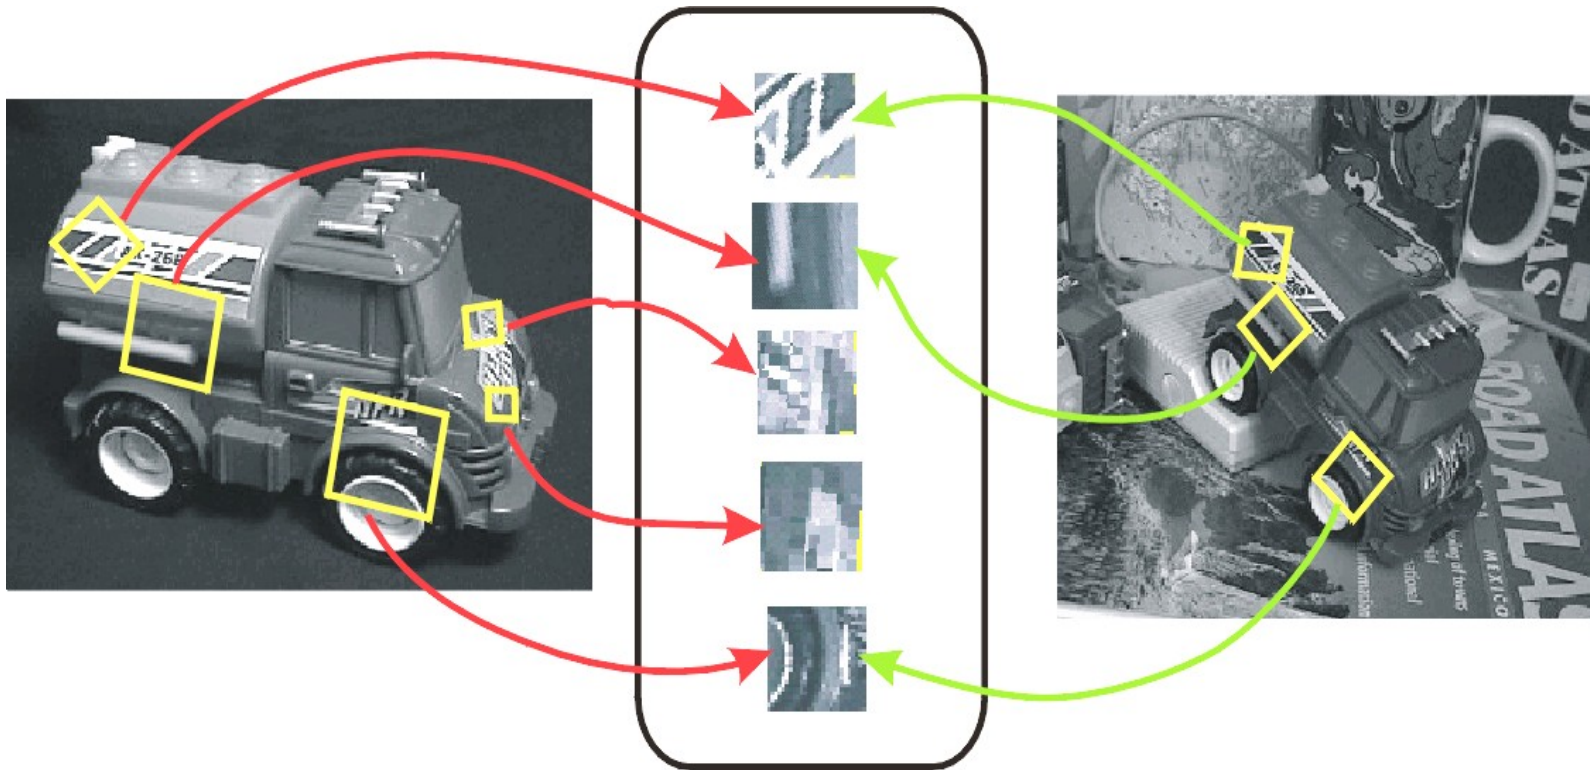


Example: Invariant Local Features

Detect points that are *repeatable* and *distinctive*.

I.E., invariant to image transformations:

- appearance variation (brightness, illumination)
- geometric variation (translation, rotation, scale).



Keypoint Descriptors

Example application


Panorama stitching

We have two images – how do we combine them?




Local features: main components



- 1) Detection:
Find a set of distinctive key points.




- 2) Description:
Extract feature descriptor around each interest point as vector.




$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$




$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$



- 3) Matching: $d(\mathbf{x}_1, \mathbf{x}_2) < T$
Compute distance between feature vectors to find correspondence.



Characteristics of good features



Repeatability

The same feature can be found in several images despite geometric and photometric transformations

Saliency

Each feature is distinctive

Compactness and efficiency

Many fewer features than image pixels

Locality

A feature occupies a relatively small area of the image; robust to clutter and occlusion

Goal: interest operator repeatability

We want to detect (at least some of) the same points in both images.




With these points, there's no chance to find true matches!

Yet we have to be able to run the detection procedure *independently* per image.

Goal: descriptor distinctiveness

We want to be able to reliably determine which point goes with which.



Must provide some invariance to geometric and photometric differences between the two views.

Local features: main components

1) Detection:

Find a set of distinctive key points.



2) Description:

Extract feature descriptor around each interest point as vector.

3) Matching:

Compute distance between feature vectors to find correspondence.

Detection: Basic Idea

We do not know which other image locations the feature will end up being matched against.


But we can compute how stable a location is in appearance with respect to small variations in position u .

Compare image patch against local neighbors.


Corner Detection: Basic Idea

We might recognize the point by looking through a small window.


We want a window shift in *any direction* to give *a large change* in intensity.



“Flat” region:
no change in
all directions



“Edge”:
no change
along the edge
direction



“Corner”:
significant
change in all
directions

Corner Detection by Auto-correlation


Change in appearance of window $w(x,y)$ for shift $[u,v]$:

$$E(u,v) = \sum_{x,y} w(x,y) \left[I(x+u, y+v) - I(x,y) \right]^2$$

Diagram illustrating the components of the equation:


- Window function**: Points to $w(x,y)$ in the equation.
- Shifted intensity**: Points to $I(x+u, y+v)$ in the equation.
- Intensity**: Points to $I(x,y)$ in the equation.

Window function $w(x,y) =$



1 in window, 0 outside

or





Gaussian

Corner Detection by Auto-correlation

Change in appearance of window $w(x,y)$ for shift $[u,v]$:



$$E(u, v) = \sum_{x,y} w(x, y) \left[I(x+u, y+v) - I(x, y) \right]^2$$



Corner Detection by Auto-correlation

Change in appearance of window $w(x,y)$ for shift $[u,v]$:

$$E(u, v) = \sum_{x,y} w(x, y) \left[I(x+u, y+v) - I(x, y) \right]^2$$






$$E(u,v) = \sum_{x,y} w(x,y) \left[I(x+u, y+v) - I(x,y) \right]^2$$

Think-Pair-Share:




Correspond the three red crosses to (b,c,d).



$E(u,v)$



$E(u,v)$



As a surface

Corner Detection by Auto-correlation

Change in appearance of window $w(x,y)$ for shift $[u,v]$:


$$E(u,v) = \sum_{x,y} w(x,y) \left[I(x+u, y+v) - I(x,y) \right]^2$$

We want to discover how E behaves for small shifts

But this is very slow to compute naively.

$O(\text{window_width}^2 * \text{shift_range}^2 * \text{image_width}^2)$

$O(11^2 * 11^2 * 600^2) = 5.2$ billion of these
14.6 thousand per pixel in your image





Corner Detection by Auto-correlation

Change in appearance of window $w(x,y)$ for shift $[u,v]$:



$$E(u,v) = \sum_{x,y} w(x,y) \left[I(x+u, y+v) - I(x,y) \right]^2$$

We want to discover how E behaves for small shifts


But we know the response in E that we are looking for – strong peak.



Can we just approximate $E(u,v)$ locally by a quadratic surface?



\approx




Recall: Taylor series expansion

A function f can be represented by an infinite series of its derivatives at a single point a :

$$f(a) + \frac{f'(a)}{1!}(x-a) + \frac{f''(a)}{2!}(x-a)^2 + \frac{f'''(a)}{3!}(x-a)^3 + \dots$$

Approximation of
 $f(x) = e^x$
centered at $f(0)$



Taylor series approximation (2D)

$$f(x+u, y+v) = f(x, y) + uf_x(x, y) + vf_y(x, y) +$$

First partial derivatives

$$\frac{1}{2!} [u^2 f_{xx}(x, y) + uv f_{xy}(x, y) + v^2 f_{yy}(x, y)] +$$

Second partial derivatives

$$\frac{1}{3!} [u^3 f_{xxx}(x, y) + u^2 v f_{xxy}(x, y) + uv^2 f_{xyy}(x, y) + v^3 f_{yyy}(x, y)]$$

Third partial derivatives

+ ... (Higher order terms)

Robert Collins
CSE486, Penn State

First order approx

$$f(x+u, y+v) \approx f(x, y) + uf_x(x, y) + vf_y(x, y)$$

Corner Detection: Mathematics

Robert Collins
CSE486, Penn State

$$\sum [I(x+u, y+v) - I(x, y)]^2$$

$$\approx \sum [I(x, y) + uI_x + vI_y - I(x, y)]^2 \quad \text{First order approx}$$

$$= \sum u^2 I_x^2 + 2uv I_x I_y + v^2 I_y^2$$

$$= \sum \begin{bmatrix} u & v \end{bmatrix} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} \quad \text{Rewrite as matrix equation}$$

$$= \begin{bmatrix} u & v \end{bmatrix} \left(\sum \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \right) \begin{bmatrix} u \\ v \end{bmatrix}$$

Corner Detection: Mathematics

The quadratic approximation simplifies to

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a *second moment matrix* computed from image derivatives:



$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y] = \sum \nabla I (\nabla I)^T$$

Corners as distinctive interest points

$$M = \sum w(x, y) \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix}$$

2 x 2 matrix of image derivatives
(averaged in neighborhood of a point)




Notation:

$$I_x \Leftrightarrow \frac{\partial I}{\partial x}$$

$$I_y \Leftrightarrow \frac{\partial I}{\partial y}$$

$$I_x I_y \Leftrightarrow \frac{\partial I}{\partial x} \frac{\partial I}{\partial y}$$




Example: 2D Gradients



Robert Collins
CSE486, Penn State

Example: 2D Gradients - Distribution




The distribution of the x and y derivatives is very different for all three types of patches



Example: 2D Gradients - Distribution

Robert Collins
CSE486, Penn State

The distribution of x and y derivatives can be characterized by the shape and size of the principal component ellipse




Interpreting the second moment matrix

The surface $E(u,v)$ is locally approximated by a quadratic form. Let's try to understand its shape.


$$E(u,v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$



Interpreting the second moment matrix


Consider a horizontal “slice” of $E(u, v)$: $[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$



Visualization of second moment matrices



Visualization of second moment matrices



Interpreting the second moment matrix


Consider a horizontal “slice” of $E(u, v)$: $[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$

This is the equation of an ellipse.


Diagonalization of M : $M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$

The axis lengths of the ellipse are determined by the eigenvalues, and the orientation is determined by a rotation matrix .

*Note inverse relationship:
larger eigenvalue =
steeper slope; smaller
ellipse in visualization*




Interpreting the second moment matrix



Make decision based on
image structure tensor

$$S_{ij} = \sum_{m=i-D}^{i+D} \sum_{n=j-D}^{j+D} w_{mn} \begin{bmatrix} h_{ij}^2 & h_{ij}v_{ij} \\ h_{ij}v_{ij} & v_{ij}^2 \end{bmatrix}$$

Classification of image points using eigenvalues of M




Classification of image points using eigenvalues of M

Cornerness

$$C = \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2$$

α : constant (0.04 to 0.06)



Classification of image points using eigenvalues of M

Cornerness

$$C = \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2$$


α : constant (0.04 to 0.06)

Remember your linear algebra:



Determinant: $\det(A) = \prod_{i=1}^n \lambda_i = \lambda_1 \lambda_2 \cdots \lambda_n$.

Trace: $\text{tr}(A) = \sum_i \lambda_i$.

$$C = \det(M) - \alpha \text{trace}(M)^2$$





Cornerness Response Example



Harris c score.


I_x , I_y computed using Sobel operator
Windowing function w = Gaussian, $\sigma=1$

CornerRadius Response Example




Threshold: $-10000 < c < 10000$
(neither edges nor corners)

Cornerness Response Example



Threshold: $c < -10000$
(edges)

Cornerness Response Example



Threshold: > 10000
(corners)

Harris corner detector

- 1) Compute M matrix for each window to recover a *cornerness* score .
 - Note: We can find M purely from the per-pixel image derivatives!
- 2) Threshold to find pixels which give large corner response ($>$ threshold).
- 3) Find the local maxima pixels, i.e., suppress non-maxima.


C.Harris and M.Stephens. "A Combined Corner and Edge Detector."
Proceedings of the 4th Alvey Vision Conference: pages 147—151, 1988.

Harris Detector: Steps




Harris Detector: Steps

Compute corner response



Harris Detector: Steps

Find points with large corner response: $> \text{threshold}$



Harris Detector: Steps

Take only the points of local maxima of




Harris Detector: Steps



Invariance and covariance

Invariance (does not change in different transformed images): e.g. the cornerness response value does not change when the image is transformed (rotated – translated)

Covariance (changes homogeneously such that we can still identify it using some process): if we have two transformed versions of the same image, the location and orientation of the feature are said to be covariant



Ellipse rotates but its shape (i.e. eigenvalues) remains the same

Corner response R is invariant w.r.t. rotation and corner location and orientation are covariant