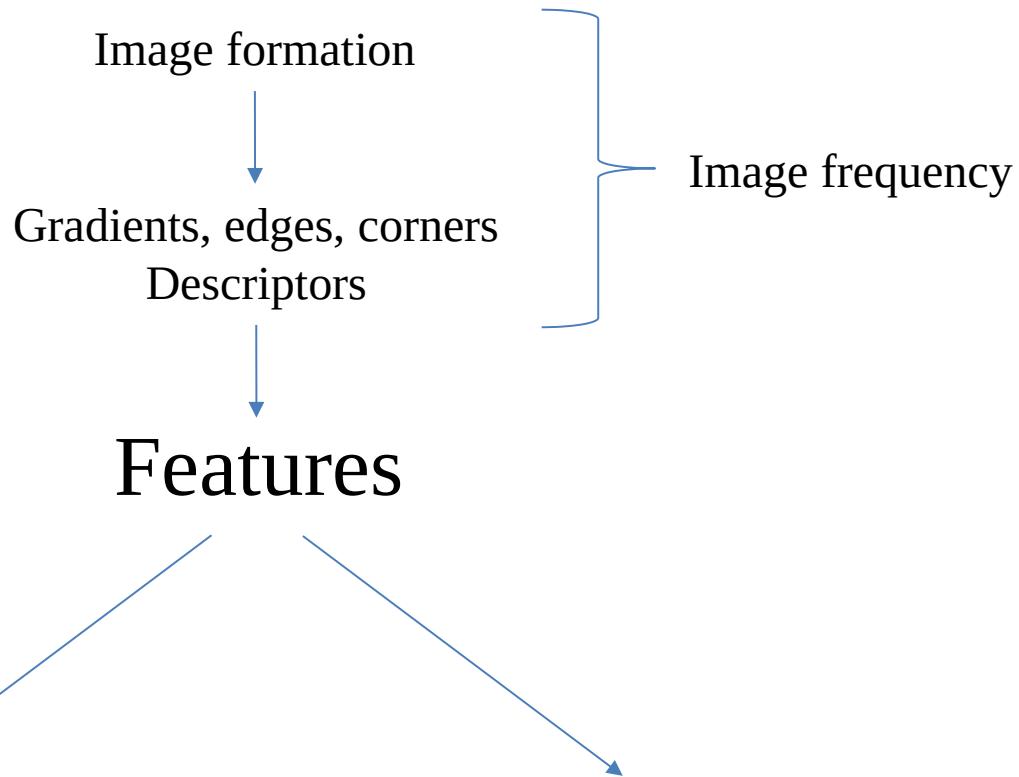


# Advanced Image Processing

## ML: Unsupervised Learning

# Where to go from our basic building block?



**Recognition**

Scenes, places, objects,  
~ 5 weeks (inc. CNNs)

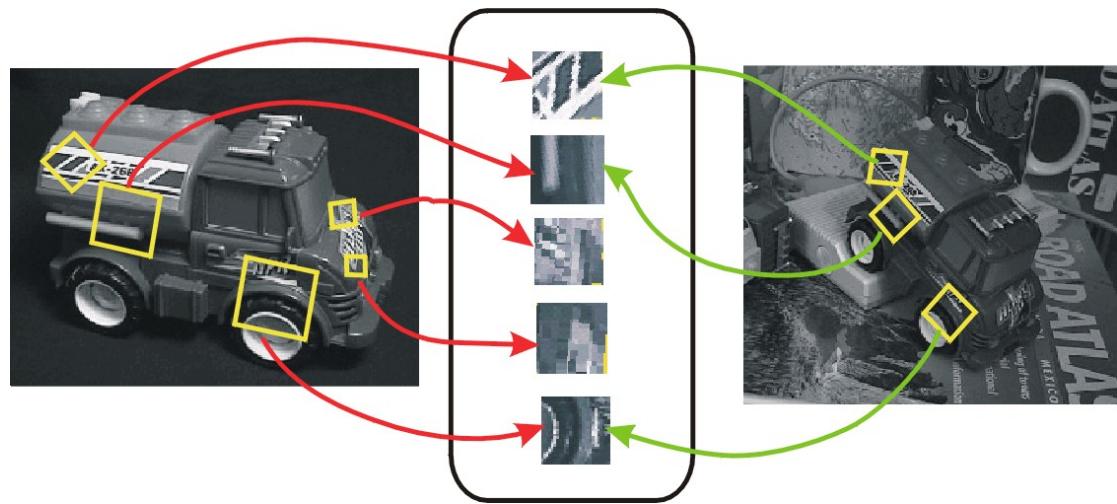
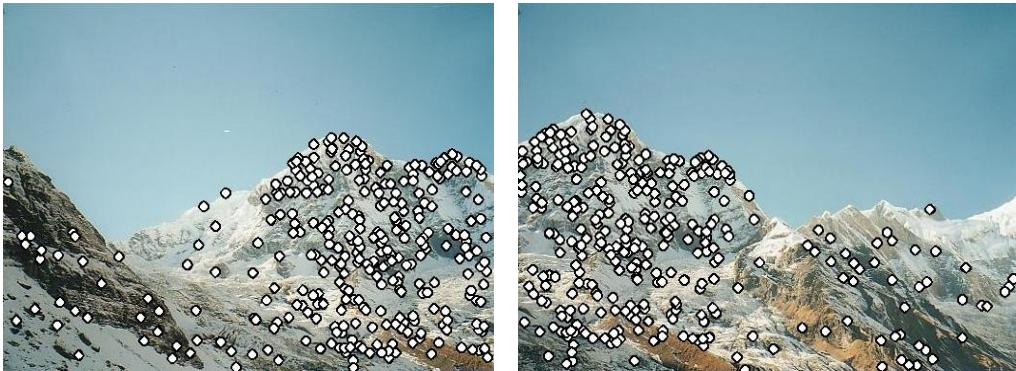
**Reconstruction**

Geometric understanding  
Last topic in class

# Panorama stitching / instance recognition

---

Needs more geometric understanding...  
...but we'll see it later on.



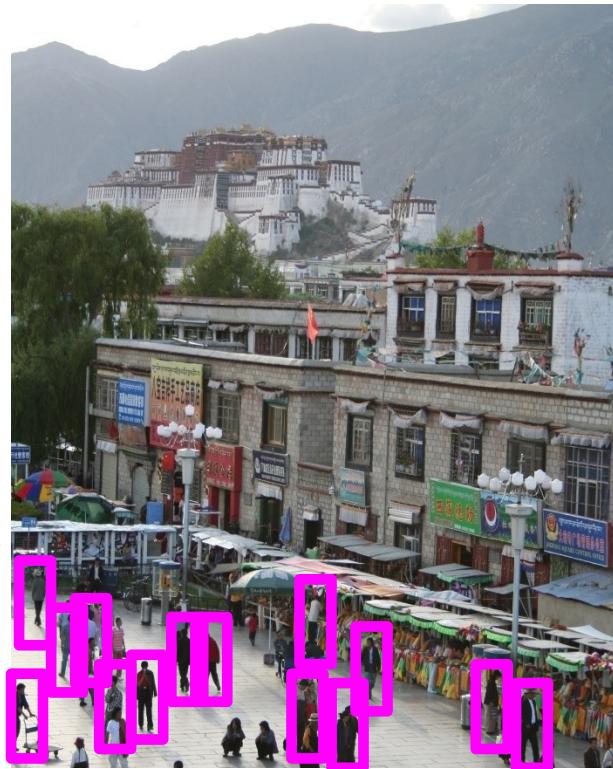
# Recognition

---

Often needs machine learning  
for compact descriptions of the visual world.



**Scene recognition**  
- City/forest/factory/...



**Find pedestrians**

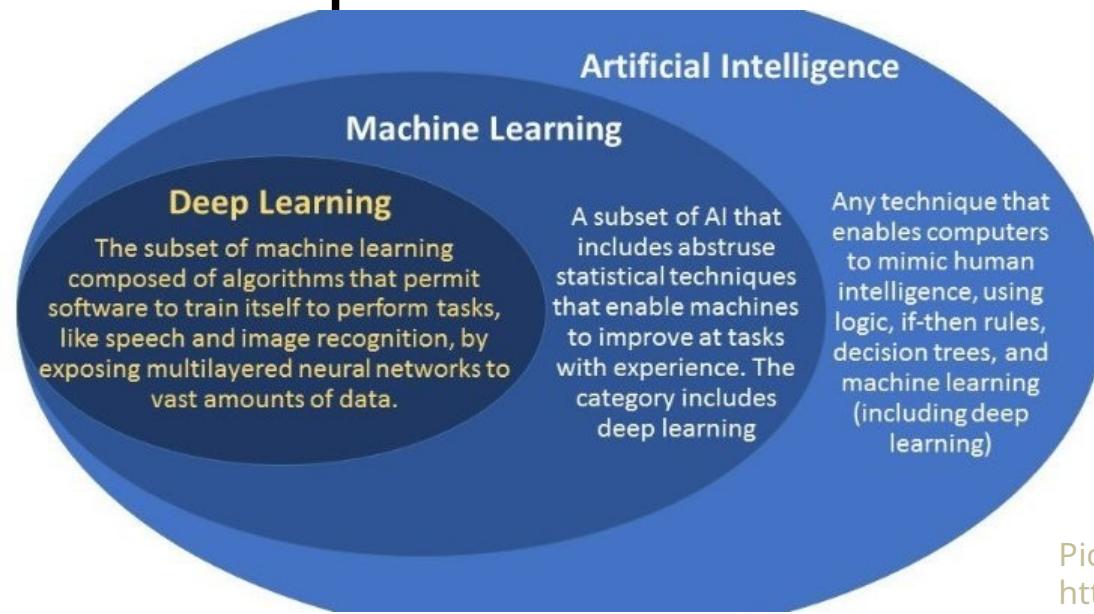
# **ML CRASH COURSE**

# Our approach

- We will look at ML as a tool. We will not detail the underpinnings of each learning method.
- Please take a machine learning course if you want to know more. Alternatively there are quite a few good reference books on the topic.

# Machine Learning

- Learn from and make predictions on data.
- Arguably the greatest export from computing to other scientific fields.
- Statisticians might disagree with Computer Scientists on the true

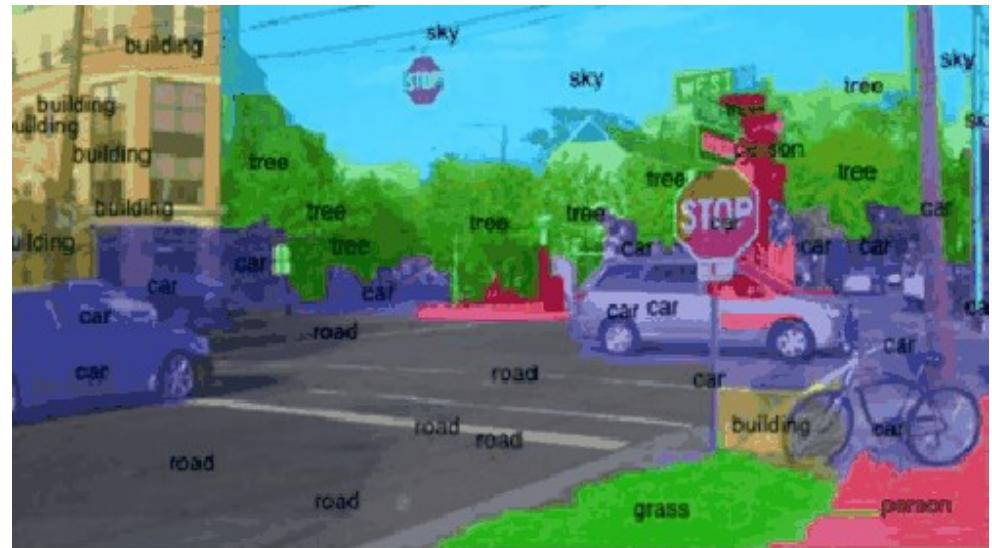
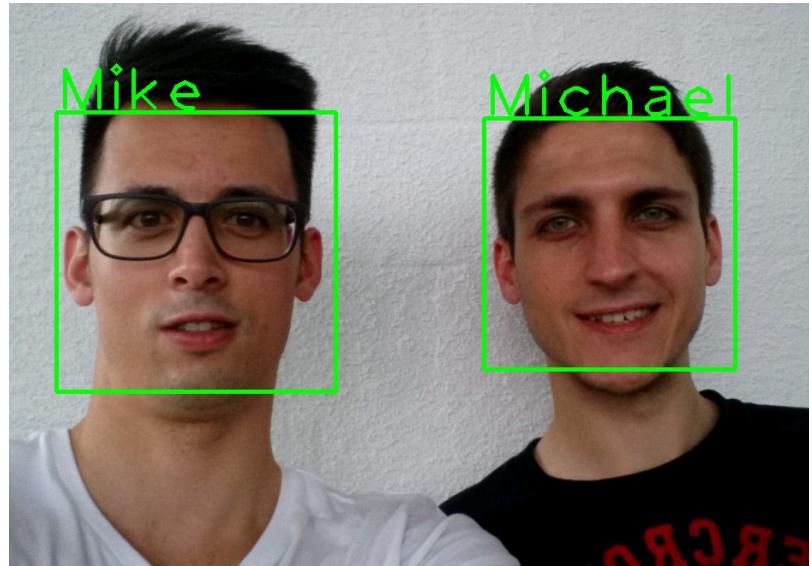


Picture credit:

<https://twitter.com/deeplearn007/status/95283364338637>

# ML for Computer Vision

- Face Recognition
- Object Classification
- Scene Segmentation



# Data, data, data!

Peter Norvig – “The Unreasonable Effectiveness of Data” (IEEE Intelligent Systems, 2009)

- “... invariably, simple models and a lot of data trump more elaborate models based on less data”

[Norvig is Director of Research @ Google, and a Brown APMA alum!]

# ImageNet

- Images for nouns in WordNet
- 1000 classes
- 1.2mil images
- 100k test
- Top 5 error

The screenshot shows the ImageNet website interface. At the top, there's a navigation bar with the IMGENET logo, a search bar containing '14,197,122 images, 21841 synsets indexed', and links for 'SEARCH', 'Home', 'About', 'Explore', and 'Download'. Below the navigation is a status message 'Not logged in. Login | Signup'.

The main content area displays a search result for 'Big cat, cat'. The title is 'Big cat, cat' with a subtitle 'Any of several large cats typically able to roar and living in the wild'. To the right, there are statistics: '1404 pictures', '93.35% Popularity Percentile', and 'Wordnet IDs'. Below this, there are three tabs: 'Treemap Visualization' (which is selected), 'Images of the Synset', and 'Downloads'.

The 'Treemap Visualization' section shows a hierarchical tree of categories under 'ImageNet 2011 Fall Release (2184)'. The tree starts with 'animal, animate being, beast, bramate (0)', then branches into 'chordate (2953)' (selected), 'vertebrate, craniate (294)', 'mammal, mammalian (367)', 'metatherian (36)', 'fossiliferous mammal (0)', 'placental, placental (362)', 'livestock, stock, hyrax, coney, co Ungulata (0)', 'bat, chiropteran (8)', 'pachyderm (8)', 'pangolin, scaly a digitigrade mami carnivorae (362)', 'bear (11)', 'musteline ma procyonid (8)', 'viverrine, vivi canine, canid fissiped mami feline, felid (5 big cat, ca sabre-t lion, kin liger (0)

The 'Images of the Synset' section displays a grid of images for various sub-categories: Tiger, Leopard, Snow, Jaguar, Lion, Cheetah, Saber-toothed, Tigon, and Liger. Each category has a grid of small images representing different sub-classes or individual instances.

[imagenet](#)

IMAGENET





mite



container ship



motor scooter



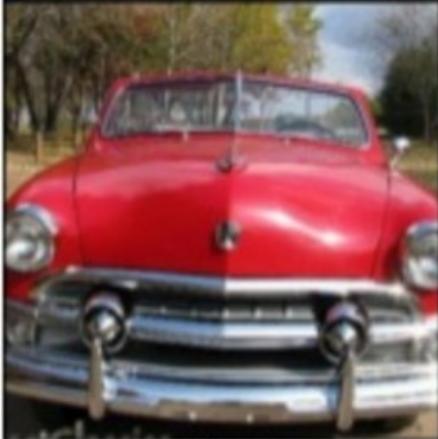
leopard

mite
black widow
cockroach
tick
starfish

container ship
lifeboat
amphibian
fireboat
drilling platform

motor scooter
go-kart
moped
bumper car
golfcart

leopard
jaguar
cheetah
snow leopard
Egyptian cat



grille



mushroom



cherry



Madagascar cat

convertible
grille
pickup
beach wagon
fire engine

agaric
mushroom
jelly fungus
gill fungus
dead-man's-fingers

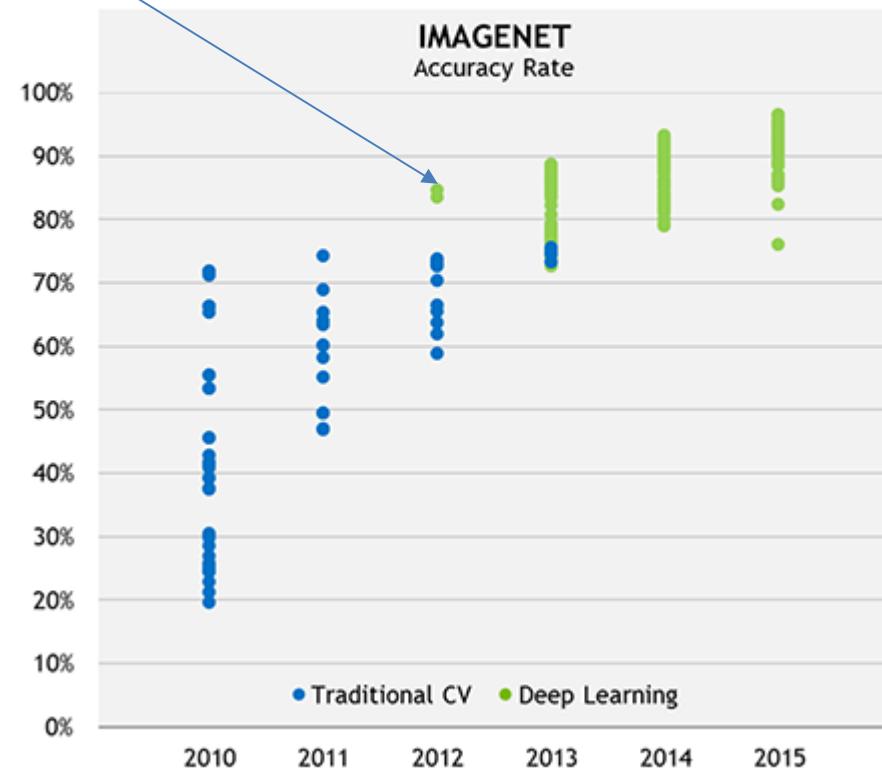
dalmatian
grape
elderberry
ffordshire bulterrier
currant

squirrel monkey
spider monkey
titi
indri
howler monkey

# ImageNet Competition

- Krizhevsky,  
2012
- Google,  
Microsoft  
2015
  - Beat the best  
human score  
in the  
ImageNet  
challenge.

2015: A MILESTONE YEAR  
IN COMPUTER SCIENCE



# Machine Learning Problems

	<i>Supervised Learning</i>	<i>Unsupervised Learning</i>
<i>Discrete</i>	classification or categorization	clustering
<i>Continuous</i>	regression	dimensionality reduction

... But not exactly

# Supervised vs Unsupervised

- Supervised
  - Input: labelled training data
  - finds relationship between input and output to predict a label/value for new data
- Unsupervised
  - Input: data
  - finds meaningful pattern in the data
- Semi-supervised?

# Machine Learning Problems

*Discrete*

*Supervised Learning*

classification or  
categorization

*Unsupervised Learning*

clustering

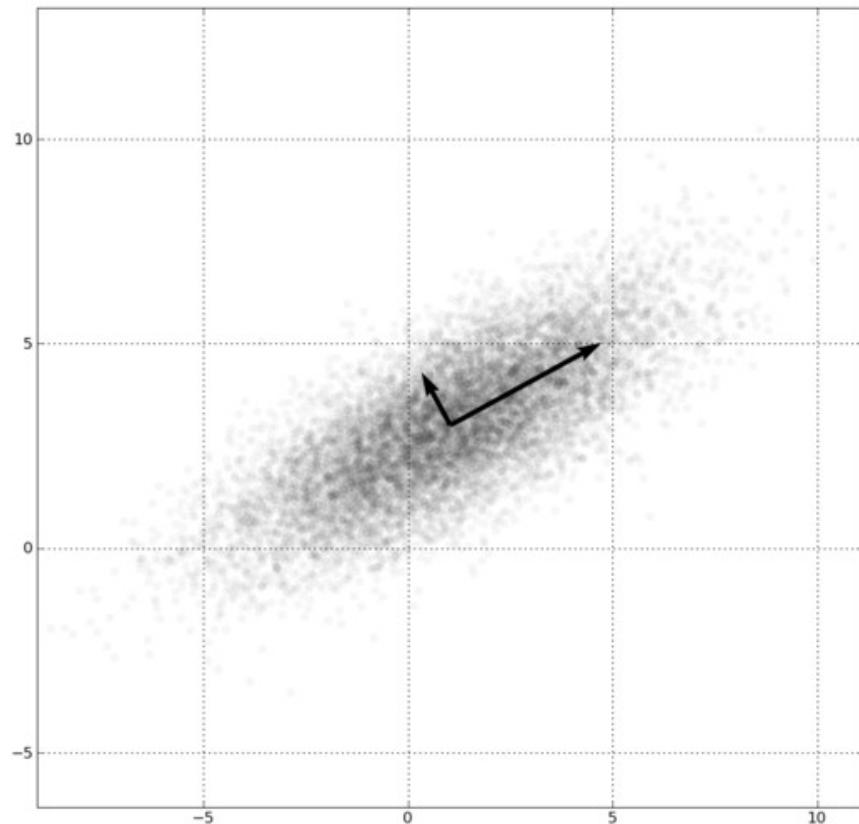
*Continuous*

regression

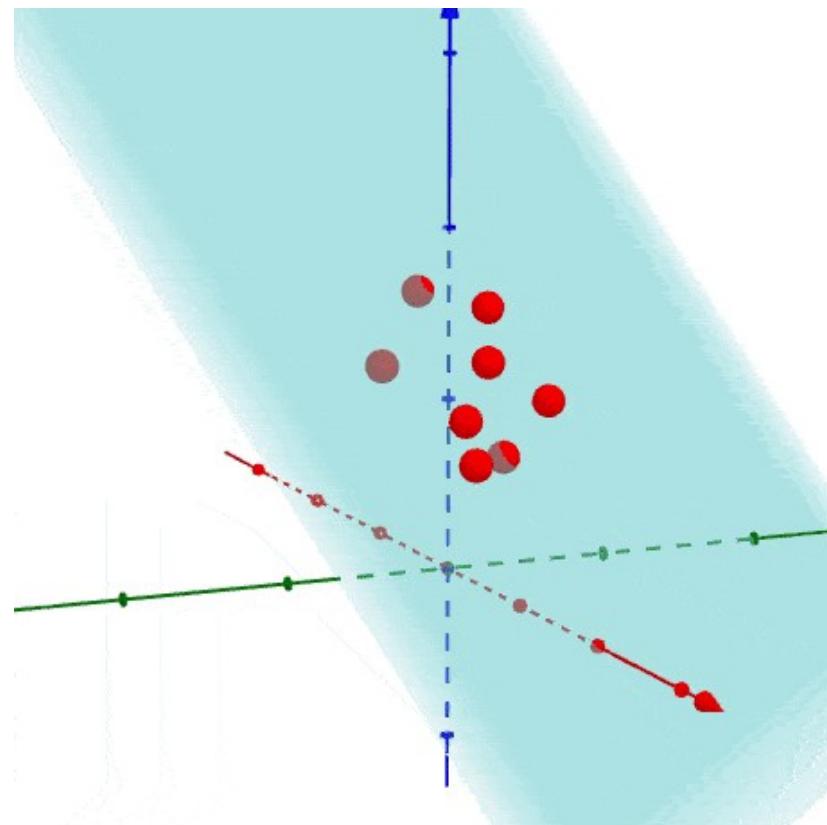
dimensionality  
reduction

# Dimensionality Reduction

- **PCA, ICA, LLE, Isomap**
- Principal component analysis
  - Creates a basis where the axes represent the dimensions of variance, from high to low.
  - Finds correlations in data dimensions to produce *best possible* lower-dimensional representation based on linear projections.



# PCA



# PCA pseudocode

1) Compute data mean.

for i in feature dimensions:

    inMean(i) = np.mean(featureVector[i,:])

2) Optional(but not really optional): align data so that mean = 0 and variance = 1

3) Compute covariance matrix: (square matrix with size FeatSizeXFeatSize

    Covar = np.cov(FeatVec[0,:], FeatVec[1,:] , ...)

4) Compute eigen value decomposition

    eigVal, eigVec = np.linalg.eig(cov\_mat)

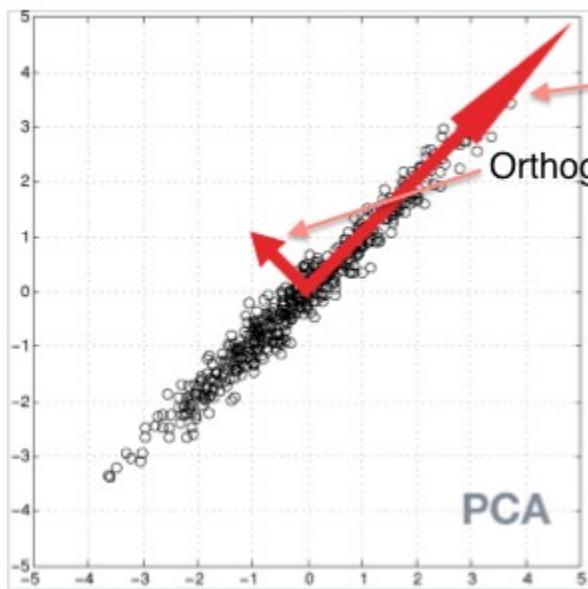
5) Sort eigVec by decreasing order of eigVal

6) Form a transformation matrix 'w' by keeping the first K eigVec such that it's a matrix of size: d x K

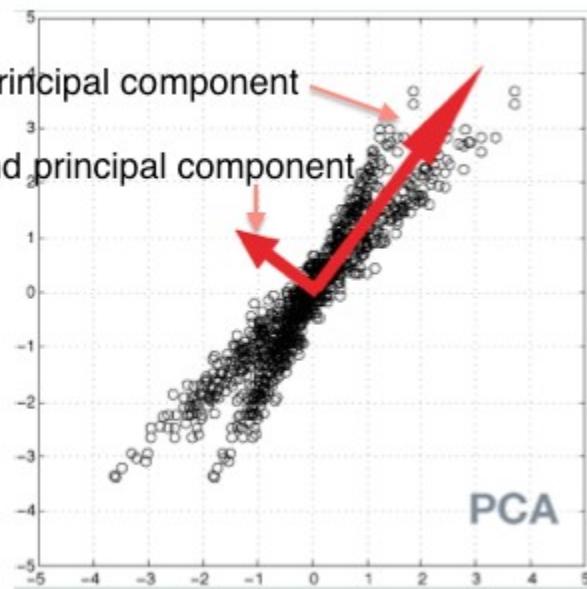
7) Transform the data using:

# PCA

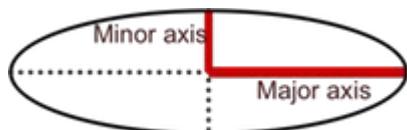
**A**



**B**



(Figure adapted from C. Beckmann, Oxford FMRIB)



# PCA

- Using sklearn:

```
from sklearn.decomposition import PCA  
sklearn_pca = sklearnPCA(k) # k=nbr dim to keep  
Sklearn_trnf = sklearn.pca.fit_transform(data)
```

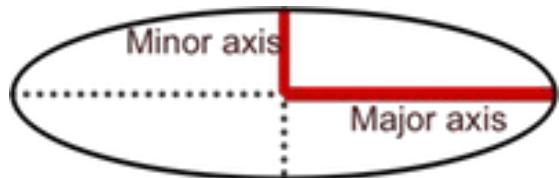
- Eigen Decomposition trouble:

$$\Sigma v = \lambda v$$

$\Sigma$ =Covariance matrix

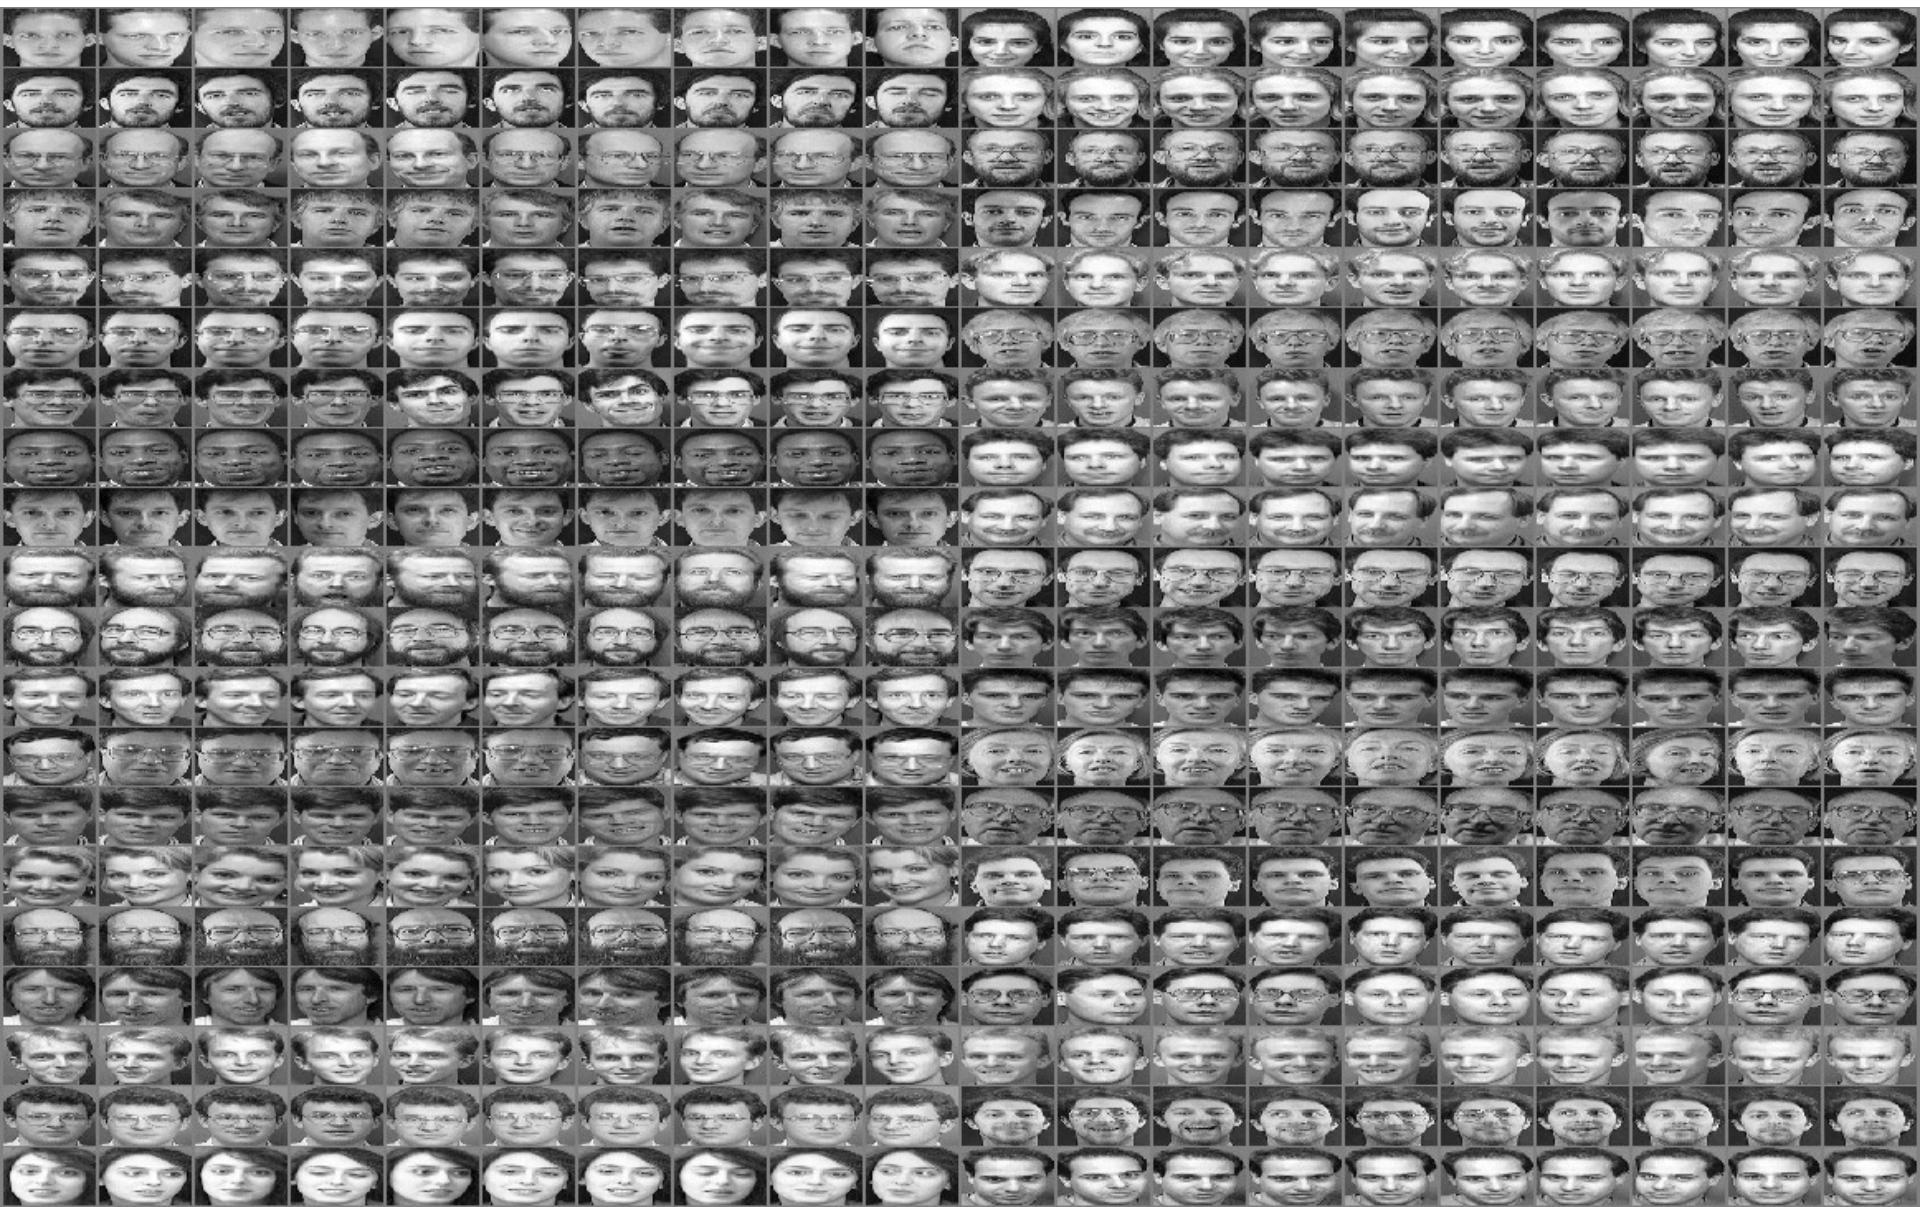
$v$ =Eigenvector

$\lambda$ =Eigenvalue



# Eigenfaces

*The ATT face database (formerly the ORL database), 10 pictures of 40 subjects each*



# Eigenfaces



Mean  
face



Basis of variance  
(eigenvectors)

M. Turk; A. Pentland (1991). ["Face recognition using eigenfaces"](#) (PDF).  
Proc. IEEE Conference on Computer Vision and Pattern Recognition. pp. 586–591.

R.P.W.

# Machine Learning Problems

*Discrete*

*Supervised Learning*

classification or  
categorization

*Unsupervised Learning*

clustering

*Continuous*

regression

dimensionality  
reduction

# Unsupervised learning Gestalt factors



Not grouped



Proximity



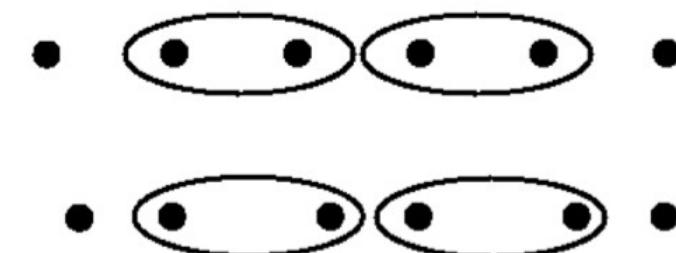
Similarity



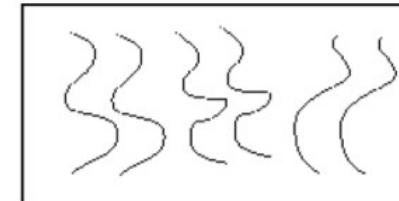
Similarity



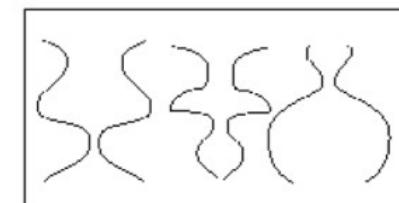
Common Fate



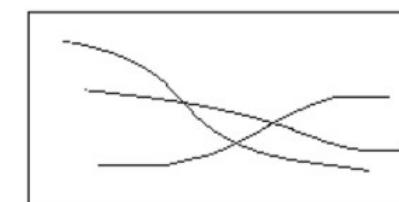
Common Region



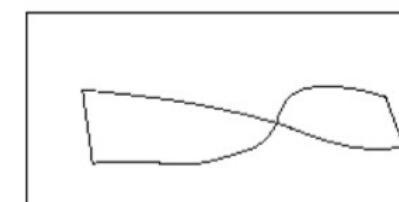
Parallelism



Symmetry



Continuity



Closure



# The United States *redrawn as* Fifty States *with Equal Population*







CANADA  
WITH THIRTEEN  
PROVINCES OF  
EQUAL POPULATION



Source:  
alasdairgunn

# Clustering example: image segmentation

Goal: Break up the image into meaningful or perceptually similar regions



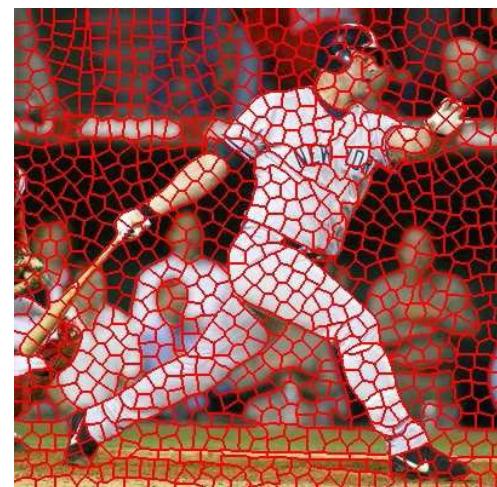
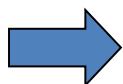
# Segmentation for feature support or efficiency



50x50  
Patch



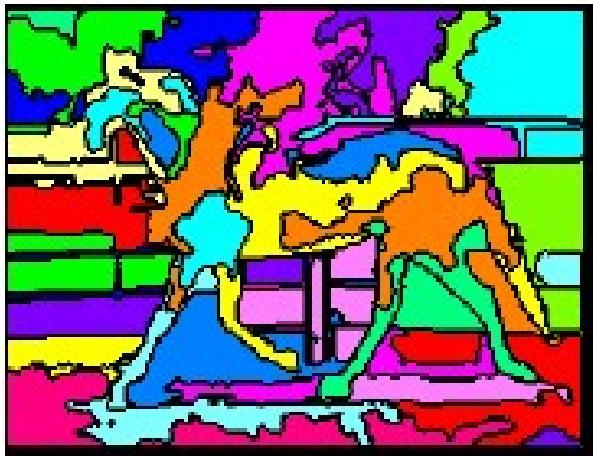
[Felzenszwalb and Huttenlocher 2004]



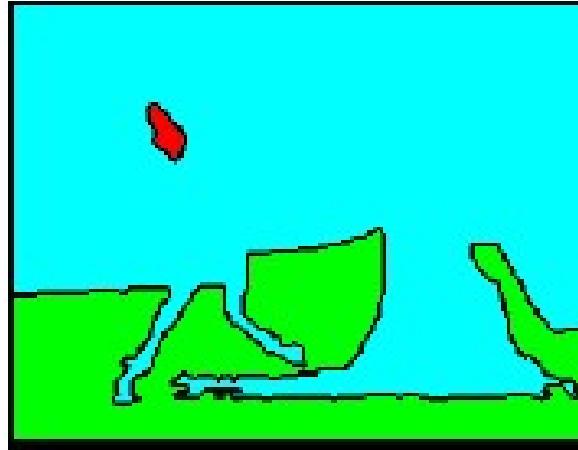
[Shi and Malik 2001]

Superpixel  
s!

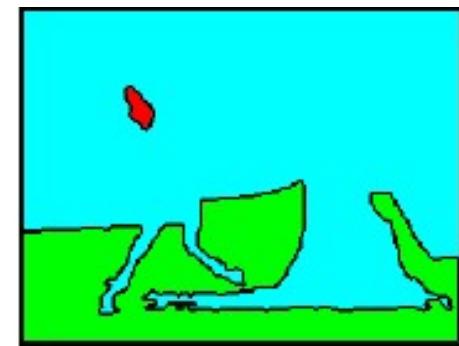
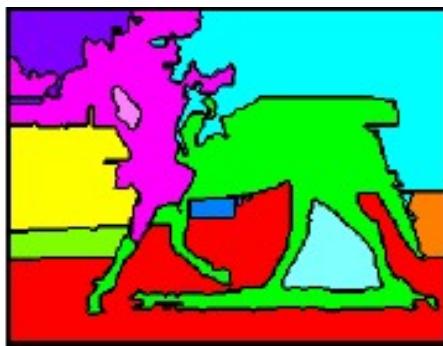
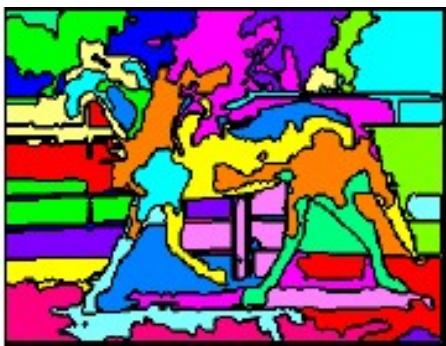
# Types of Segmentations



Oversegmentation



Undersegmentation



Hierarchical Segmentations

# Clustering

Group together similar ‘points’ and represent them with a single token.

Key Challenges:

- 1) What makes two points/images/patches similar?
- 2) How do we compute an overall grouping from pairwise similarities?

# Why do we cluster?

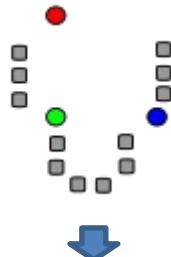
- **Summarizing data**
  - Look at large amounts of data
  - Patch-based compression or denoising
  - Represent a large continuous vector with the cluster number
- **Counting**
  - Histograms of texture, color, SIFT vectors
- **Segmentation**
  - Separate the image into different regions
- **Prediction**
  - Images in the same cluster may have the same labels

# How do we cluster?

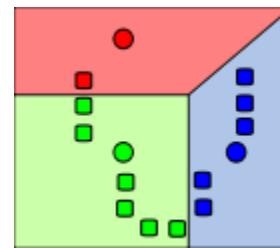
- K-means
  - Iteratively re-assign points to the nearest cluster center
- Agglomerative clustering
  - Start with each point as its own cluster and iteratively merge the closest clusters
- Mean-shift clustering
  - Estimate modes of pdf
- Spectral clustering
  - Split the nodes in a graph based on assigned links with similarity weights

# K-means algorithm

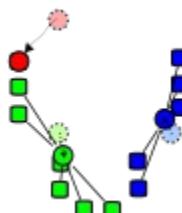
1. Randomly select K centers



2. Assign each point to nearest center

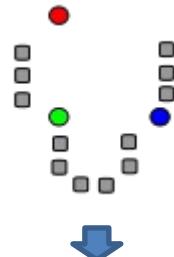


3. Compute new center (mean) for each cluster

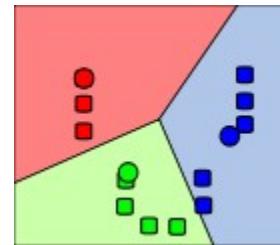


# K-means algorithm

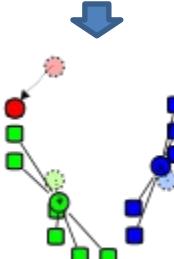
1. Randomly select K centers



2. Assign each point to nearest center



3. Compute new center (mean) for each cluster



Back to 2



# K-means

1. Initialize cluster centers:  $\mathbf{c}^0$  ; t=0
2. Assign each point to the closest center

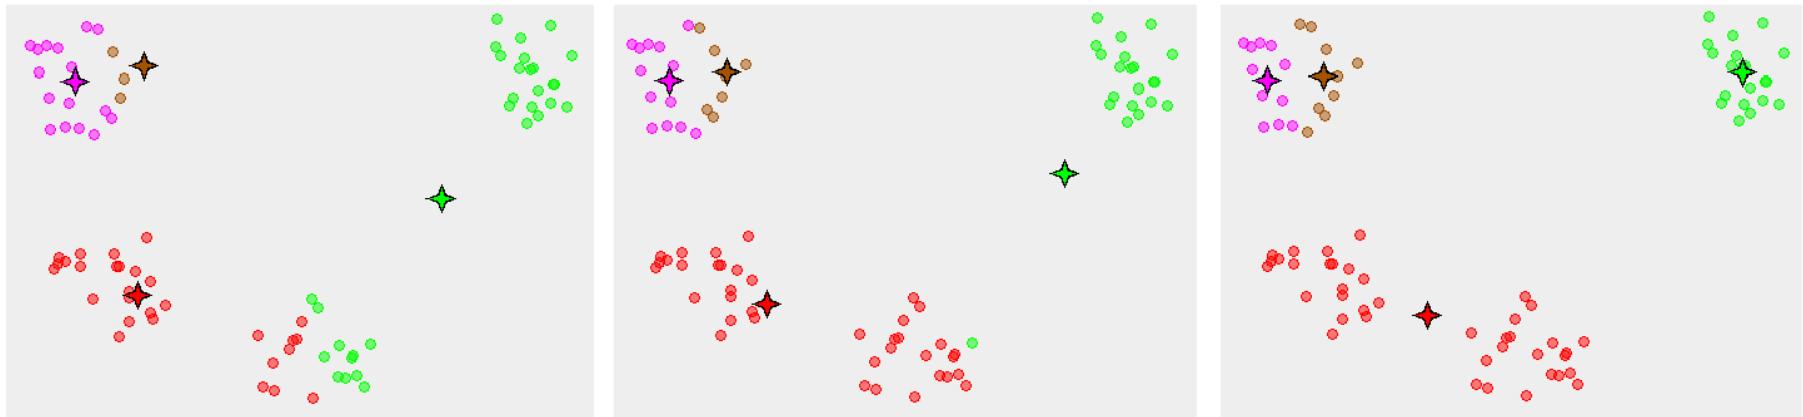
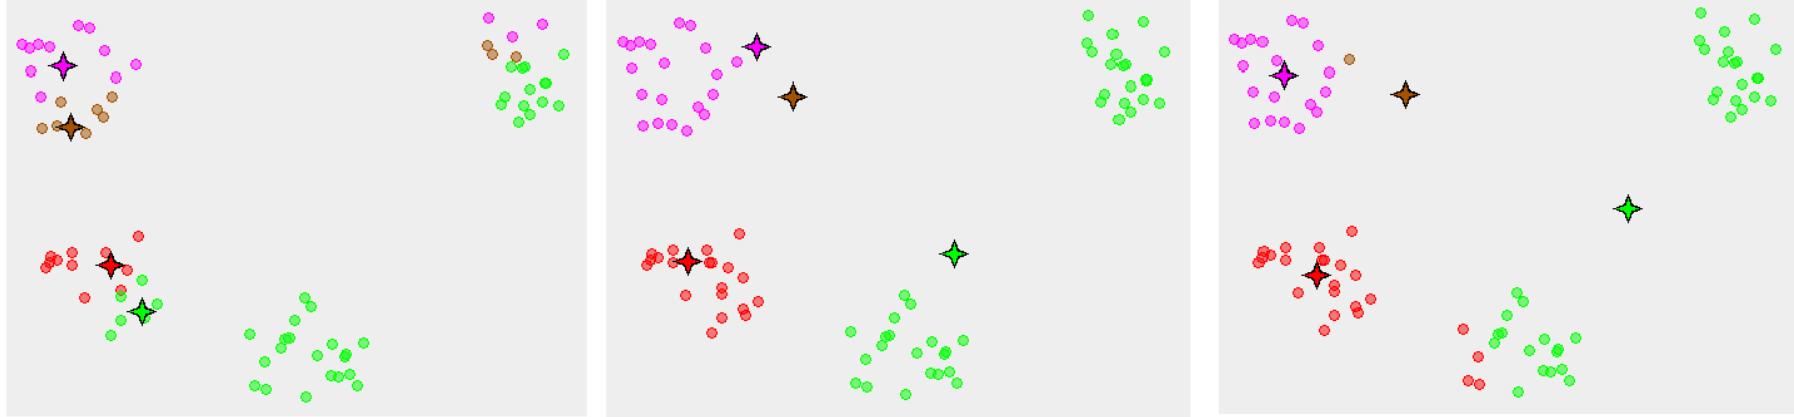
$$\delta^t = \underset{\mathbf{c}}{\operatorname{argmin}} \frac{1}{N} \sum_j \sum_i \delta_{ij} (\mathbf{c}_i^{t-1} - \mathbf{x}_j)^2$$

3. Update cluster centers as the mean of the points

$$\mathbf{c}^t = \underset{\mathbf{c}}{\operatorname{argmin}} \frac{1}{N} \sum_j \sum_i \delta_{ij}^t (\mathbf{c}_i - \mathbf{x}_j)^2$$

4. Repeat 2-3 until no points are re-assigned (t=t+1)

# K-means convergence



# K-means: design choices

- Initialization
  - Randomly select K points as initial cluster center
  - Or greedily choose K points to minimize residual
- Distance measures
  - Traditionally Euclidean, could be others
- Optimization
  - Will converge to a *local minimum*

# K-means clustering using intensity or color

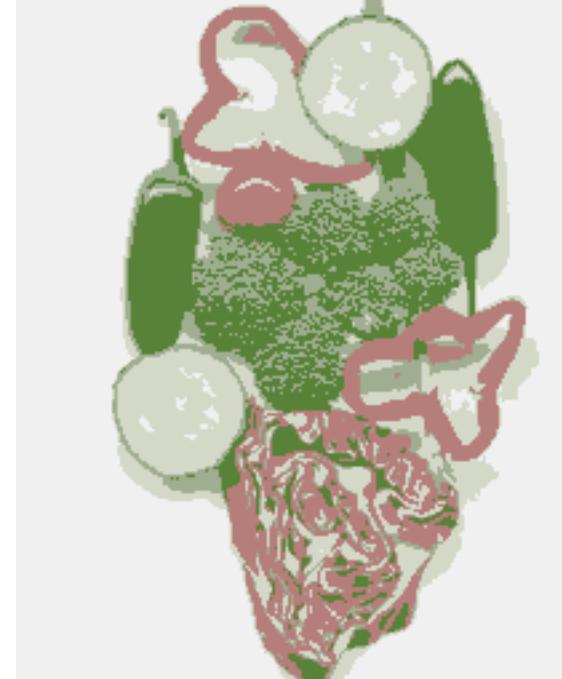
Image



Clusters on intensity

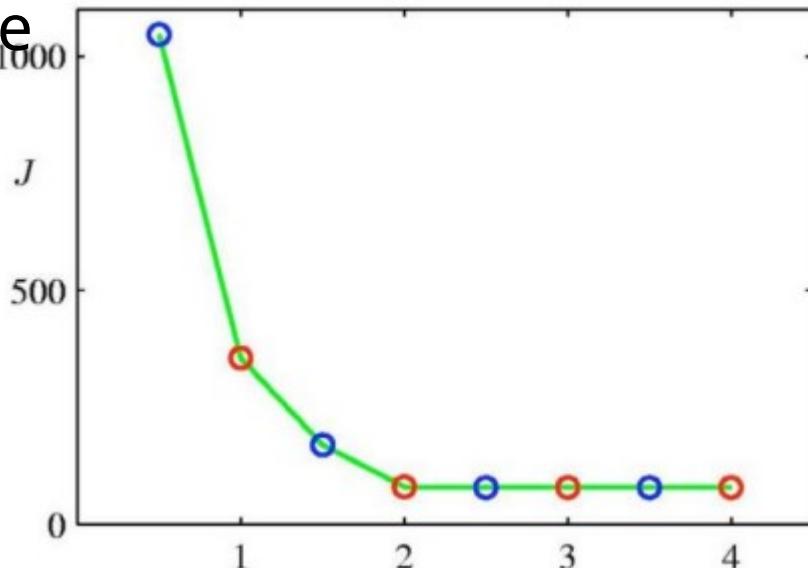


Clusters on color



# How to choose the number of clusters?

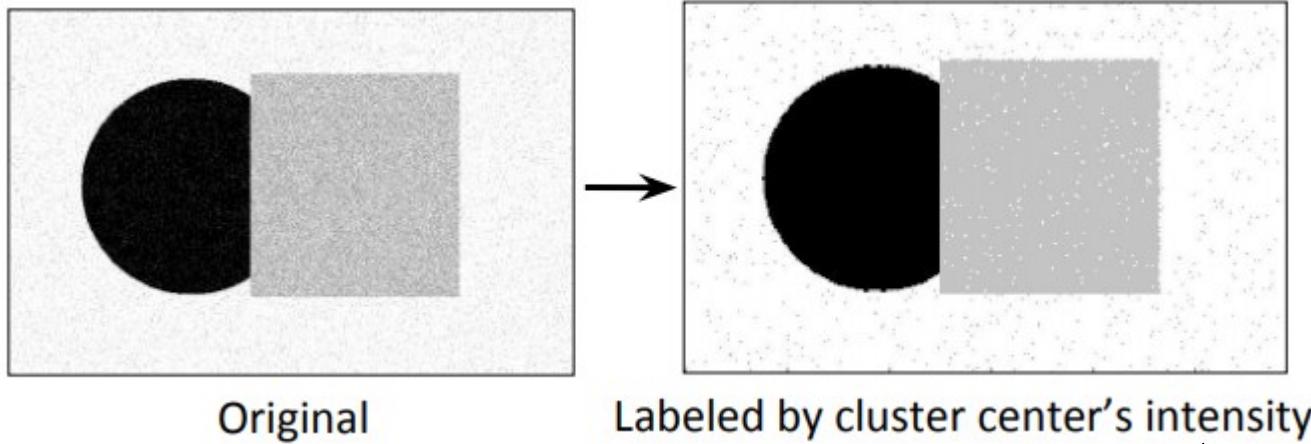
- Validation set
  - Try different numbers of clusters and look at performance
    - When building dictionaries (discussed later), more clusters typically work better.
- Plot k-means objective vs. K, and pick k at the elbow (in this example its k=2)
  - $J$  here is the total sum of the distances of all points to their associated centers



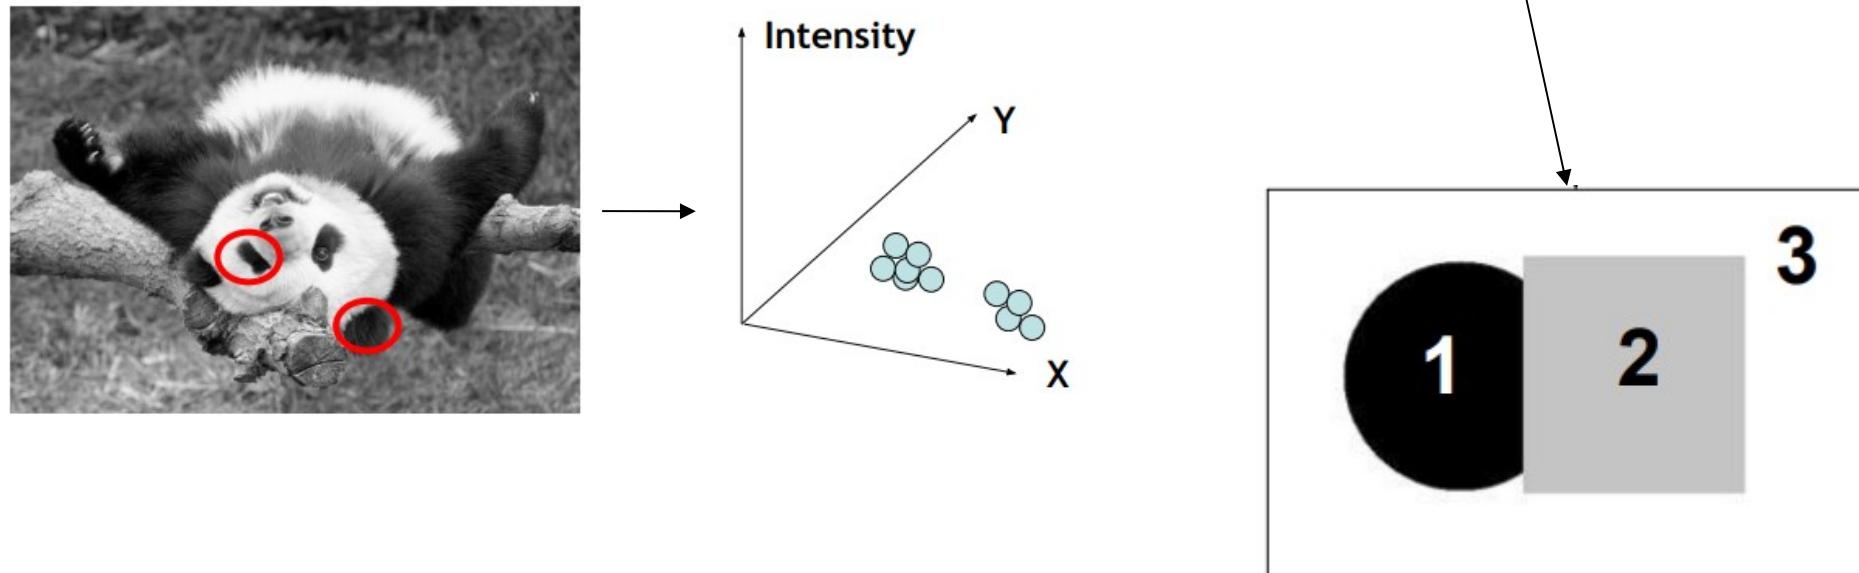
# Distance measures and termination

- Choice of “distance” measures
  - Euclidean (commonly used)
  - Cosine similarity
  - Non-linear (called kernel k-means)
- Termination criteria:
  - Maximum number of iterations reaches
  - No change during the re-assignment step  
**(convergence)**

# Smoothing cluster assignments



Group pixels based on **intensity + position**

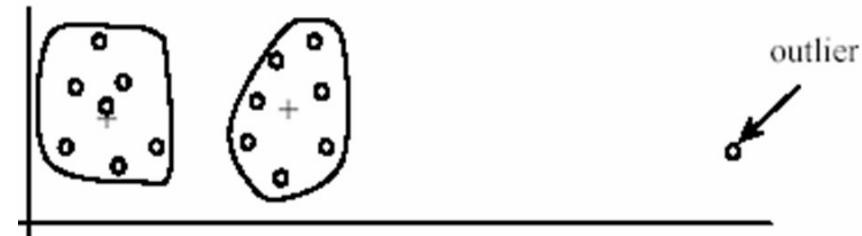


# How to initialize the clusters?

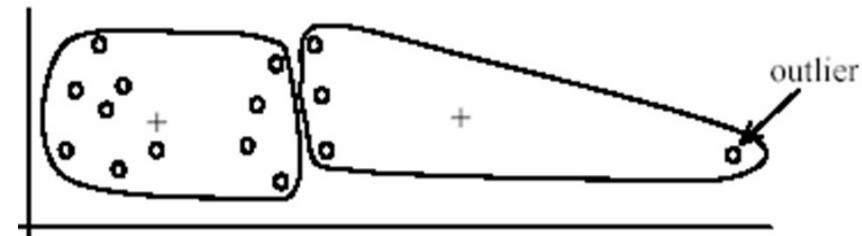
- k-means++ initialization
    - Make the initial cluster centers span the space
1. Choose one center uniformly at random from all data points.
  2. For each point  $x$ , compute the distance  $D(x)$  between  $x$  and the nearest center that has already been chosen.
  3. Choose one new data point at random as a new center, using a weighted probability distribution where a point  $x$  is chosen with probability proportional to  $D(x)^2$ .
  4. Repeat Steps 2 and 3 until  $k$  centers have been chosen.
  5. Proceed using standard k-means clustering.

# K-Means pros and cons

- Pros
  - Finds cluster centers that minimize conditional variance (good representation of data)
  - Simple and fast\*
  - Easy to implement
- Cons
  - Need to choose K
  - Sensitive to outliers
  - Prone to local minima
  - All clusters have the same parameters (e.g., distance measure is non-adaptive)
  - \*Can be slow: each iteration is  $O(KNd)$  for N d-dimensional points
- Usage
  - Cluster features to build visual dictionaries



(B): Ideal clusters



# Building Visual Dictionaries

1. Sample features from a database

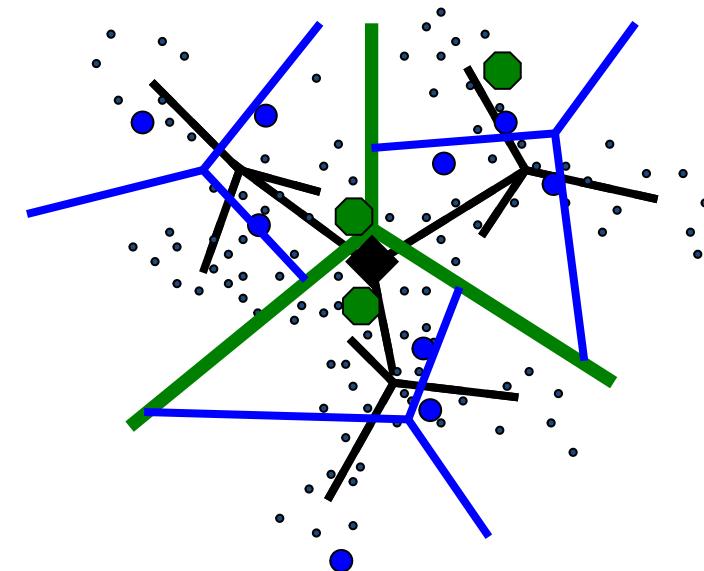
- E.g., 128 dimensional SIFT vectors



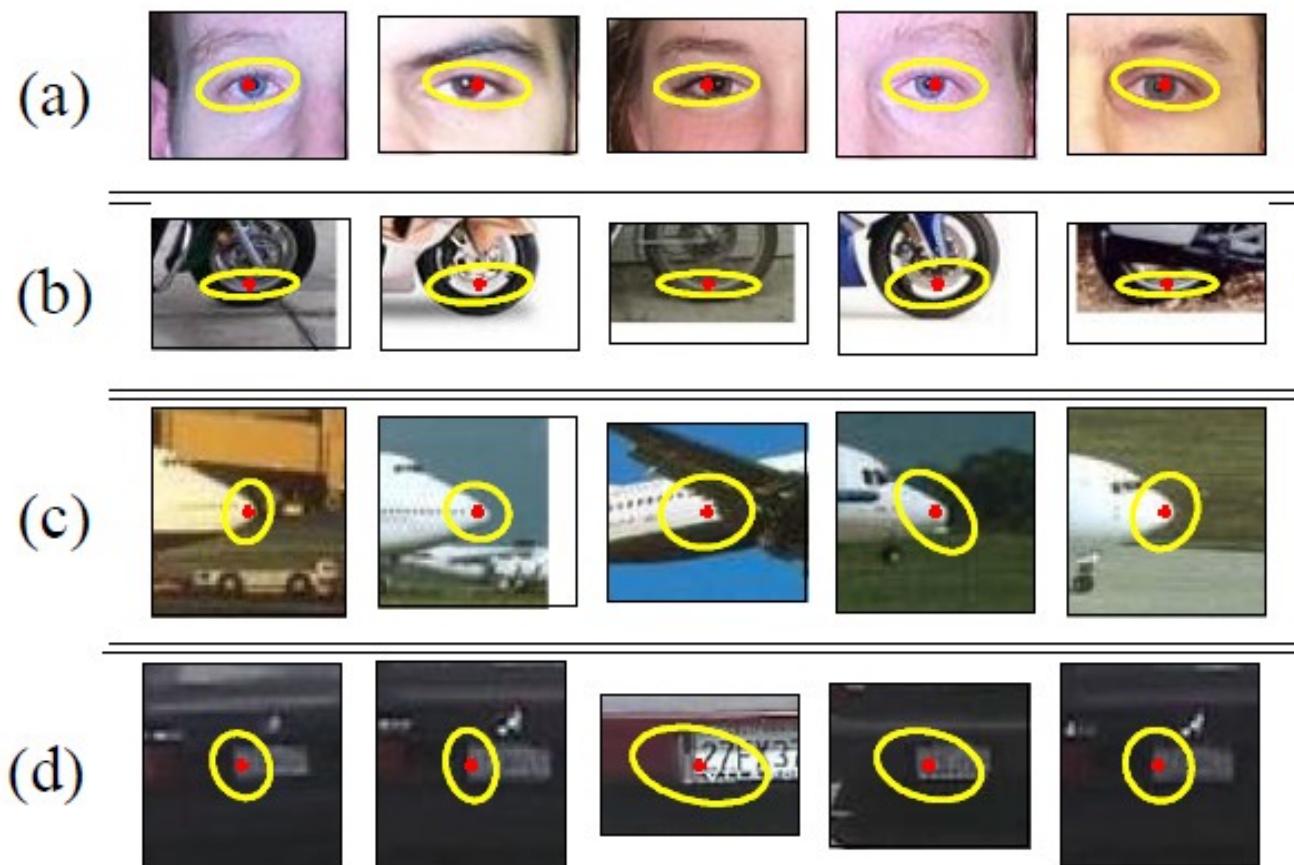
2. Cluster to build dictionary

- Cluster centers are the dictionary words

3. To match new features, assign to the nearest cluster to save rebuilding dictionary

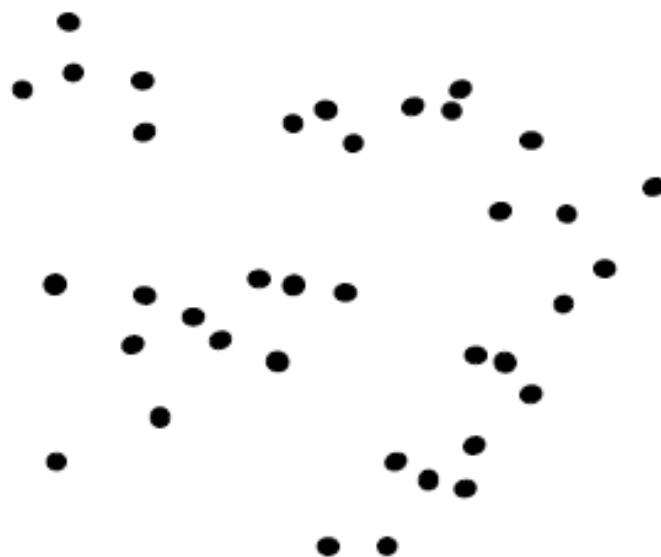


# Examples of learned codewords



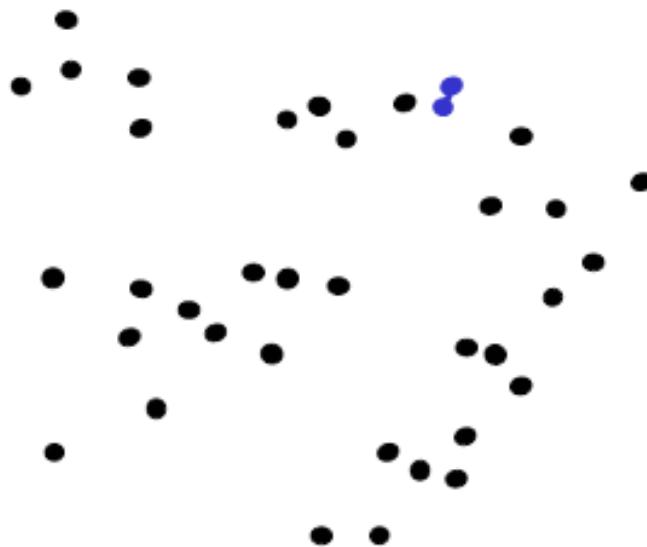
Most likely codewords for 4 learned “topics”  
EM with multinomial (problem 3) to get topics

# Agglomerative clustering



1. Say "Every point is its own cluster"

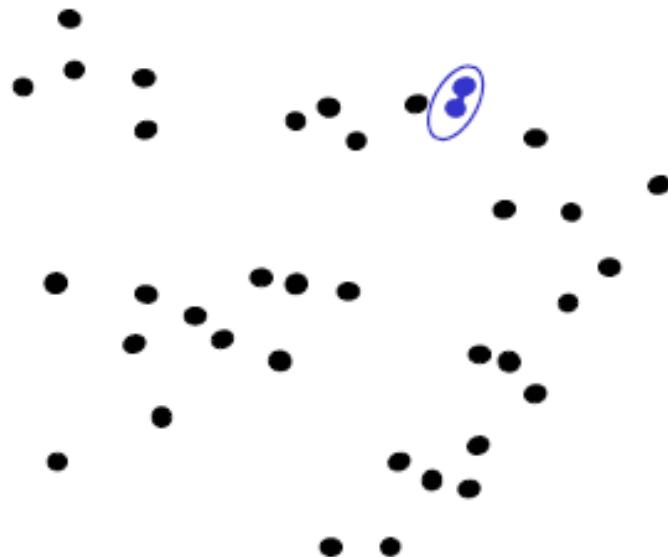
# Agglomerative clustering



1. Say "Every point is its own cluster"
2. Find "most similar" pair of clusters



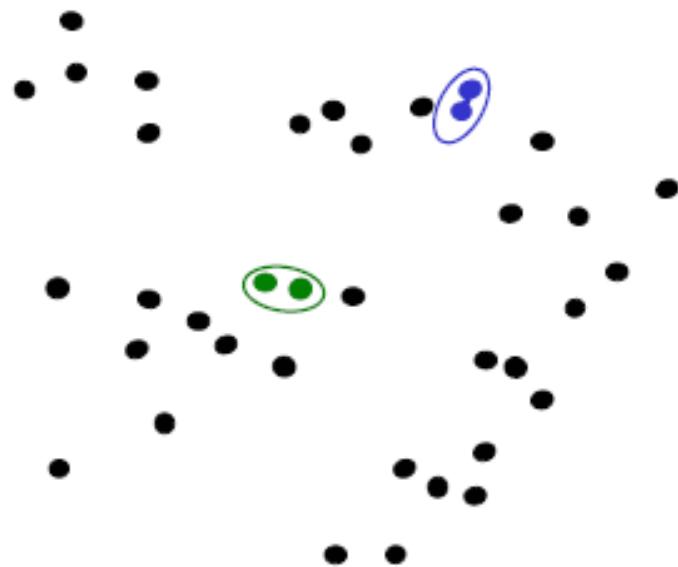
# Agglomerative clustering



1. Say "Every point is its own cluster"
2. Find "most similar" pair of clusters
3. Merge it into a parent cluster



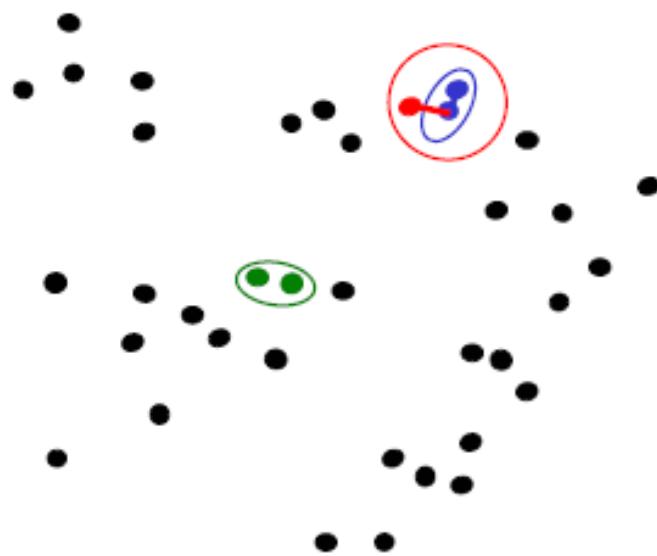
# Agglomerative clustering



1. Say "Every point is its own cluster"
2. Find "most similar" pair of clusters
3. Merge it into a parent cluster
4. Repeat



# Agglomerative clustering



1. Say "Every point is its own cluster"
2. Find "most similar" pair of clusters
3. Merge it into a parent cluster
4. Repeat

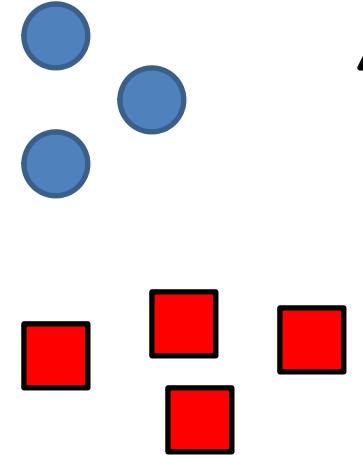


# Agglomerative clustering



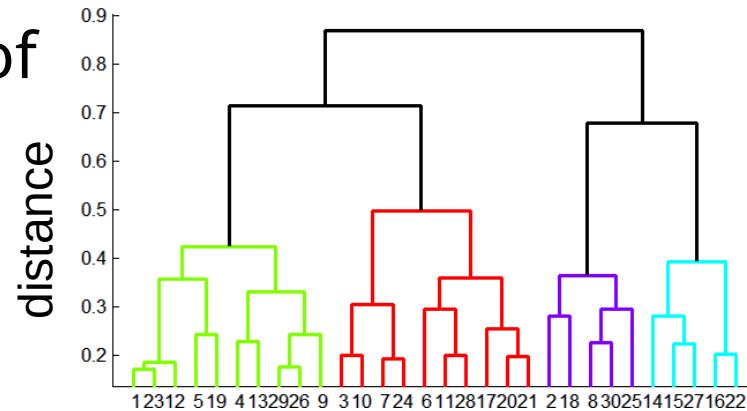
How to define cluster similarity?

- Average distance between points, maximum distance, minimum distance
- Distance between means or medoids



How many clusters?

- Clustering creates a dendrogram (a tree)
- Threshold based on max number of clusters or based on distance between merges



# Conclusions: Agglomerative Clustering

## Good

- Simple to implement, widespread application
- Clusters have adaptive shapes
- Provides a hierarchy of clusters

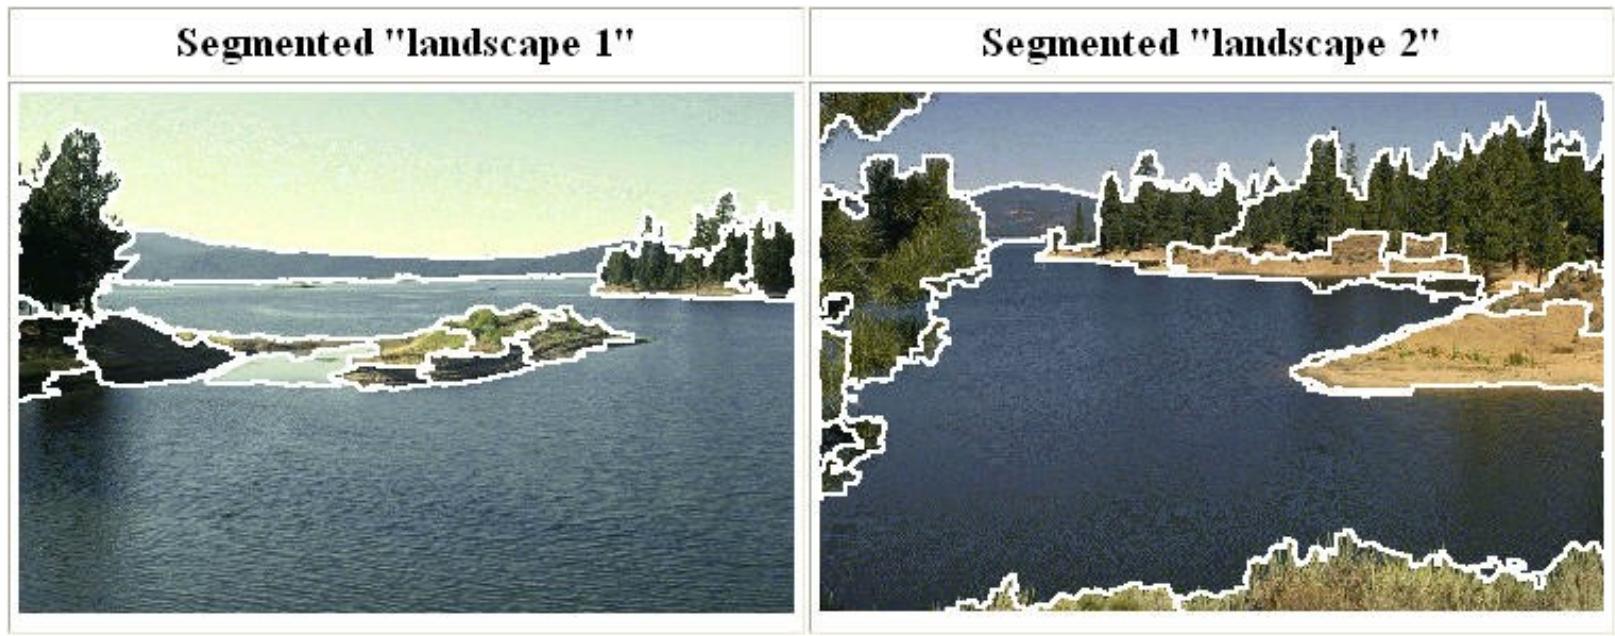
## Bad

- May have imbalanced clusters
- Still have to choose number of clusters or threshold
- Need to use an “ultrametric” to get a meaningful hierarchy
- Can be very slow  $O(n^2d + n^3)$

# Mean shift segmentation

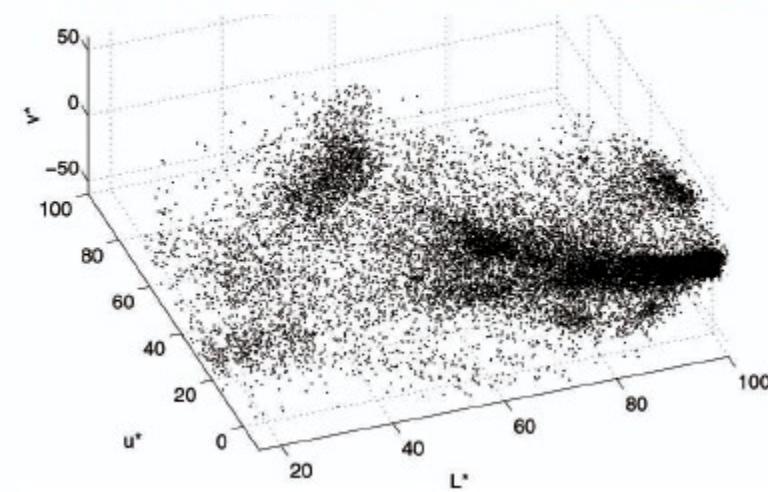
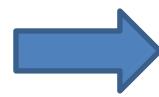
D. Comaniciu and P. Meer, Mean Shift: A Robust Approach toward Feature Space Analysis, PAMI 2002.

- Versatile technique for clustering-based segmentation

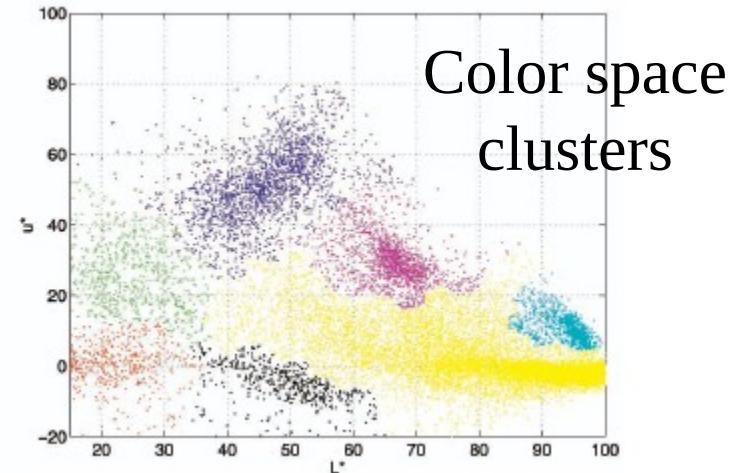
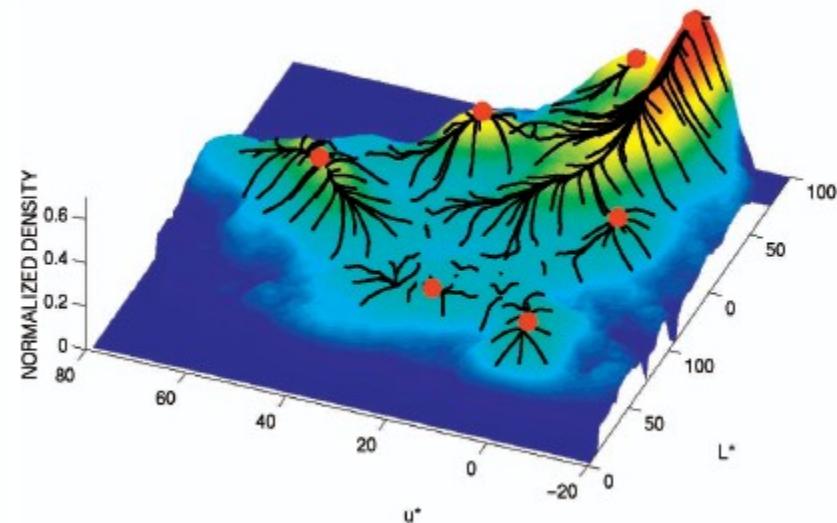


# Mean shift algorithm

Try to find *modes* of a non-parametric density.



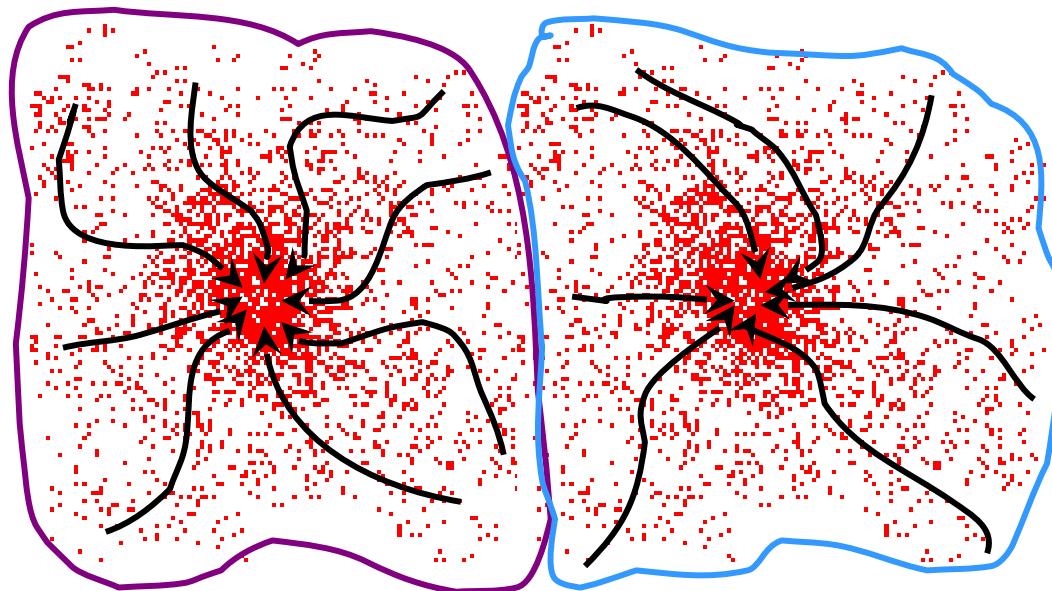
Color  
space



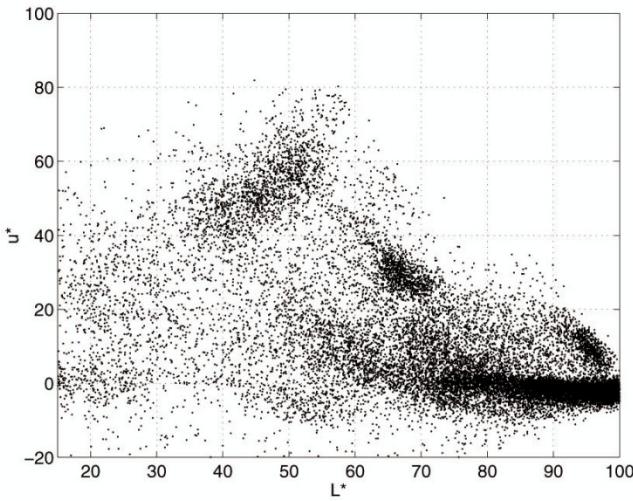
Color space  
clusters

# Attraction basin

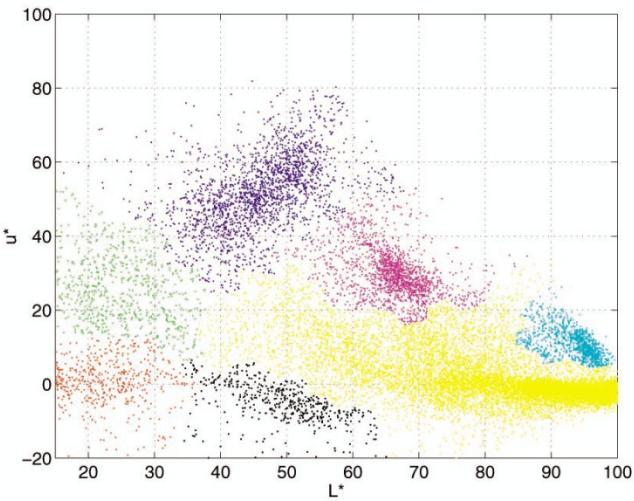
- **Attraction basin:** the region for which all trajectories lead to the same mode
- **Cluster:** all data points in the attraction basin of a mode



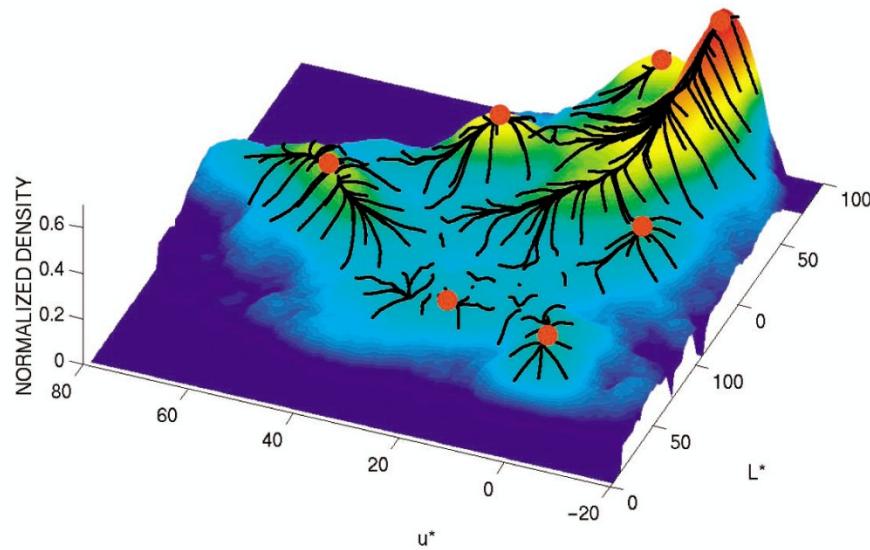
# Attraction basin



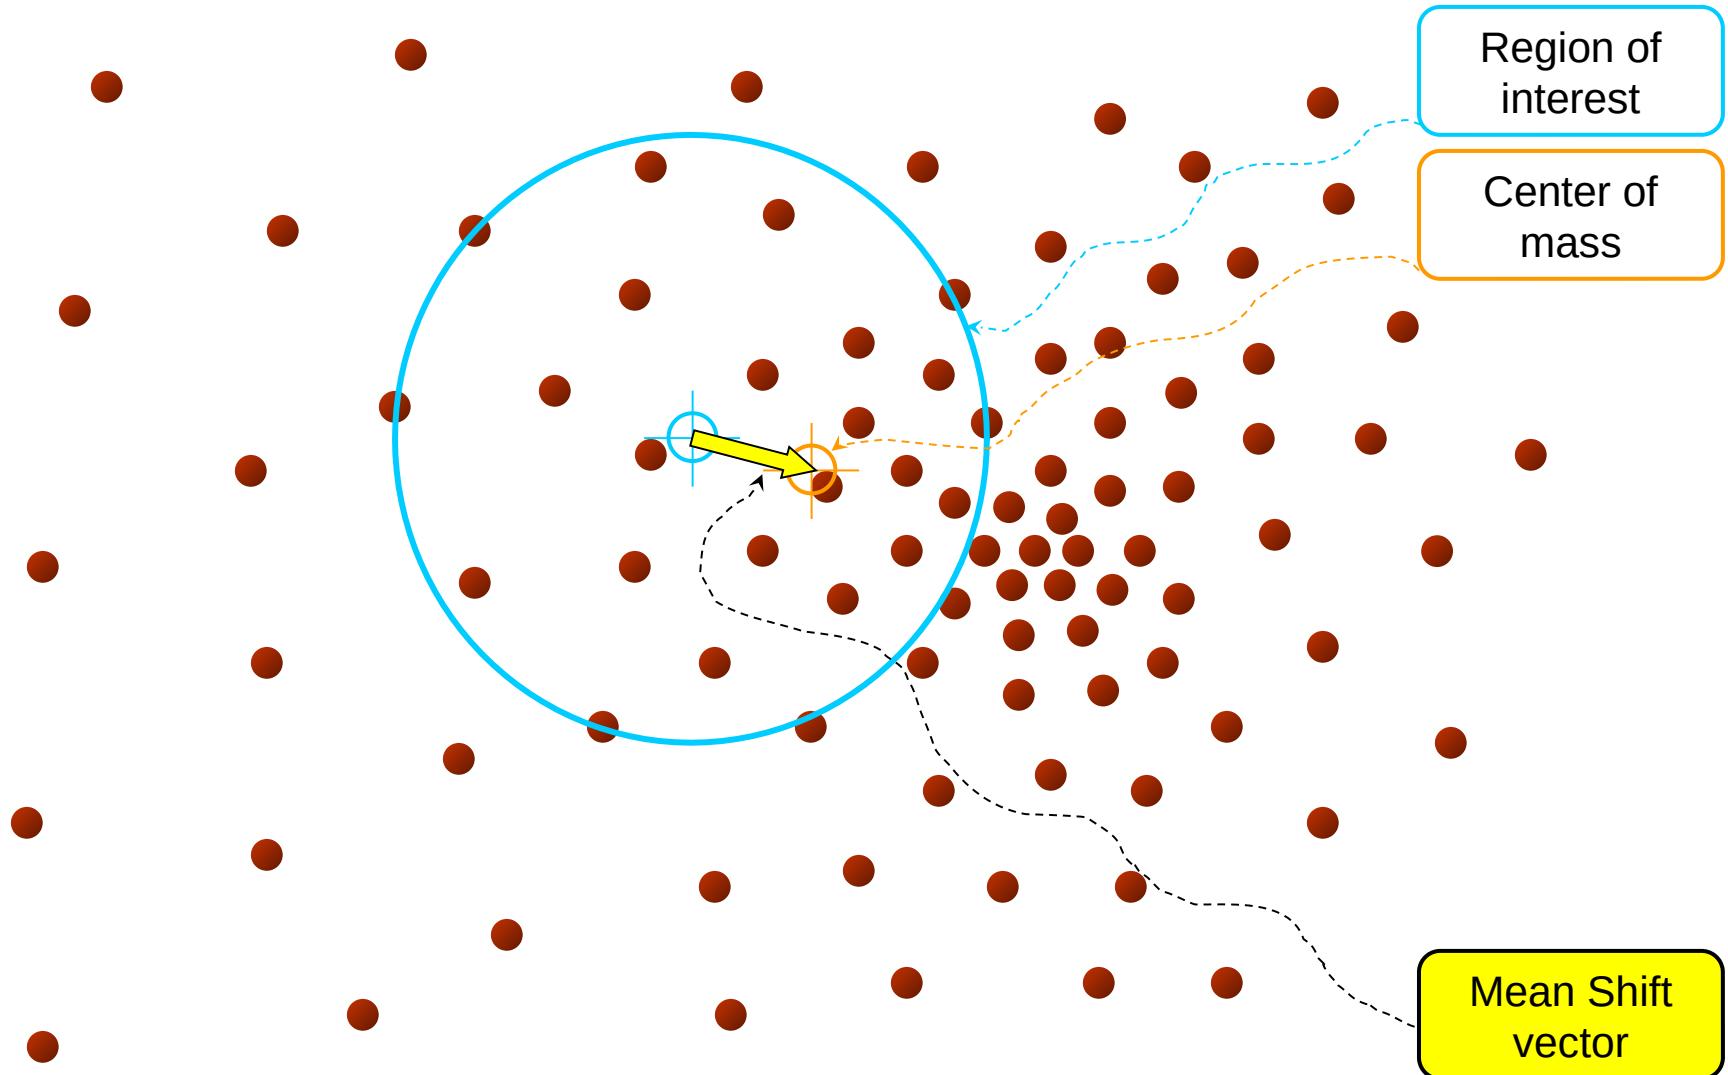
(a)



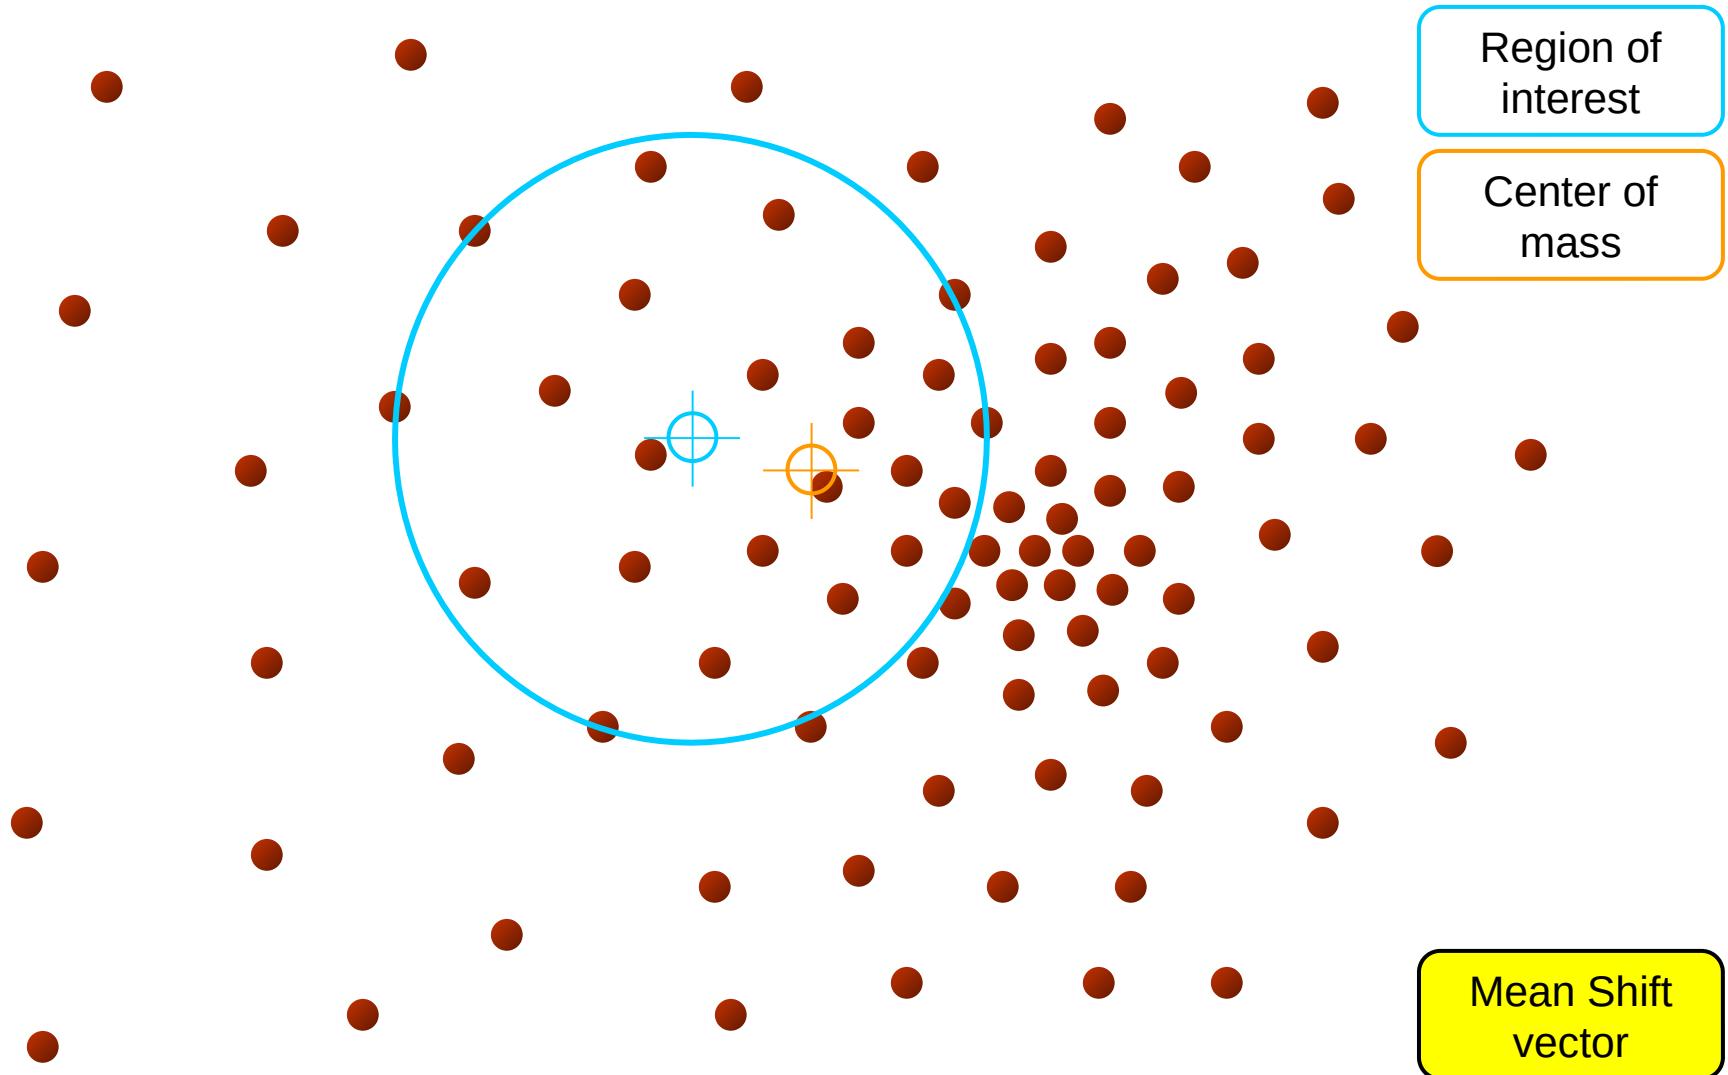
(b)



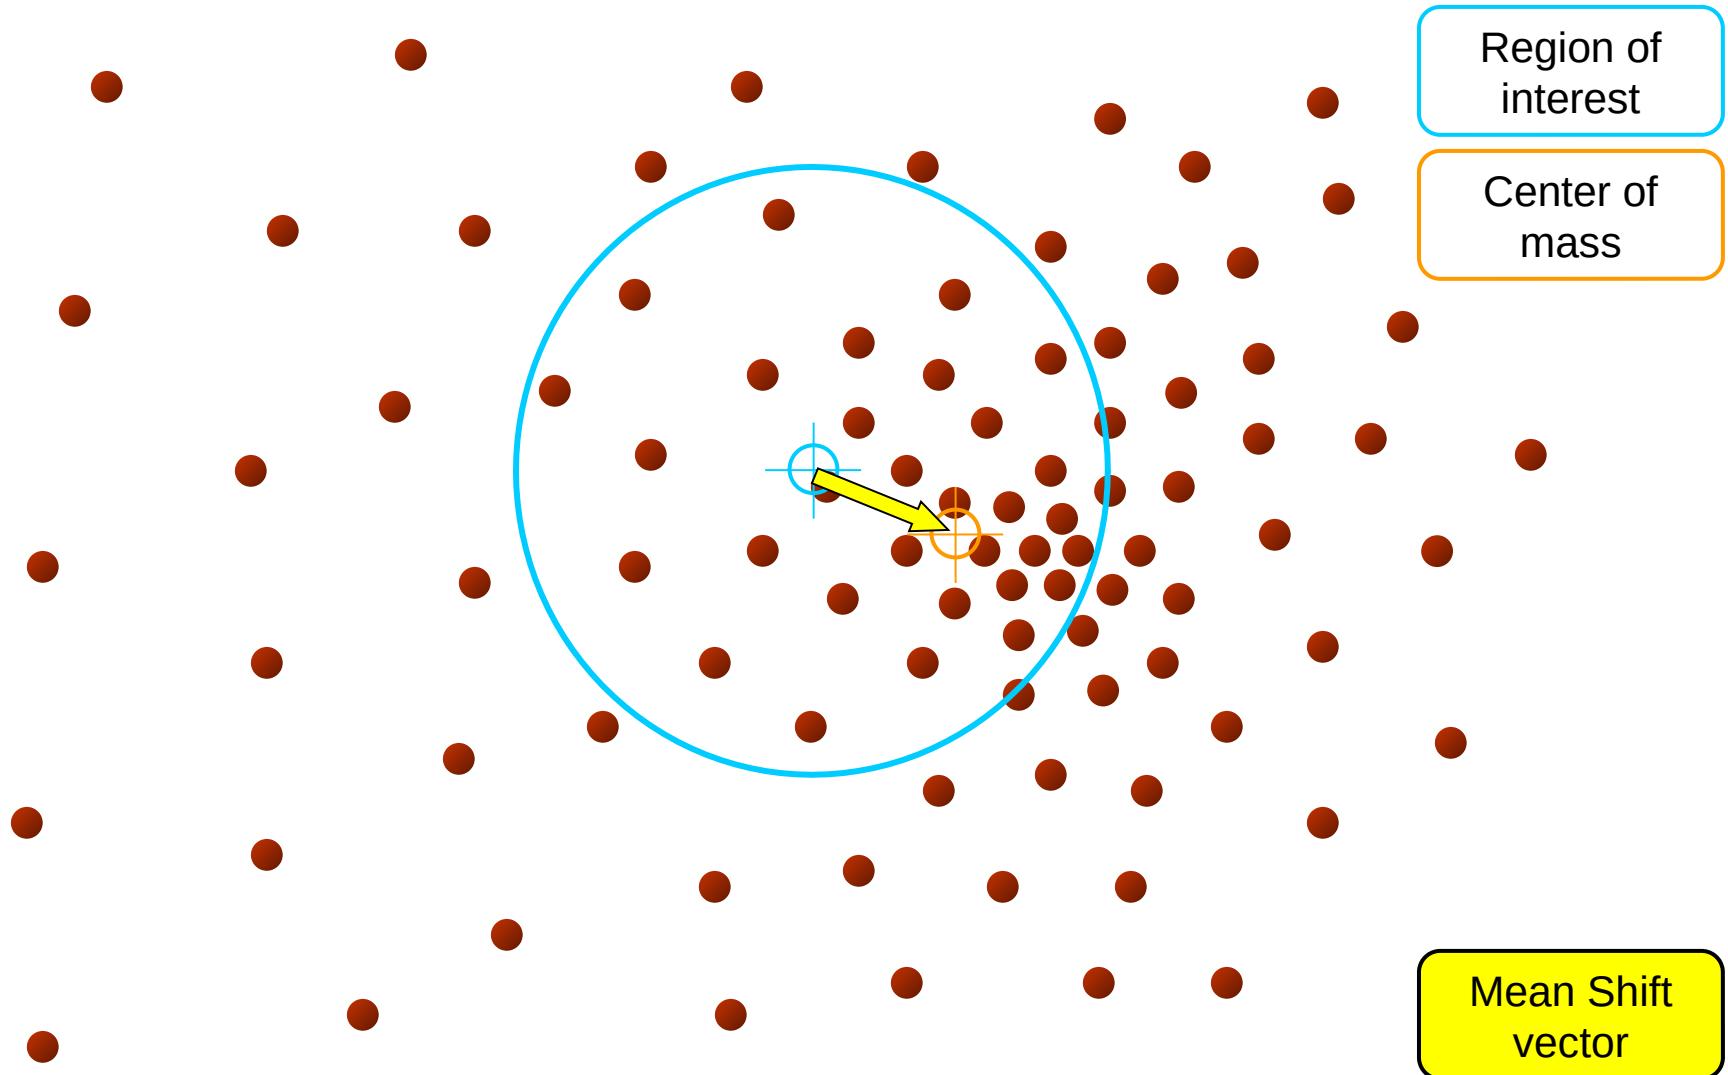
# Mean shift



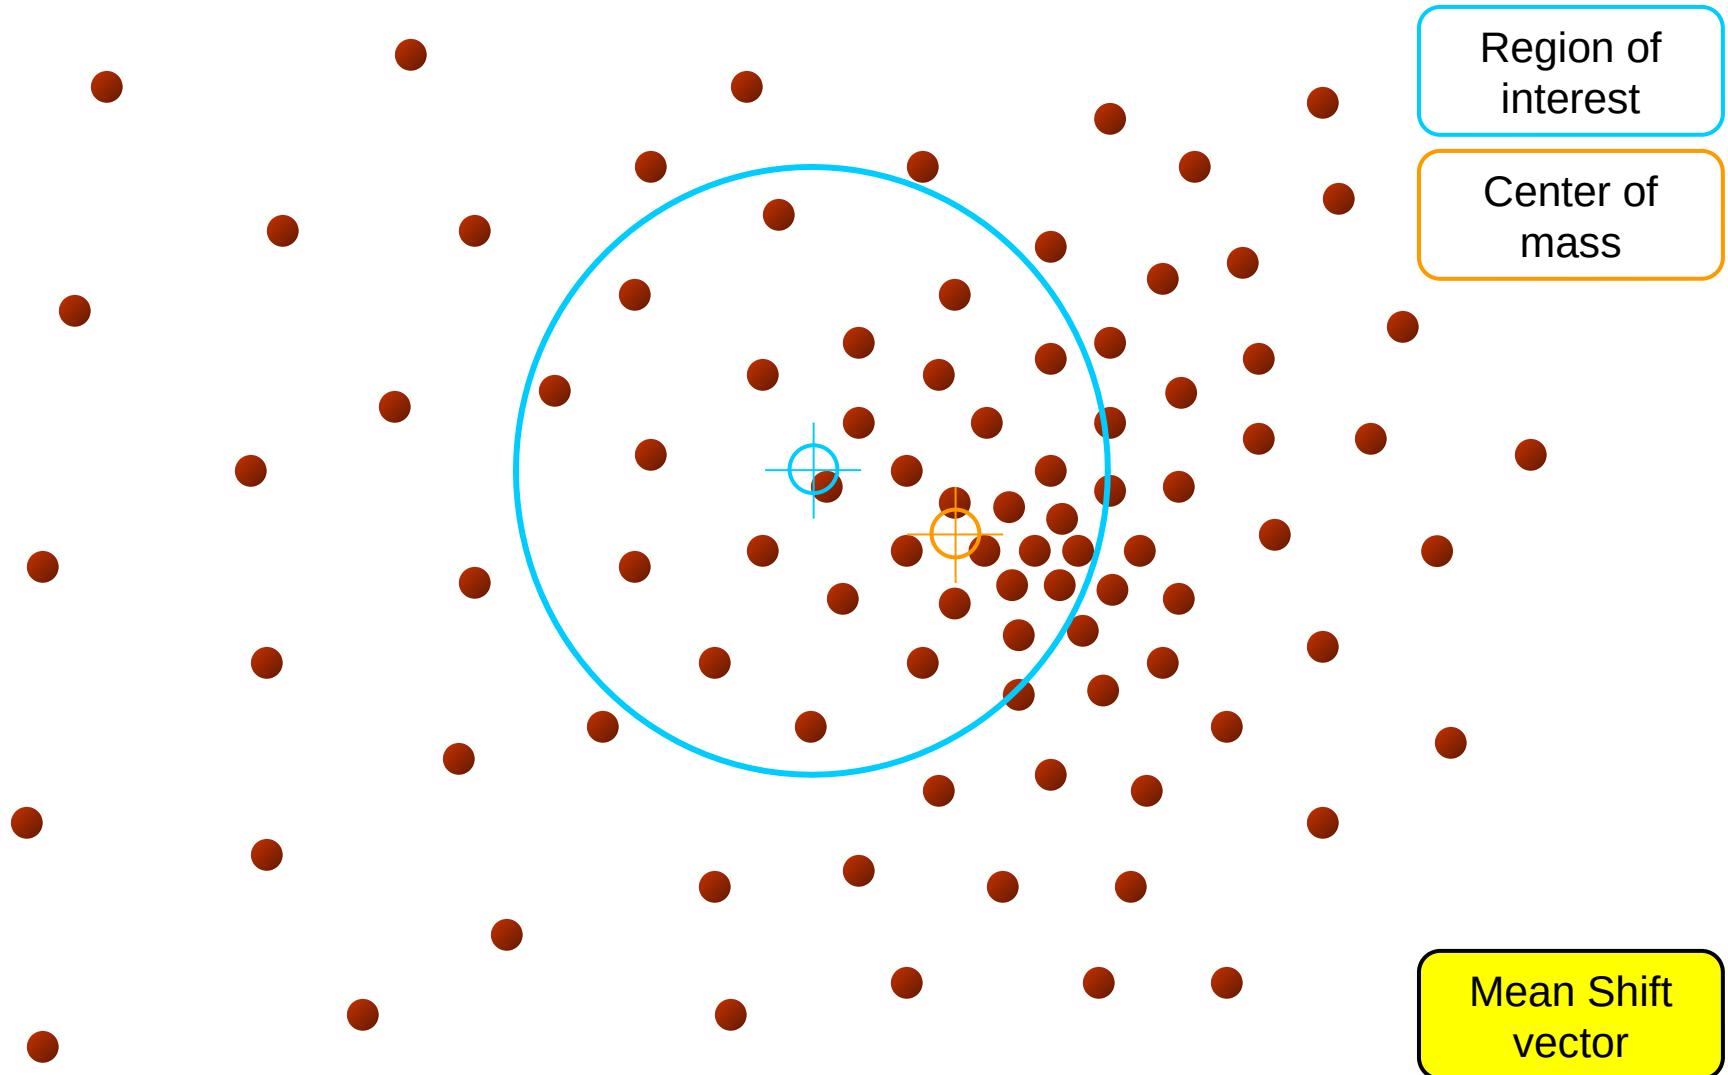
# Mean shift



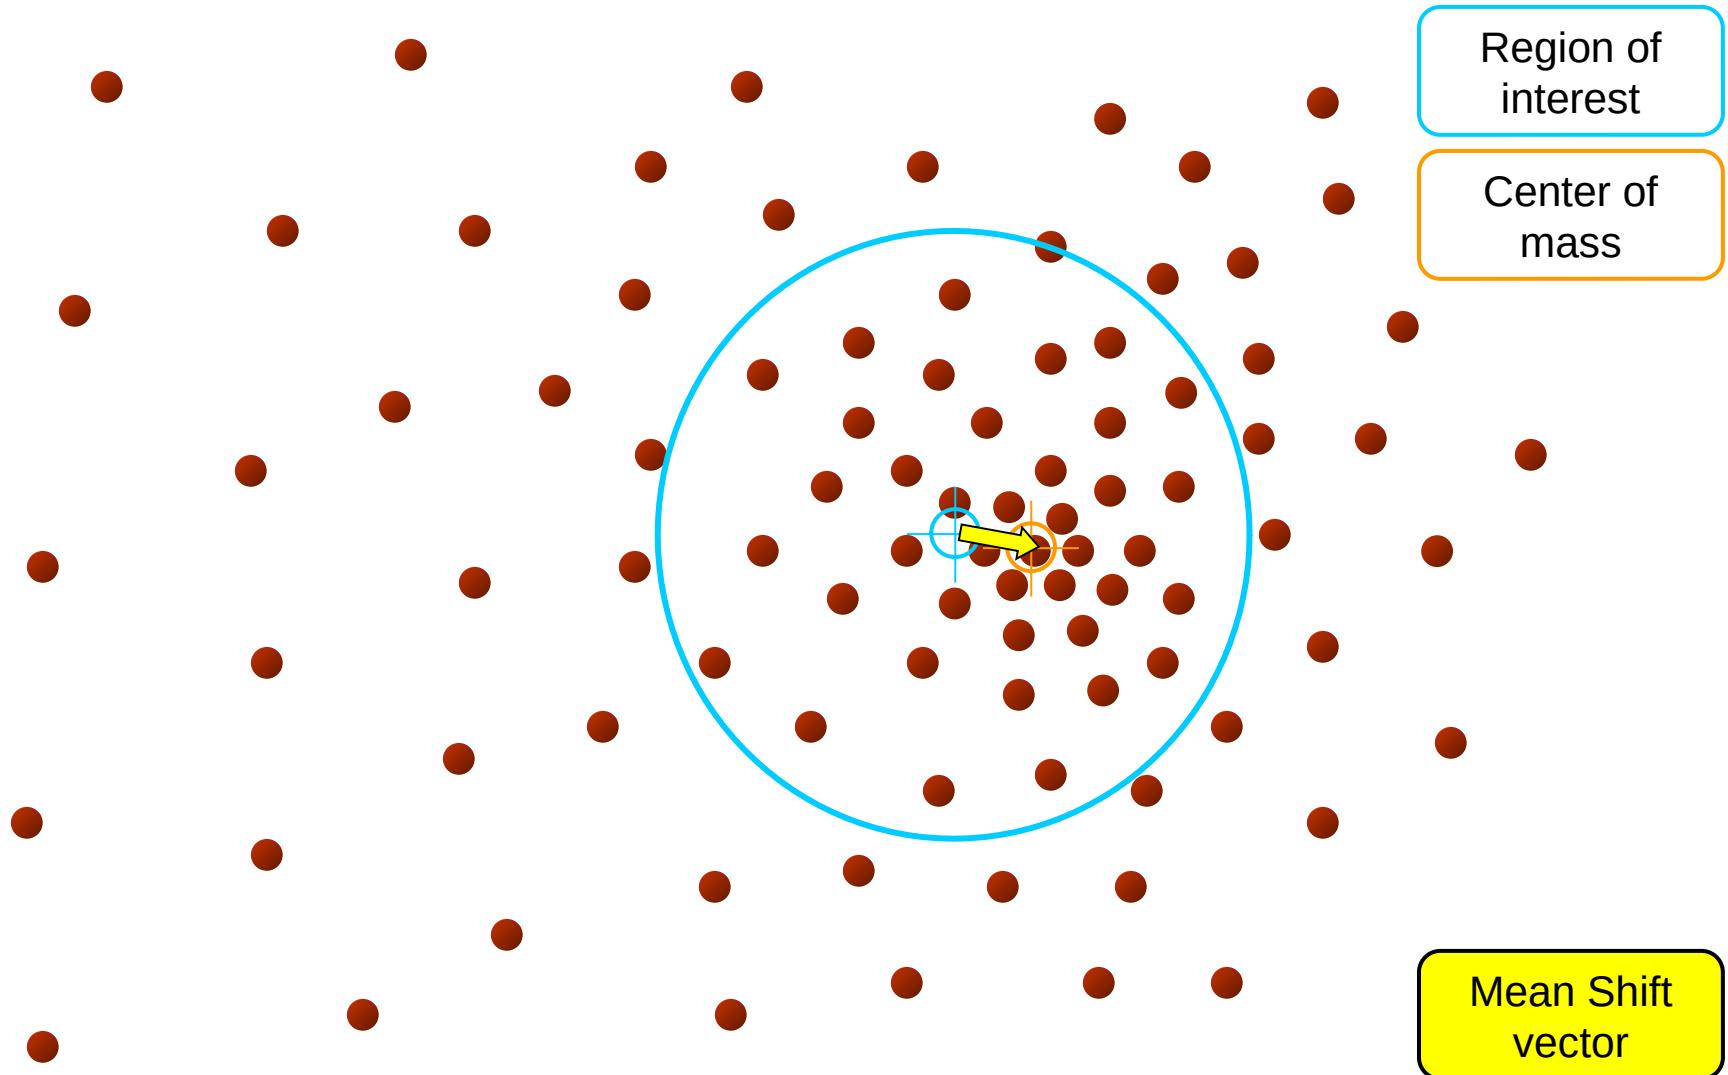
# Mean shift



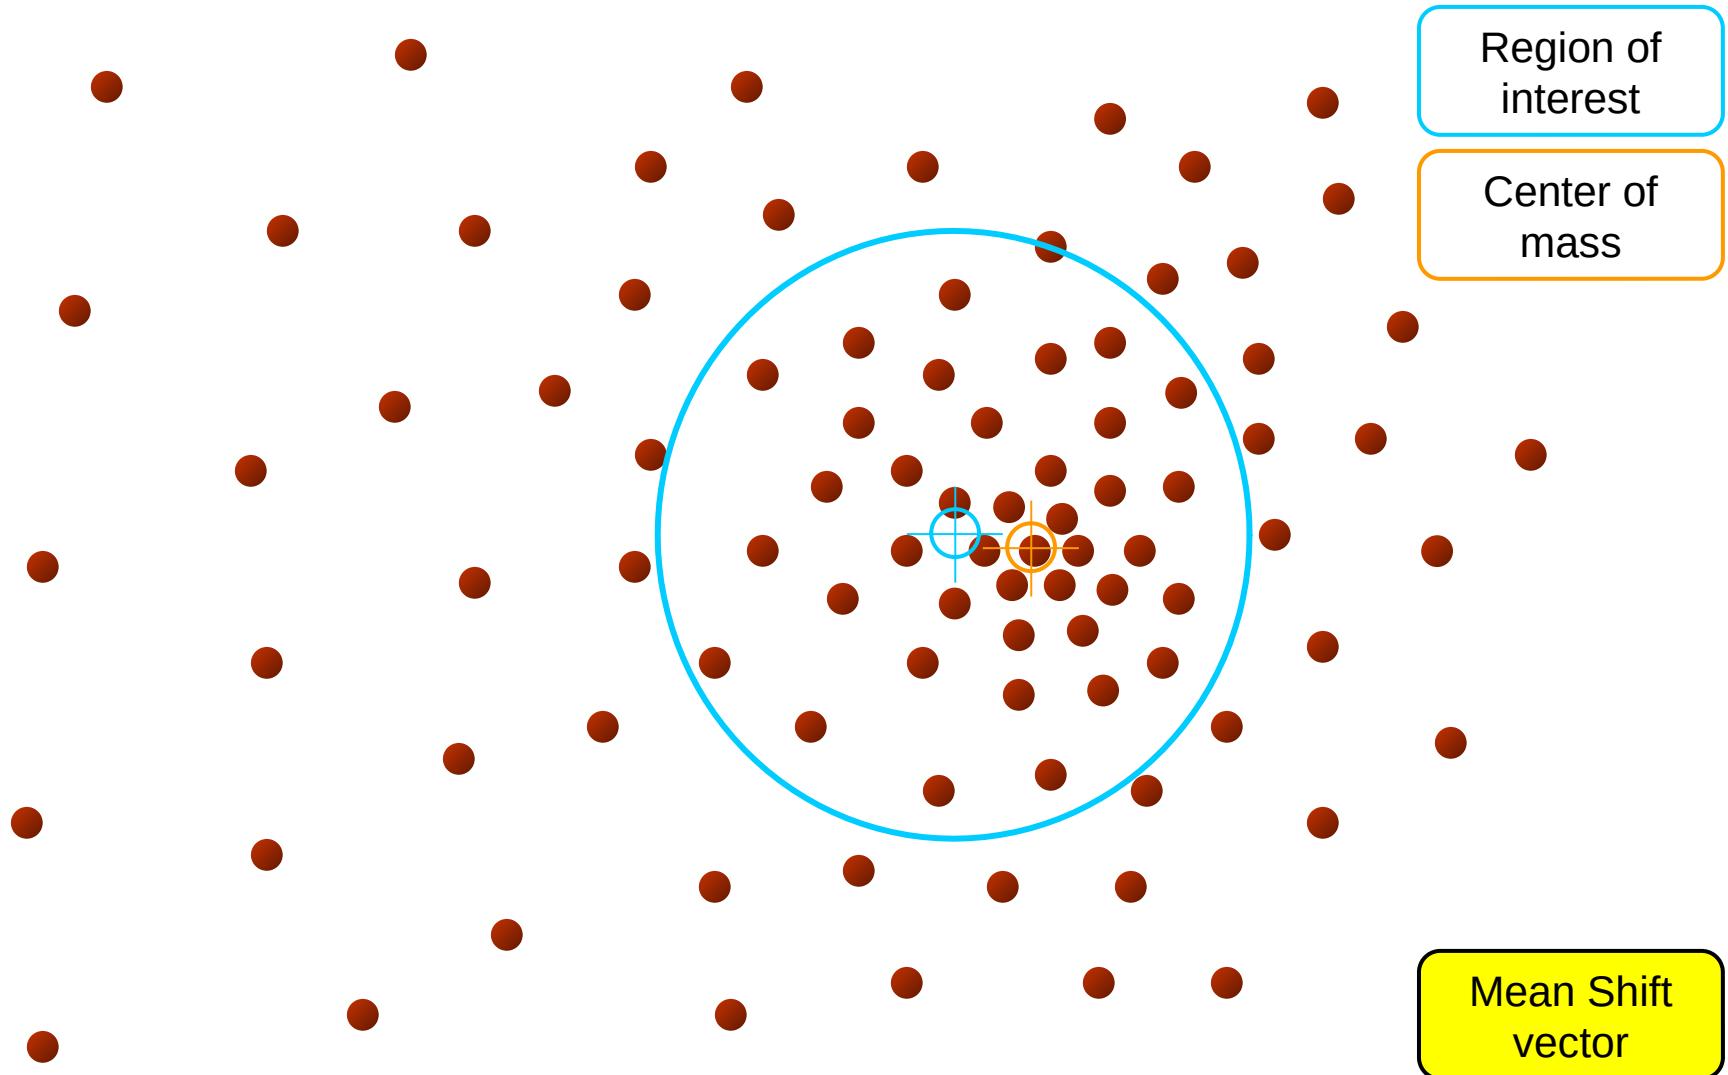
# Mean shift



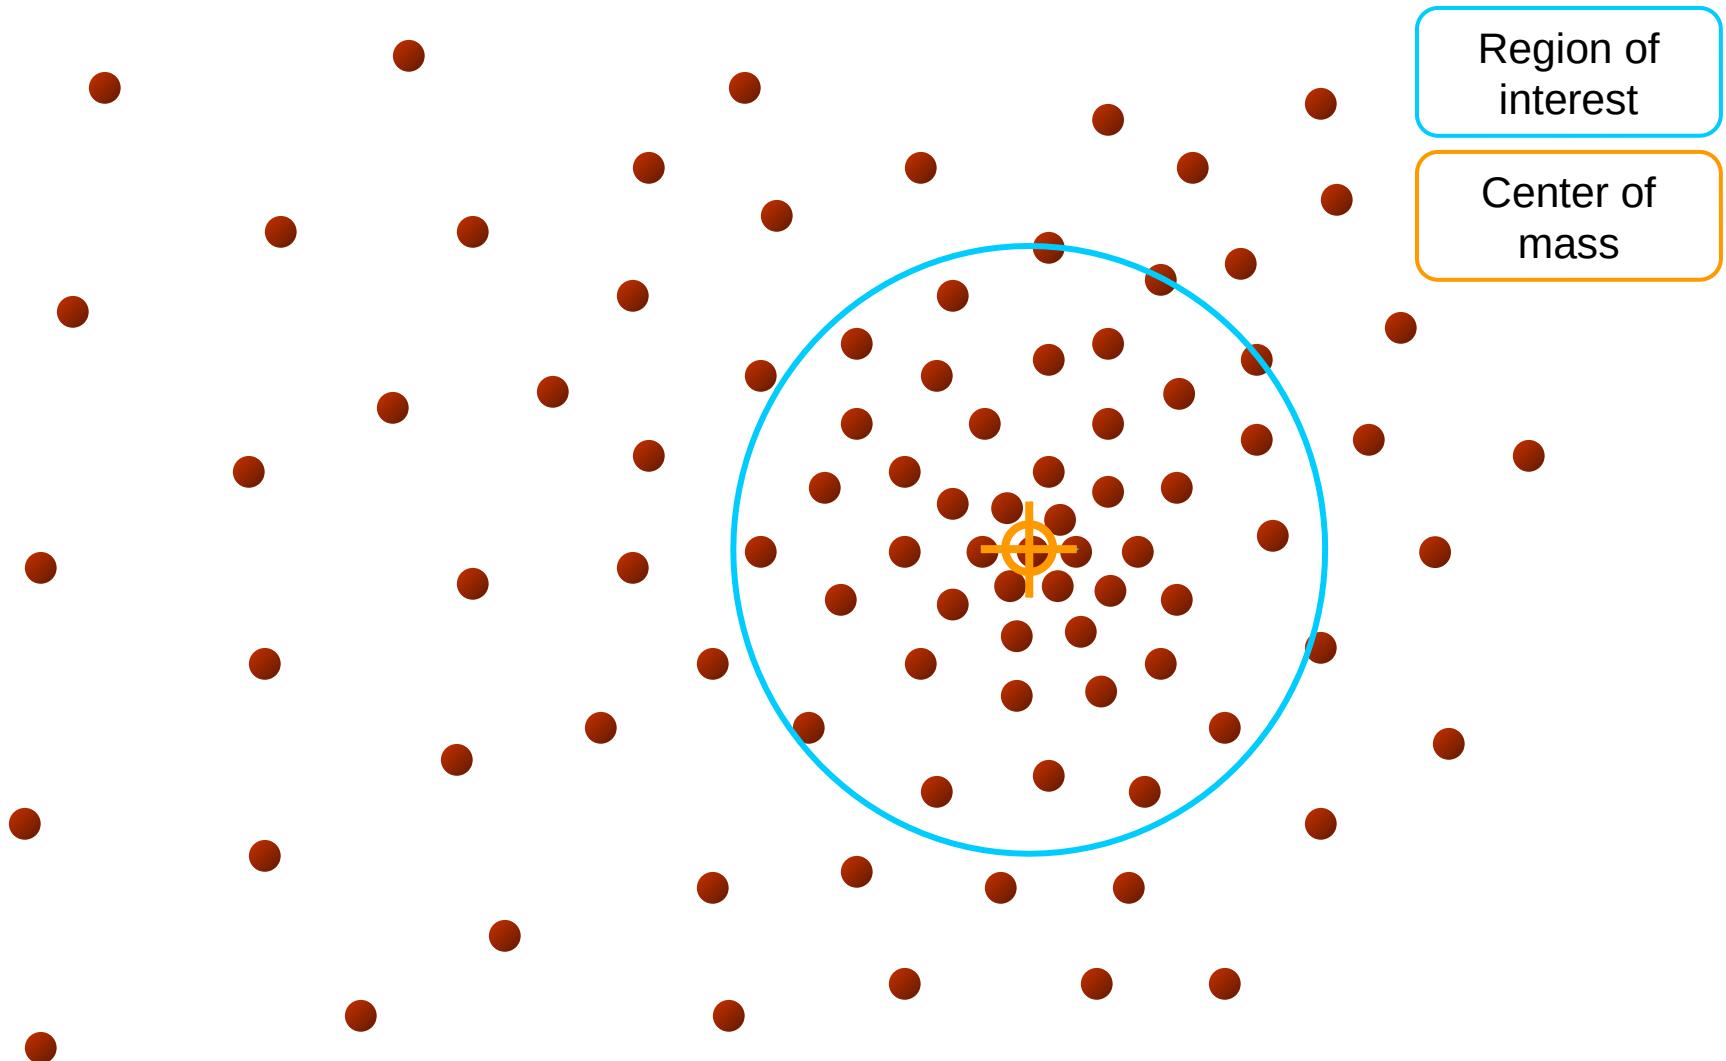
# Mean shift



# Mean shift



# Mean shift



# Kernel density estimation

Kernel density estimation function

$$\hat{f}_h(x) = \frac{1}{nh} \sum_{i=1}^n K\left(\frac{x - x_i}{h}\right)$$

n = number of  $x_i \in \mathbb{R}^d$

h = radius of the kernel  
(bandwidth parameter)

Gaussian kernel

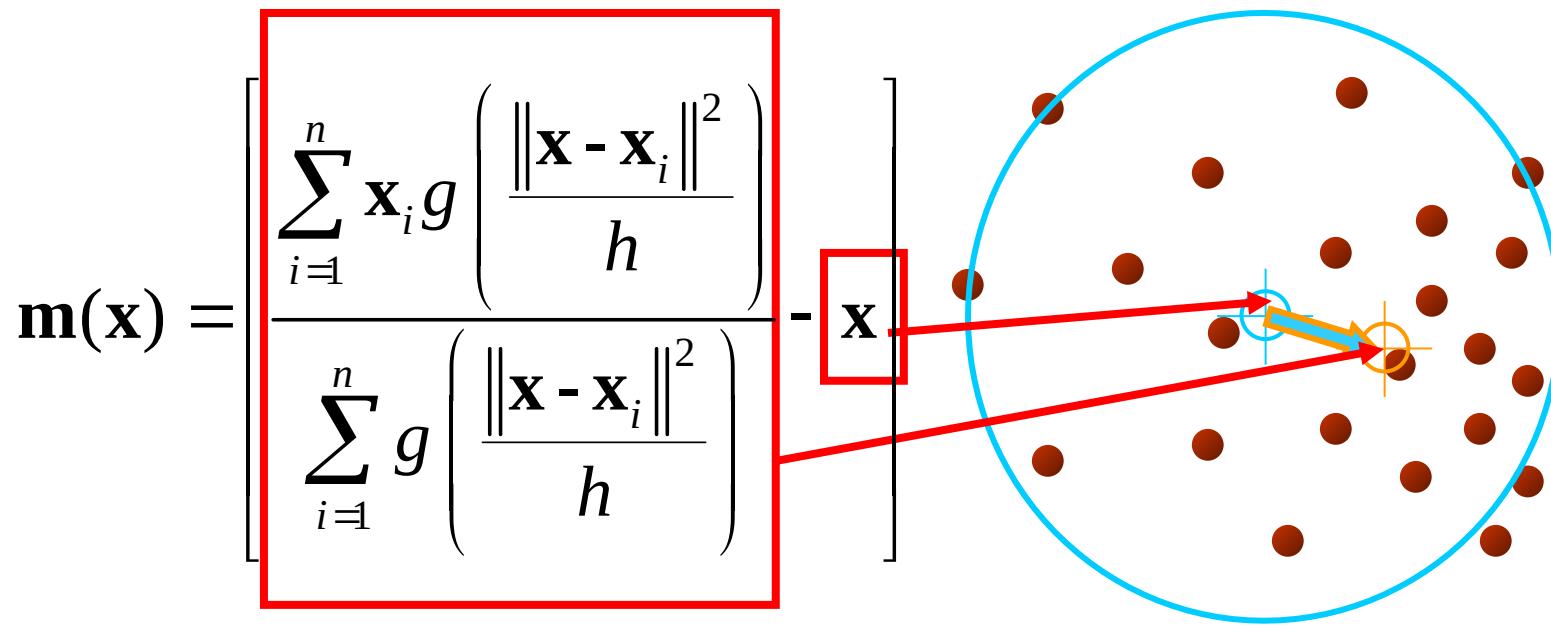
$$K\left(\frac{x - x_i}{h}\right) = \frac{1}{\sqrt{2\pi}} e^{-\frac{(x-x_i)^2}{2h^2}}.$$

$$g(x) = -K'(x)$$

# Computing the Mean Shift

Simple Mean Shift procedure:

- Compute mean shift vector  $\mathbf{m}(\mathbf{x})$
- Iteratively translate the kernel window by  $\mathbf{m}(\mathbf{x})$  until convergence.



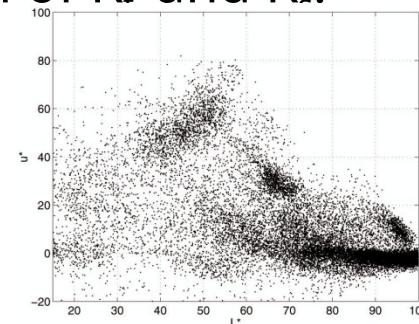
# Mean shift clustering

The mean shift algorithm seeks *modes* of the given set of points

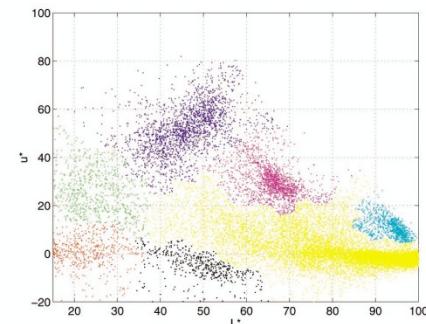
1. Choose kernel and bandwidth
2. For each point:
  - a) Center a window on that point
  - b) Compute the mean of the data in the search window
  - c) Center the search window at the new mean location
  - d) Repeat (b,c) until convergence
3. Assign points that lead to nearby modes to the same cluster

# Segmentation by Mean Shift

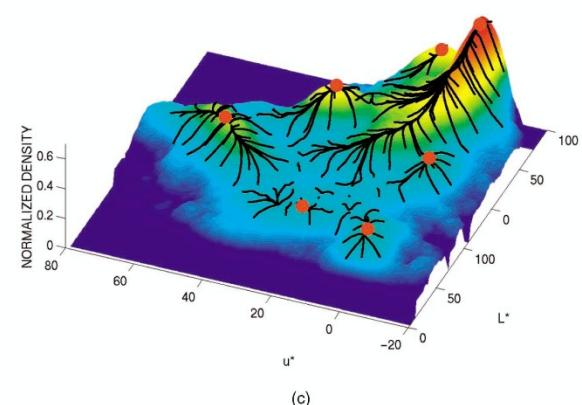
- Compute features for each pixel (color, gradients, texture, etc.).
- Set kernel size for features  $K_f$  and position  $K_s$ .
- Initialize windows at individual pixel locations.
- Perform mean shift for each window until convergence.
- Merge windows that are within width of  $K_f$  and  $K_s$ .



(a)



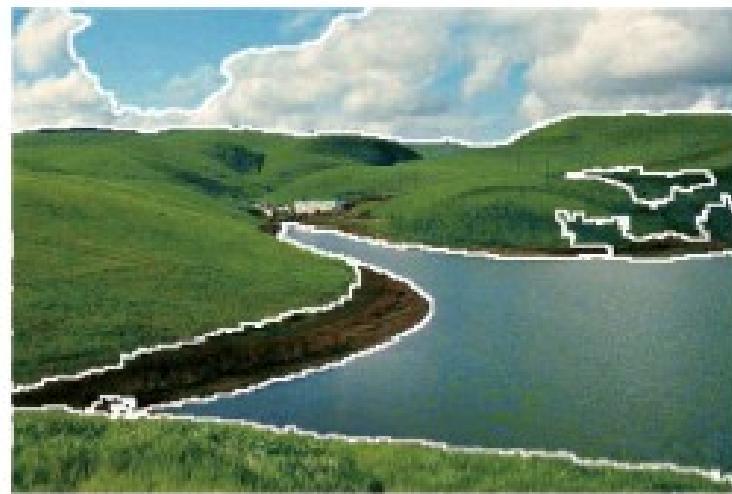
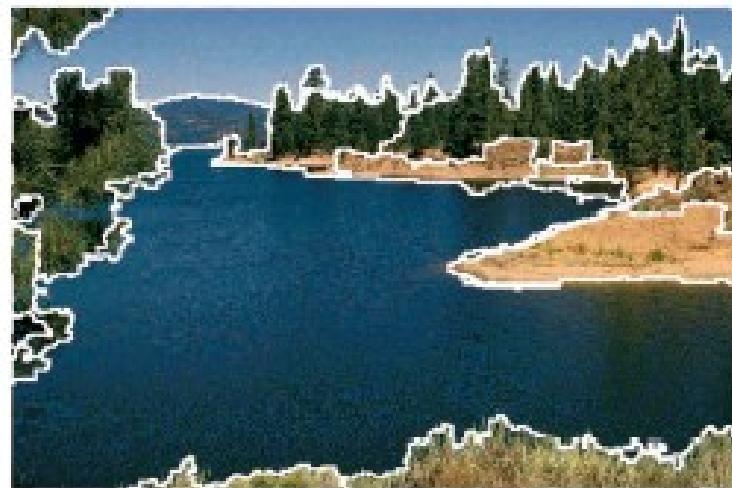
(b)



(c)

# Mean shift segmentation results





Comaniciu and Meer 2002

# Mean shift pros and cons

- Pros
  - Good general-practice segmentation
  - Flexible in number and shape of regions
  - Robust to outliers
- Cons
  - Have to choose kernel size in advance
  - Not suitable for high-dimensional features
- When to use it
  - Oversegmentation
  - Multiple segmentations
  - Tracking, clustering, filtering applications

# Spectral clustering

Group points based on graph structure & edge costs.

Captures “neighborhood-ness” or local smoothness.

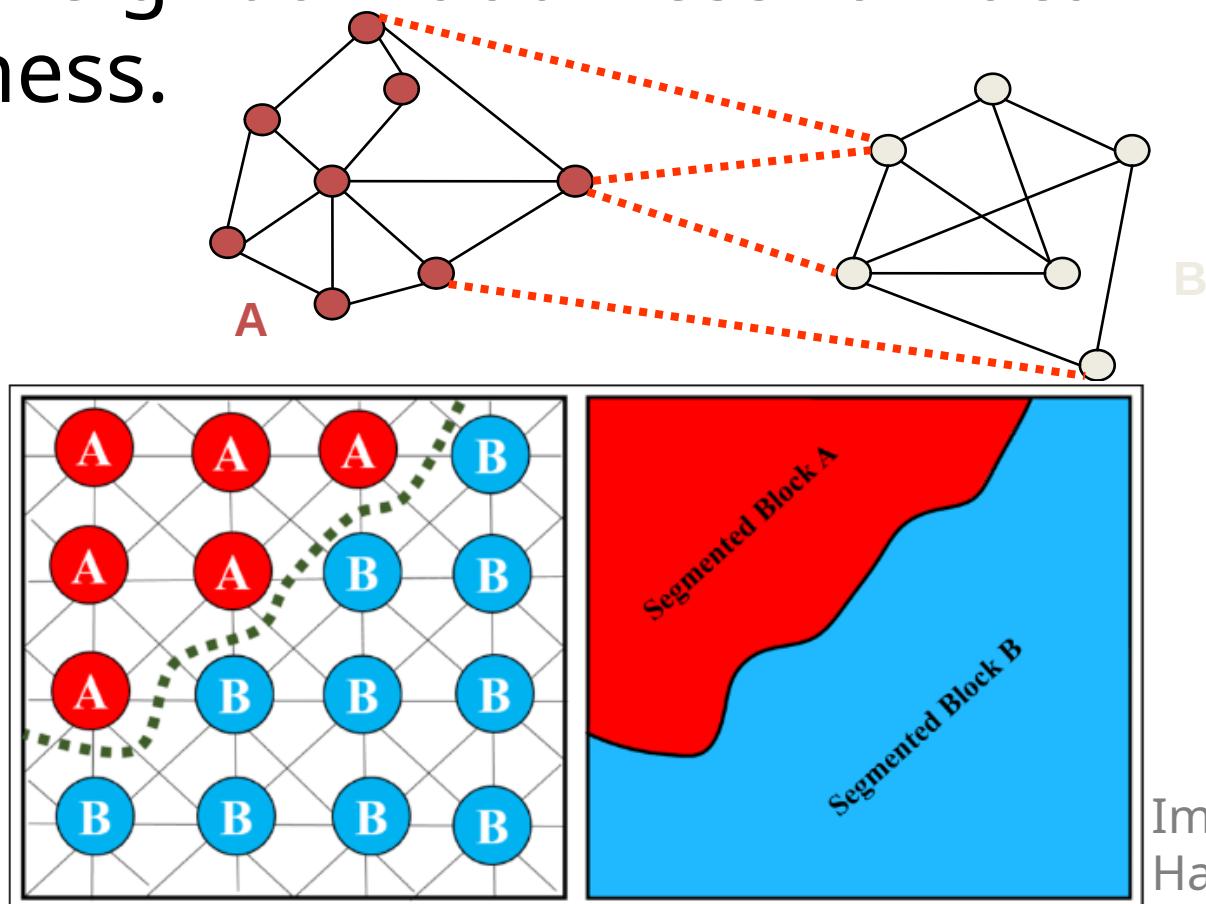
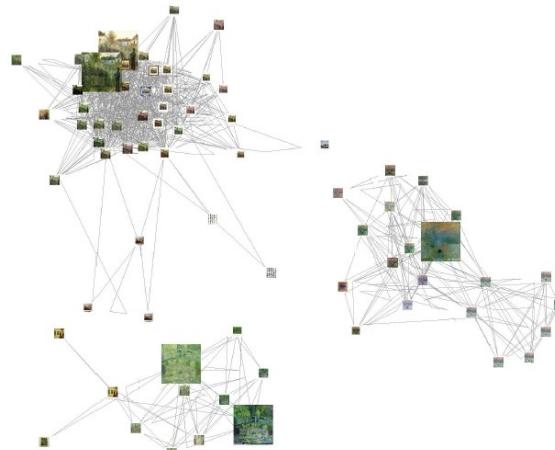
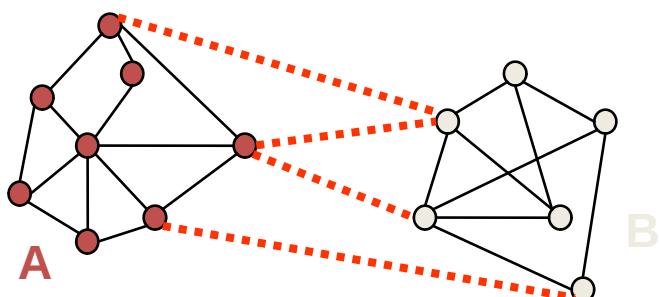


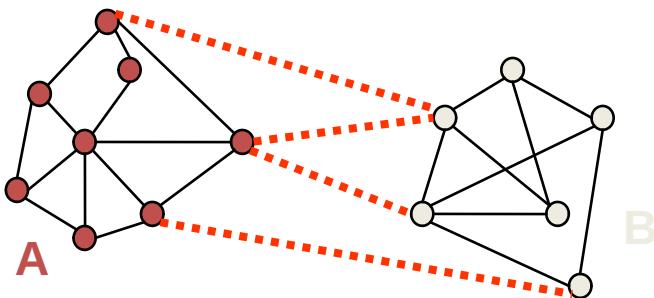
Image:  
Hassan et al.

# Spectral clustering

Group points based on links in a graph



# Cuts in a graph



## Normalized Cut

- a cut penalizes large segments
- fix by normalizing for size of segments

$$Ncut(A, B) = \frac{cut(A, B)}{volume(A)} + \frac{cut(A, B)}{volume(B)}$$

- $volume(A)$  = sum of costs of all edges that touch A

# Normalized cuts for segmentation

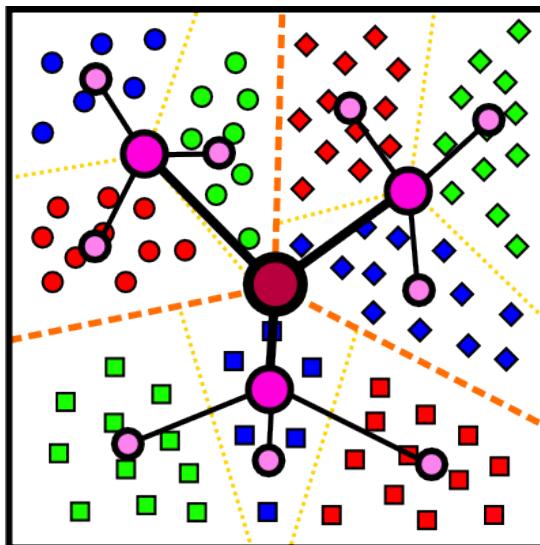


# Segmentation as a result

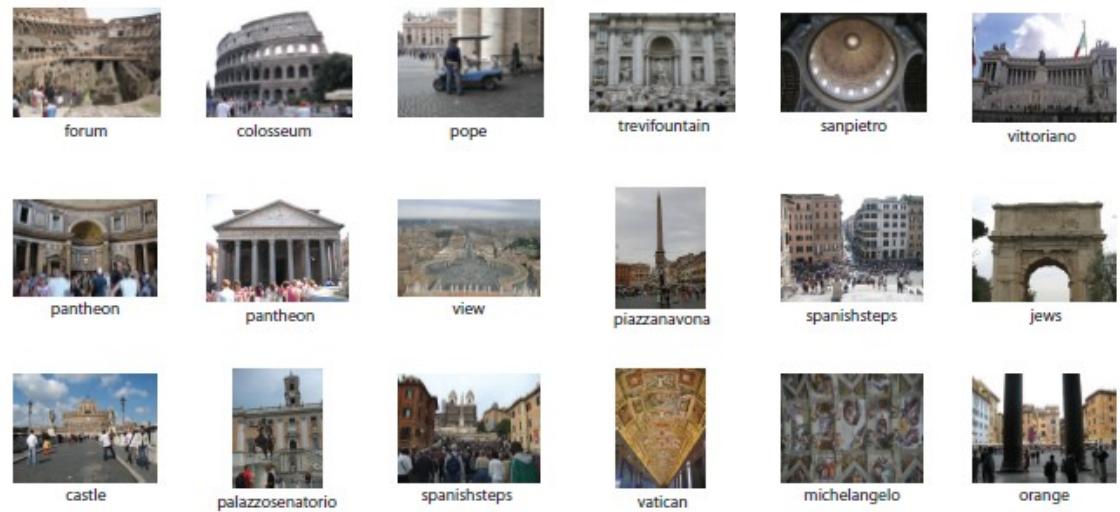


# Which algorithm to use?

- Quantization/Summarization: K-means
  - Aims to preserve variance of original data
  - Can easily assign new point to a cluster



Quantization for  
computing histograms



Summary of 20,000 photos of Rome using  
“greedy k-means”

<http://grail.cs.washington.edu/projects/canonview/>

# Which algorithm to use?

- Image segmentation: agglomerative clustering
  - More flexible with distance measures (e.g., can be based on boundary prediction)
  - Adapts better to specific data
  - Hierarchy can be useful

