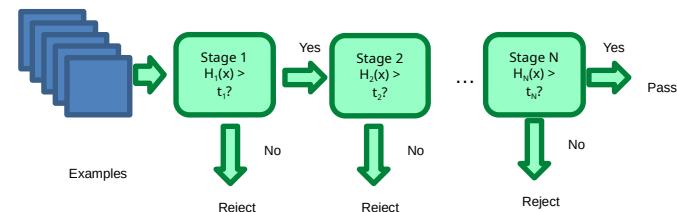
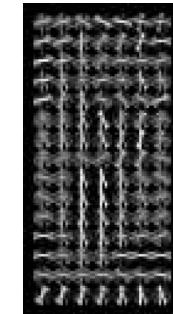


# Advanced Image Processing

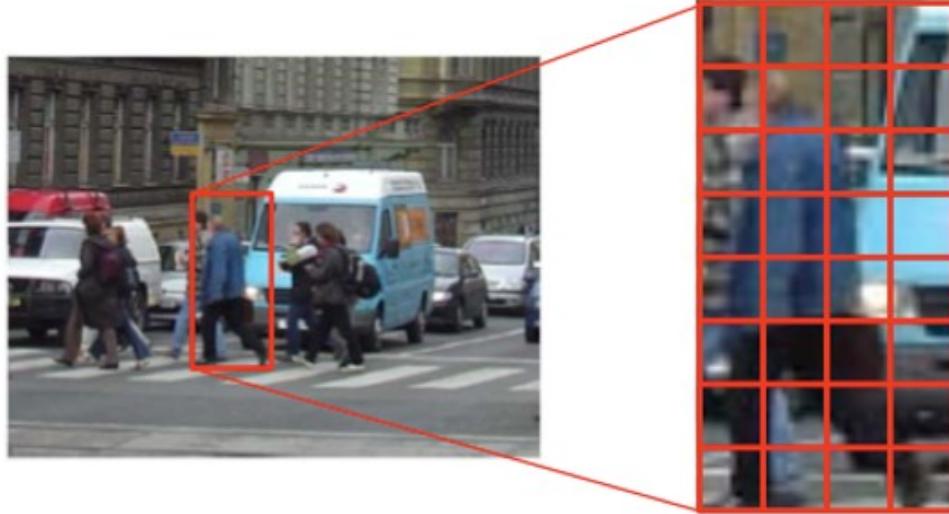
## Big Data

# Object detection

- Sliding window for search
- Features based on differences of intensity (gradient, wavelet, etc.)
- Boosting for feature selection
- Integral images, cascade for speed
- Bootstrapping to deal with many, many negative examples



# Starting point: sliding window classifiers

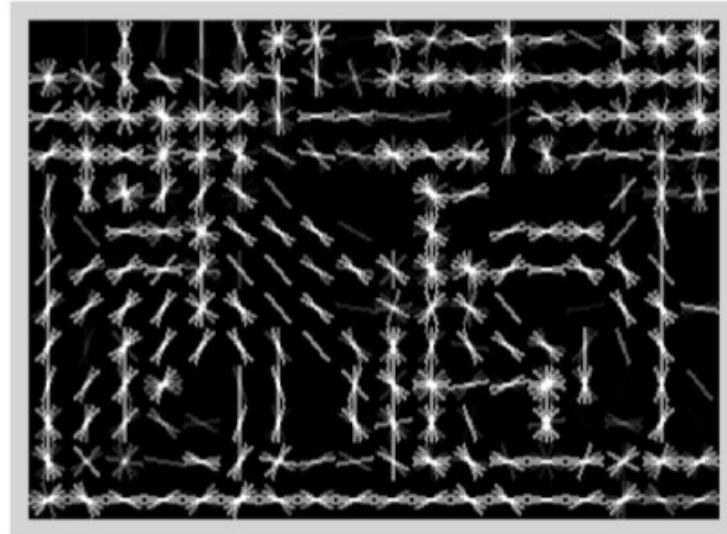


Feature vector

$$x = [\dots, \dots, \dots, \dots]$$

- Detect objects by testing each subwindow
  - Reduces object detection to binary classification
  - Dalal & Triggs: HOG features + linear SVM classifier
  - Previous state of the art for detecting people

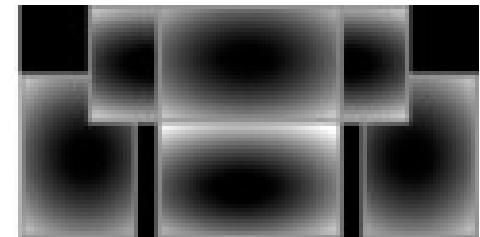
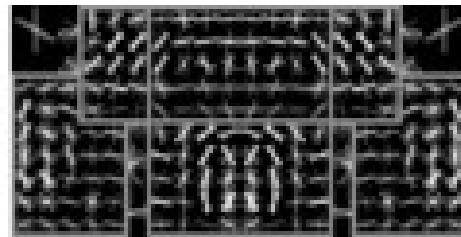
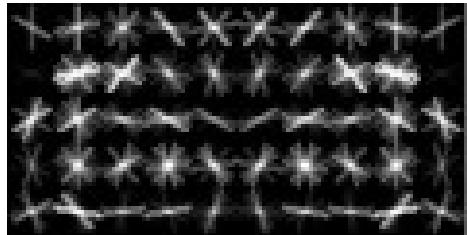
# Histogram of Gradient (HOG) features



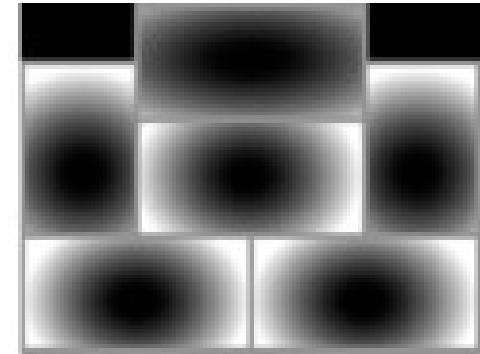
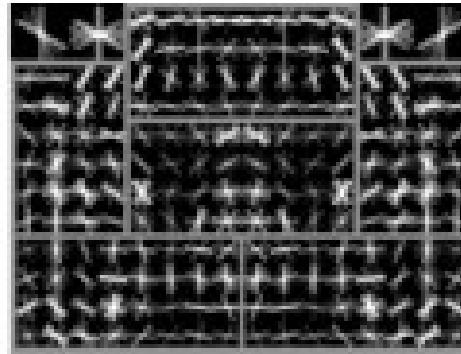
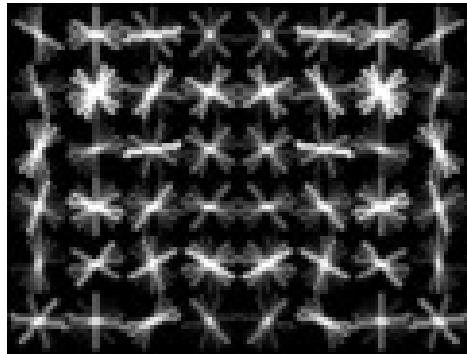
- Image is partitioned into 8x8 pixel blocks
- In each block we compute a histogram of gradient orientations
  - **Invariant** to changes in lighting, small deformations, etc.
- Compute features at different resolutions (pyramid)

# Car model

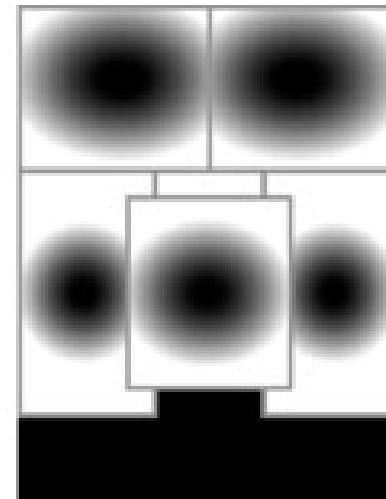
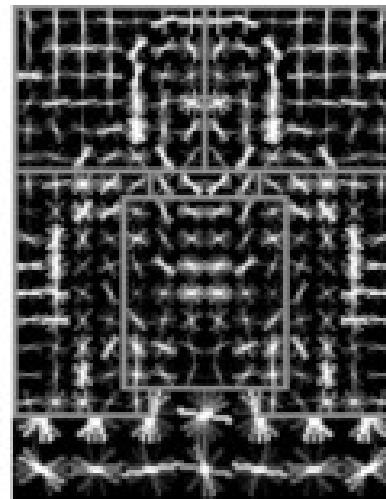
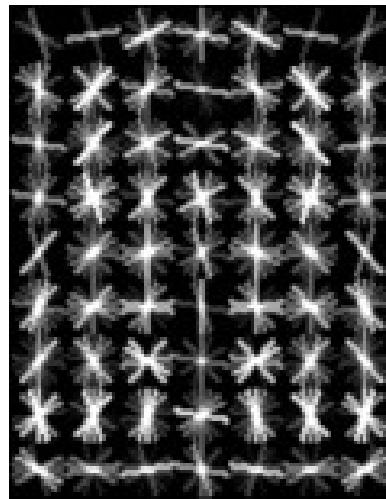
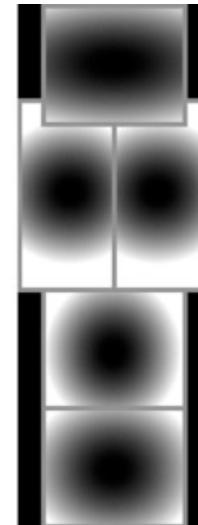
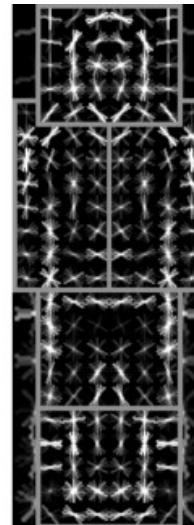
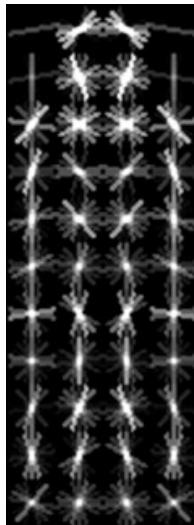
Component 1



Component 2

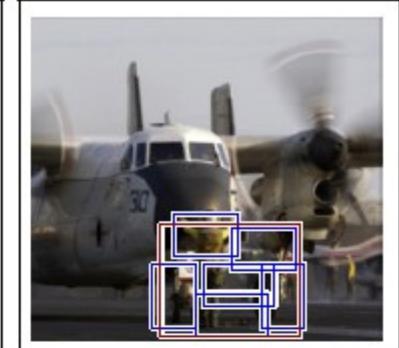
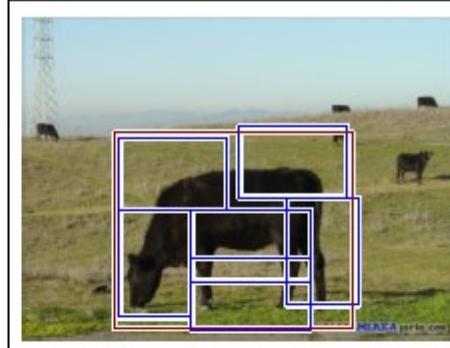
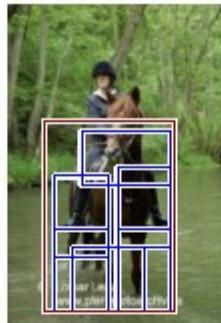
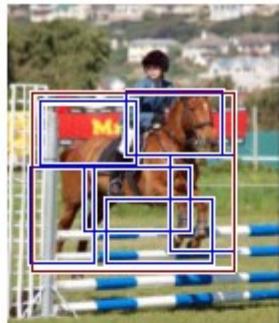
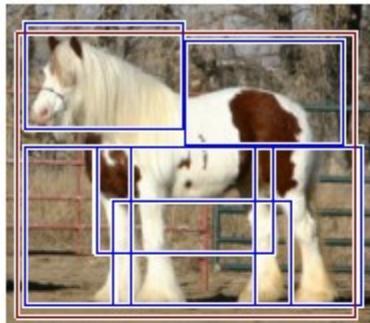


# Person model

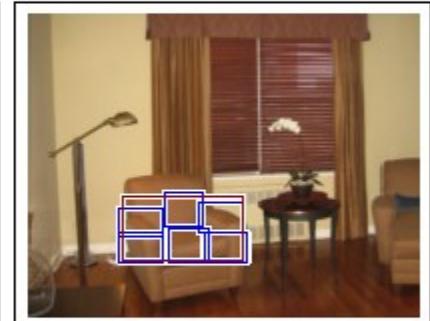
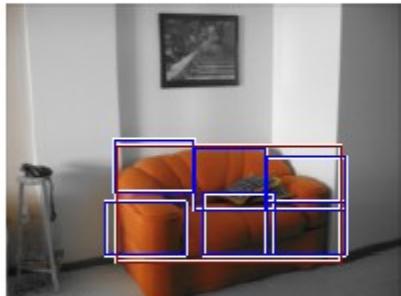


# Good detections?

horse



sofa

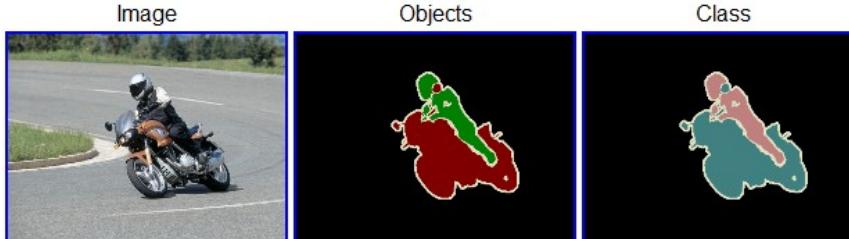


bottle



# The PASCAL Visual Object Classes Challenge 2009 (VOC2009)

- Twenty object categories (aeroplane to TV/monitor)
- Three challenges:
  - Classification challenge (is there an X in this image?)
  - Detection challenge (draw a box around every X)
  - Segmentation challenge



# Dataset: Collection

---

- Images downloaded from **flickr**
  - 500,000 images downloaded and random subset selected for annotation
  - Queries
    - Keyword e.g. “car”, “vehicle”, “street”, “downtown”
    - Date of capture e.g. “taken 21-July”
      - Removes “recency” bias in flickr results
    - Images selected from random page of results
      - Reduces bias toward particular flickr users

# Dataset: Annotation

---

- “Complete” annotation of all objects
- Annotated over web with written guidelines
  - High quality (?)

20 classes.

- Train / validation data has 11,530 images containing 27,450 ROI annotated objects and 6,929 segmentations.

# Examples

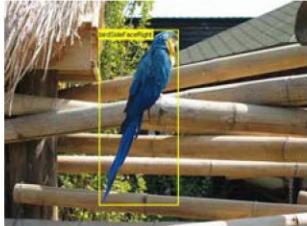
Aeroplane



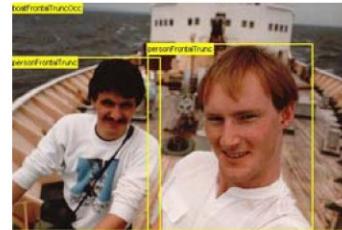
Bicycle



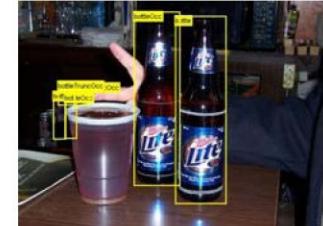
Bird



Boat



Bottle



Bus



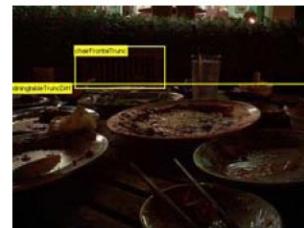
Car



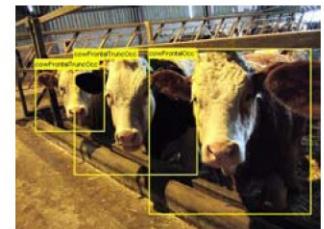
Cat



Chair



Cow



# Examples

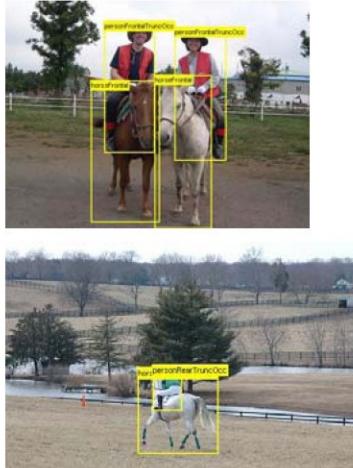
Dining Table



Dog



Horse



Motorbike



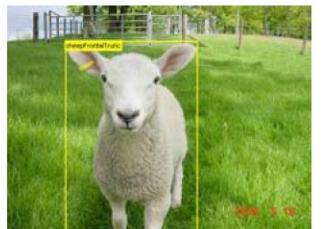
Person



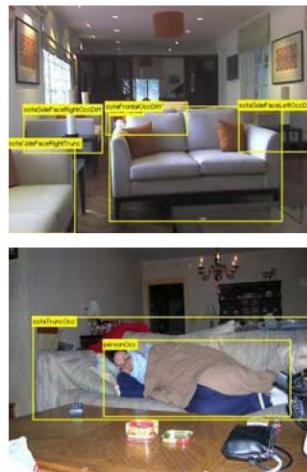
Potted Plant



Sheep



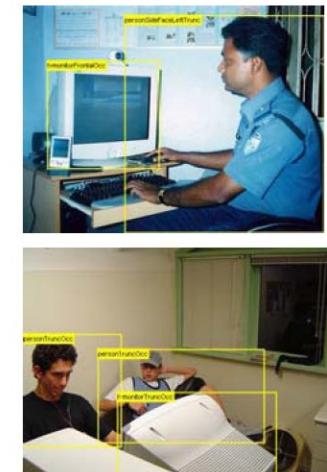
Sofa



Train



TV/Monitor



# Classification Challenge

- Predict whether at least one object of a given class is present in an image



is there a cat?

# Participation

---

- 48 Methods, 20 Groups
- VOC2008: 21 Methods, 11 Groups
- Overwhelmingly “bag of visual words” methods with multiple features e.g. SIFT, color
- Multiple submissions of methods with small variations e.g. different features or classifier architectures

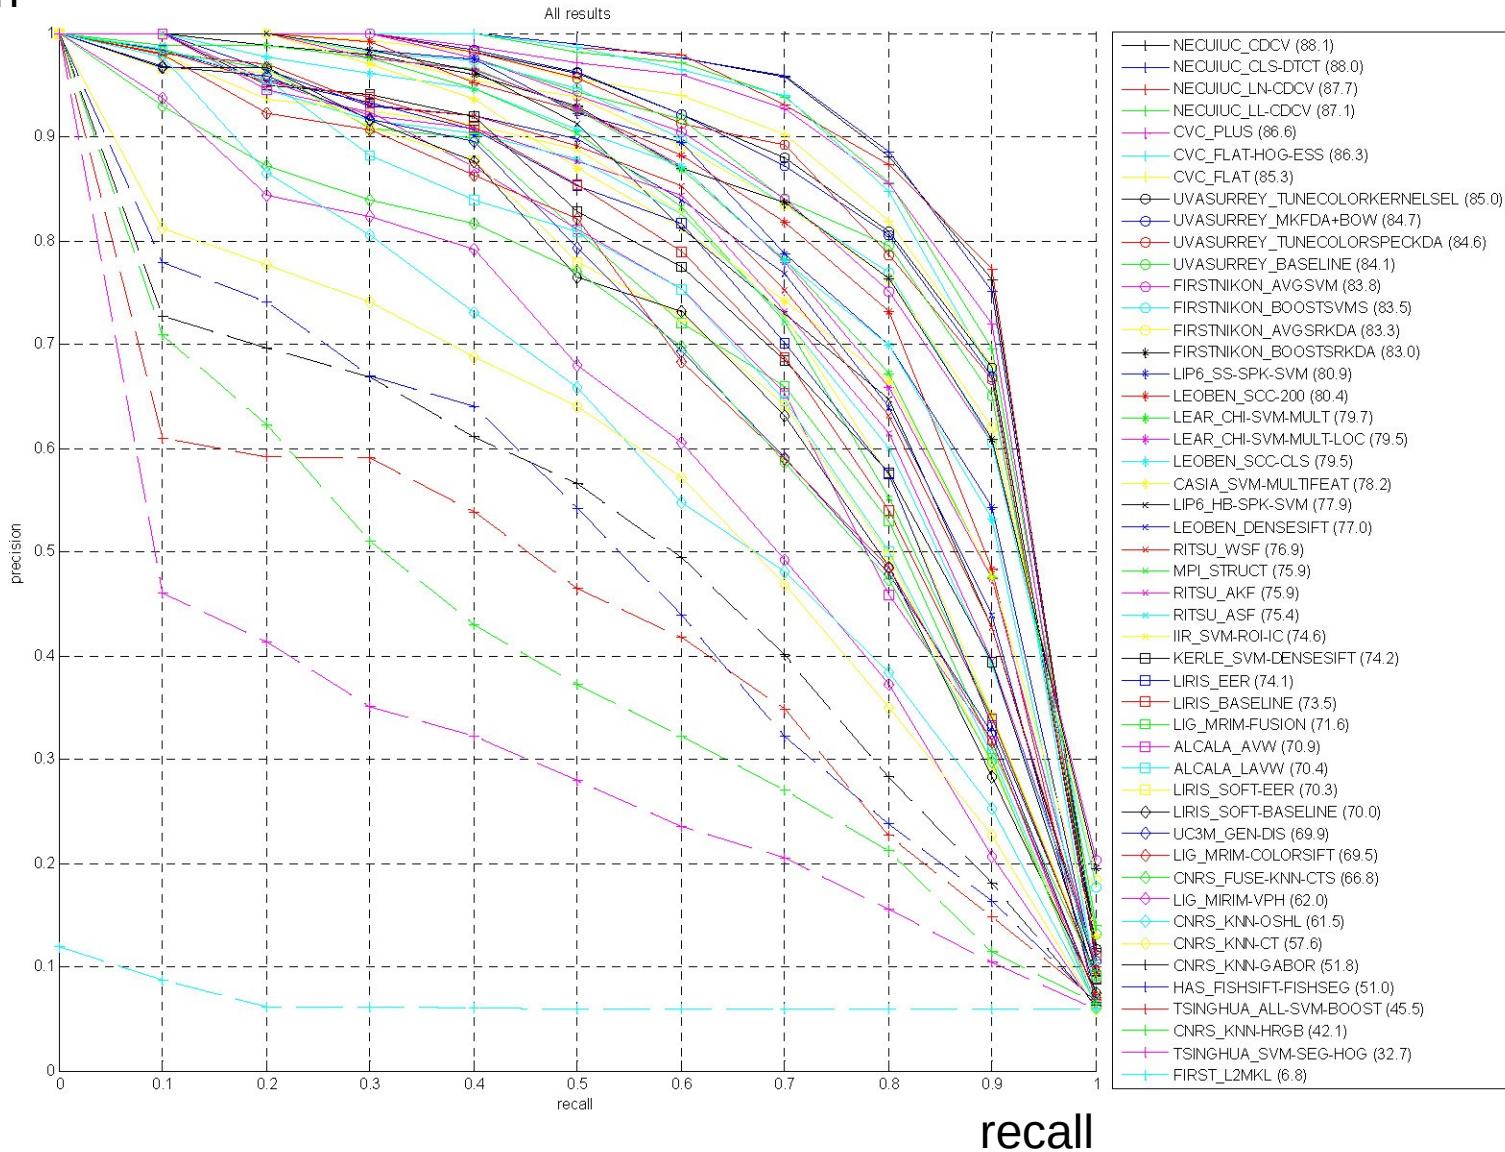
# Results: AP by Method and Class

|                                 | aero<br>plane | bicycle | bird | boat | bottle | bus  | car  | ccf  | chair | cow  | dining<br>table | dog  | horse | motor<br>bike | person | potted<br>plant | sheep | sofa | train | tv/<br>monitor |
|---------------------------------|---------------|---------|------|------|--------|------|------|------|-------|------|-----------------|------|-------|---------------|--------|-----------------|-------|------|-------|----------------|
| CVC_FLAT                        | 85.3          | 57.8    | 66.0 | 66.1 | 36.2   | 70.6 | 60.6 | 63.5 | 55.1  | 44.6 | 53.4            | 49.1 | 64.4  | 66.8          | 84.8   | 37.4            | 44.1  | 47.9 | 81.9  | 67.5           |
| CVC_FLAT-HOG-ESS                | 86.3          | 60.7    | 66.4 | 65.3 | 41.0   | 71.7 | 64.7 | 63.9 | 55.5  | 40.1 | 51.3            | 45.9 | 65.2  | 68.9          | 85.0   | 40.8            | 49.0  | 49.1 | 81.8  | 68.6           |
| CVC_PLUS                        | 86.6          | 58.4    | 66.7 | 67.3 | 34.8   | 70.4 | 60.0 | 64.2 | 52.5  | 43.0 | 50.8            | 46.5 | 64.1  | 66.8          | 84.4   | 37.5            | 45.1  | 45.4 | 82.1  | 67.0           |
| FIRSTNIKON_AVGSRKDA             | 83.3          | 59.3    | 62.7 | 65.3 | 30.2   | 71.6 | 58.2 | 62.2 | 54.3  | 40.7 | 49.2            | 50.0 | 66.6  | 62.9          | 83.3   | 34.2            | 48.2  | 46.1 | 83.4  | 65.5           |
| FIRSTNIKON_AVGSVM               | 83.8          | 58.2    | 62.6 | 65.2 | 32.0   | 69.8 | 57.7 | 61.1 | 54.5  | 44.0 | 50.3            | 49.6 | 64.6  | 61.7          | 83.2   | 33.4            | 46.5  | 48.0 | 81.6  | 65.3           |
| FIRSTNIKON_BOOSTSRKDA           | 83.0          | 59.2    | 61.4 | 64.6 | 33.2   | 71.1 | 57.5 | 61.0 | 54.8  | 40.7 | 48.3            | 50.0 | 65.5  | 63.4          | 82.8   | 32.8            | 47.0  | 47.1 | 83.3  | 64.6           |
| FIRSTNIKON_BOOSTSVMS            | 83.5          | 56.8    | 61.8 | 65.5 | 33.2   | 69.7 | 57.3 | 60.5 | 54.6  | 43.1 | 48.3            | 50.3 | 64.3  | 62.4          | 82.3   | 32.9            | 46.9  | 48.4 | 82.0  | 64.2           |
| LEAR_CHI-SVM-MULT-LOC           | 79.5          | 55.5    | 54.5 | 63.9 | 43.7   | 70.3 | 66.4 | 56.5 | 54.4  | 38.8 | 44.1            | 46.2 | 58.5  | 64.2          | 82.2   | 39.1            | 41.3  | 39.8 | 73.6  | 66.2           |
| NECUIUC_CDCV                    | 88.1          | 68.0    | 68.0 | 72.5 | 41.0   | 78.9 | 70.4 | 70.4 | 58.1  | 53.4 | 55.7            | 59.3 | 73.1  | 71.3          | 84.5   | 32.3            | 53.3  | 56.7 | 86.0  | 66.8           |
| NECUIUC_CLS-DTCT                | 88.0          | 68.6    | 67.9 | 72.9 | 44.2   | 79.5 | 72.5 | 70.8 | 59.5  | 53.6 | 57.5            | 59.0 | 72.6  | 72.3          | 85.3   | 36.6            | 56.9  | 57.9 | 85.9  | 68.0           |
| NECUIUC_LL-CDCV                 | 87.1          | 67.4    | 65.8 | 72.3 | 40.9   | 78.3 | 69.7 | 69.7 | 58.5  | 50.1 | 55.1            | 56.3 | 71.8  | 70.8          | 84.1   | 31.4            | 51.5  | 55.1 | 84.7  | 65.2           |
| NECUIUC_LN-CDCV                 | 87.7          | 67.8    | 68.1 | 71.1 | 39.1   | 78.5 | 70.6 | 70.7 | 57.4  | 51.7 | 53.3            | 59.2 | 71.6  | 70.6          | 84.0   | 30.9            | 51.7  | 55.9 | 85.9  | 66.7           |
| UVASURREY_BASELINE              | 84.1          | 59.2    | 62.7 | 65.4 | 35.7   | 70.6 | 59.8 | 61.3 | 56.7  | 45.3 | 52.4            | 50.6 | 66.1  | 66.6          | 83.7   | 34.8            | 47.2  | 47.7 | 80.8  | 65.9           |
| UVASURREY_MKFDA+BOW             | 84.7          | 63.9    | 66.1 | 67.3 | 37.9   | 74.1 | 63.2 | 64.0 | 57.1  | 46.2 | 54.7            | 53.5 | 68.1  | 70.6          | 85.2   | 38.5            | 47.2  | 49.3 | 83.2  | 68.1           |
| UVASURREY_TUNE COLOR KERNEL SEL | 85.0          | 62.8    | 65.1 | 66.5 | 37.6   | 73.5 | 62.1 | 62.0 | 57.4  | 45.1 | 54.5            | 52.5 | 67.7  | 69.8          | 84.8   | 39.1            | 46.8  | 49.9 | 82.9  | 68.1           |
| UVASURREY_TUNE COLOR SPEC KDA   | 84.6          | 62.4    | 65.6 | 67.2 | 39.4   | 74.0 | 63.4 | 62.8 | 56.7  | 43.8 | 54.7            | 52.7 | 67.3  | 70.6          | 85.0   | 38.8            | 46.9  | 50.0 | 82.2  | 66.2           |

- Only methods in 1st, 2nd or 3rd place by group shown
- Groups: CVC, FIRST/Nikon, NEC/UIUC, UVA/Surrey

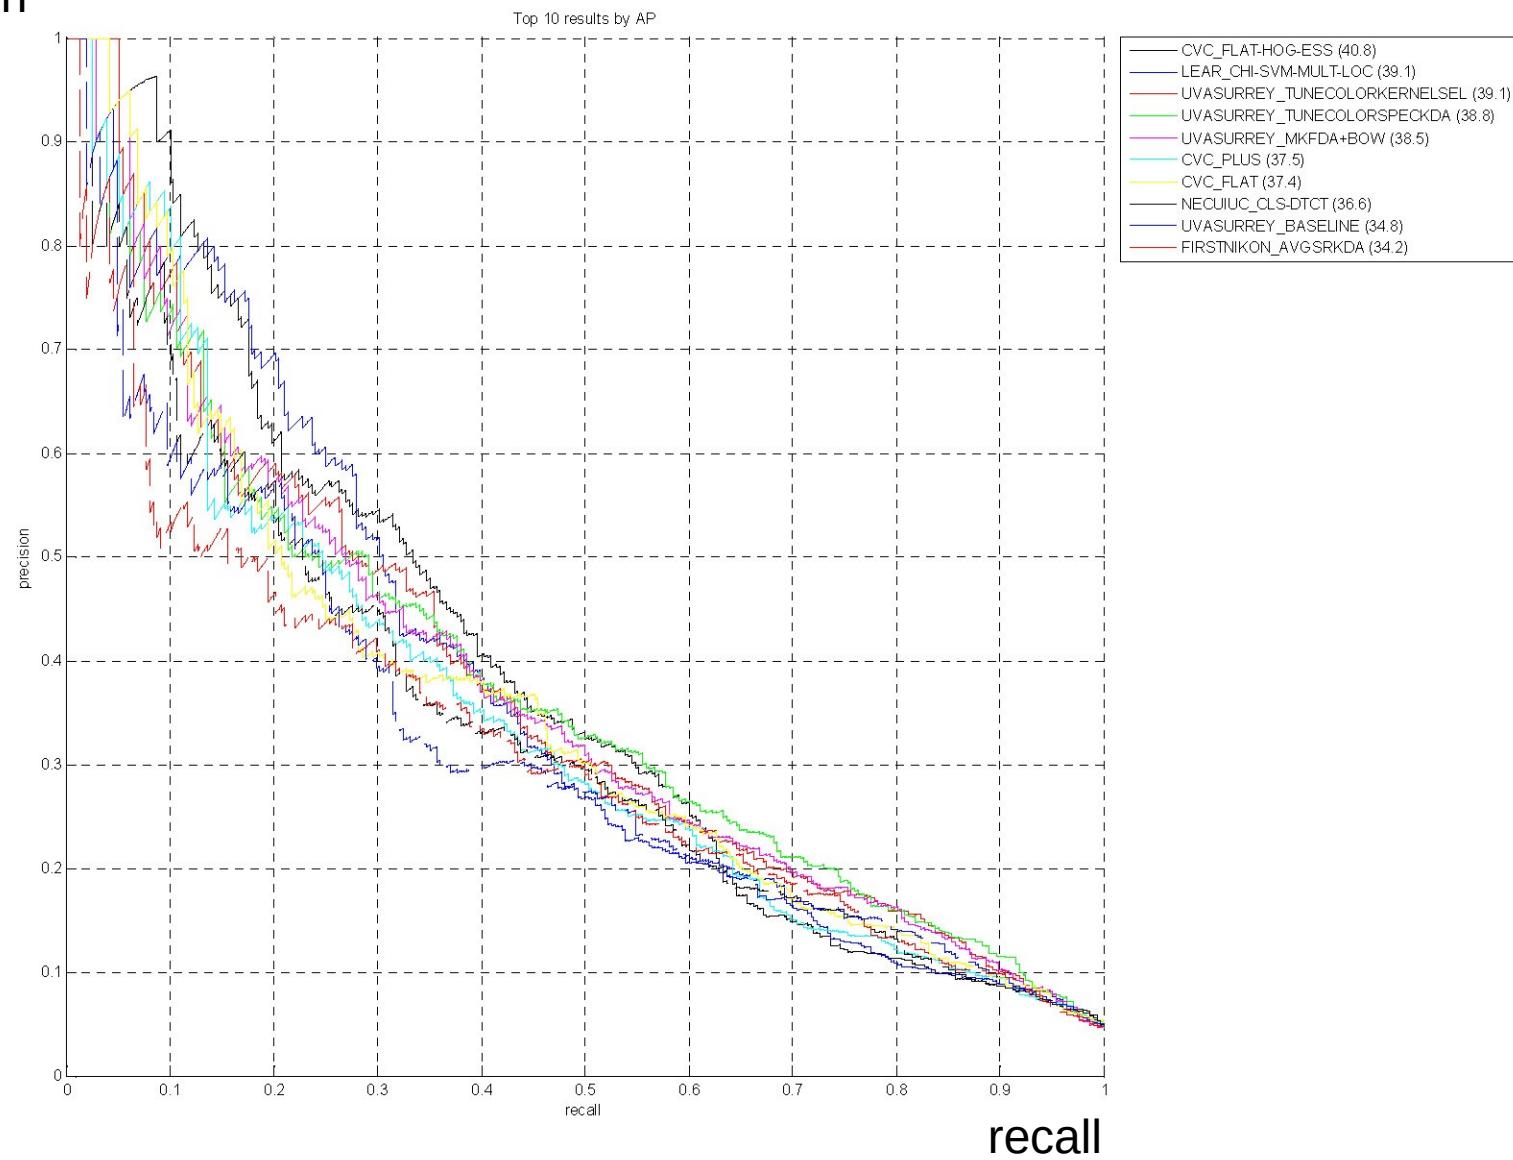
# Precision/Recall: Aeroplane (All)

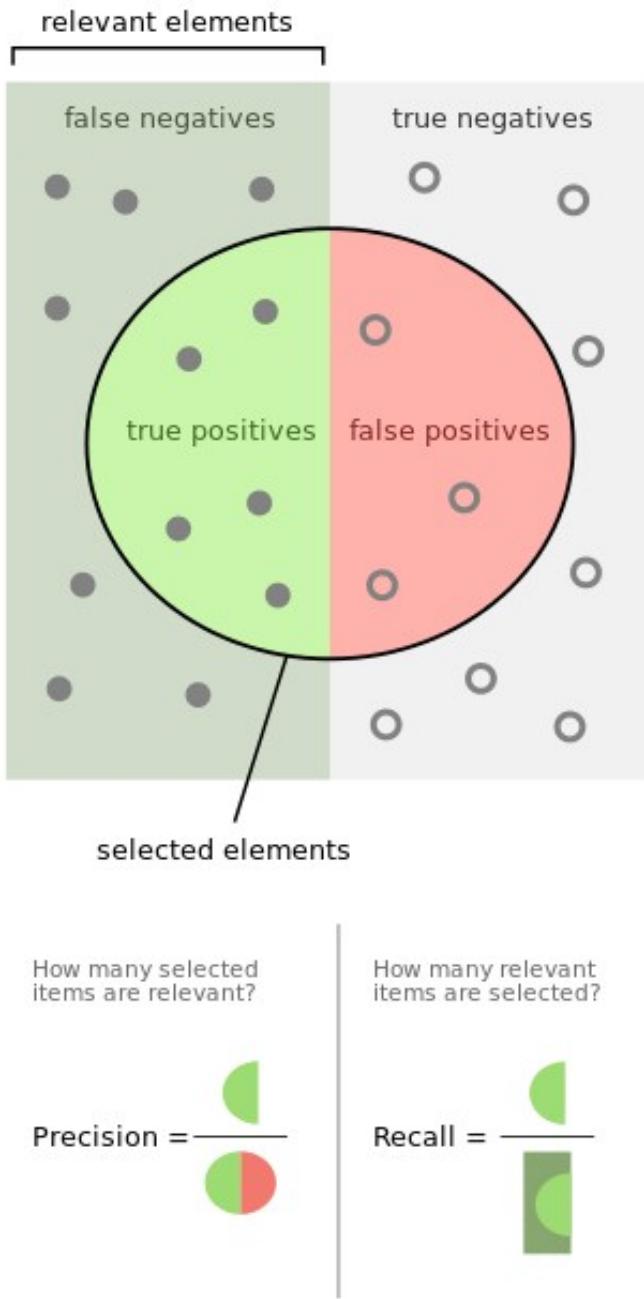
precision



# Precision/Recall: Potted plant (Top 10 by AP)

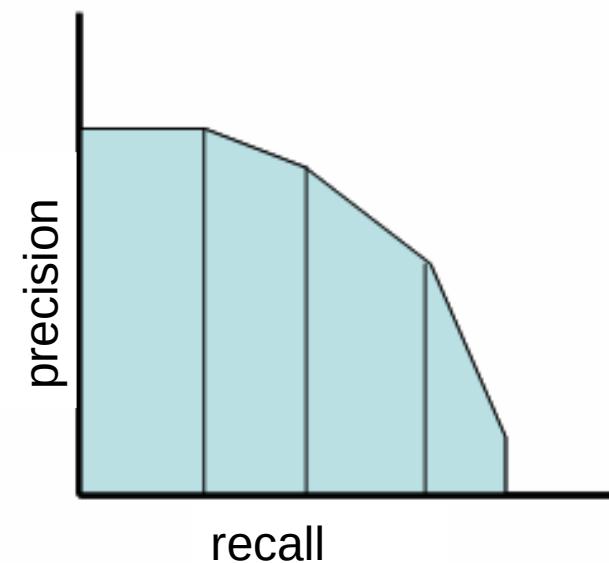
precision





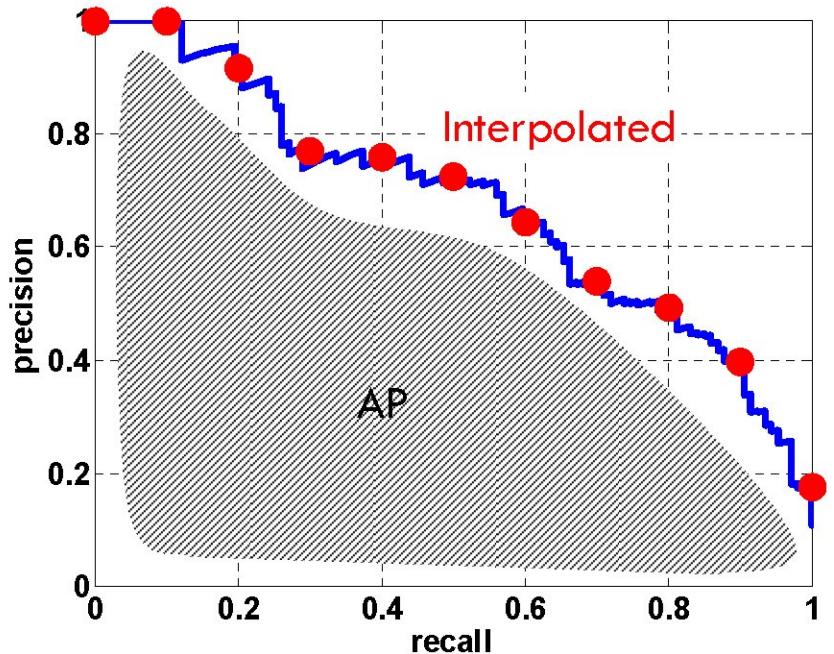
Set threshold on ‘detection’ to create one pair of precision / recall values.

Vary threshold across all values to generate precision / recall curves:



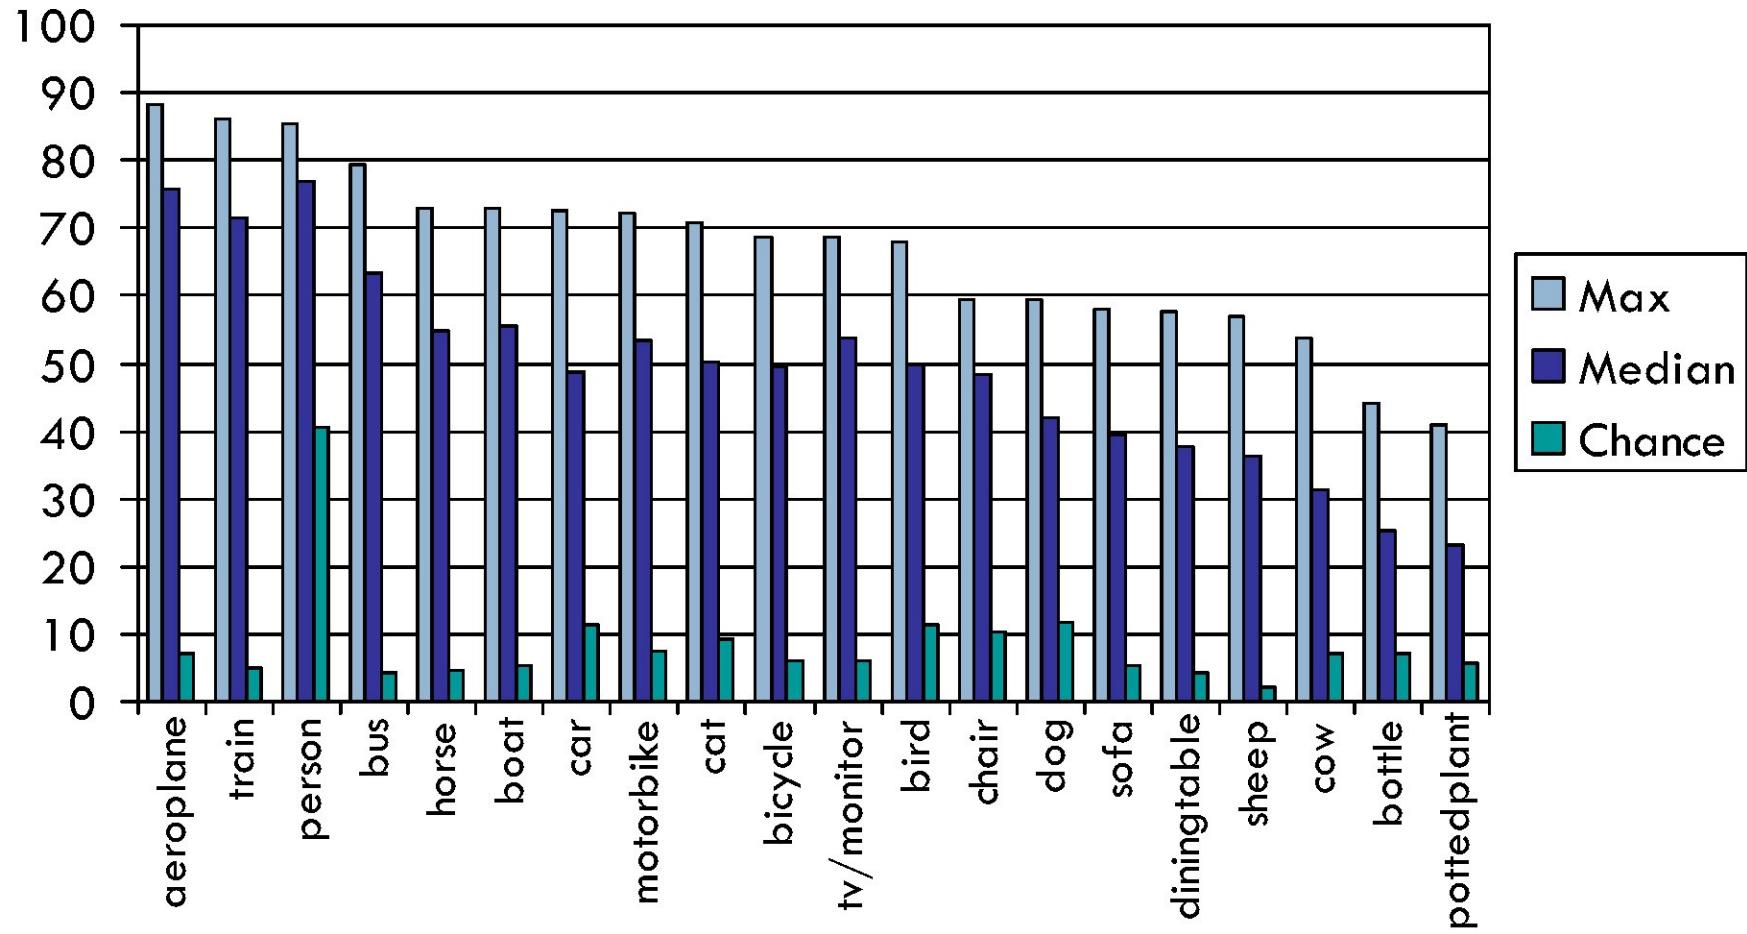
# Evaluation

- Average Precision [TREC] averages precision over the entire range of recall
  - Curve interpolated to reduce influence of “outliers”



- A good score requires both high recall **and** high precision
- Application-independent
- Penalizes methods giving high precision but low recall

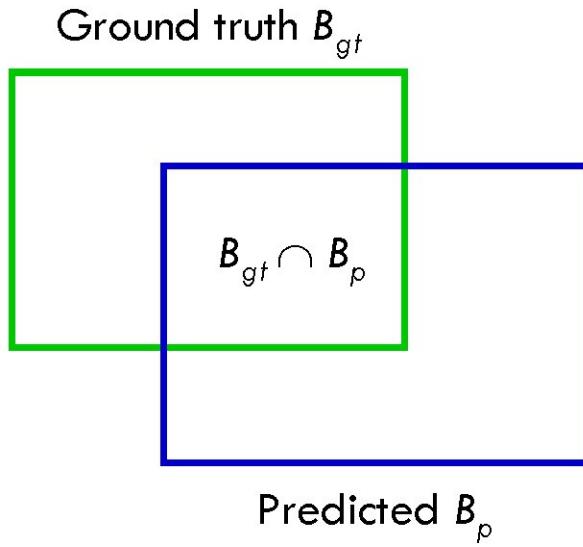
# AP for Image classification



- Max AP: 88.1% (aeroplane) ... 40.8% (potted plant)

# Evaluating Bounding Boxes

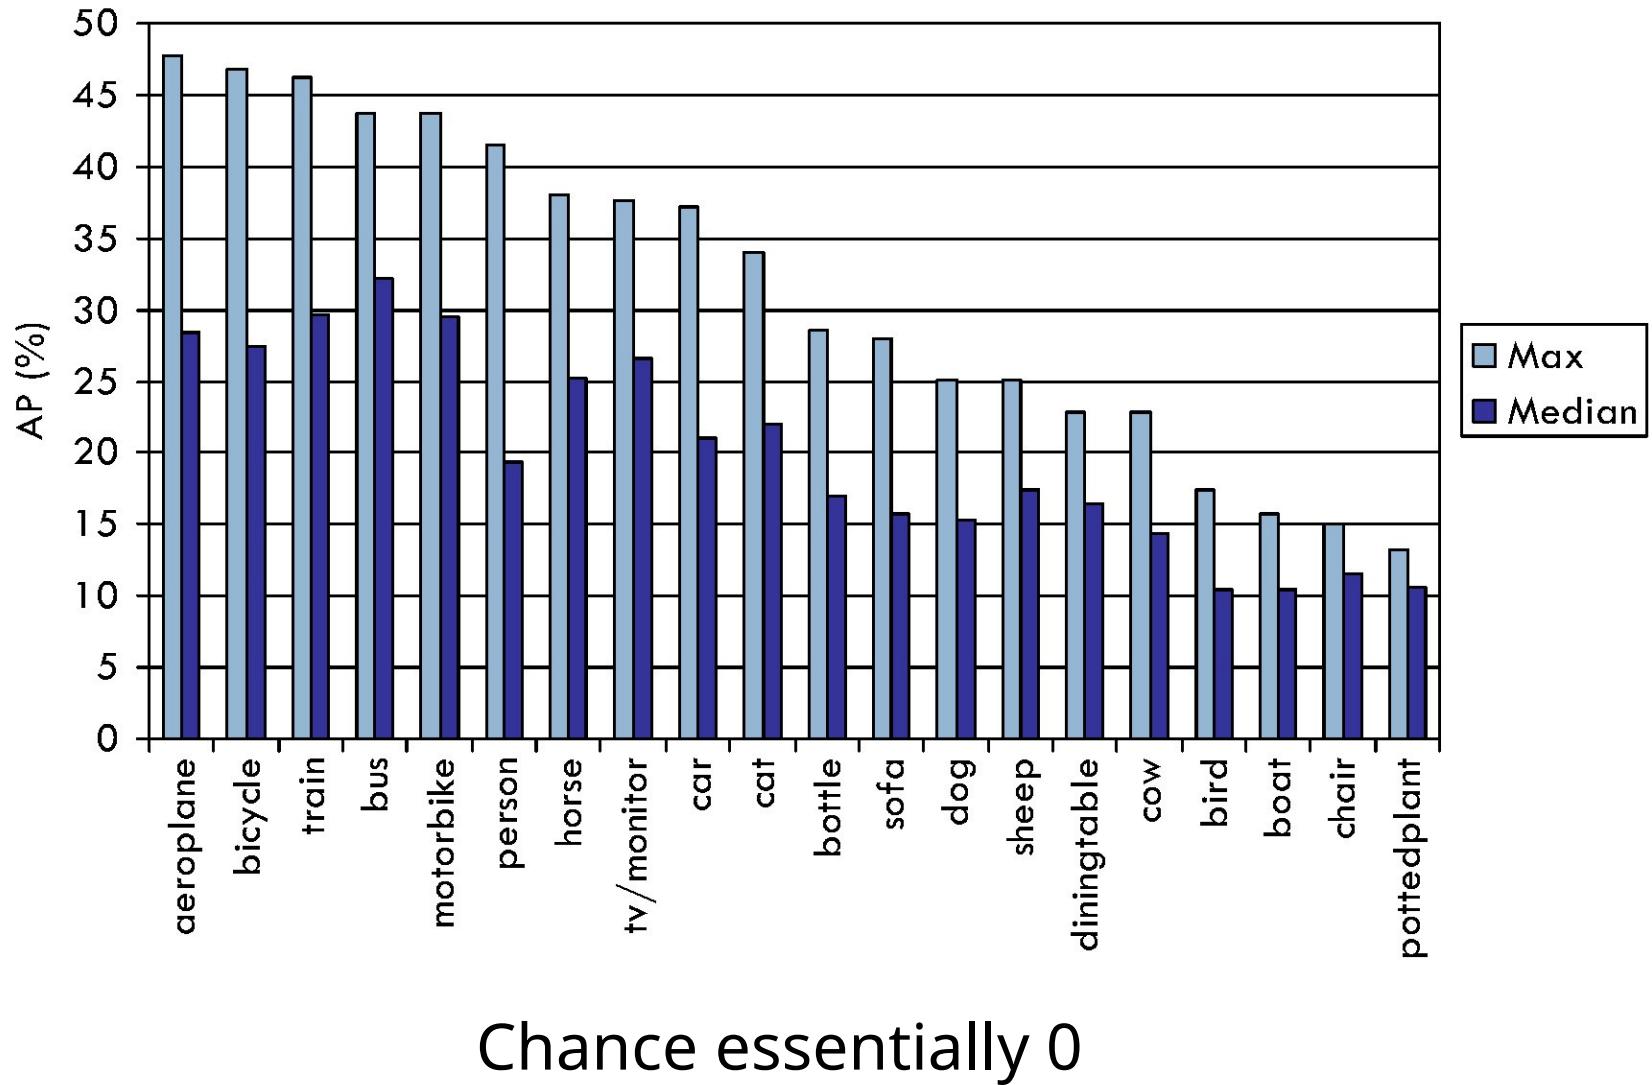
- Area of Overlap (AO) Measure



$$AO(B_{gt}, B_p) = \frac{|B_{gt} \cap B_p|}{|B_{gt} \cup B_p|}$$

- Need to define a threshold  $t$  such that  $AO(B_{gt}, B_p)$  implies a correct detection: 50%

# AP for Object detection

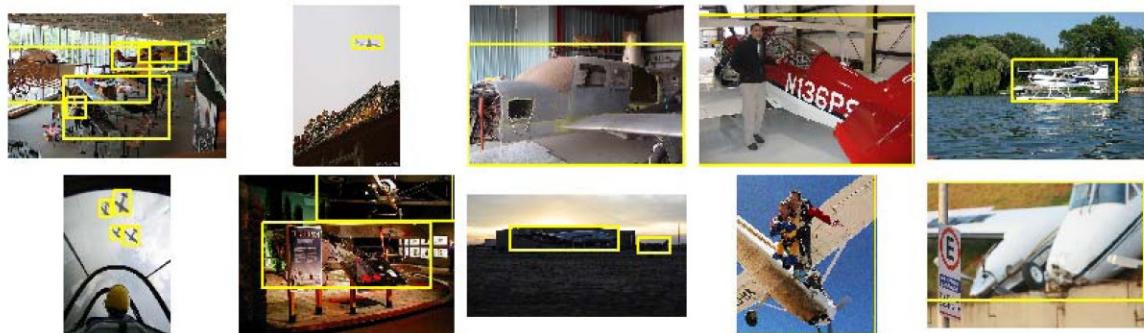


# Ranked Images: Aeroplane

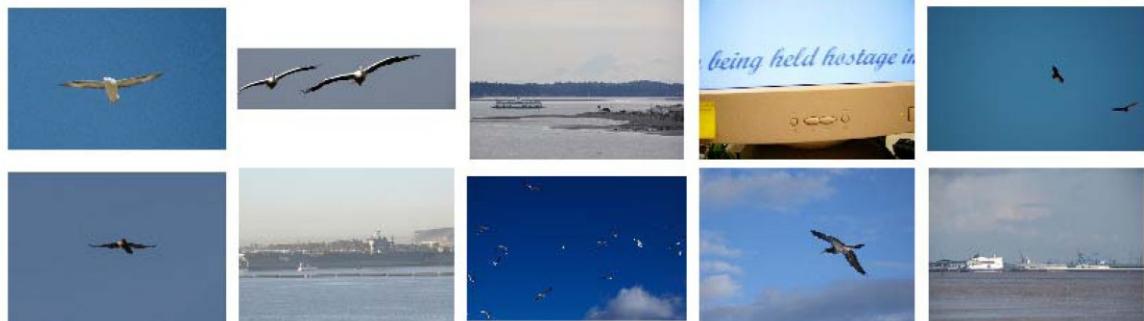
- Class images:  
Highest ranked



- Class images:  
Lowest ranked



- Non-class images:  
Highest ranked



- Context?

# True Positives - Person

UoCTTI\_LSVM-MDPM



MIZZOU\_DEF-HOG-LBP



NECUIUC\_CLS-DTCT



# False Positives - Person

UoCTTI\_LSVM-MDPM



MIZZOU\_DEF-HOG-LBP



NECUIUC\_CLS-DTCT



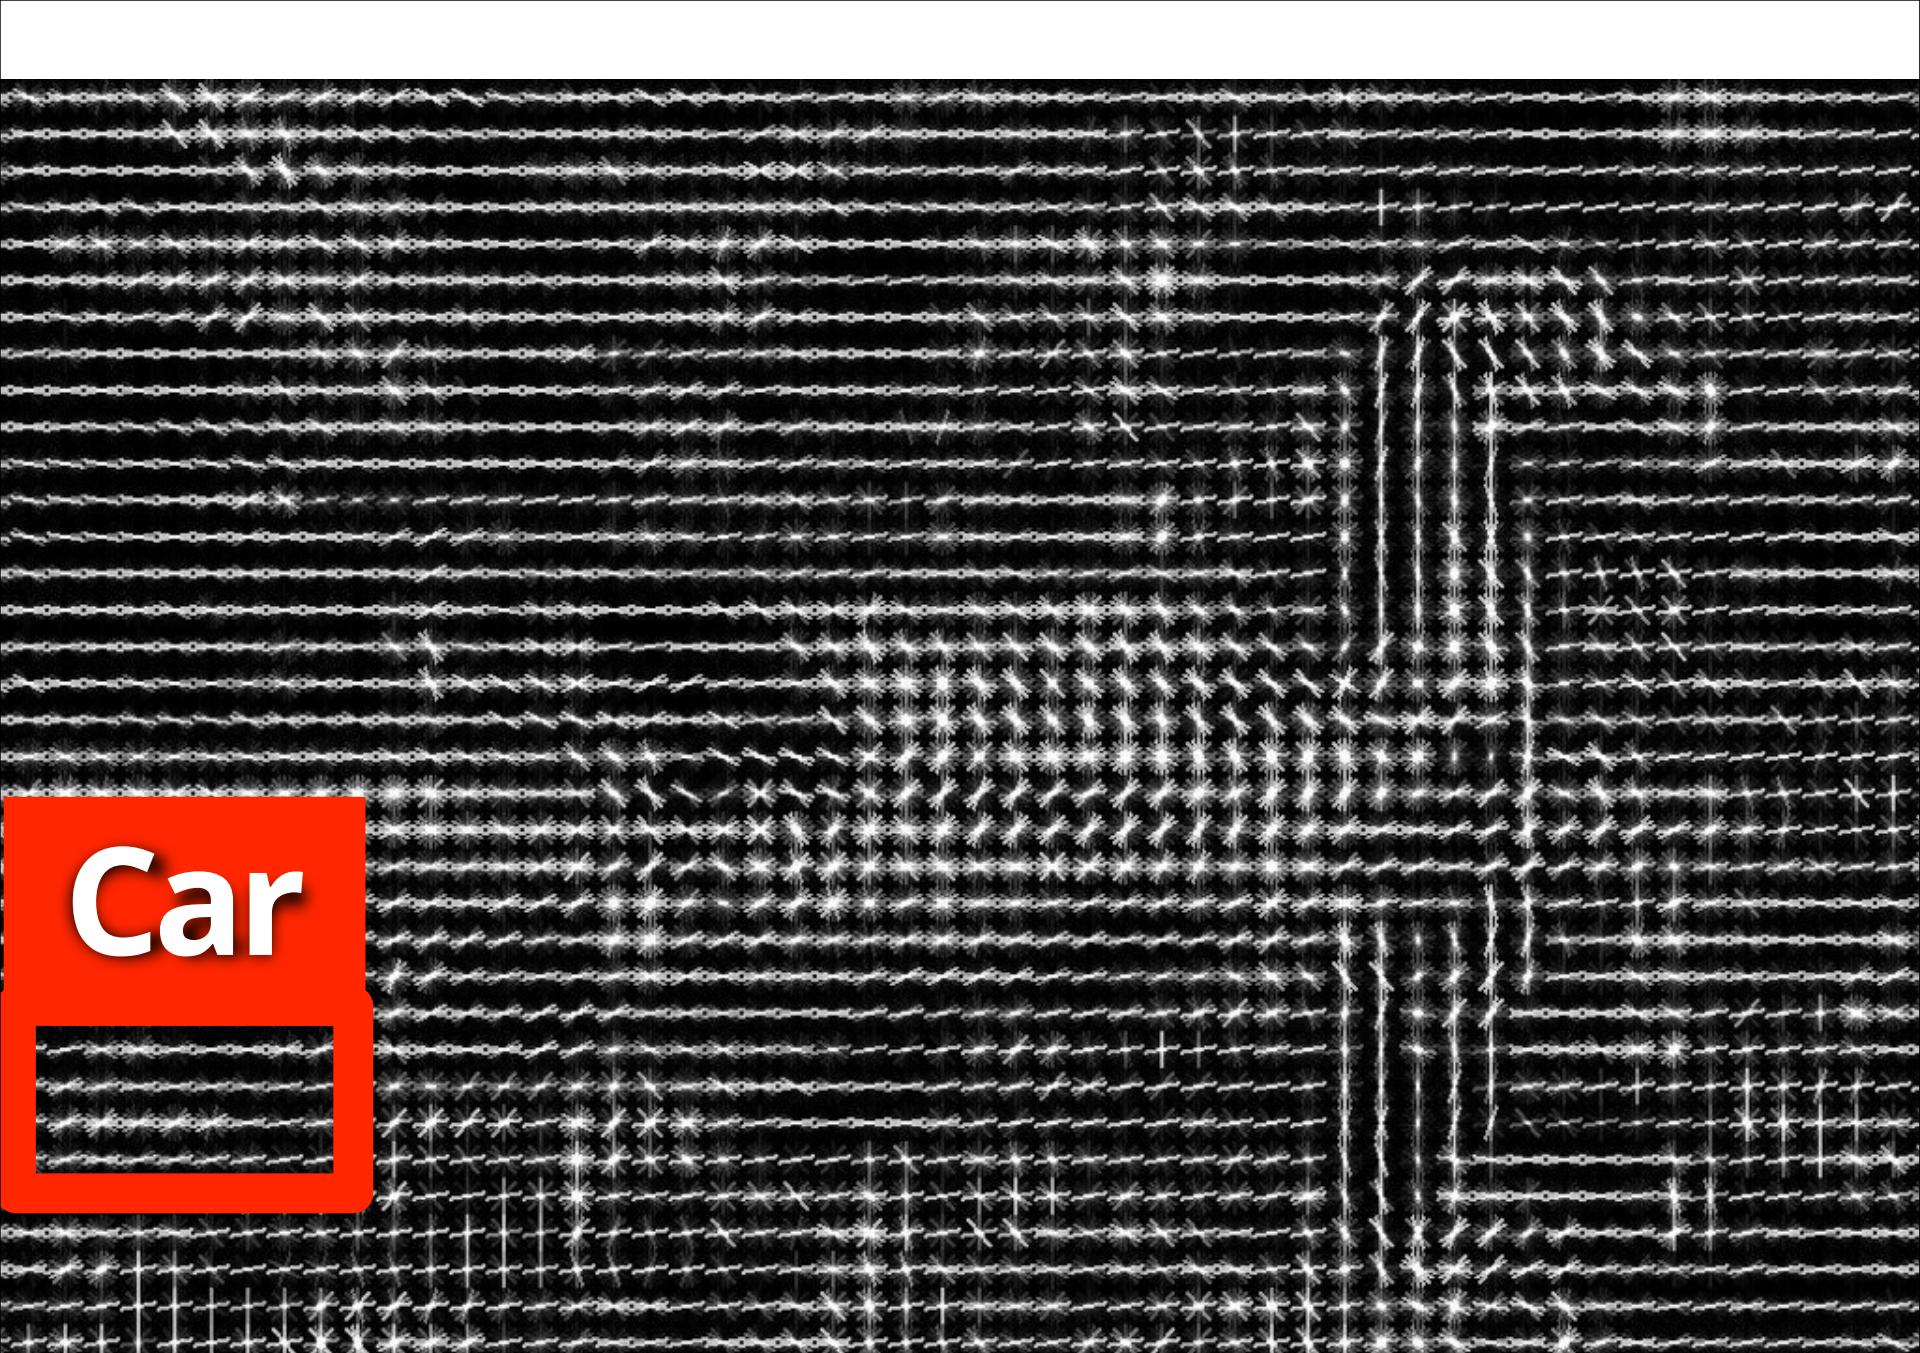
HOGgles (Vondrick et al. ICCV 2013)



Vondrick et

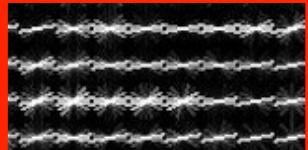
A Canada goose is swimming in a body of water. The bird has a dark brown neck and head, a white patch on its forehead, and a light-colored body with brown wing feathers. It is facing towards the right of the frame. The water is slightly rippled.

Car

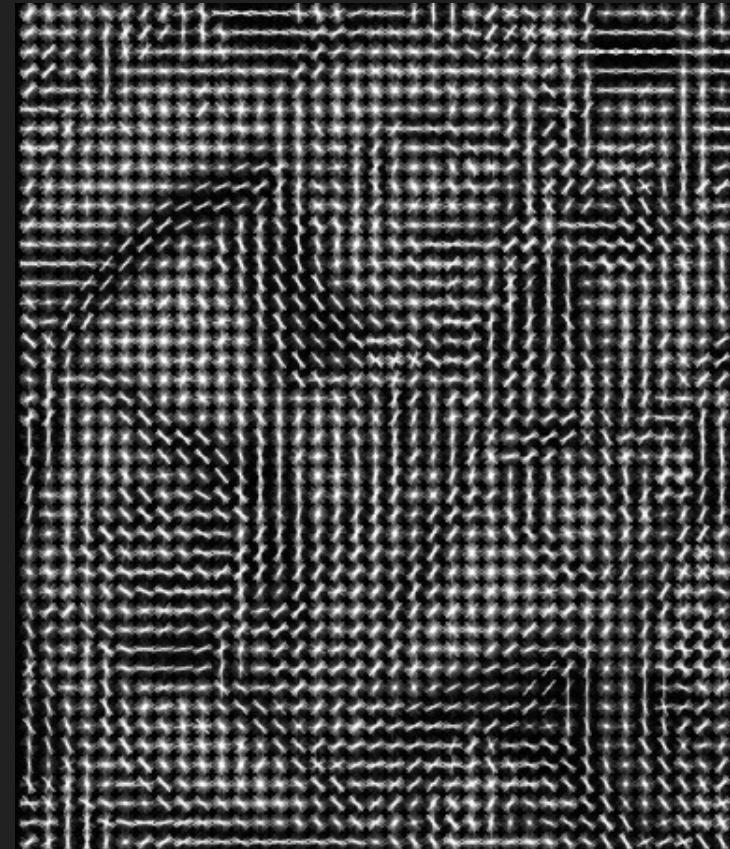


Car

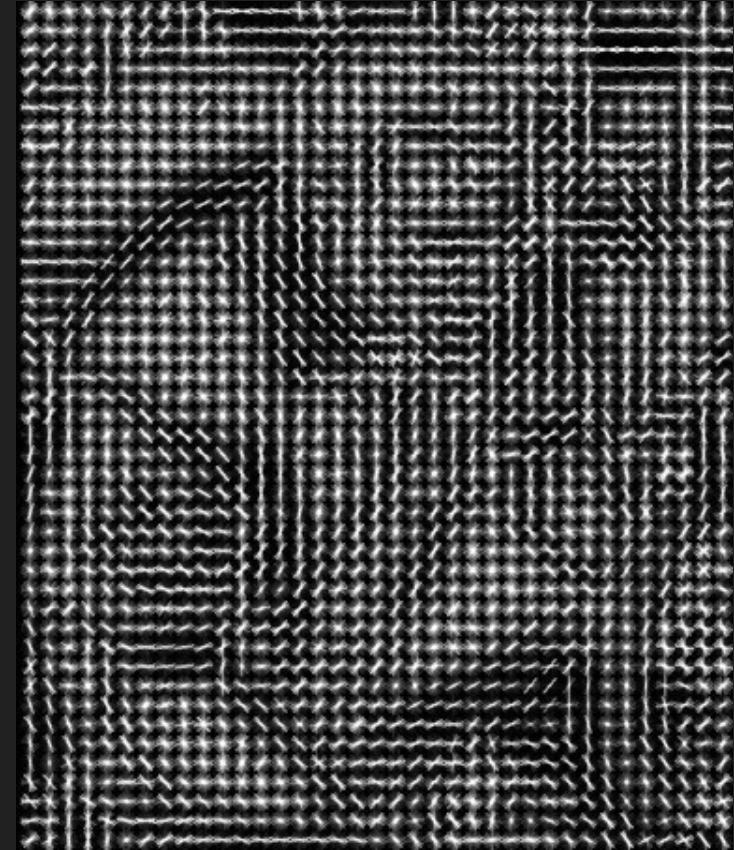
Car



# What information is lost?



# What information is lost?



# How can we ‘invert’ lossy HOG?

- Gradient computation
- Oriented magnitude sum (via bins)
  - Loss of precision
  - Loss of specificity – any number of values can sum to the same total
- Normalization
  - No way to unnormalize without knowing normalization factors

*Many different image patches translate to the same HOG feature  
:(*

# What information is lost?

$x = \text{input patch}$

$y = \text{HOG descriptor}$

$\phi(x) = \text{HOG transform}$

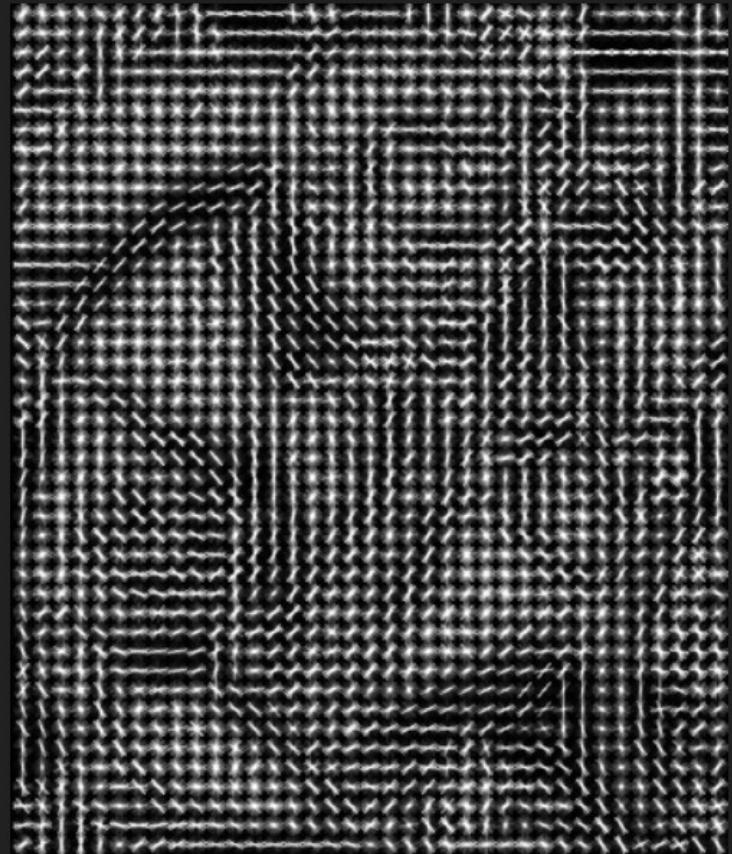
$$\min_{x \in \mathbb{R}^d} \|\phi(x) - y\|_2^2$$

Hard to optimize!

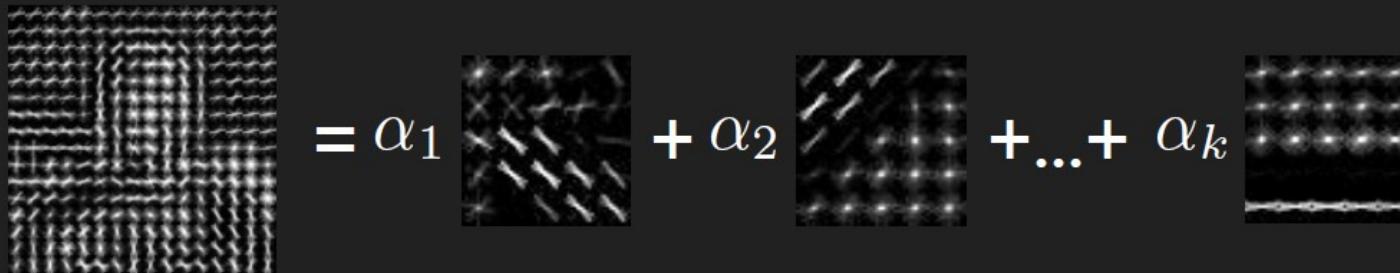
Many-to-one = unconstrained!

The paper tries five approaches:

- Gradient descent
- Exemplar LDA
- Ridge regression
- Direct optimization
- Paired dictionary learning



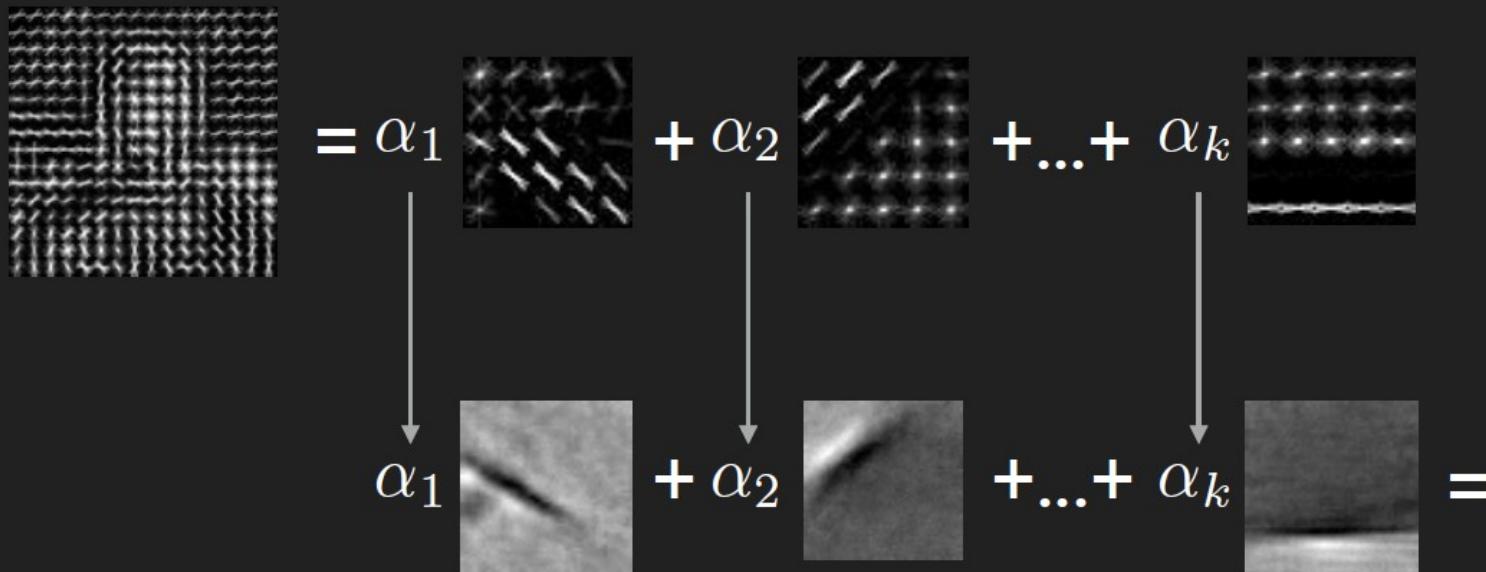
# Method: Paired Dictionary


$$\text{[Complex Image Patch]} = \alpha_1 \text{ [Image Patch with Diagonal Lines]} + \alpha_2 \text{ [Image Patch with Horizontal Lines]} + \dots + \alpha_k \text{ [Image Patch with Vertical Lines]}$$

How to constrain (two parts):

1. Learn a basis over HOG windows

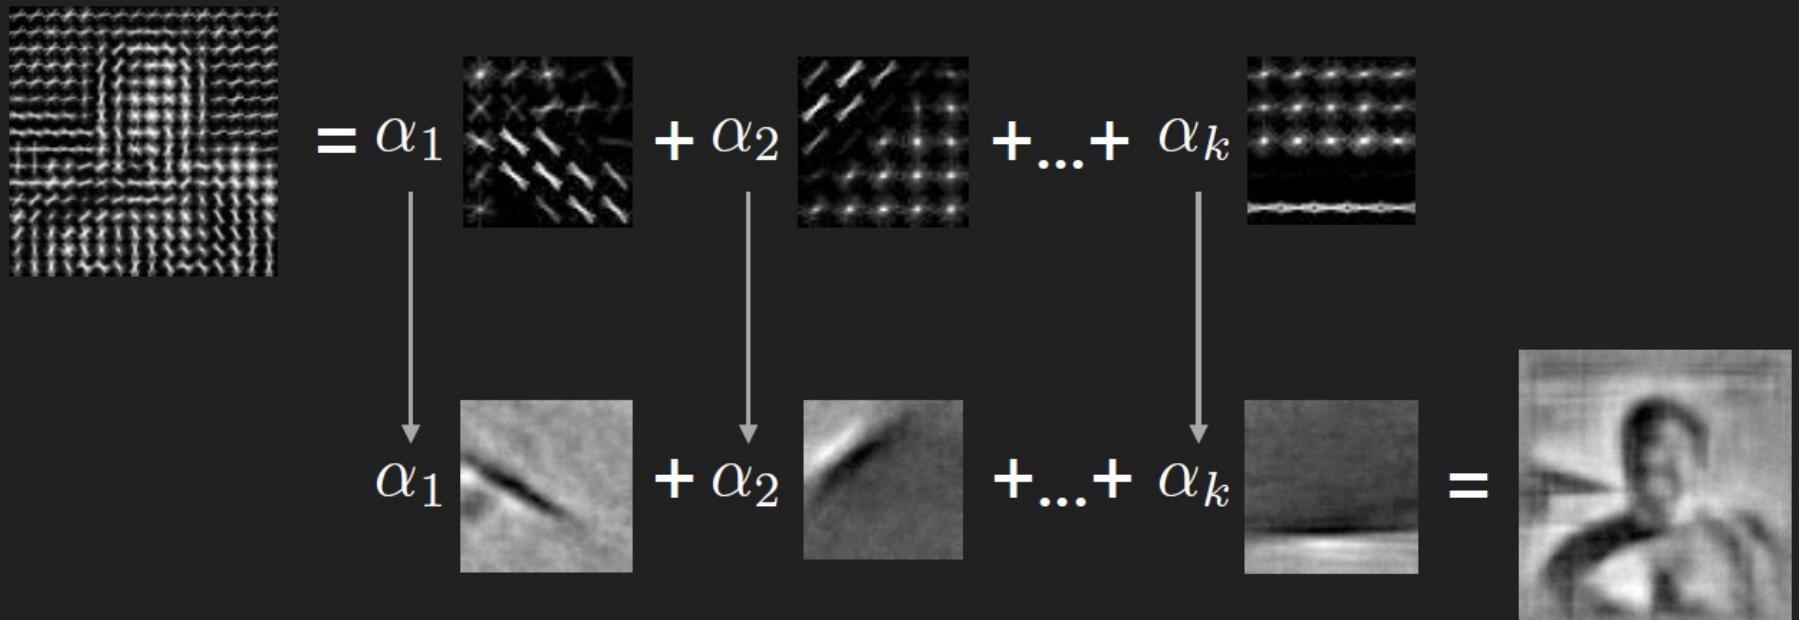
# Method: Paired Dictionary



How to constrain (two parts):

1. Learn a basis over HOG windows
2. Simultaneously learn a basis over input windows,  
and *share the weights over the training data*

# Method: Paired Dictionary



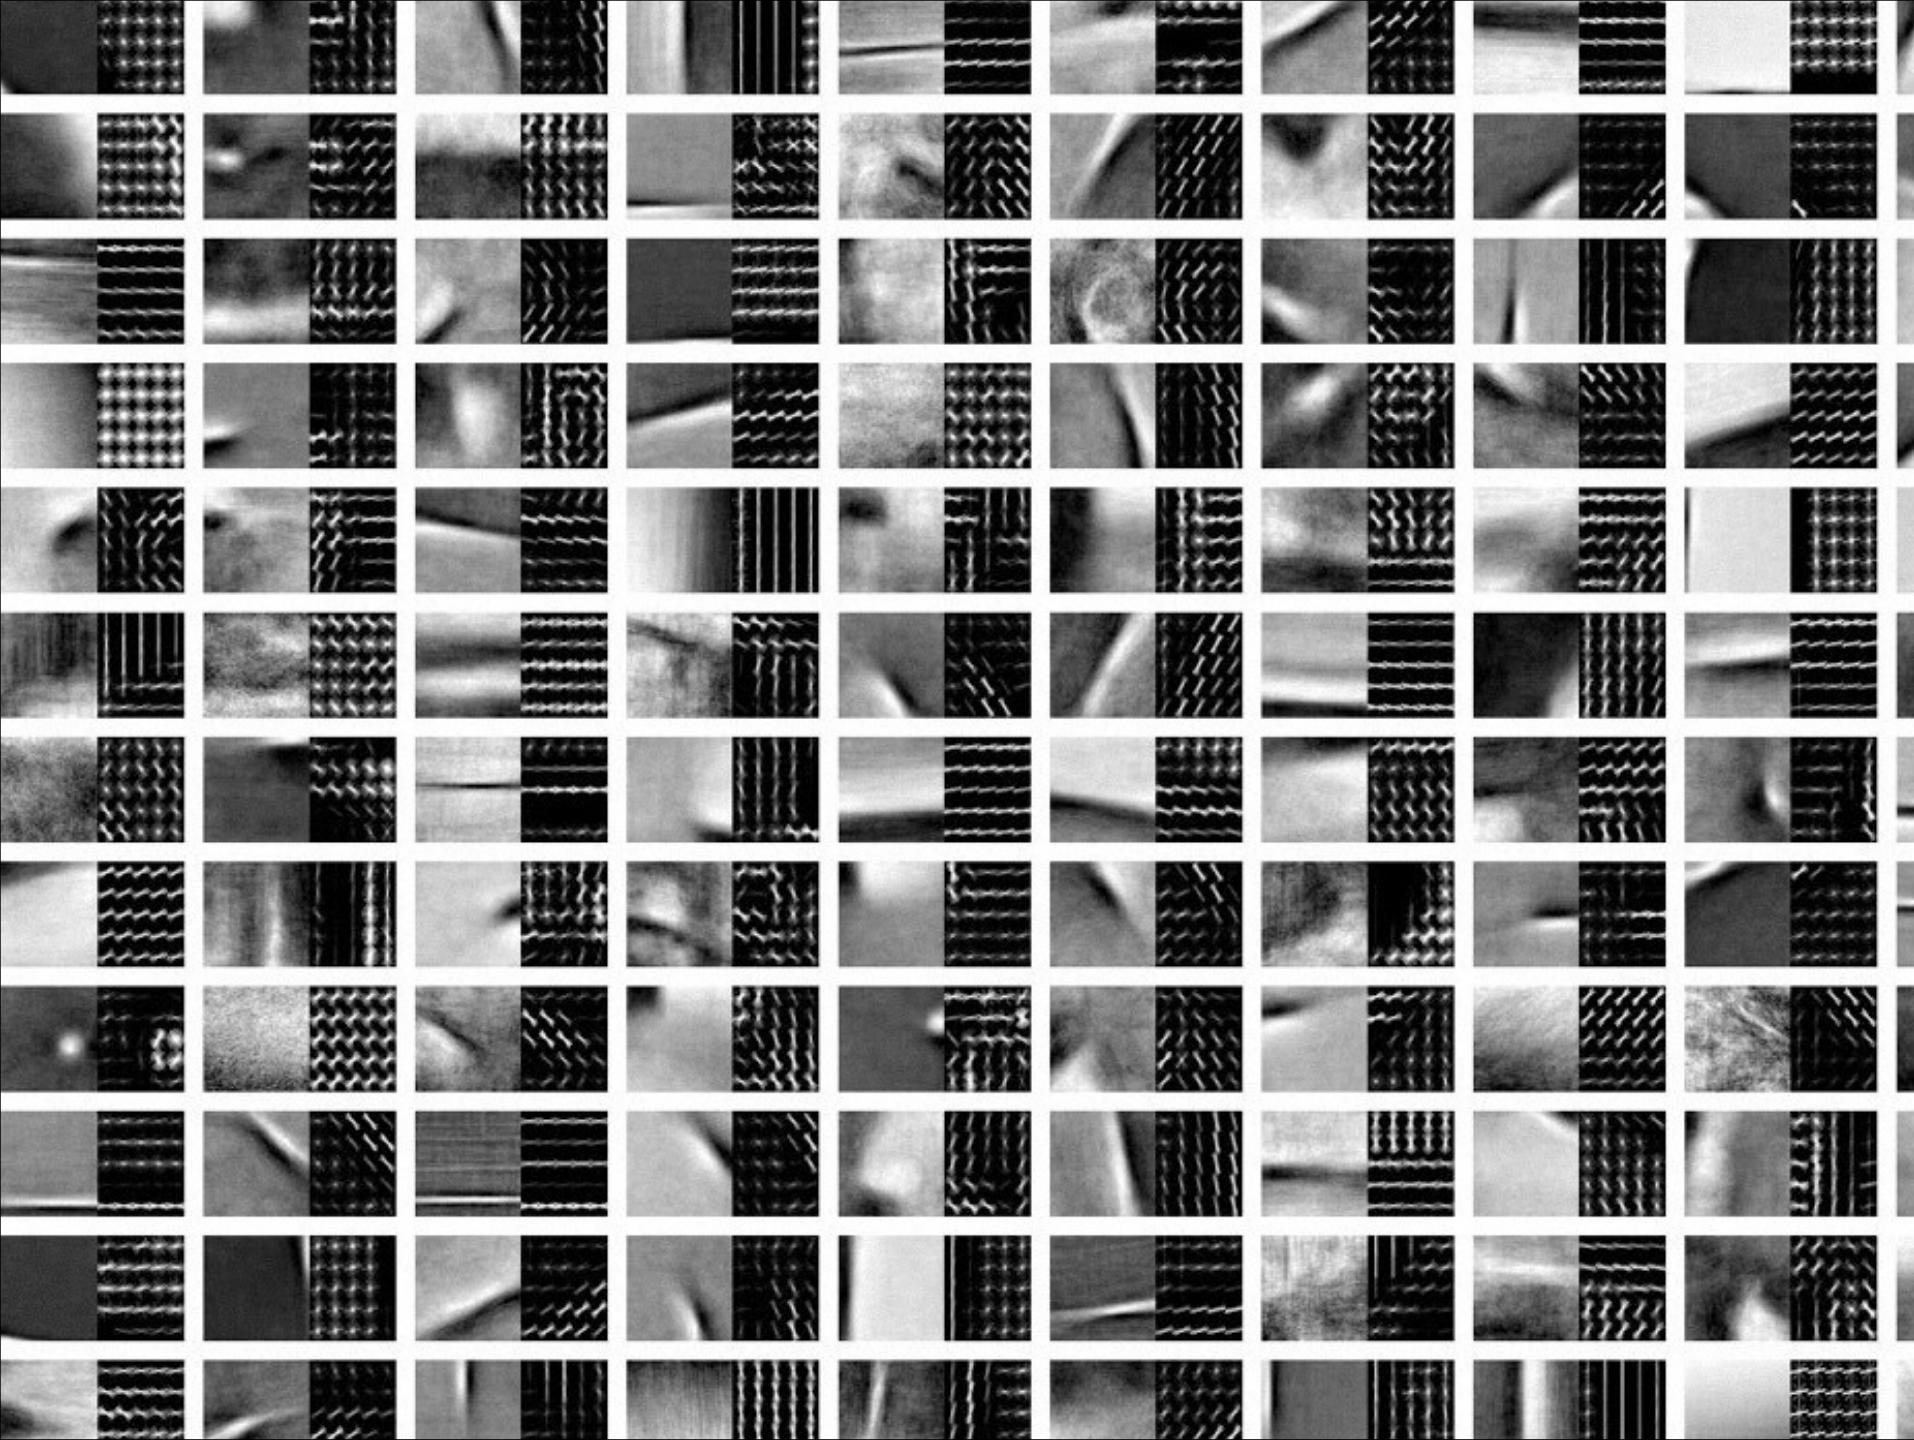
Inference to invert HOG:

1. Transform HOG patch into basis vectors
2. Take weights and apply to input basis

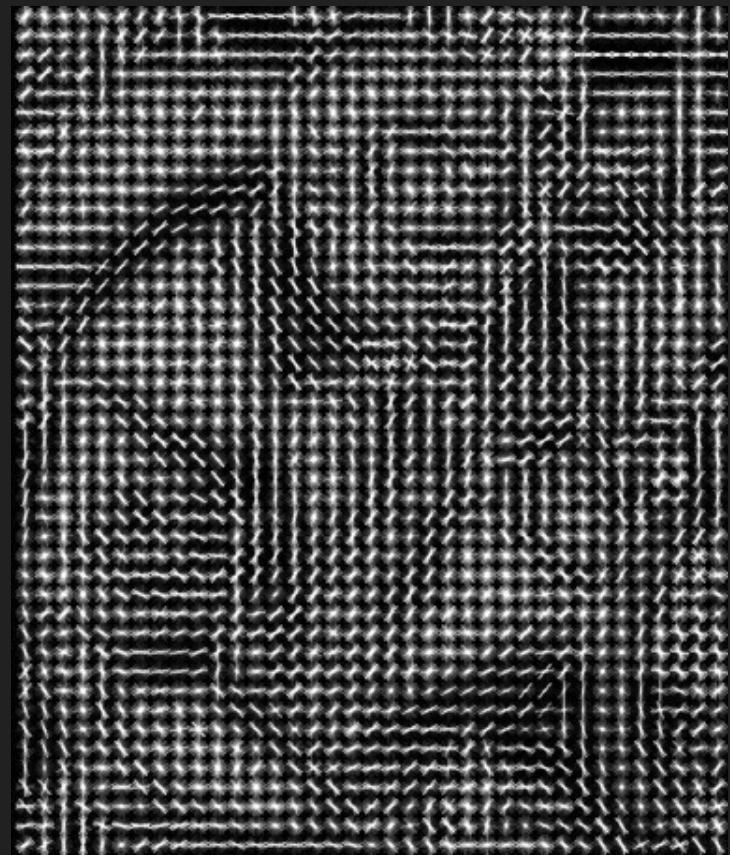
# Paired Dictionary Learning

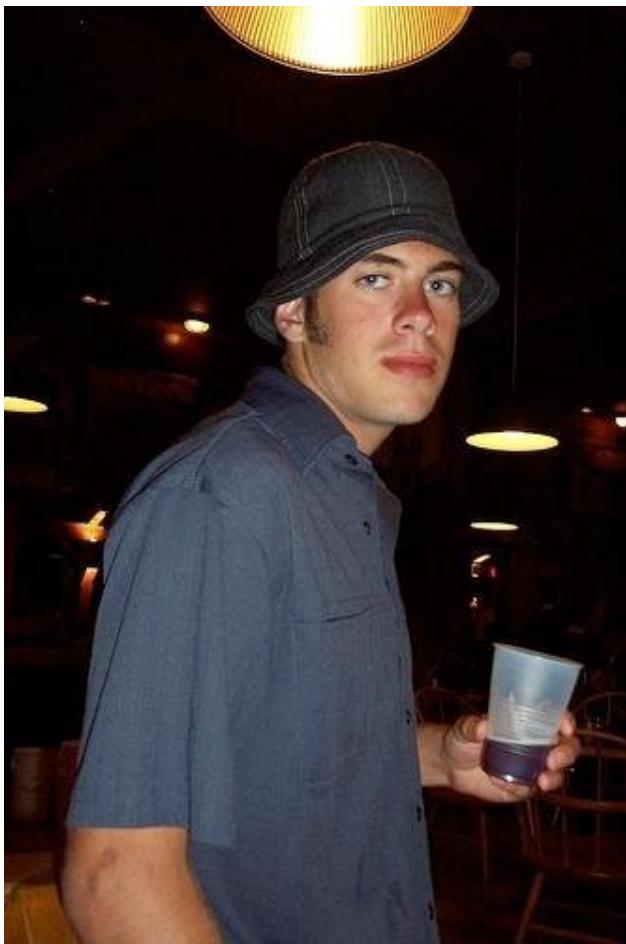
$$\begin{aligned} & \arg \min_{U, V, \alpha} \sum_{i=1}^N \|x_i - U\alpha_i\|_2^2 + \|y_i - V\alpha_i\|_2^2 \\ \text{s.t. } & \|\alpha_i\|_1 \leq \lambda, \|U\|_2^2 \leq \gamma_1, \|V\|_2^2 \leq \gamma_2 \end{aligned}$$

**Just sparse coding!**  
(Optimize using off the shelf solvers)



# What information is lost?





Human  
Vision

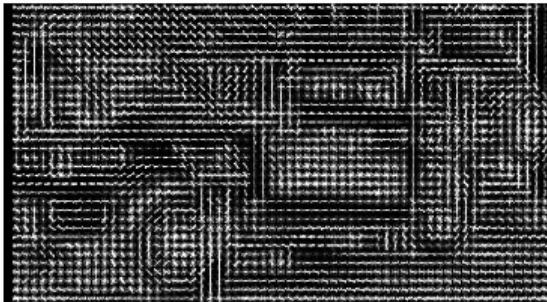


HOG  
Vision

Vondrick et al.

# HOGgles (Vondrick et al. ICCV 2013)

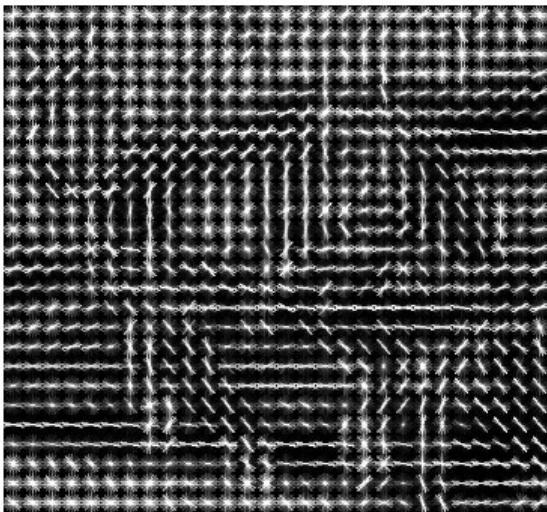
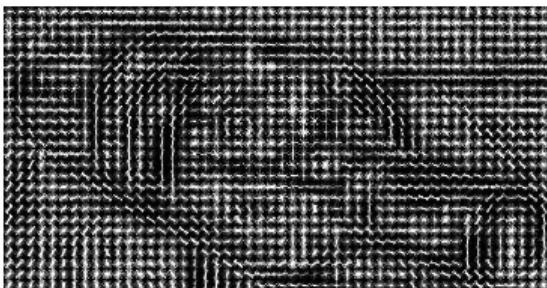
HOG [1]



Inverse ( $U_s$ )



Original

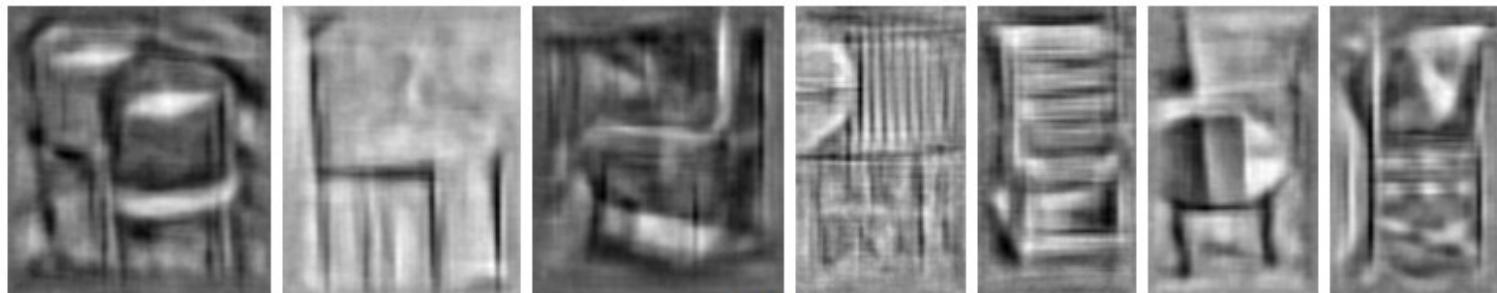


## Visualizing Top Detections

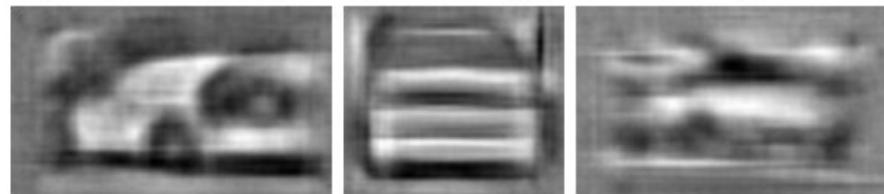
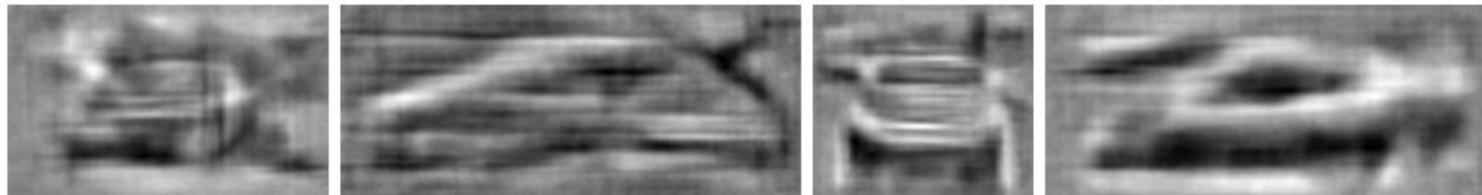
We have visualized some high scoring detections from the deformable parts model. Can you guess which are false alarms? Click on the images below to reveal the corresponding RGB patch. You might be surprised!



Person



Chair

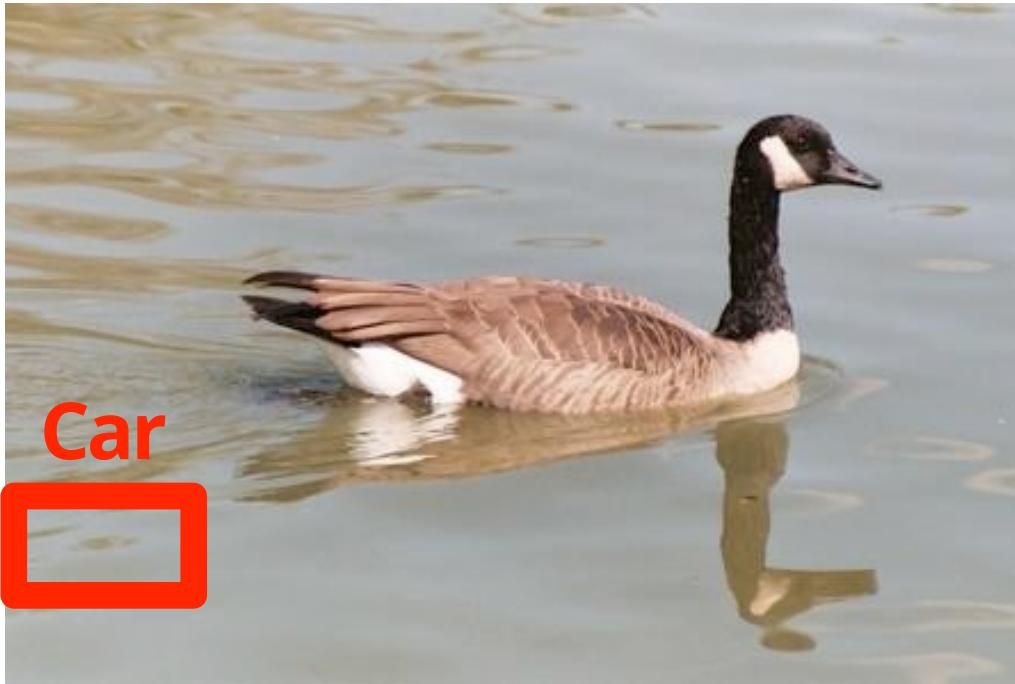


Car

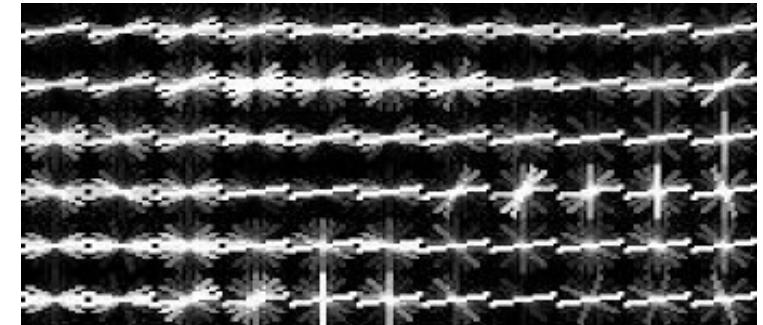
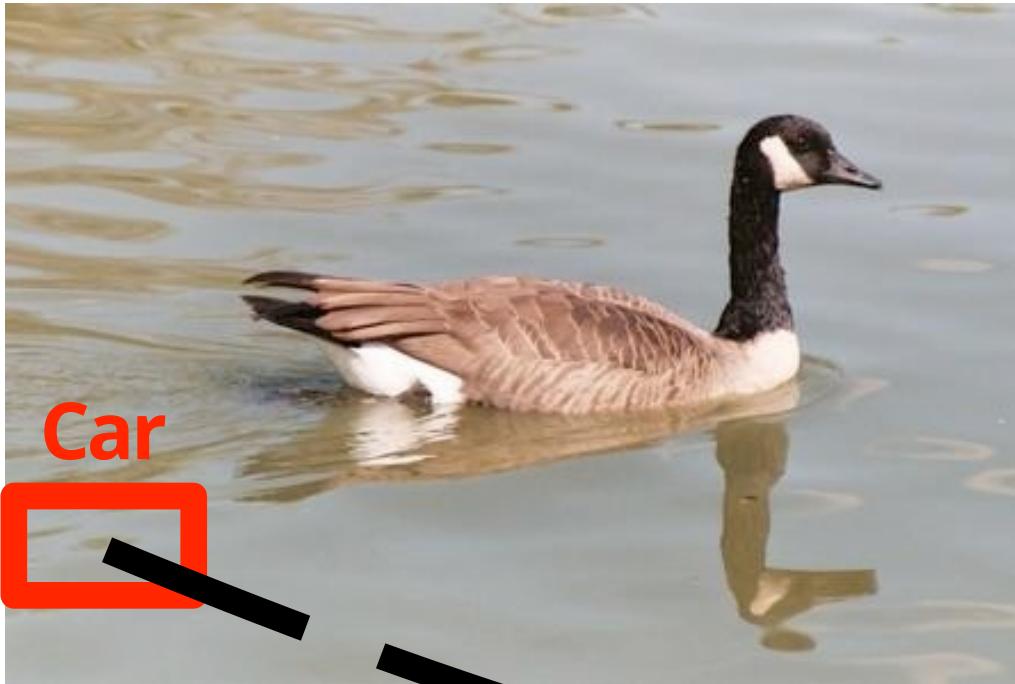
[carlvondrick.com/ihog/](http://carlvondrick.com/ihog/)

Vondrick et

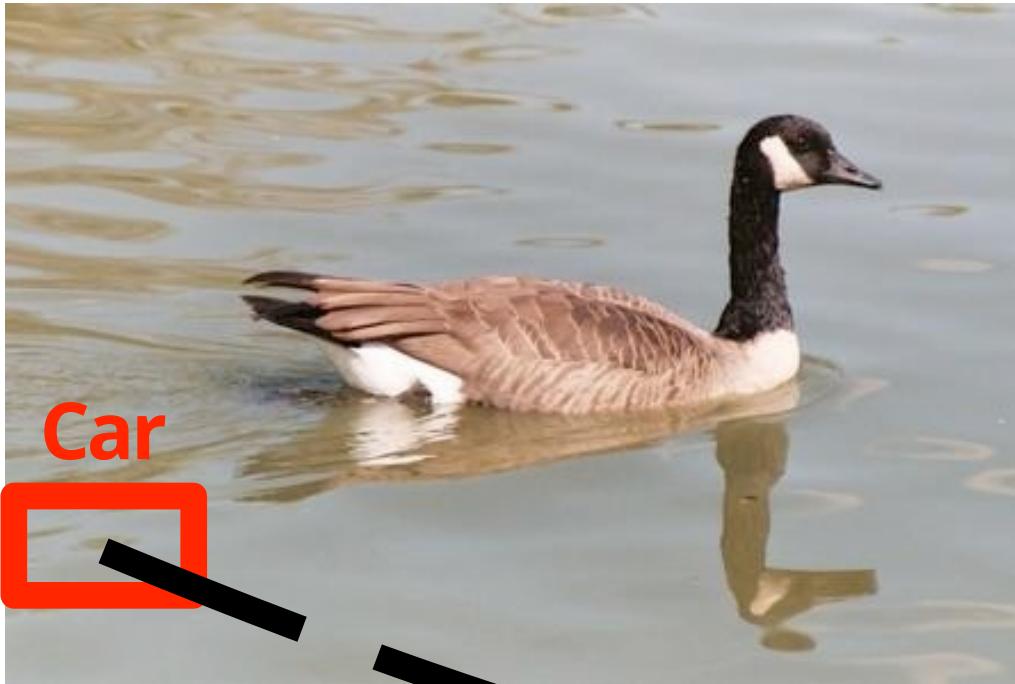
# Why did the detector fail?



# Why did the detector fail?



# Why did the detector fail?



# Recursive HOG!



Original  $x$



$x' = \phi^{-1}(\phi(x))$



$x'' = \phi^{-1}(\phi(x'))$

Figure 11: We recursively compute HOG and invert it with a paired dictionary. While there is some information loss, our visualizations still do a good job at accurately representing HOG features.  $\phi(\cdot)$  is HOG, and  $\phi^{-1}(\cdot)$  is the inverse.

# Code Available

Try it on your projects!

```
ihog = invertHOG(feat);
```



<https://github.com/cvondrick/ihog/blob/master/README.md>

# Computer Vision so far

- The geometry of image formation
  - Ancient / Renaissance
- Signal processing / Convolution
  - 1800, but really the 50's and 60's
- Hand-designed Features for recognition, either instance-level or categorical
  - 1999 (SIFT), 2003 (Video Google), 2005 (Dalal-Triggs), 2006 (spatial pyramid)
- Learning from Data
  - 1991 (EigenFaces) but late 90's to now especially

# What has changed in the last decade?

- The Internet
- Crowdsourcing
- Learning representations from the data these sources provide (deep learning)

# Big Idea

- Do we need computer vision systems to have strong AI-like reasoning about our world?
- What if invariance / generalization isn't actually the core difficulty of computer vision?
- What if we can perform high level reasoning with brute-force, data-driven algorithms?

# Short cuts to AI

With billions of images on the web, it's often possible to find a close nearest neighbor.

We can shortcut hard problems by “looking up” the answer, stealing the labels from our nearest neighbor.

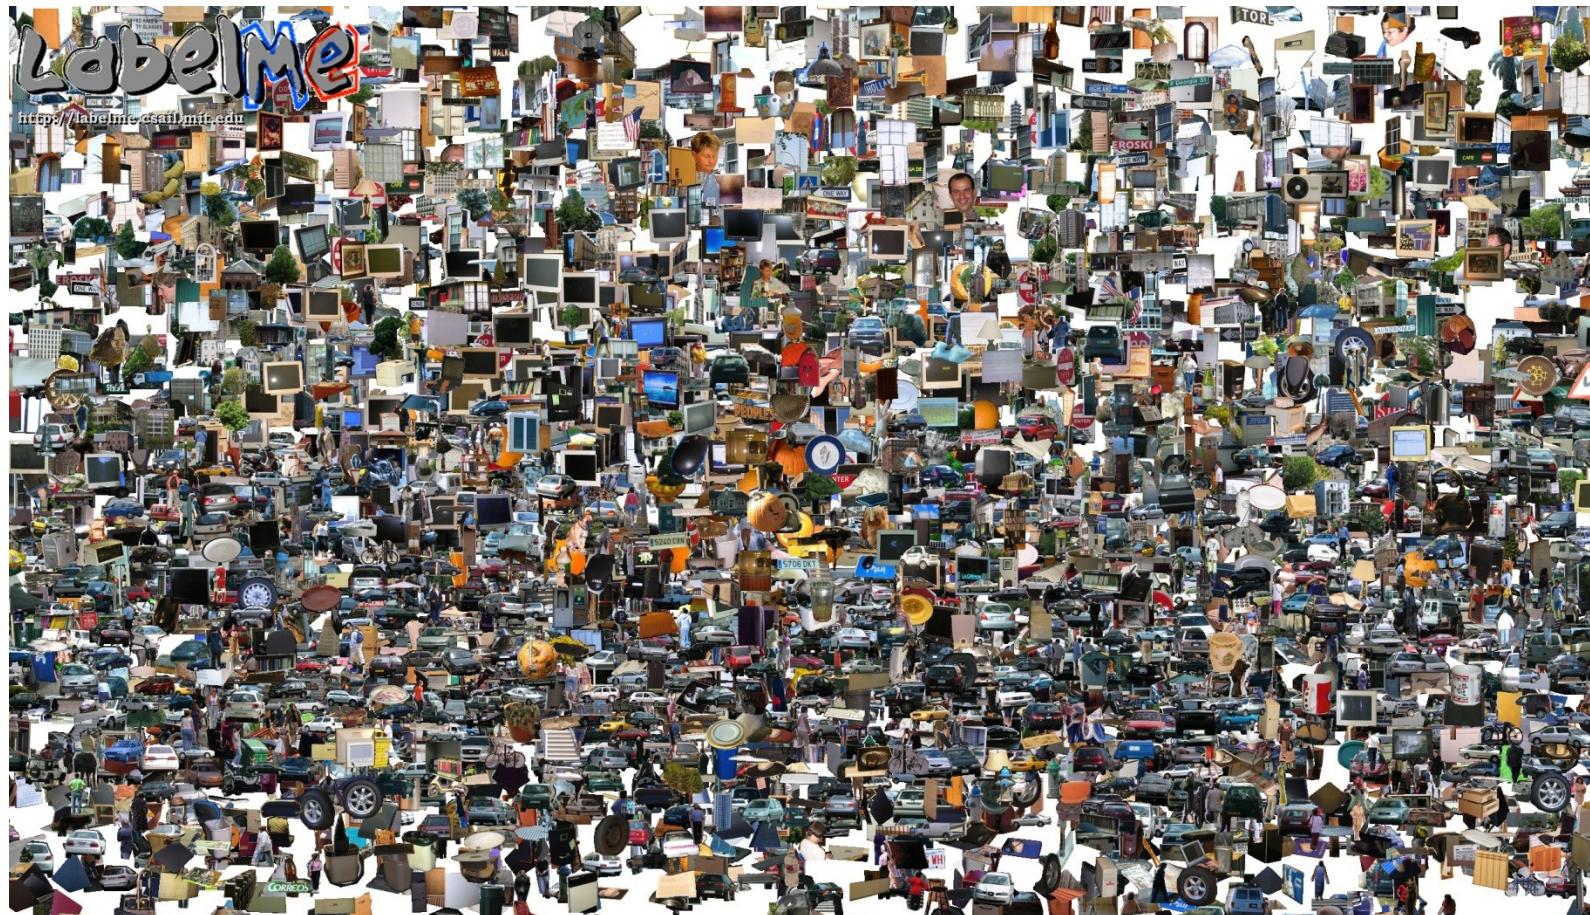


# So what is intelligence?

Weak AI: The simulation of a ‘mind’ is a model for the ‘mind’.

Strong AI: The simulation of a ‘mind’ is a ‘mind’.

# Opportunities of Scale



# Powers of 10

Number of images on my hard drive:  $10^6$



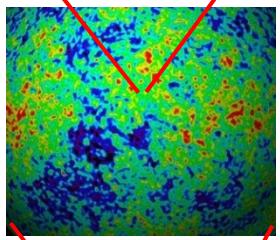
Number of images seen during my first 10 years:  $10^8$   
(3 images/second \* 60 \* 60 \* 16 \* 365 \* 10 = 630,720,000)



Number of images seen by all humanity:  $10^{20}$   
 $106,456,367,669 \text{ humans}^1 * 60 \text{ years} * 3 \text{ images/second} * 60 * 60 * 16 * 365 =$   
1 from <http://www.prb.org/Articles/2002/HowManyPeopleHaveEverLivedonEarth.aspx>



Number of photons in the universe:  $10^{88}$



Number of all 32x32 images:  $10^{7373}$   
 $256^{32*32*3} \sim 10^{7373}$



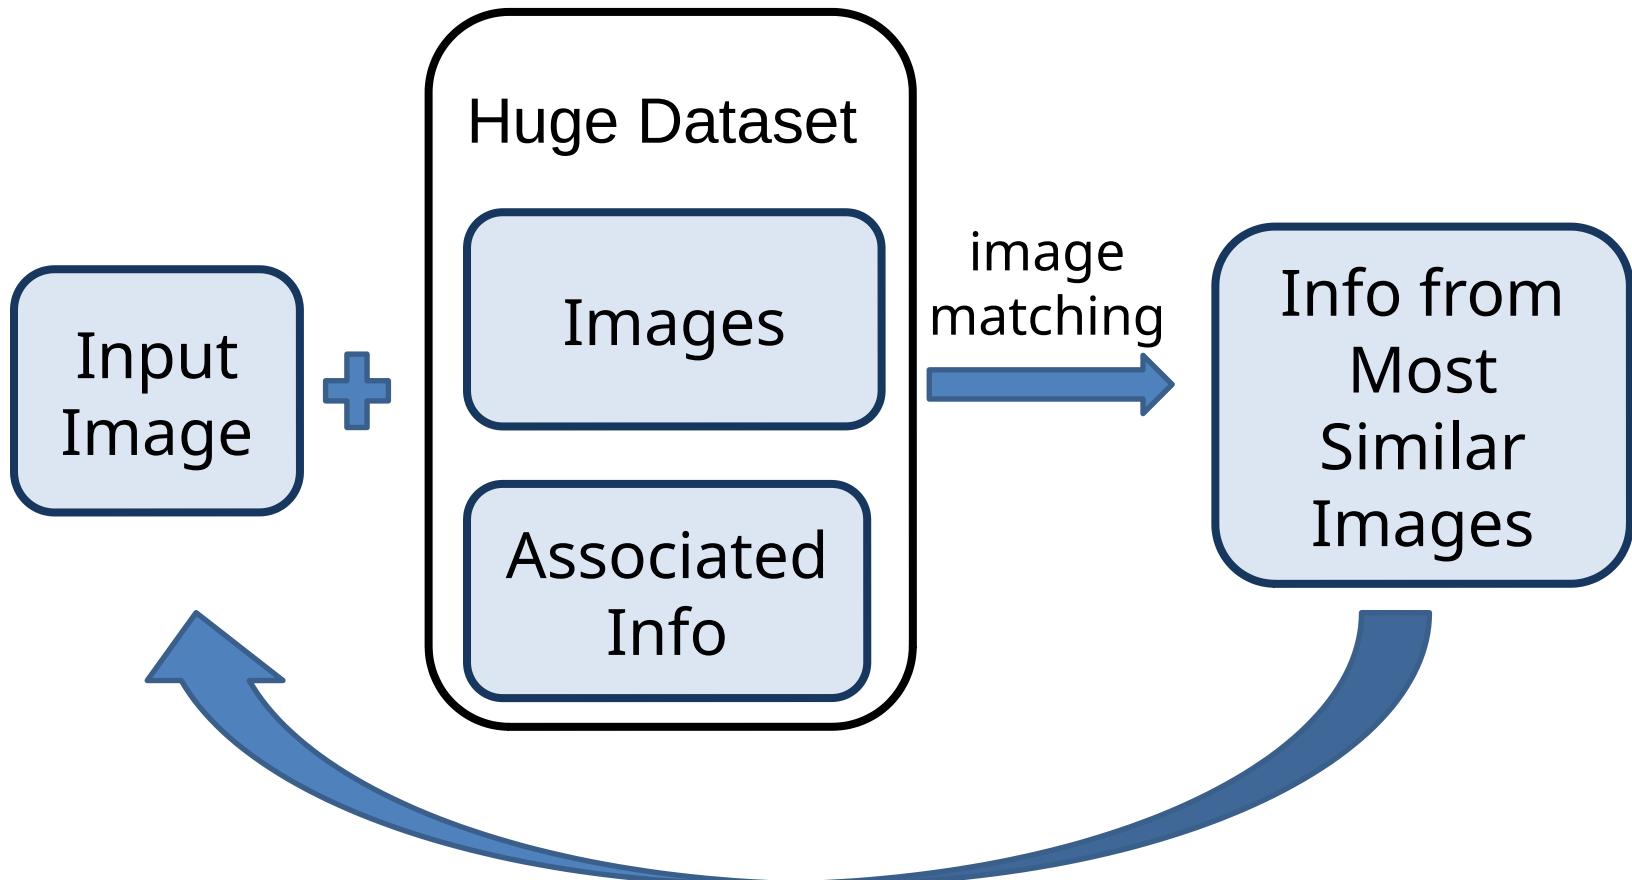
# Understanding scenes encompasses all kinds of knowledge



# But not all scenes are so original

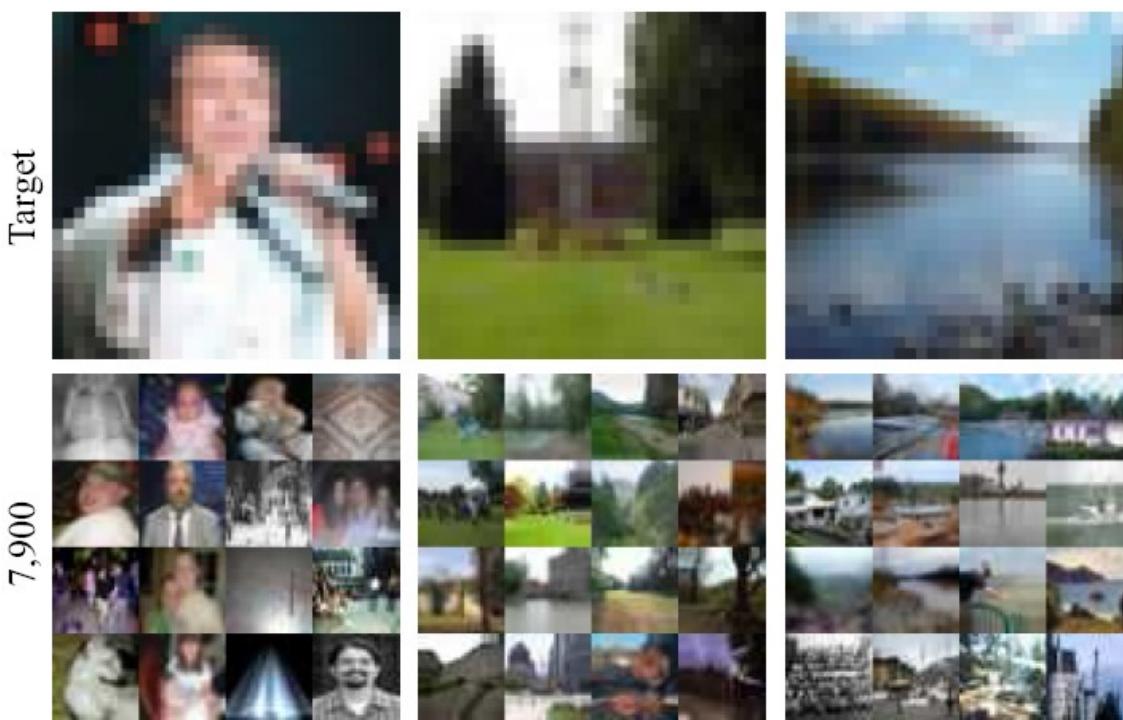


# General Principal



Hopefully, If you have enough images, the dataset will contain very similar images that you can find with simple matching methods.

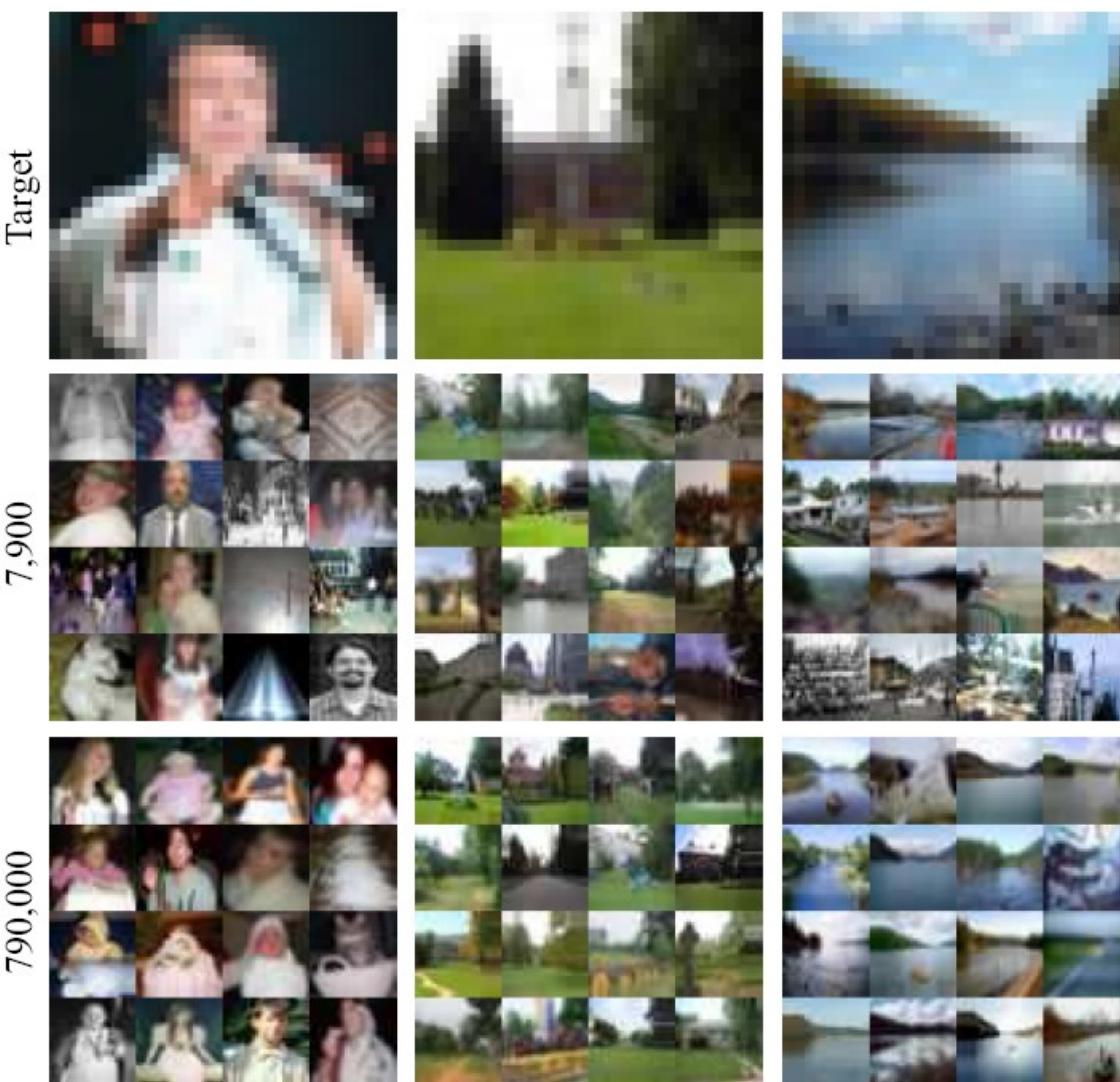
# Lots Of Images



# Lots

# Of

# Images



# Lots

# Of

# Images

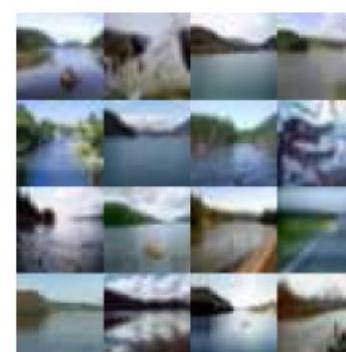
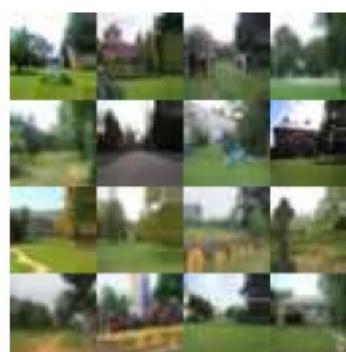
Target



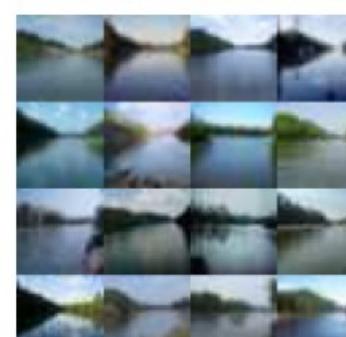
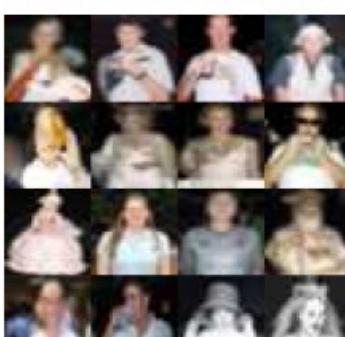
7,900



790,000



79,000,000



# Application: Automatic Colorization



Input



Color Transfer



Color Transfer



Matches (gray)



Matches (w/ color)



Avg Color of Match

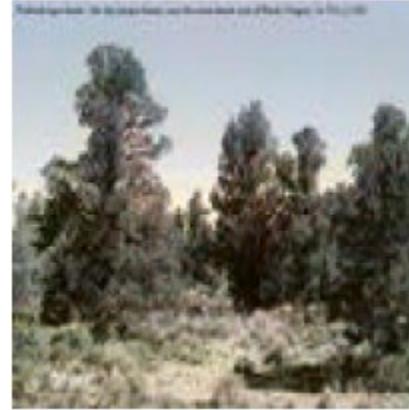
# Application: Automatic Colorization



Input



Color Transfer



Color Transfer



Matches (gray)



Matches (w/ color)



Avg Color of Match

# How much can an image tell about its geographic location?



# How much can an image tell about its geographic location?

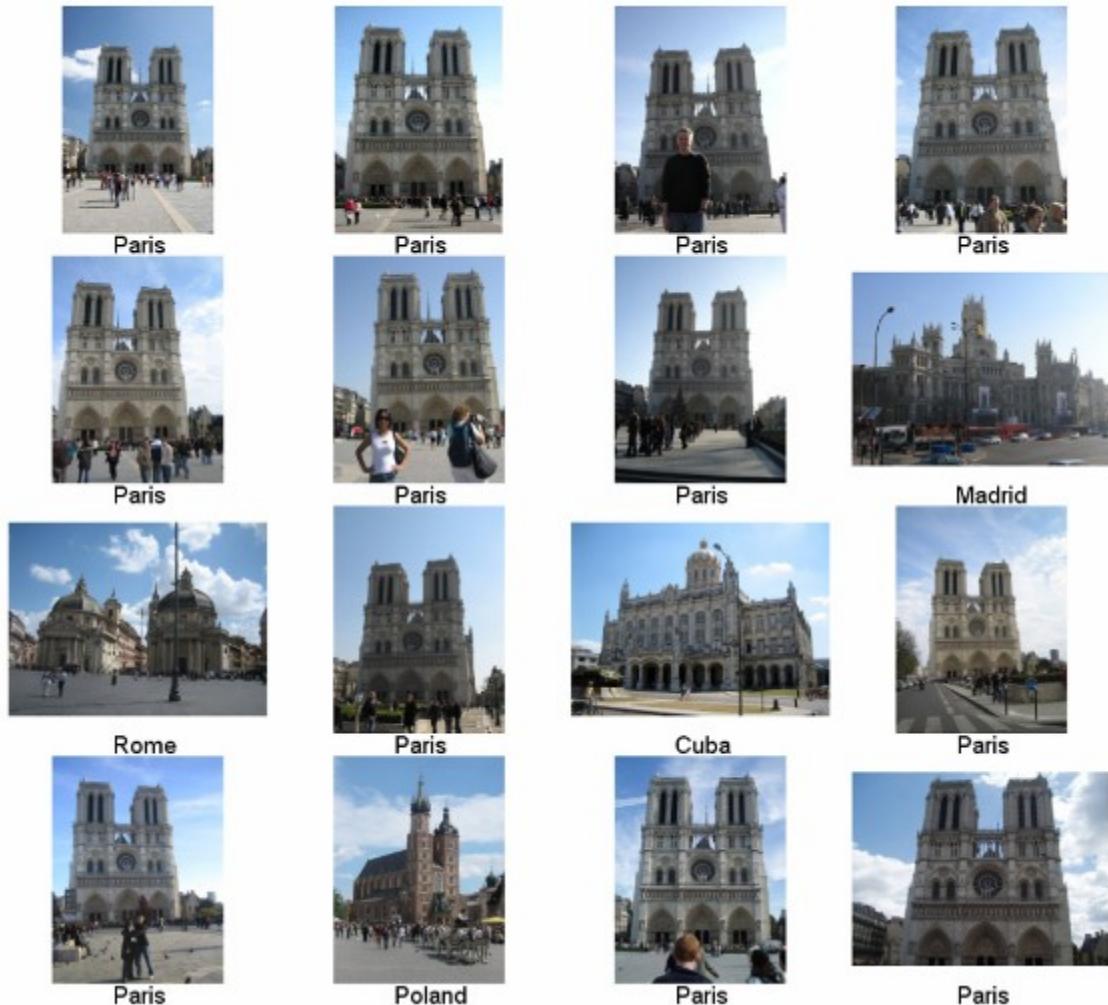


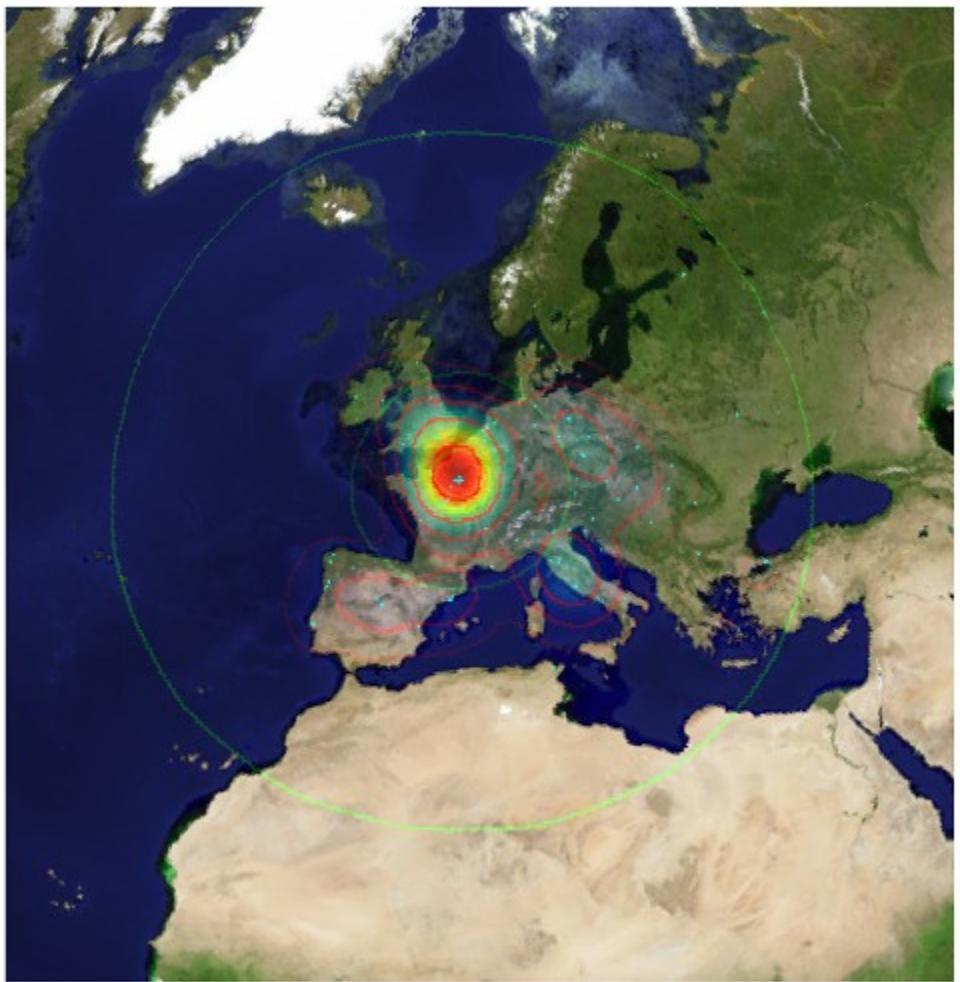
6 million geo-tagged Flickr images

<http://graphics.cs.cmu.edu/projects/im2gps/>

im2gps (Hays & Efros, CVPR 2008)

## Nearest Neighbors according to gist + bag of SIFT + color histogram + a few others

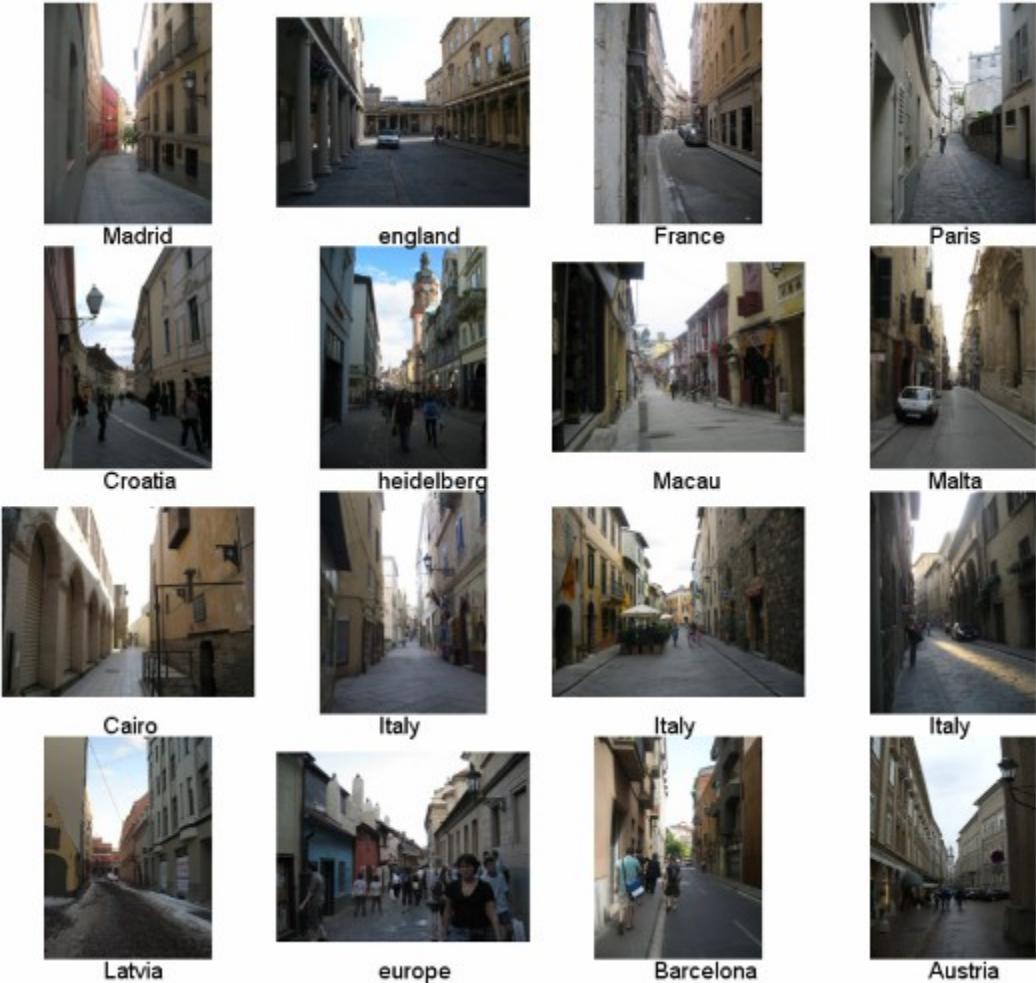




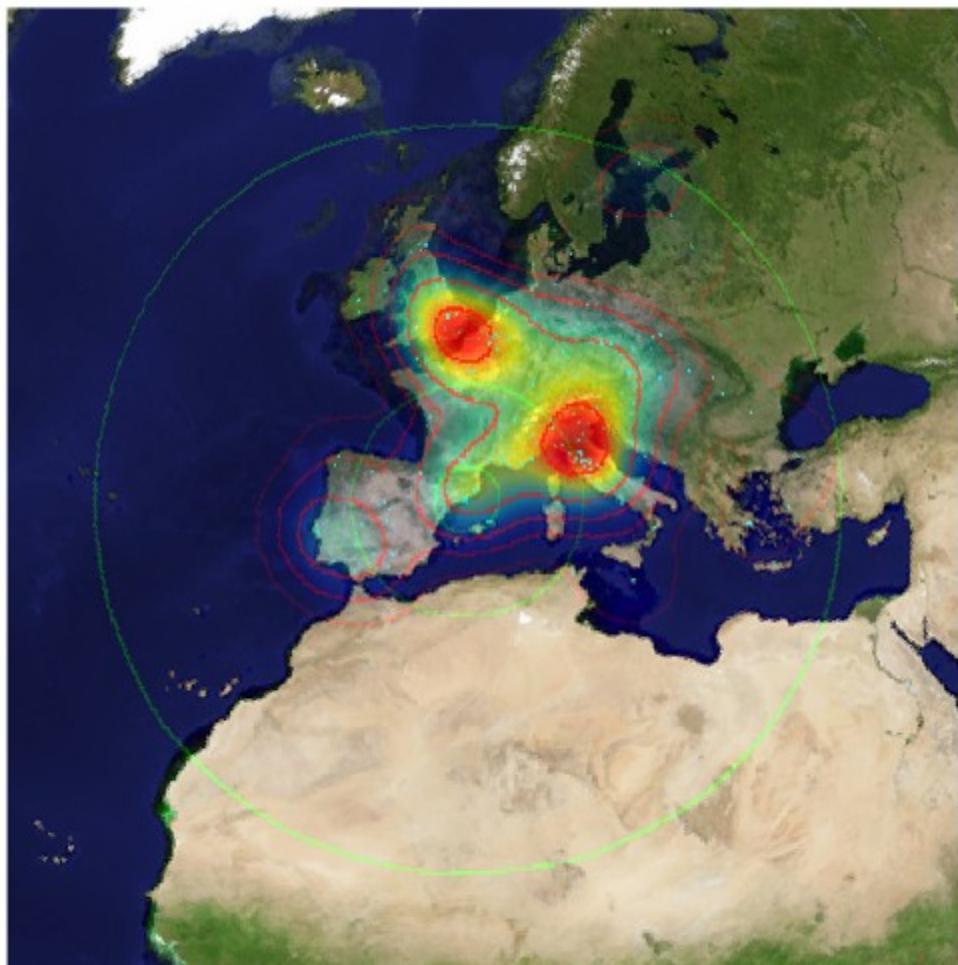
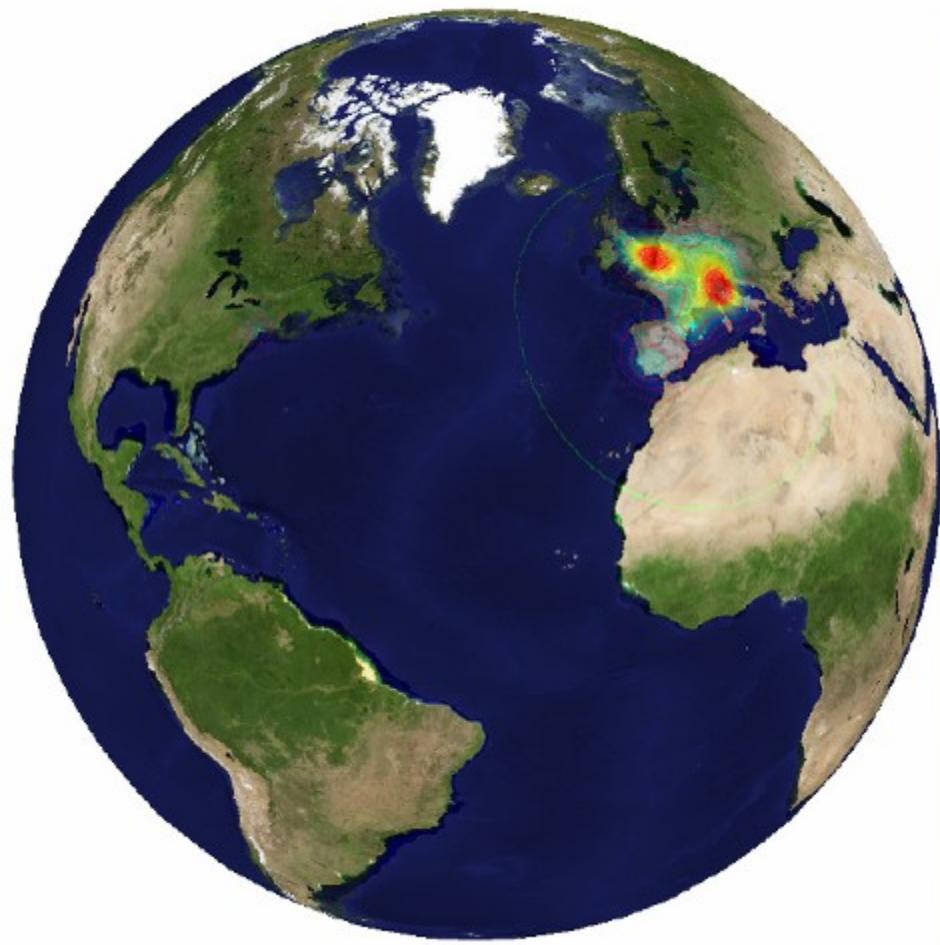
# Im2gps



# Example Scene Matches

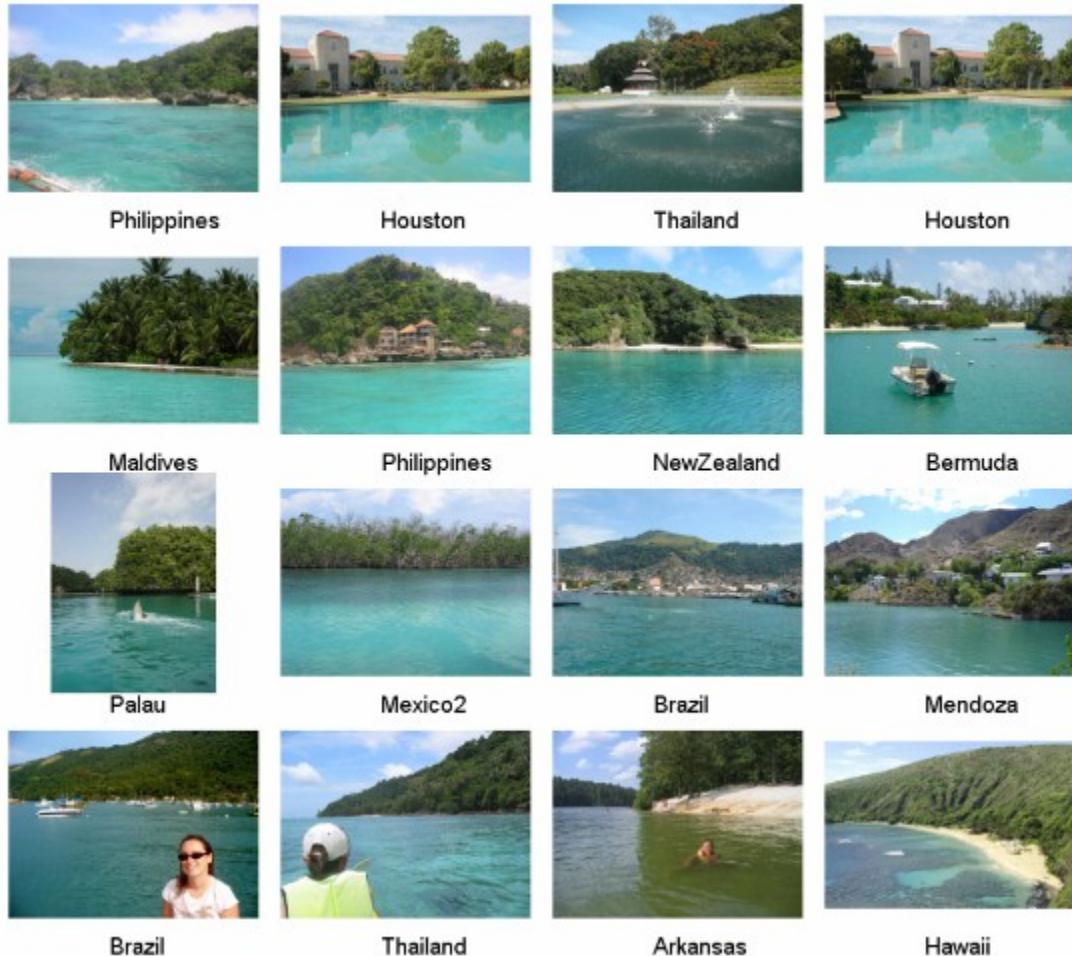


# Voting Scheme



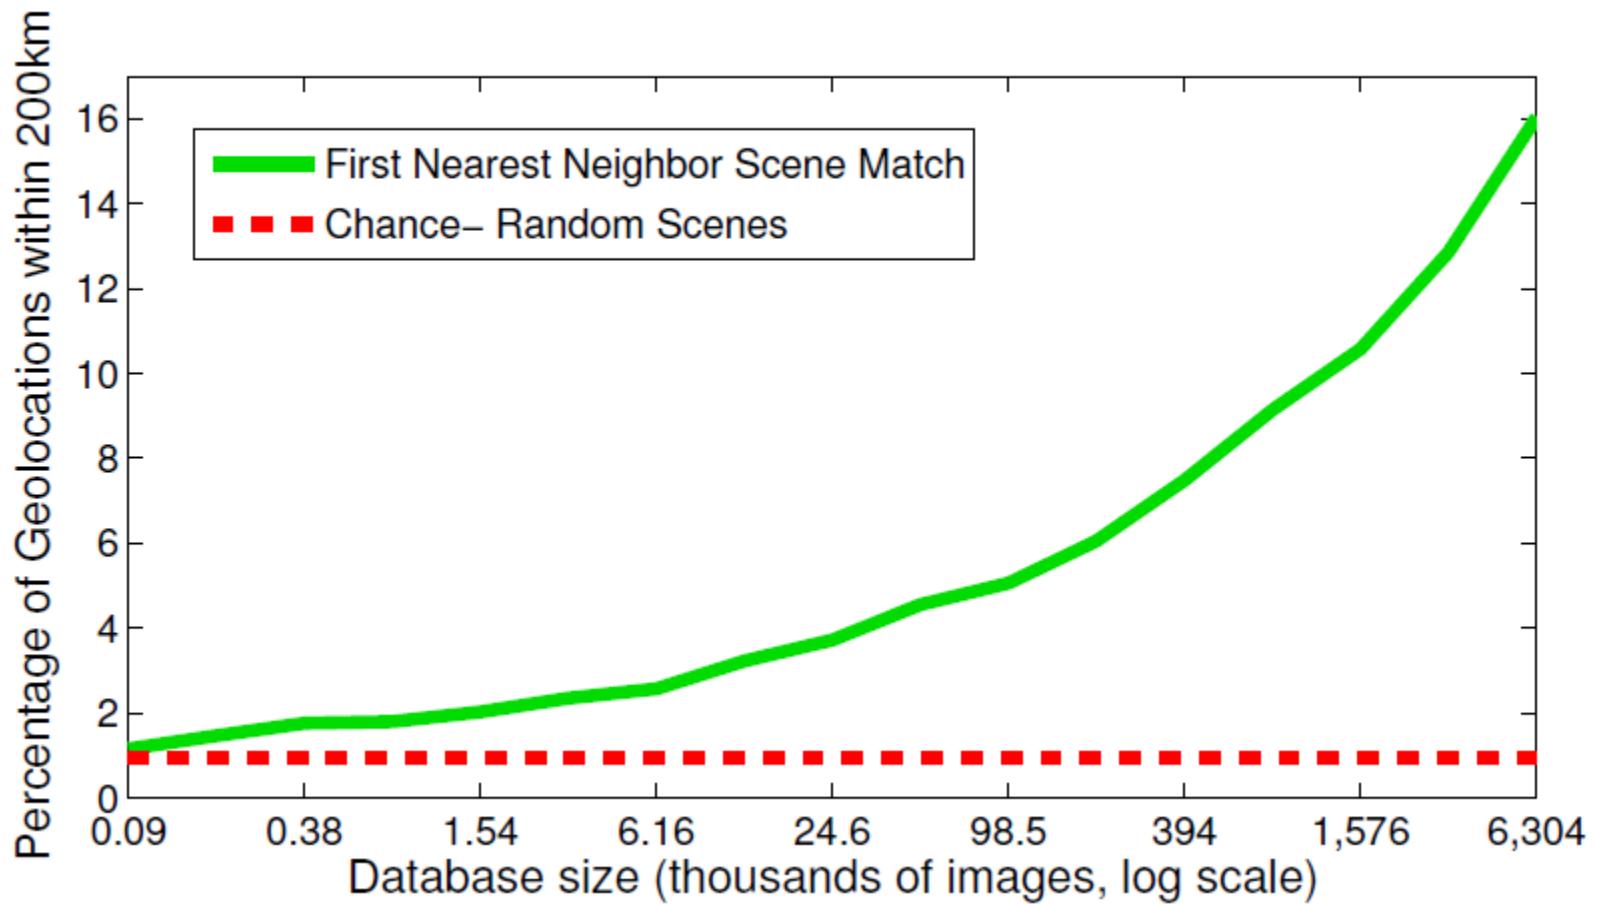
# im2gps



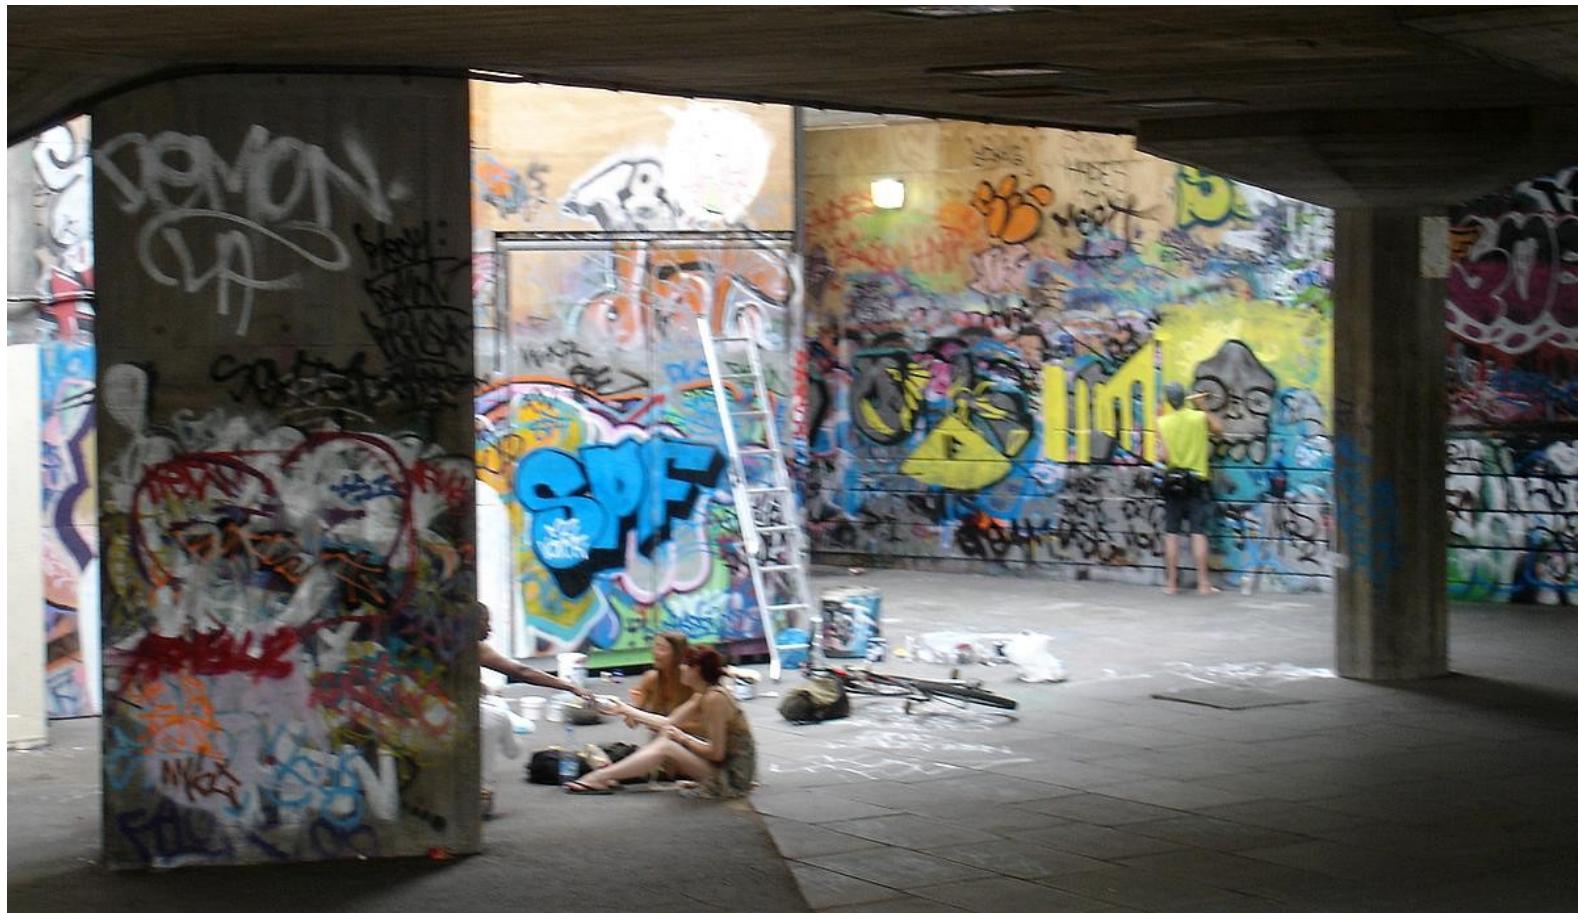




# Effect of Dataset Size



# Where is This?



[Vesselova, Kalogerakis, Hertzmann, Hays, Efros. Image Sequence Geolocation. ICCV'09]

# Where is This?



# Where are These?



15:14,  
June 18<sup>th</sup>, 2006



16:31,  
June 18<sup>th</sup>, 2006

# Where are These?



15:14,  
June 18<sup>th</sup>, 2006

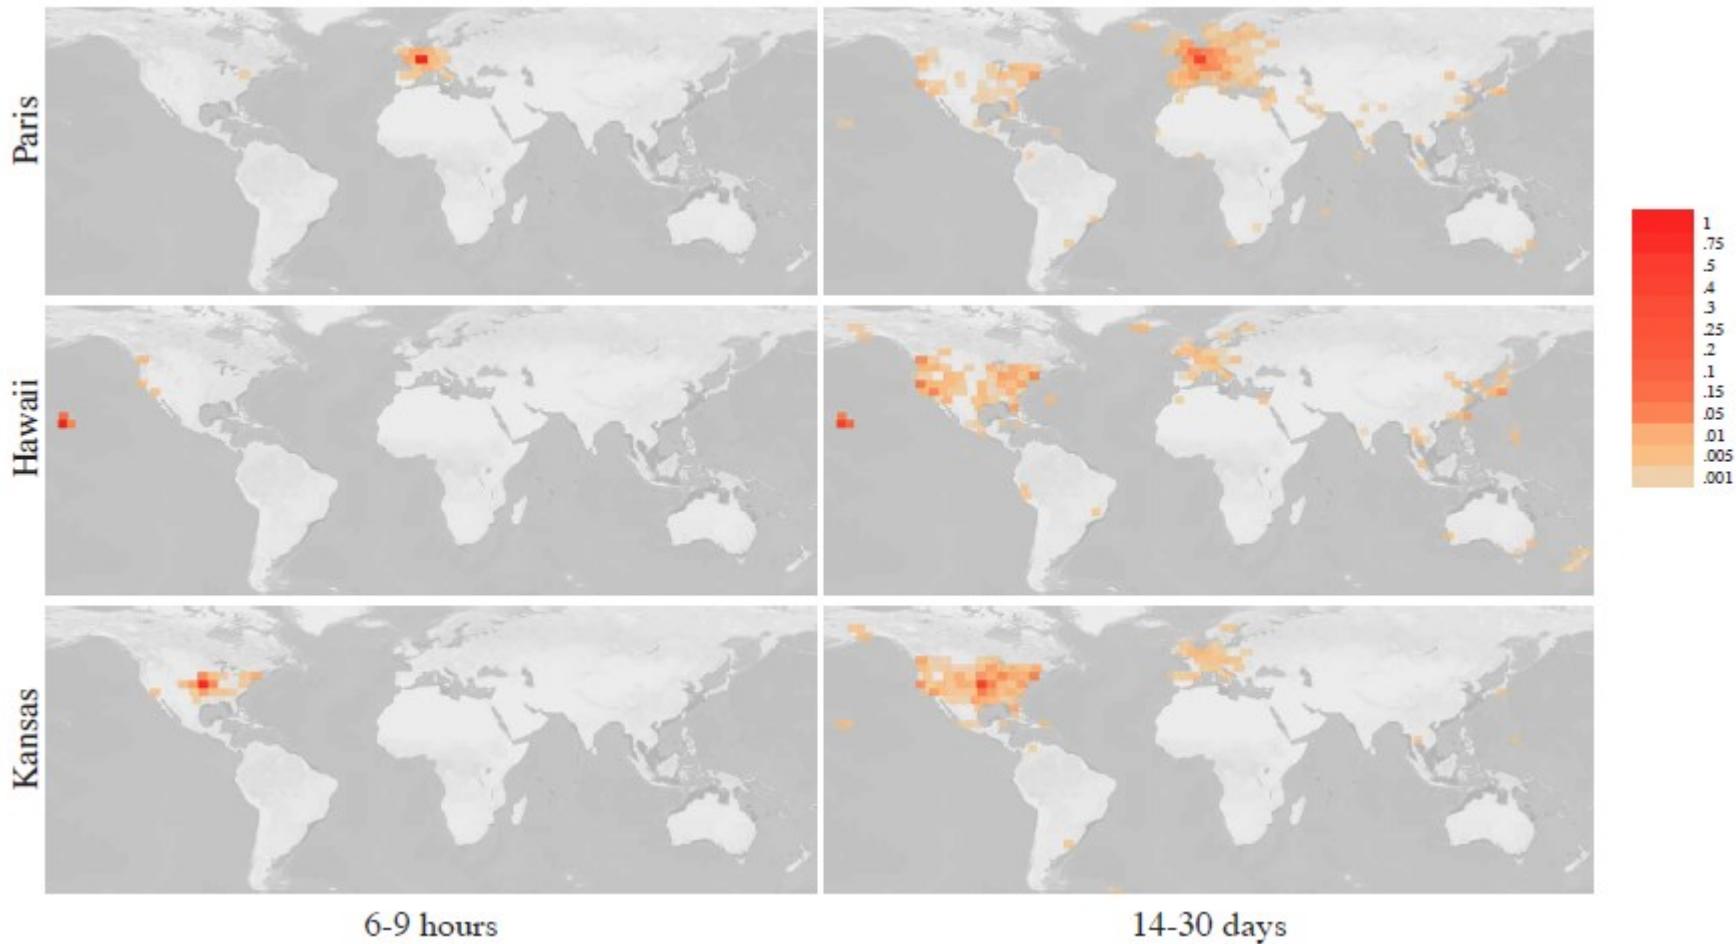
16:31,  
June 18<sup>th</sup>, 2006

17:24,  
June 19<sup>th</sup>, 2006

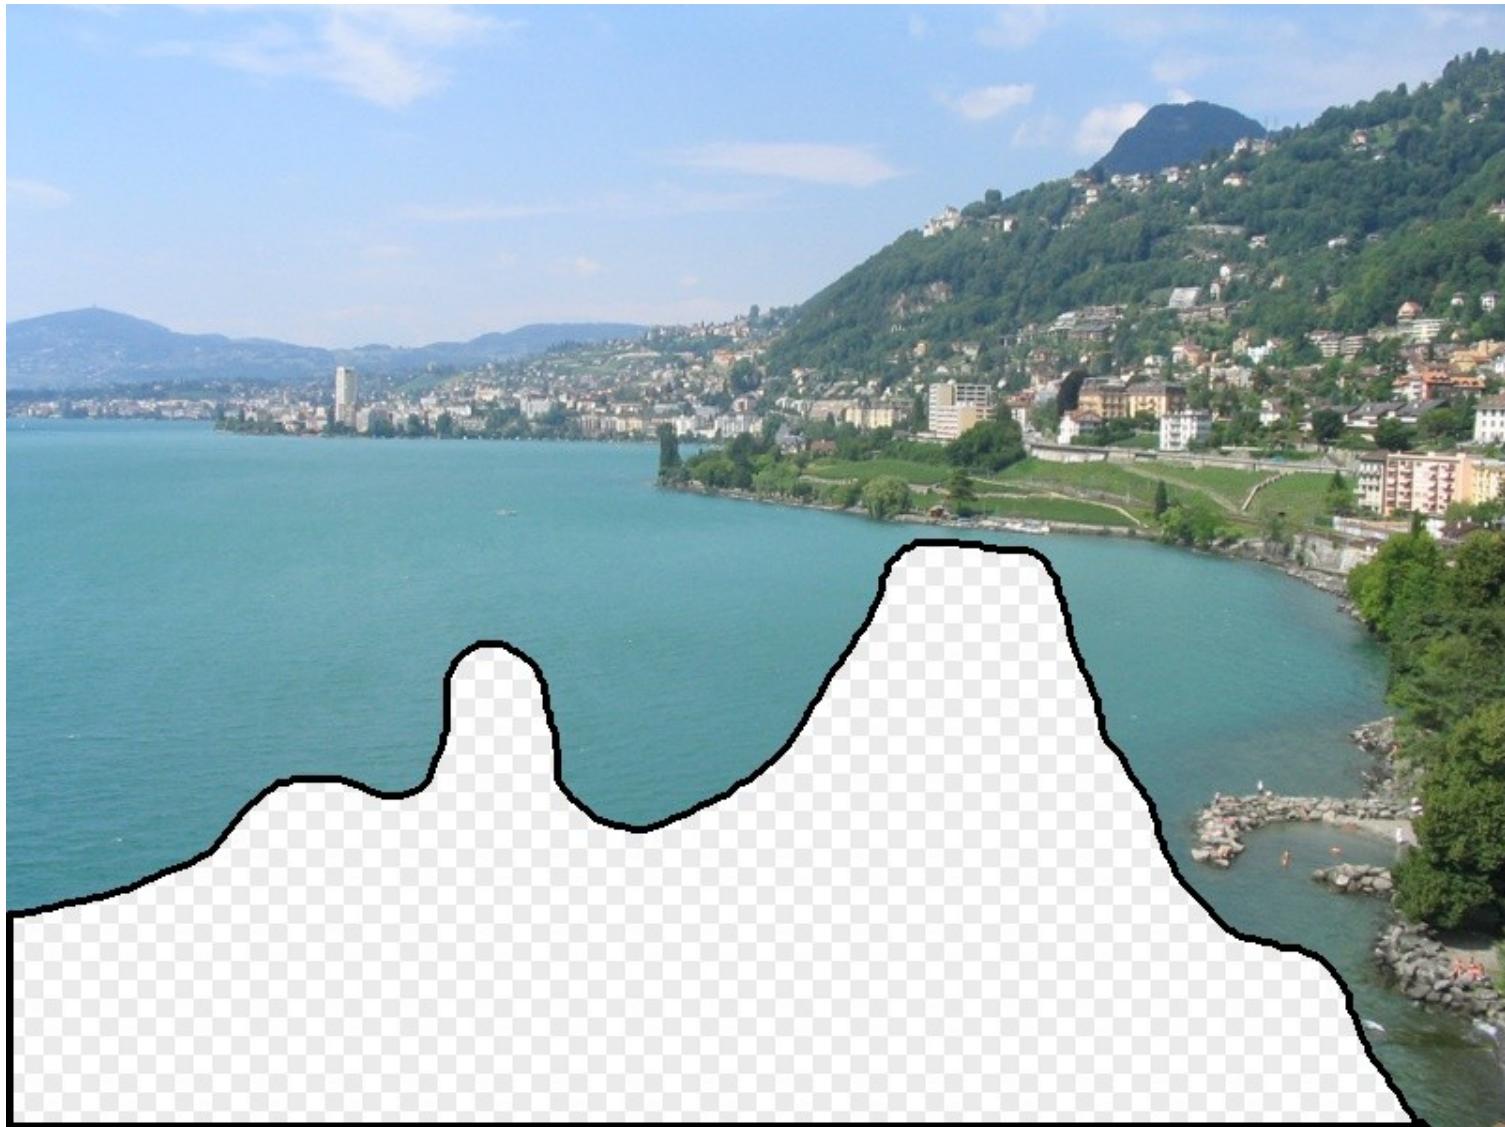
# Results

<https://github.com/lugiavn/revisiting-im2gps>

- im2gps – 10% (geo-loc within 400 km)
- temporal im2gps – 56%



# What should the missing region contain?









# Which is the original?



(a)



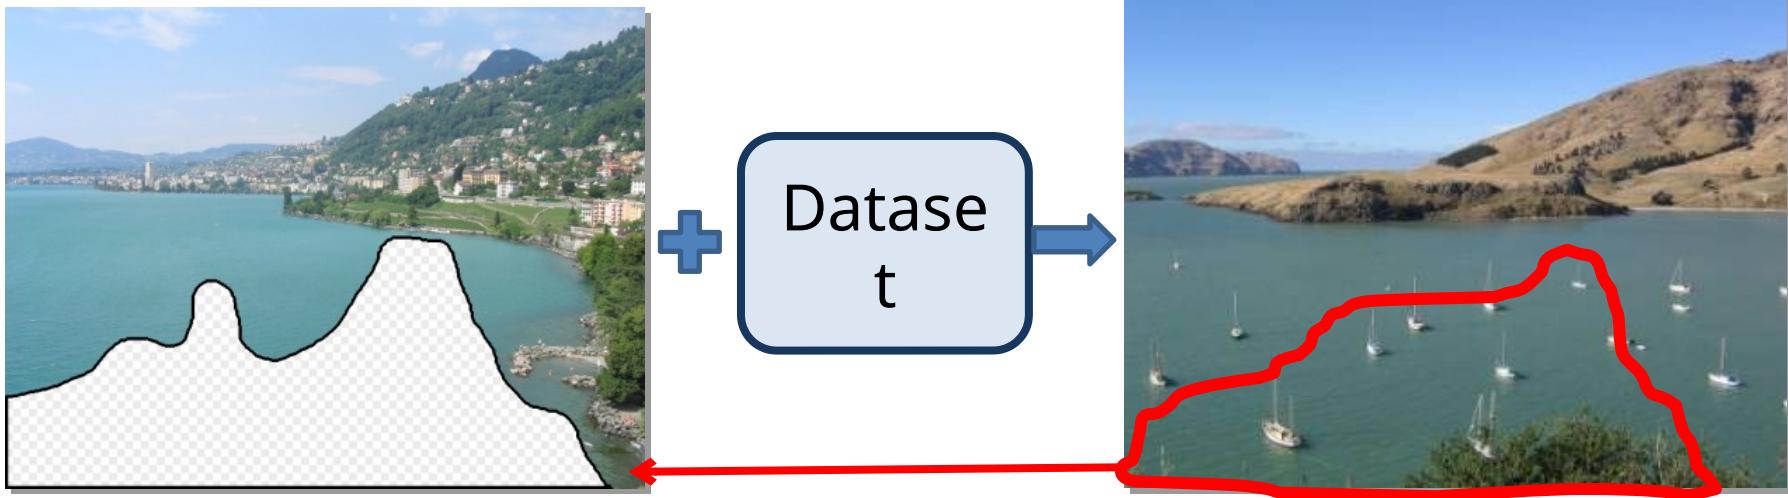
(b)



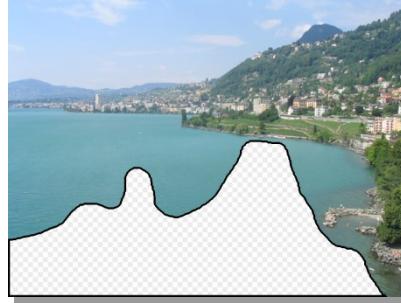
(c)

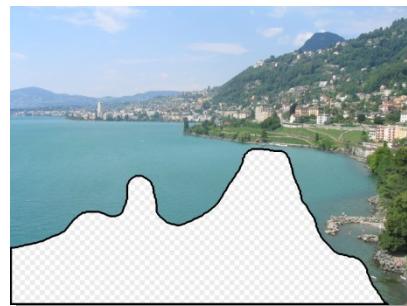
# How it works

- Find a similar image from a large dataset
- Blend a region from that image into the hole

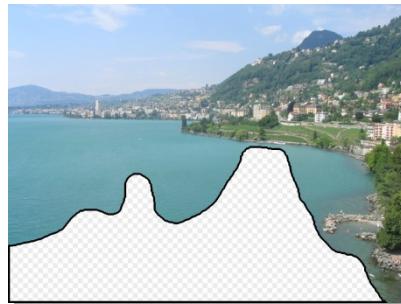


# How many images is enough?





Nearest neighbors from a collection of 20 thousand images



Nearest neighbors from a  
collection of 2 million images

# Image Data on the Internet

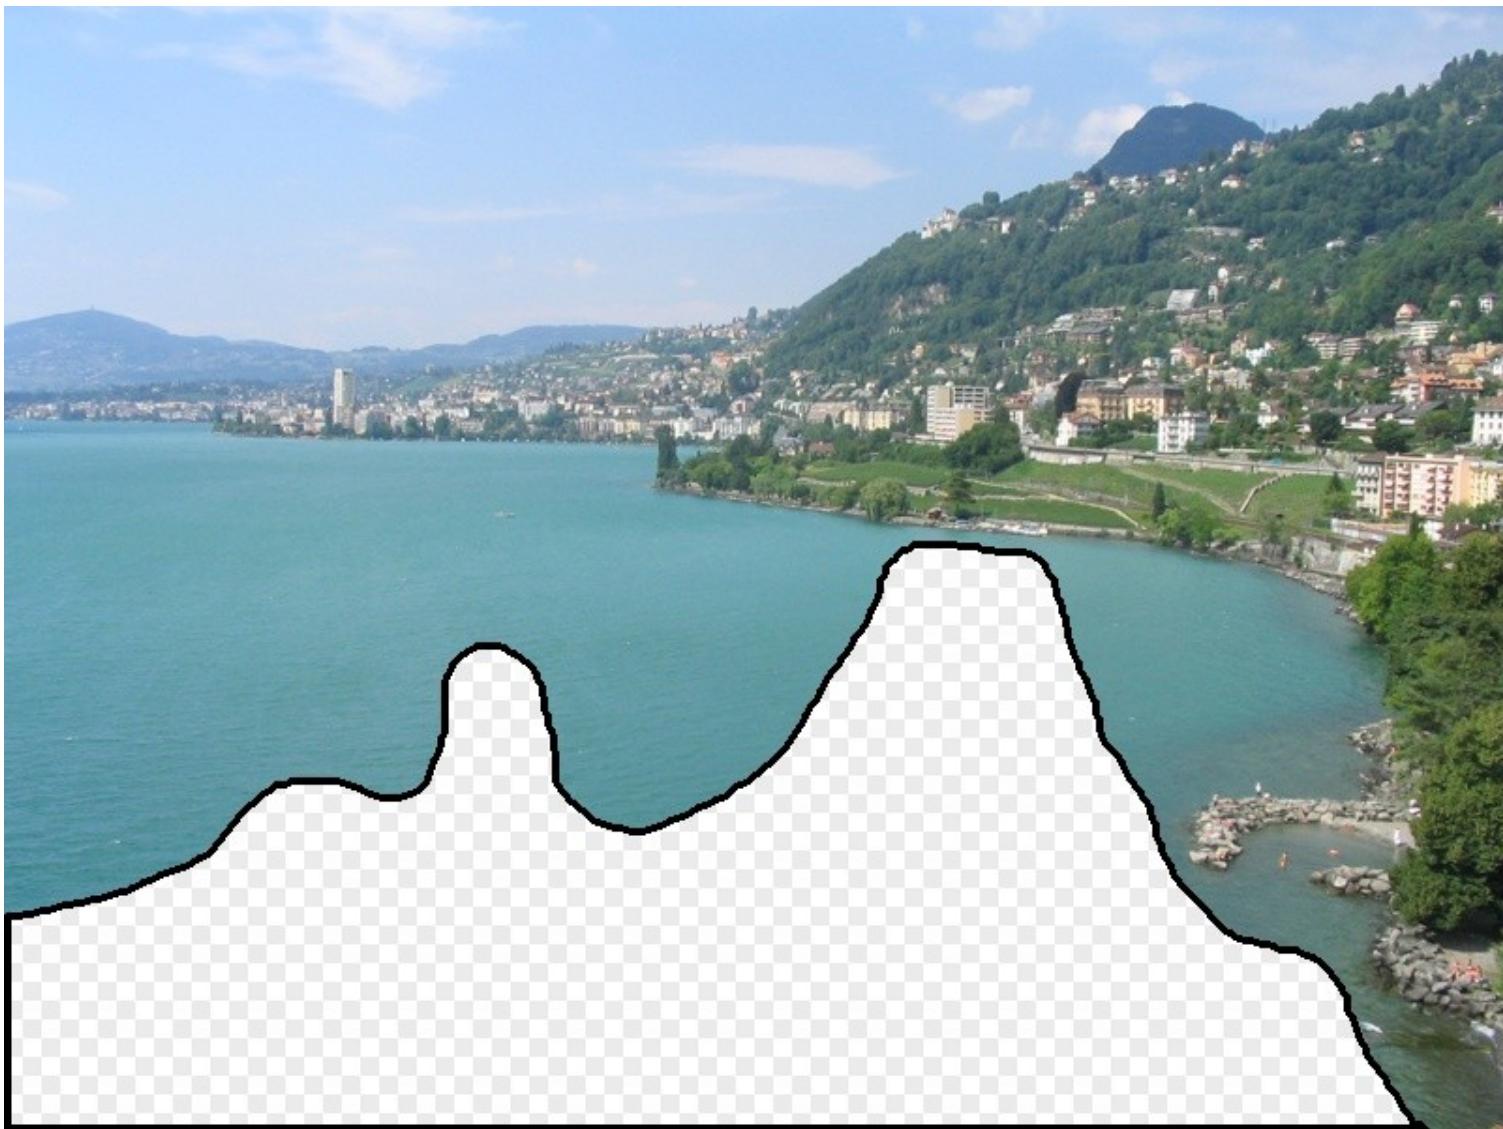
- Flickr (as of Nov 2015)
  - 10 billion photographs
  - 100+ million geotagged images
  - 3.5 million a day
- Facebook (as of Sept 2013)
  - 250 billion+
  - 300 million a day
- Instagram
  - 93 million a day

# The Algorithm:

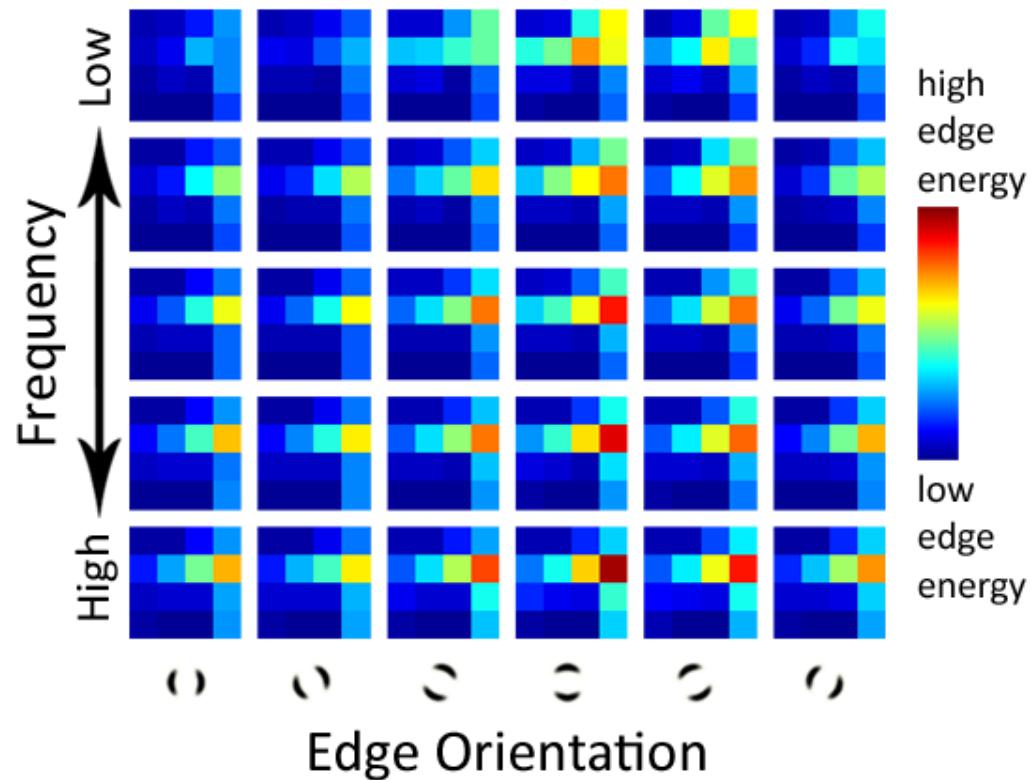
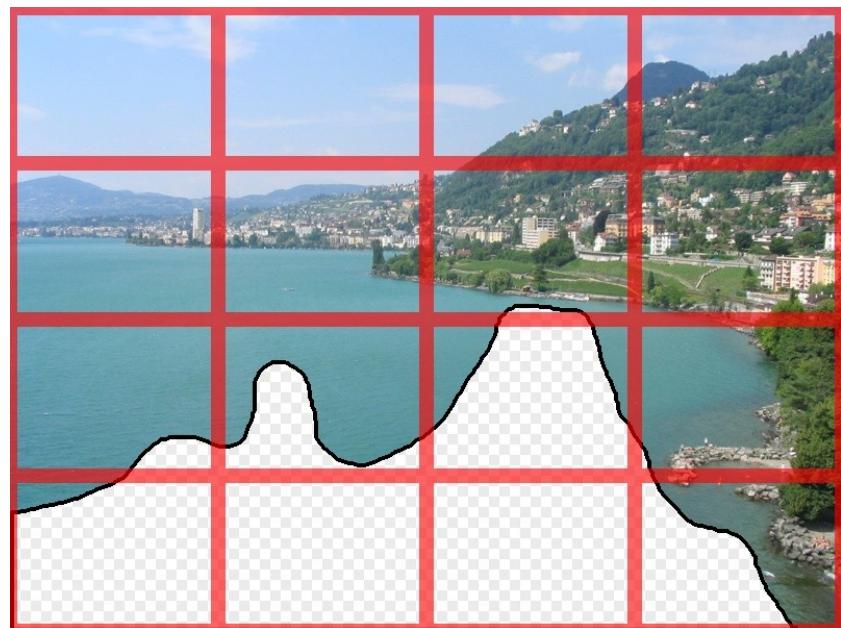
[Hays and Efros. Scene Completion Using Millions of Photographs. SIGGRAPH 2007 and CACM October 2008.]



# Scene Matching

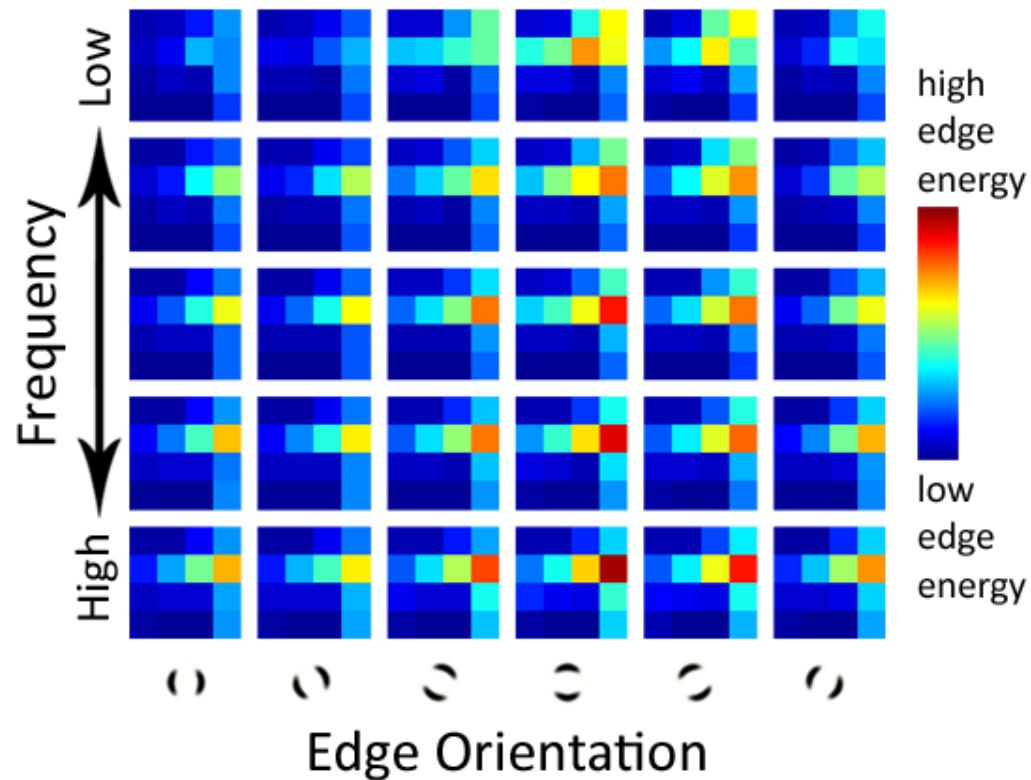
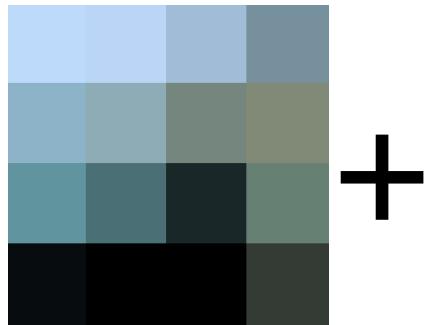


# Scene Descriptor



Scene Gist Descriptor  
(Oliva and Torralba 2001)

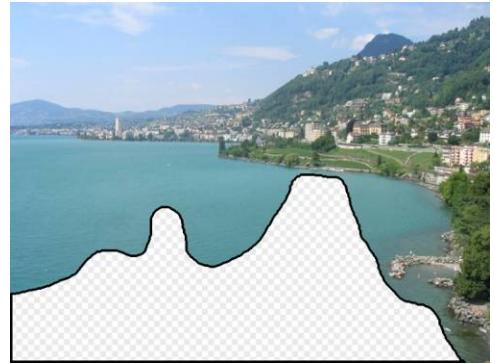
# Scene Descriptor



Scene Gist Descriptor  
(Oliva and Torralba 2001)

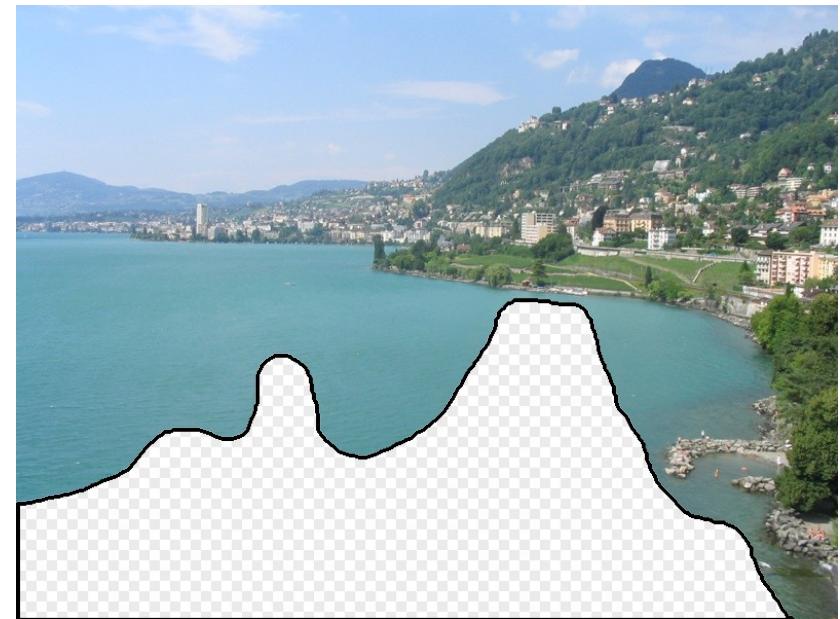
# 2 Million Flickr Images

The background of the image is a dense, uniform grid composed of numerous small, square thumbnail images. These thumbnails represent a vast collection of photographs from Flickr, showing a wide variety of subjects and colors. The overall effect is a visual representation of the scale and diversity of user-generated content.



... 200 total

# Context Matching

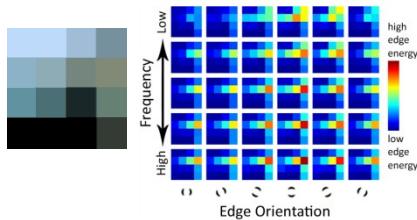




Graph cut + Poisson

# Result Ranking

We assign each of the 200 results a score which is the sum of:



The scene matching distance



The context matching distance  
(color + texture)



The graph cut cost

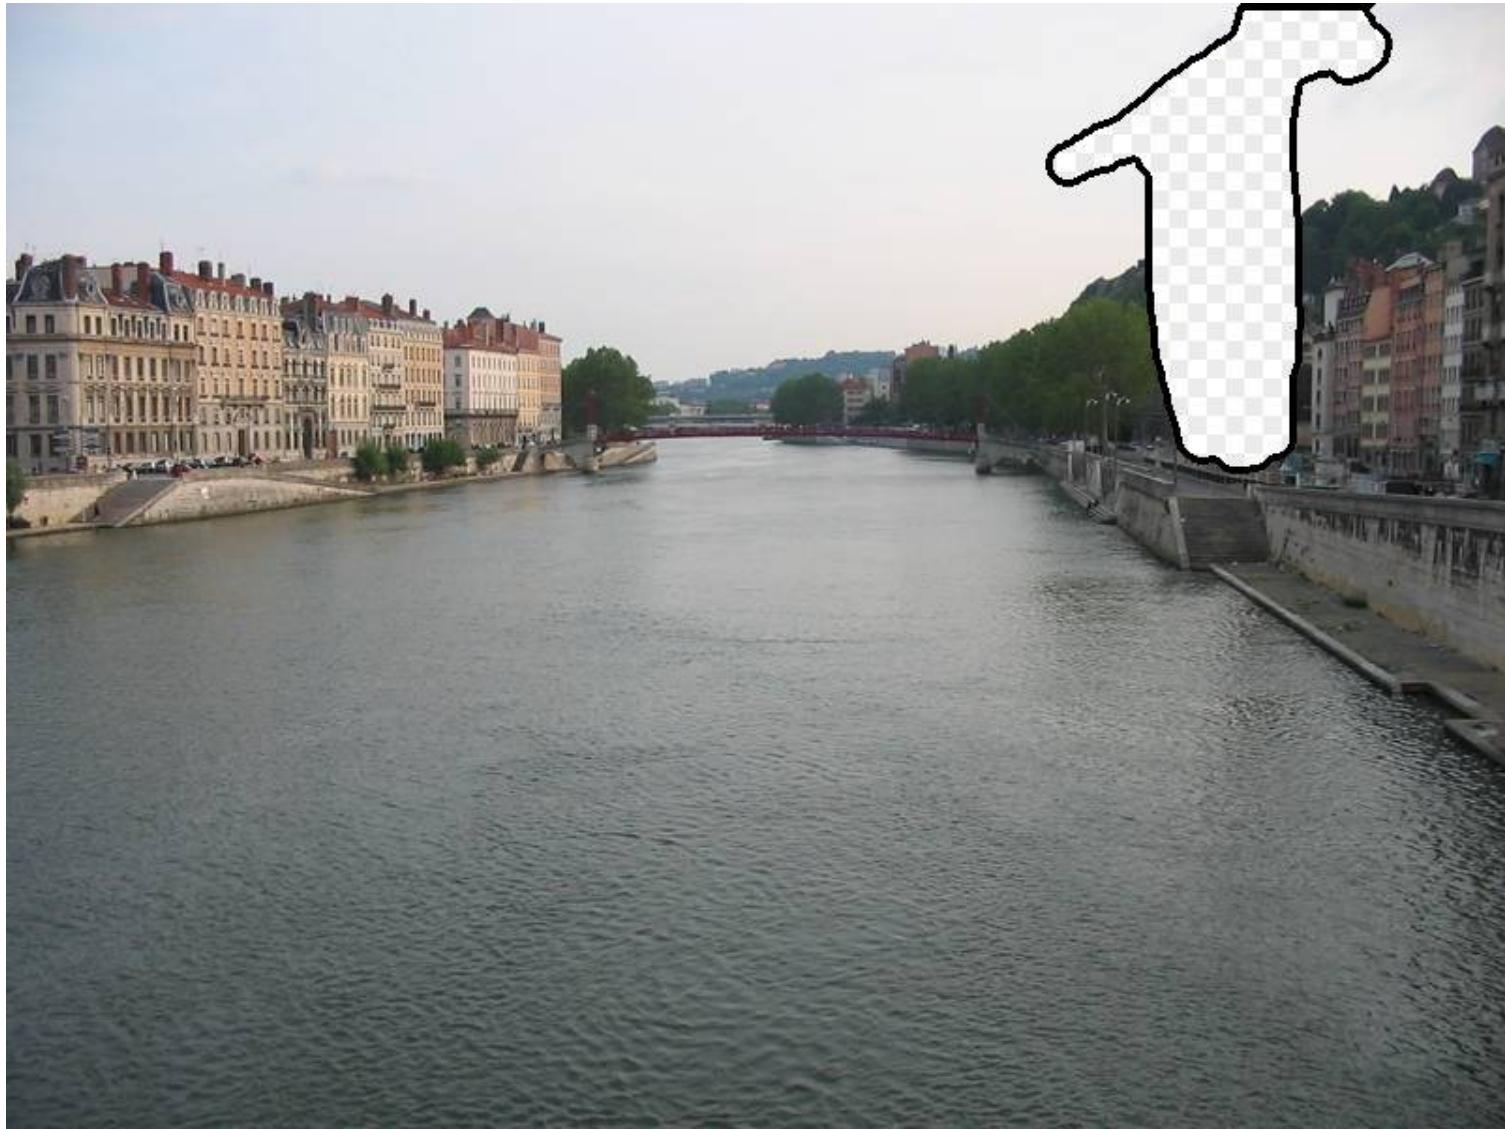




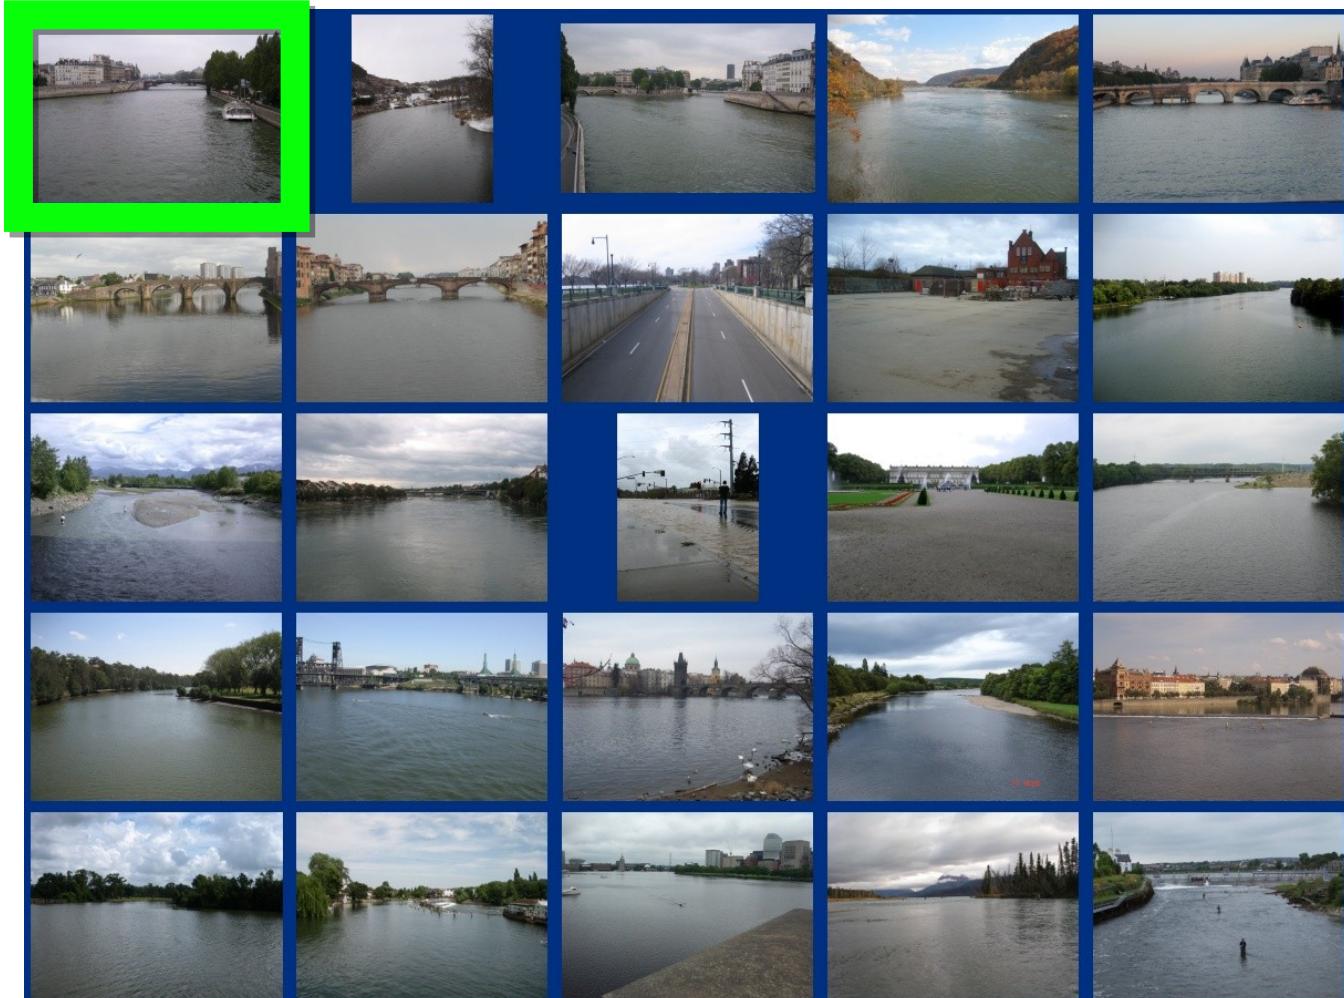








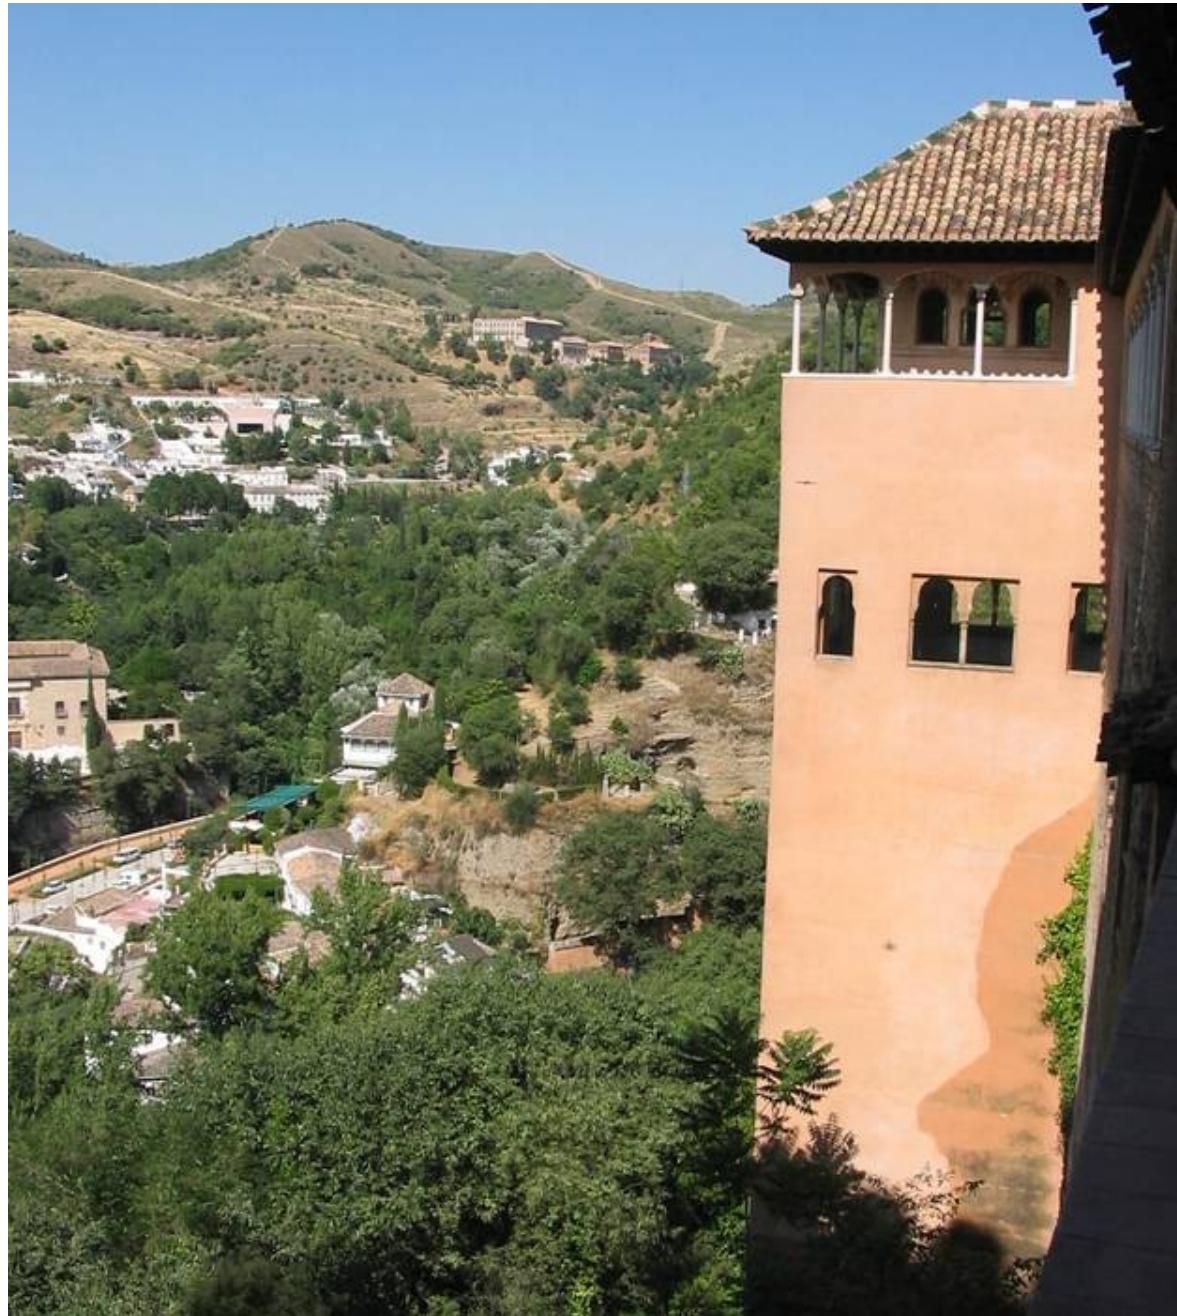


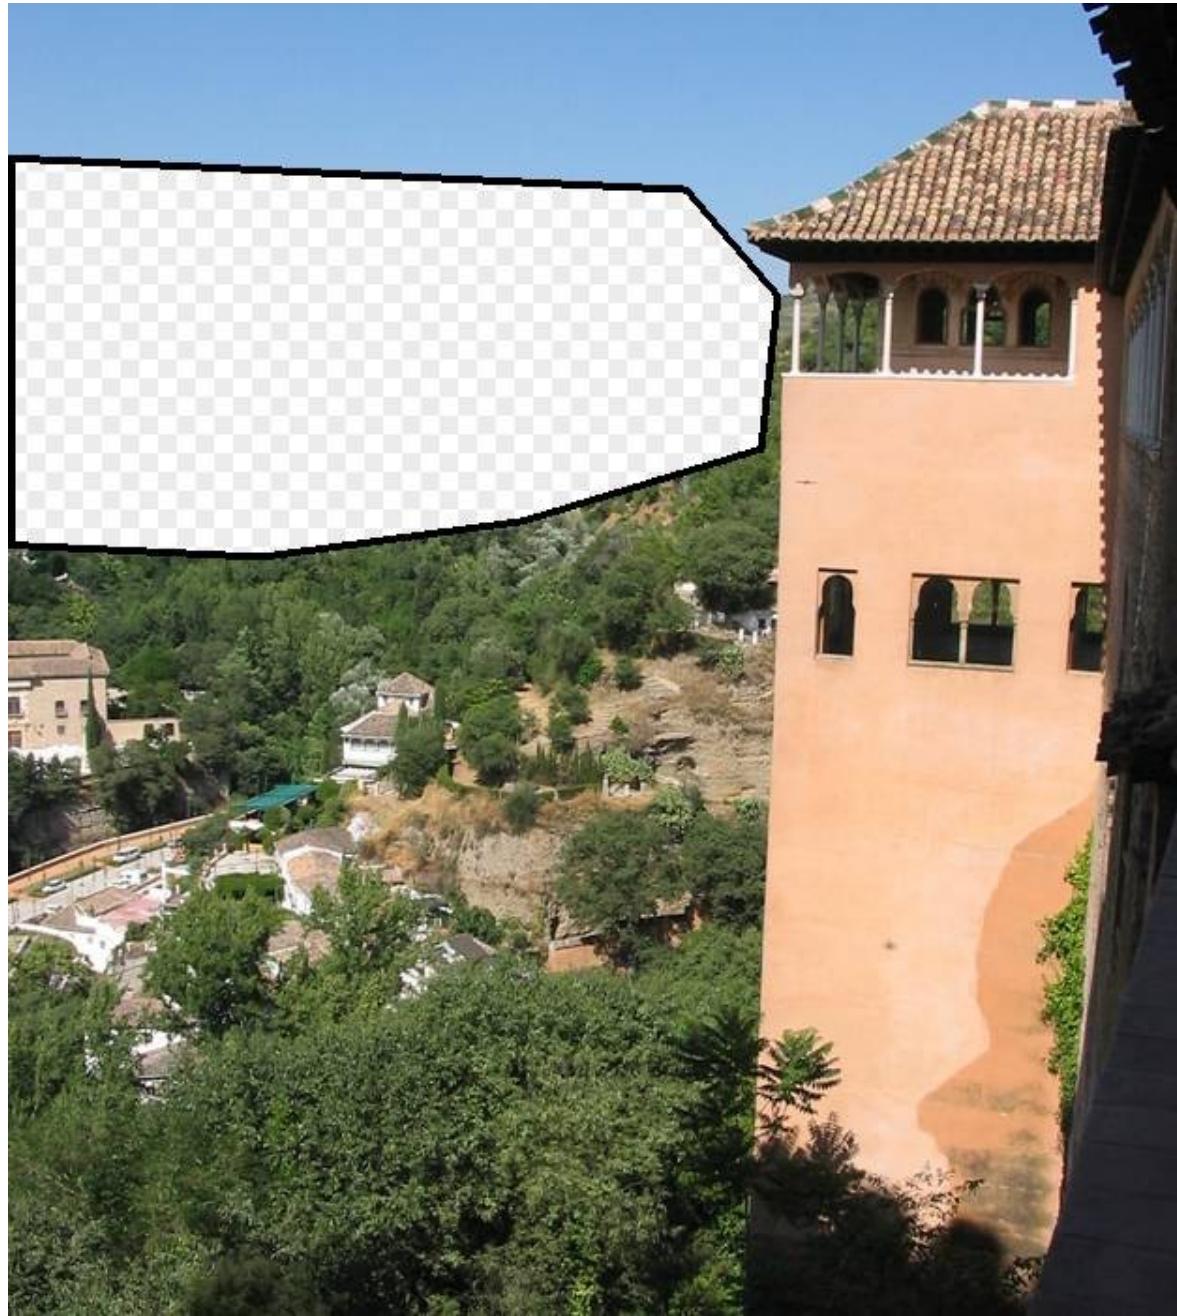


... 200 scene  
matches











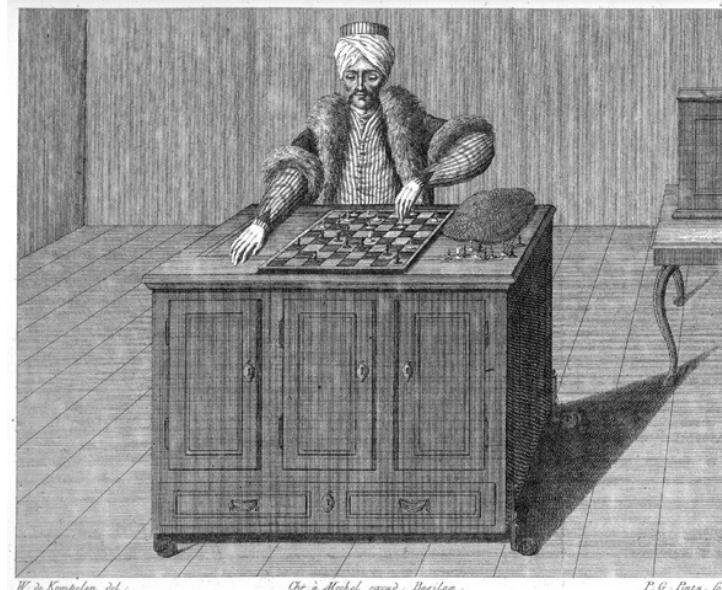
# Which is the original?





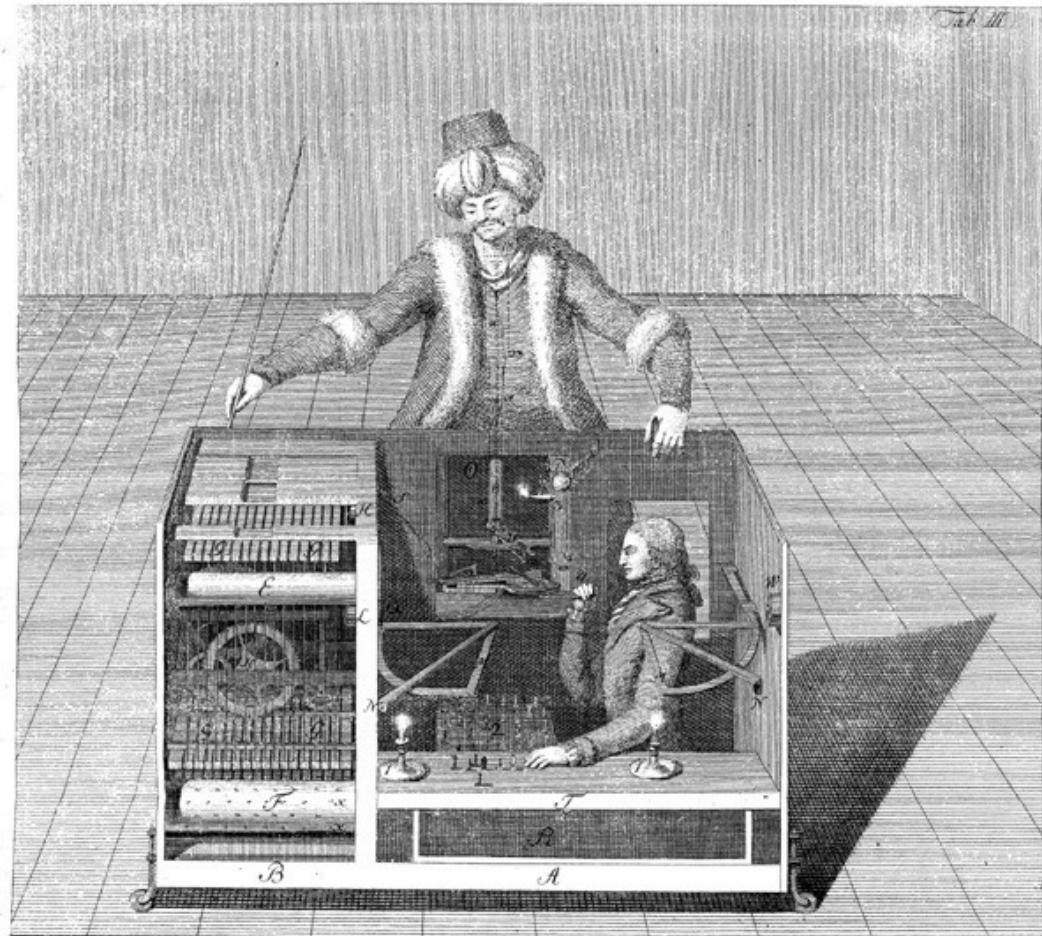
# Mechanical Turk

- von Kempelen, 1770 to 1854.
- Robotic chess player.
- Clockwork routines.
- Magnetic induction (not vision)
- Toured the world; played Napoleon Bonaparte and Benjamin Franklin.



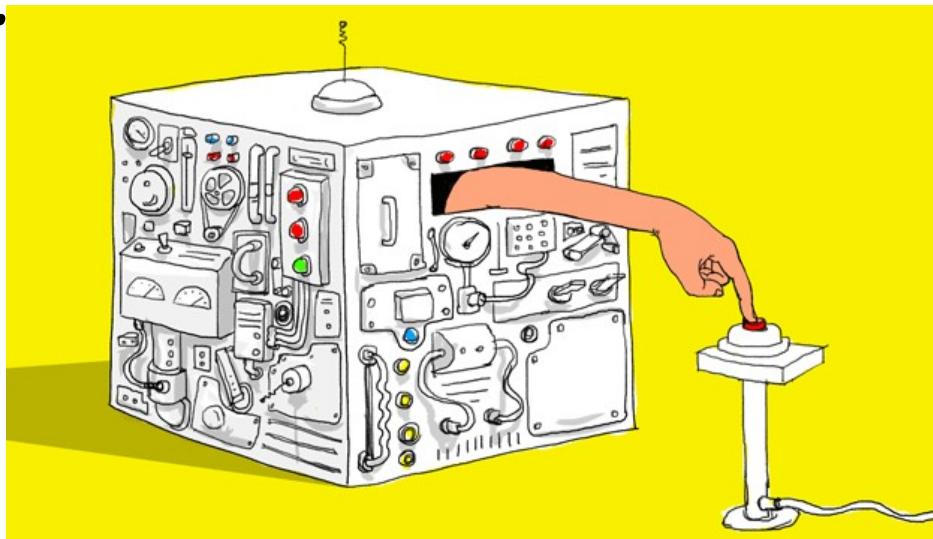
# Mechanical Turk

- It was all a ruse!
- Ho ho ho.



# Amazon Mechanical Turk

*Artificial artificial  
intelligence.*

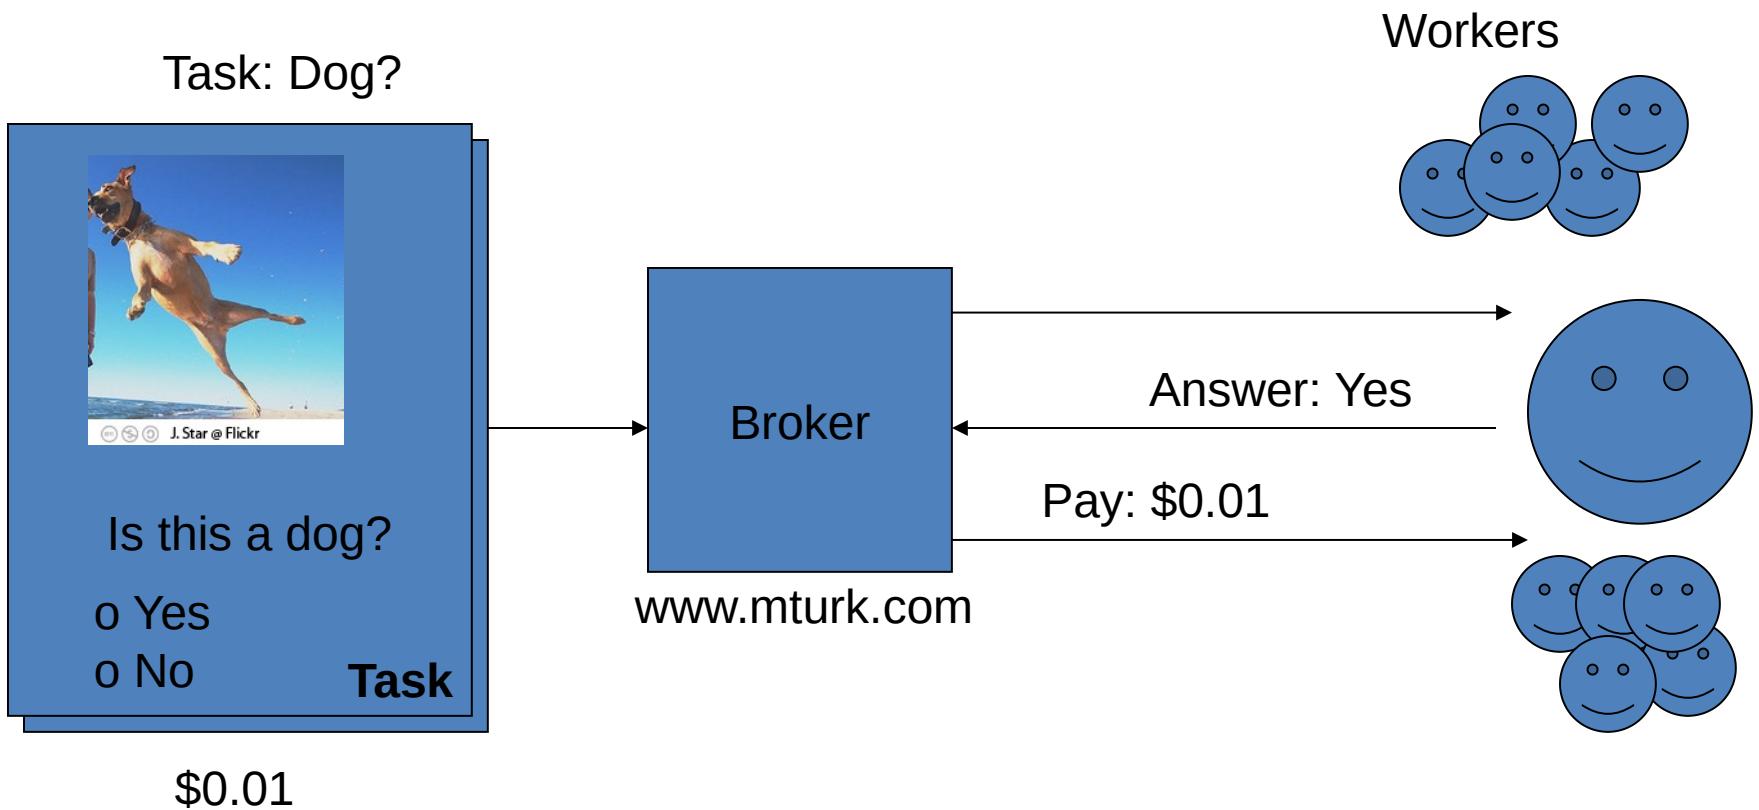


Launched 2005.  
Small tasks, small  
pay.  
Used extensively  
in data collection.

## Amazon Mechanical Turk

Access a global, on-demand, 24x7 workforce

# Amazon Mechanical Turk

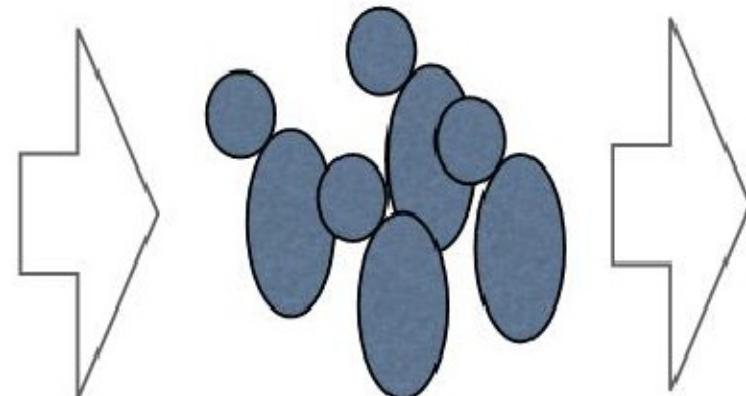


6000 images  
from flickr.com



# Building datasets

## Annotators



**amazon**mechanical turk  
beta Artificial Intelligence

Is there an Indigo bunting in the image?

100s of  
training images



# Typical payments via Amazon Mechanical Turk

- To transcribe a receipt (\$0.01)
- Summarize a block of text (\$0.35)
- Take a behavioral economics survey (\$1)

| HIT Groups (1-20 of 640) |                                     |       |        |         |         |               |
|--------------------------|-------------------------------------|-------|--------|---------|---------|---------------|
| Requester                | Title                               | HITS  | Reward | Created | Actions |               |
| ScoutIt                  | Classify Receipt                    | 151   | \$0.03 | 14s ago | Preview | Qualify       |
| Crowdsurf Support        | Full Text Review - Earn up to \$... | 53    | \$0.17 | 3m ago  | Preview | Qualify       |
| Laura A. King            | Personality, Information Proce...   | 1     | \$0.15 | 4m ago  | Preview | Accept & Work |
| Crowdsurf Support        | Review, edit, and score the tra...  | 1,091 | \$0.02 | 5m ago  | Preview | Qualify       |
| Erica Fissel             | Quick Demographic Survey! (~...     | 1     | \$0.01 | 6m ago  | Preview | Accept & Work |
| ScoutIt                  | Extract summary information fr...   | 1     | \$0.05 | 9m ago  | Preview | Accept & Work |
| Crowdsurf Support        | Transcribe up to 35 Seconds o...    | 1,042 | \$0.05 | 10m ago | Preview | Qualify       |
| Ben Stevens              | Help Pick a Book Cover!             | 1     | \$0.10 | 12m ago | Preview | Accept & Work |
| ScoutIt                  | Extract summary information fr...   | 1     | \$0.05 | 12m ago | Preview | Accept & Work |
| Michael Busseri, PhD     | Answer survey (10 minutes) a...     | 1     | \$1.00 | 12m ago | Preview | Qualify       |
| Amy Minnikin             | Feedback Seeking Motives Pr...      | 111   | \$0.25 | 15m ago | Preview | Qualify       |
| SEO BrainTrust           | Summarize and write three ke...     | 23    | \$0.35 | 25m ago | Preview | Qualify       |

<https://thehustle.co/making-money-on-amazon-mechanical-turk/>

# Vision (Segmentation): LabelMe

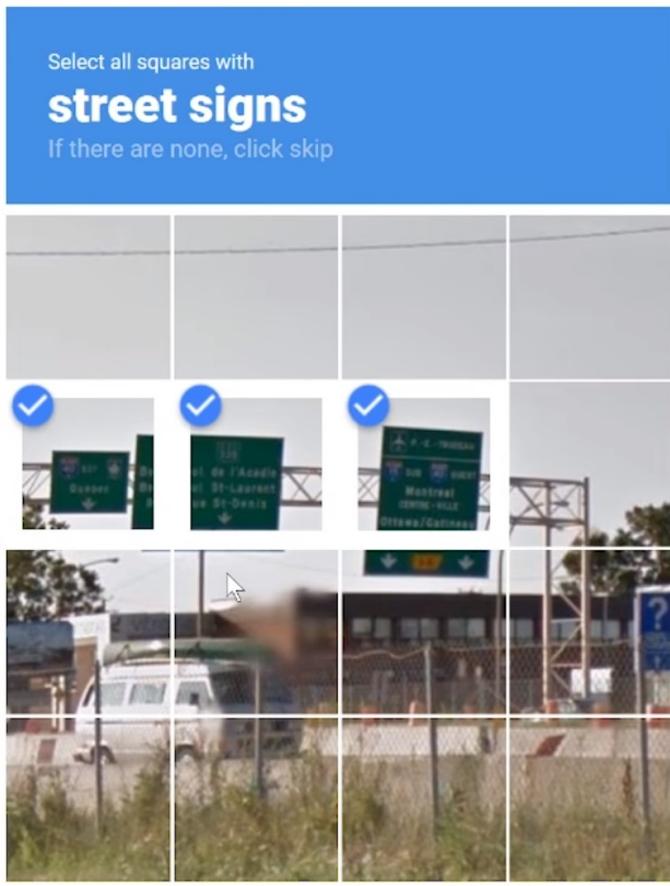
<http://labelme.csail.mit.edu>

“Open world” database annotated by the  
community\*

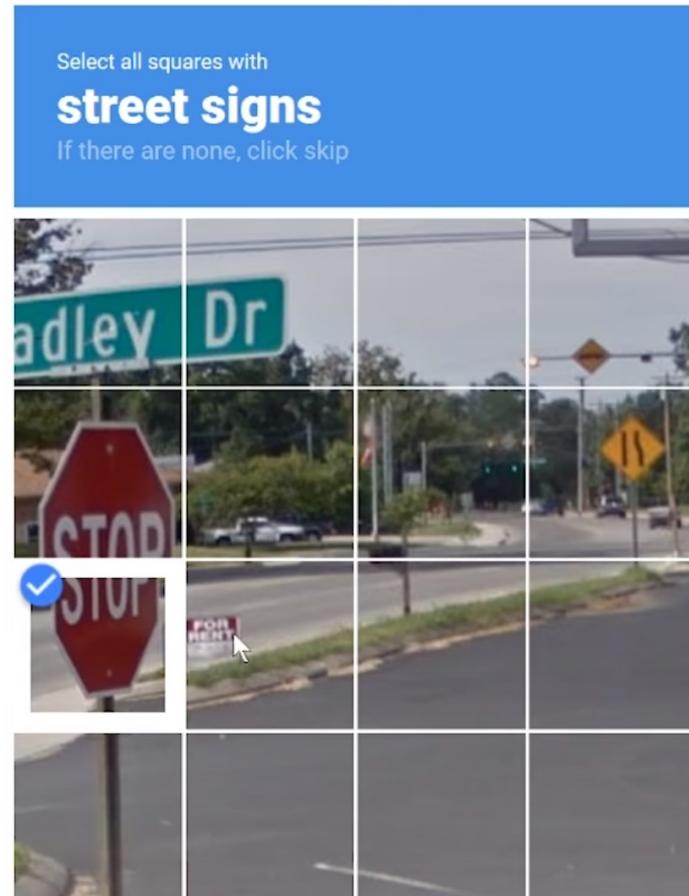
**Notes on Image Annotation**, Barriuso  
and Torralba 2012.

<http://arxiv.org/abs/1210.3448>

# “Security checks”



NEXT



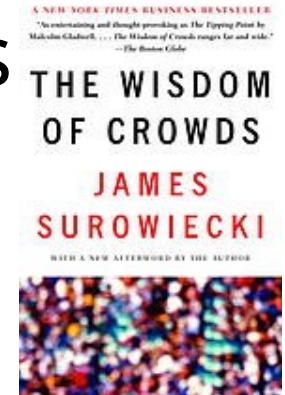
VERIFY

# Issues

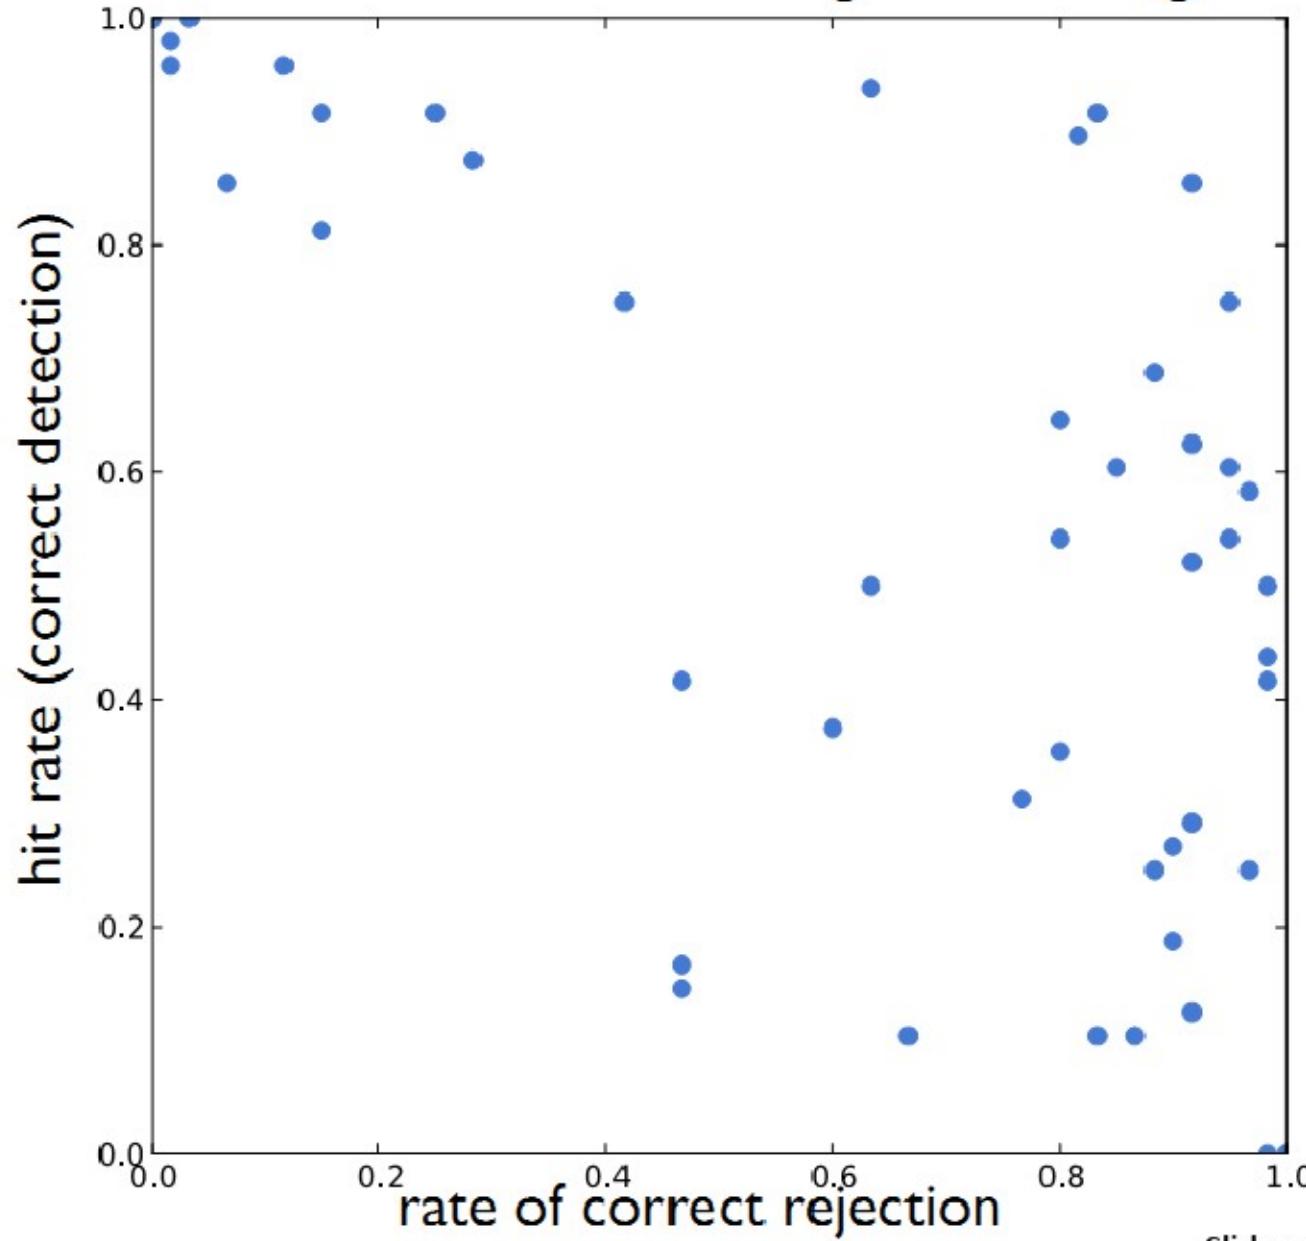
- Quality?
  - How good is it?
  - How to be sure?
- Price?
  - Trade off between throughput and cost
  - Higher pay can attract scammers

# Ensuring Annotation Quality

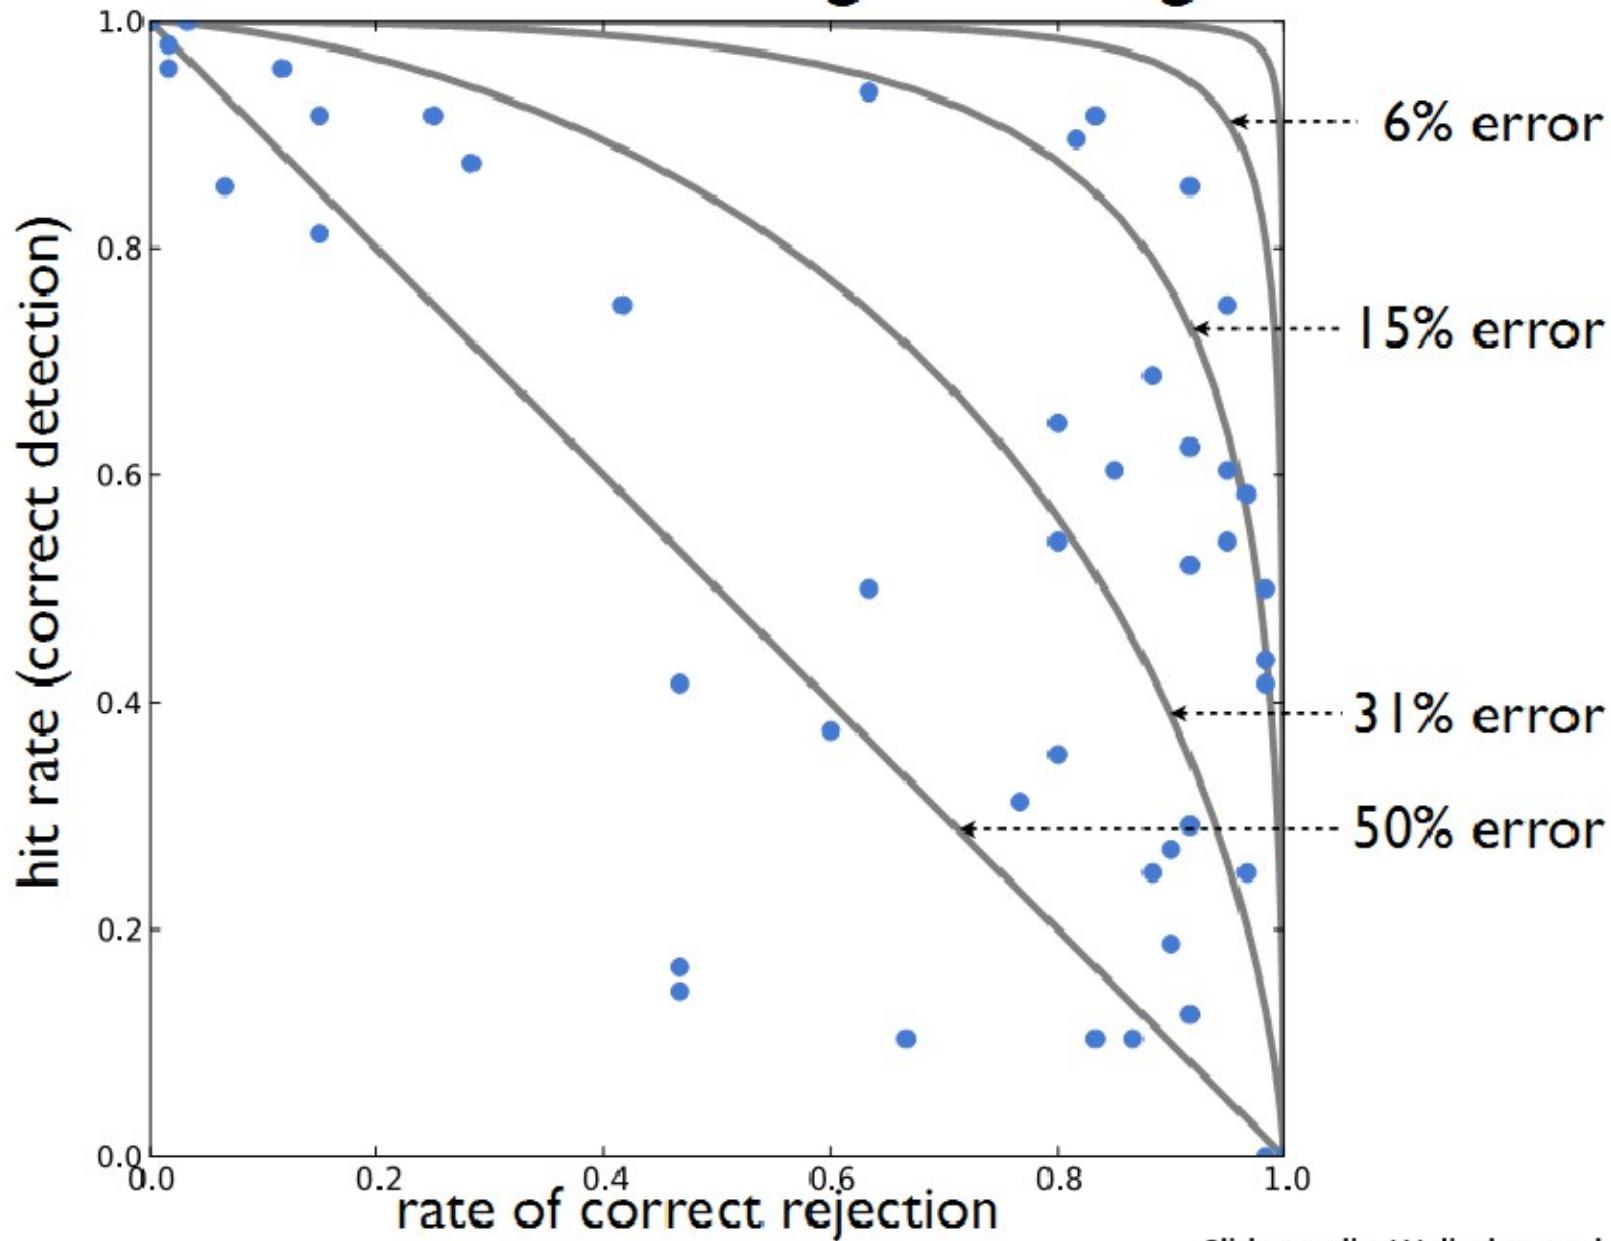
- Consensus in multiple annotations  
“Wisdom of the Crowds”  
*Not enough on its own, but widely used*
- Gold Standard / Sentinel
  - Special case: qualification exam  
*Widely used & important. Find good annotators; keep them honest.*
- Grading Tasks
  - A second tier of workers who grade others  
*Not widely used*



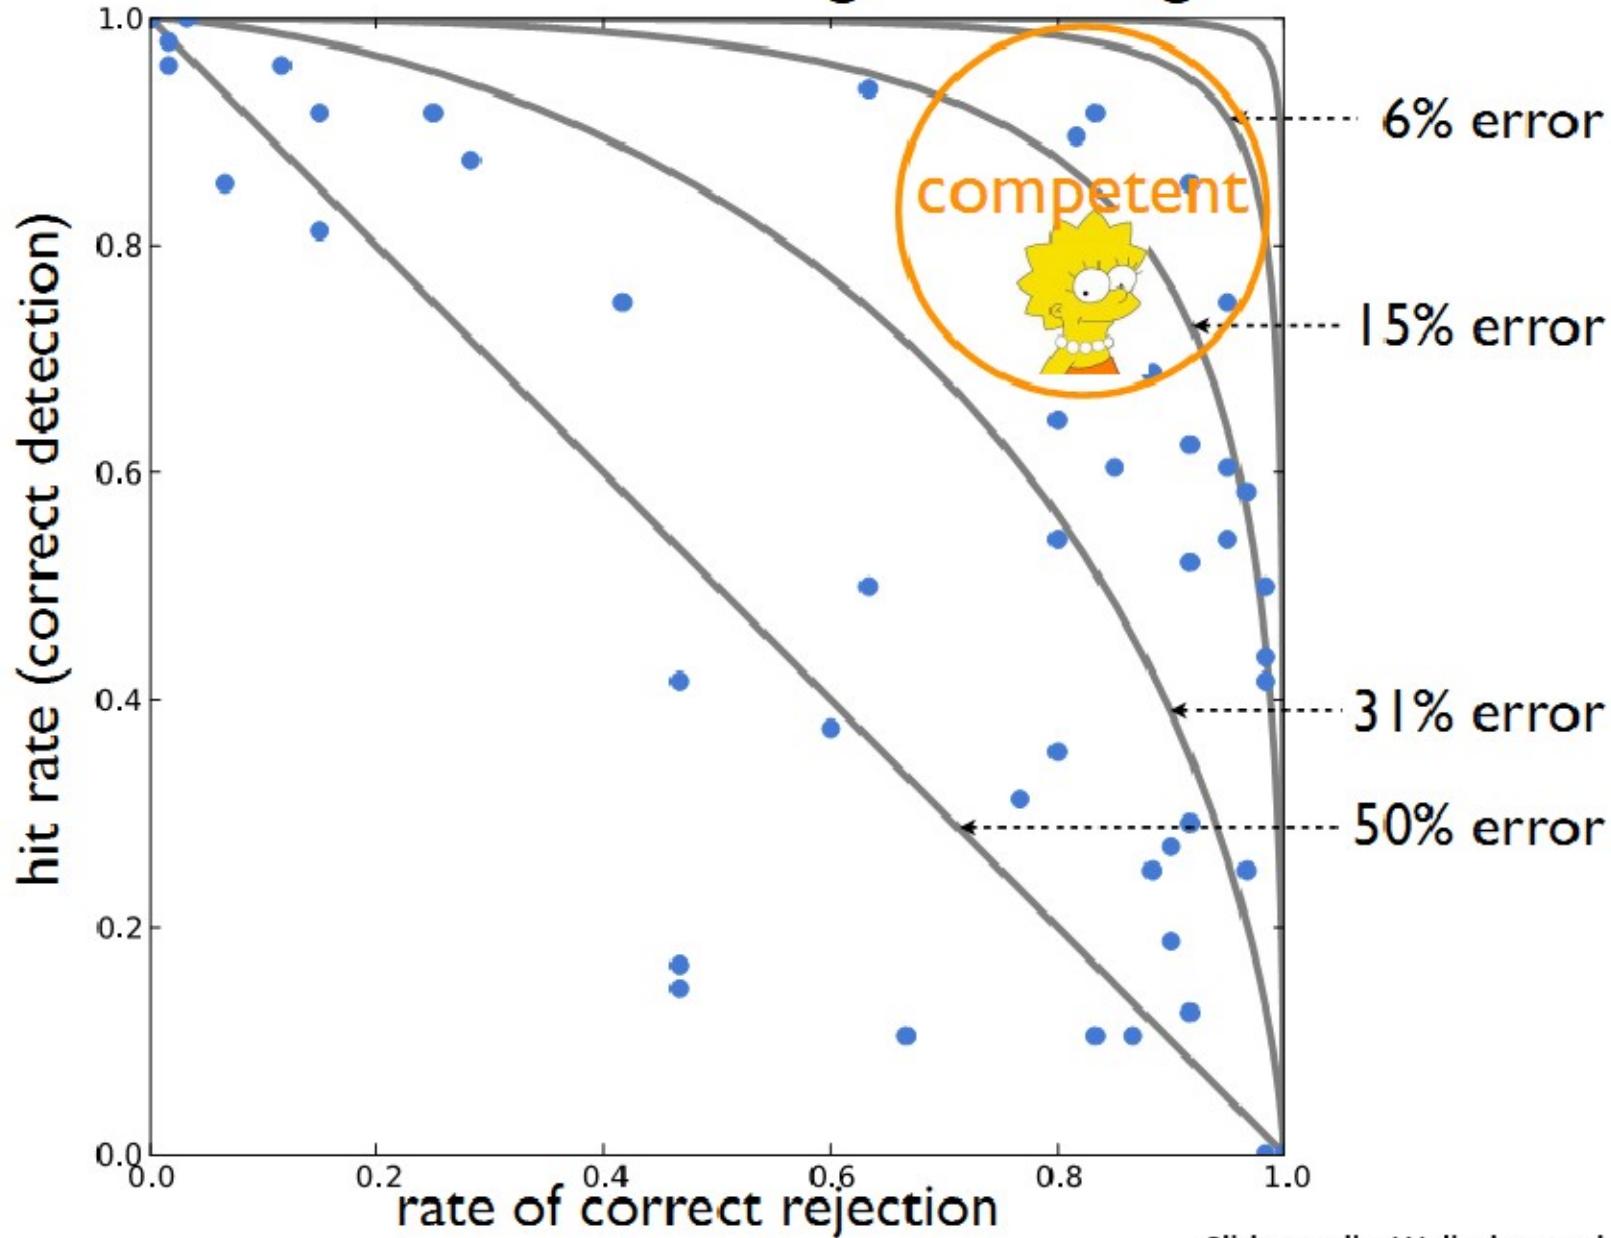
# Task: Find the Indigo Bunting



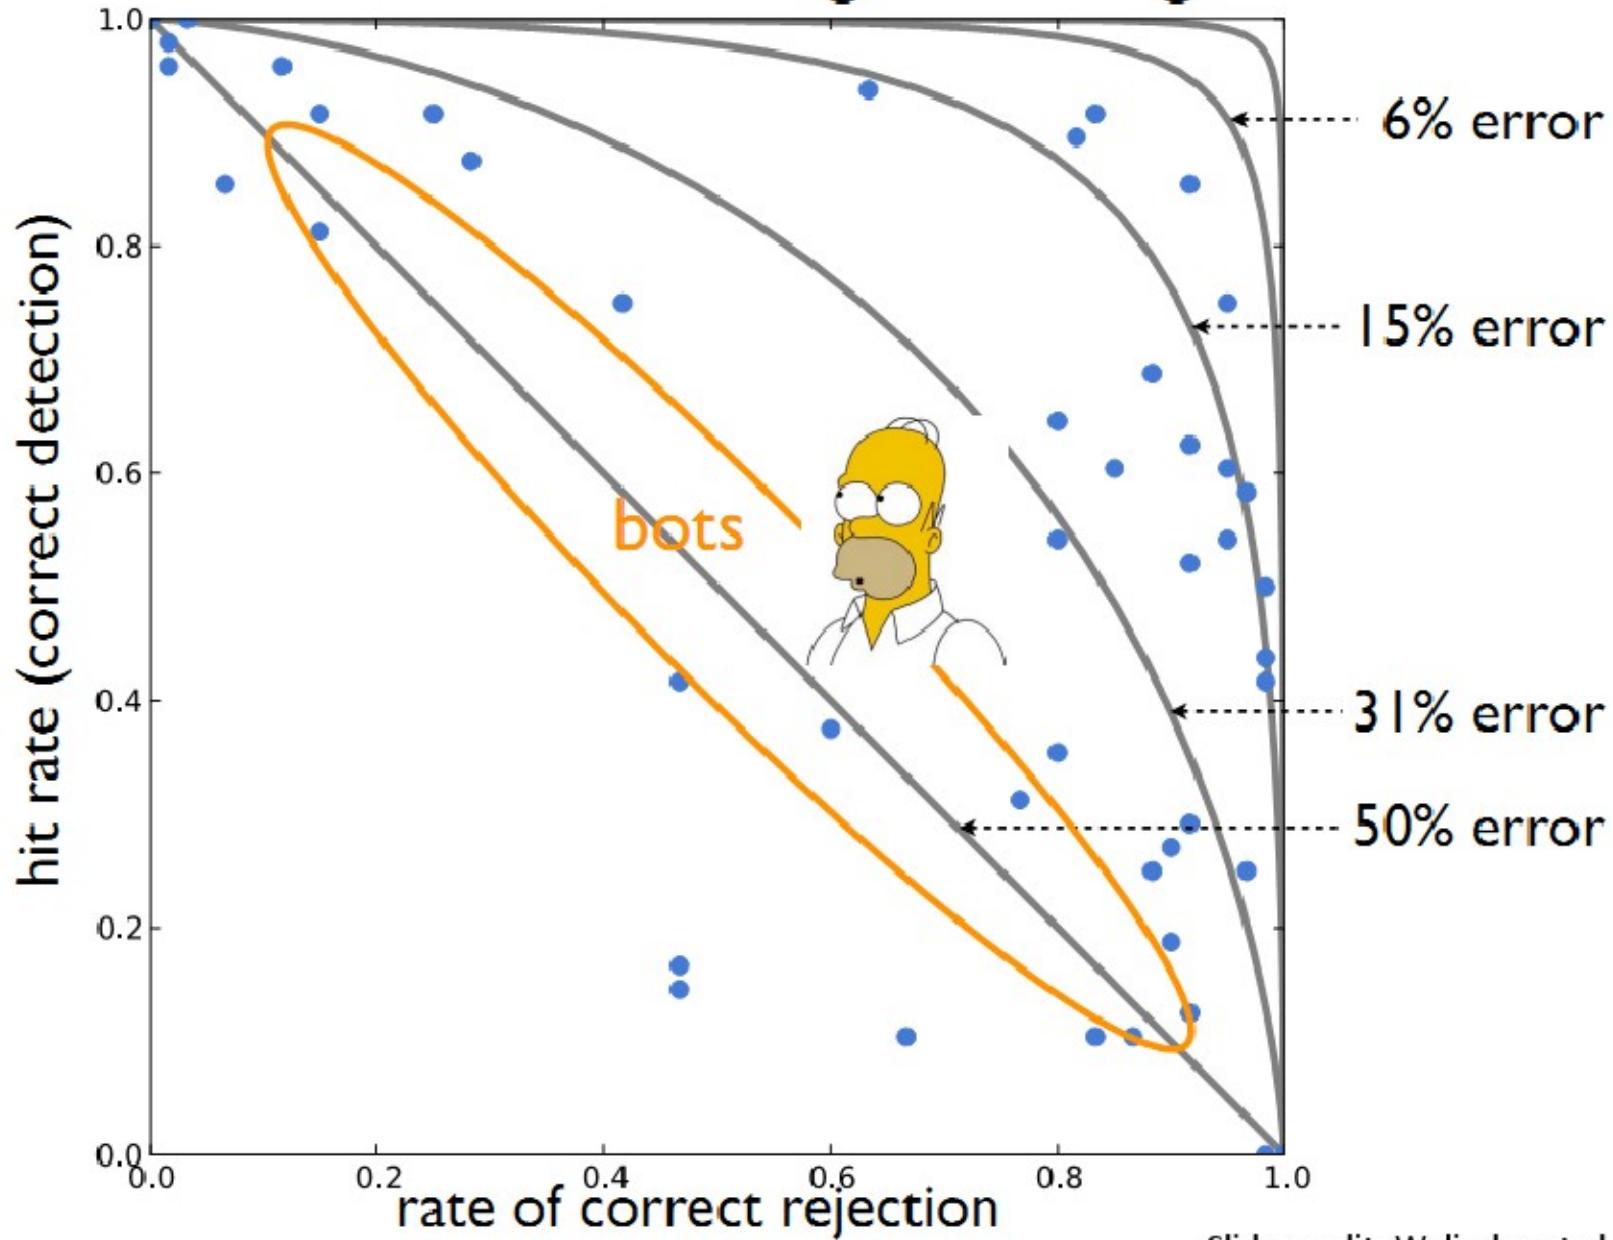
# Task: Find the Indigo Bunting



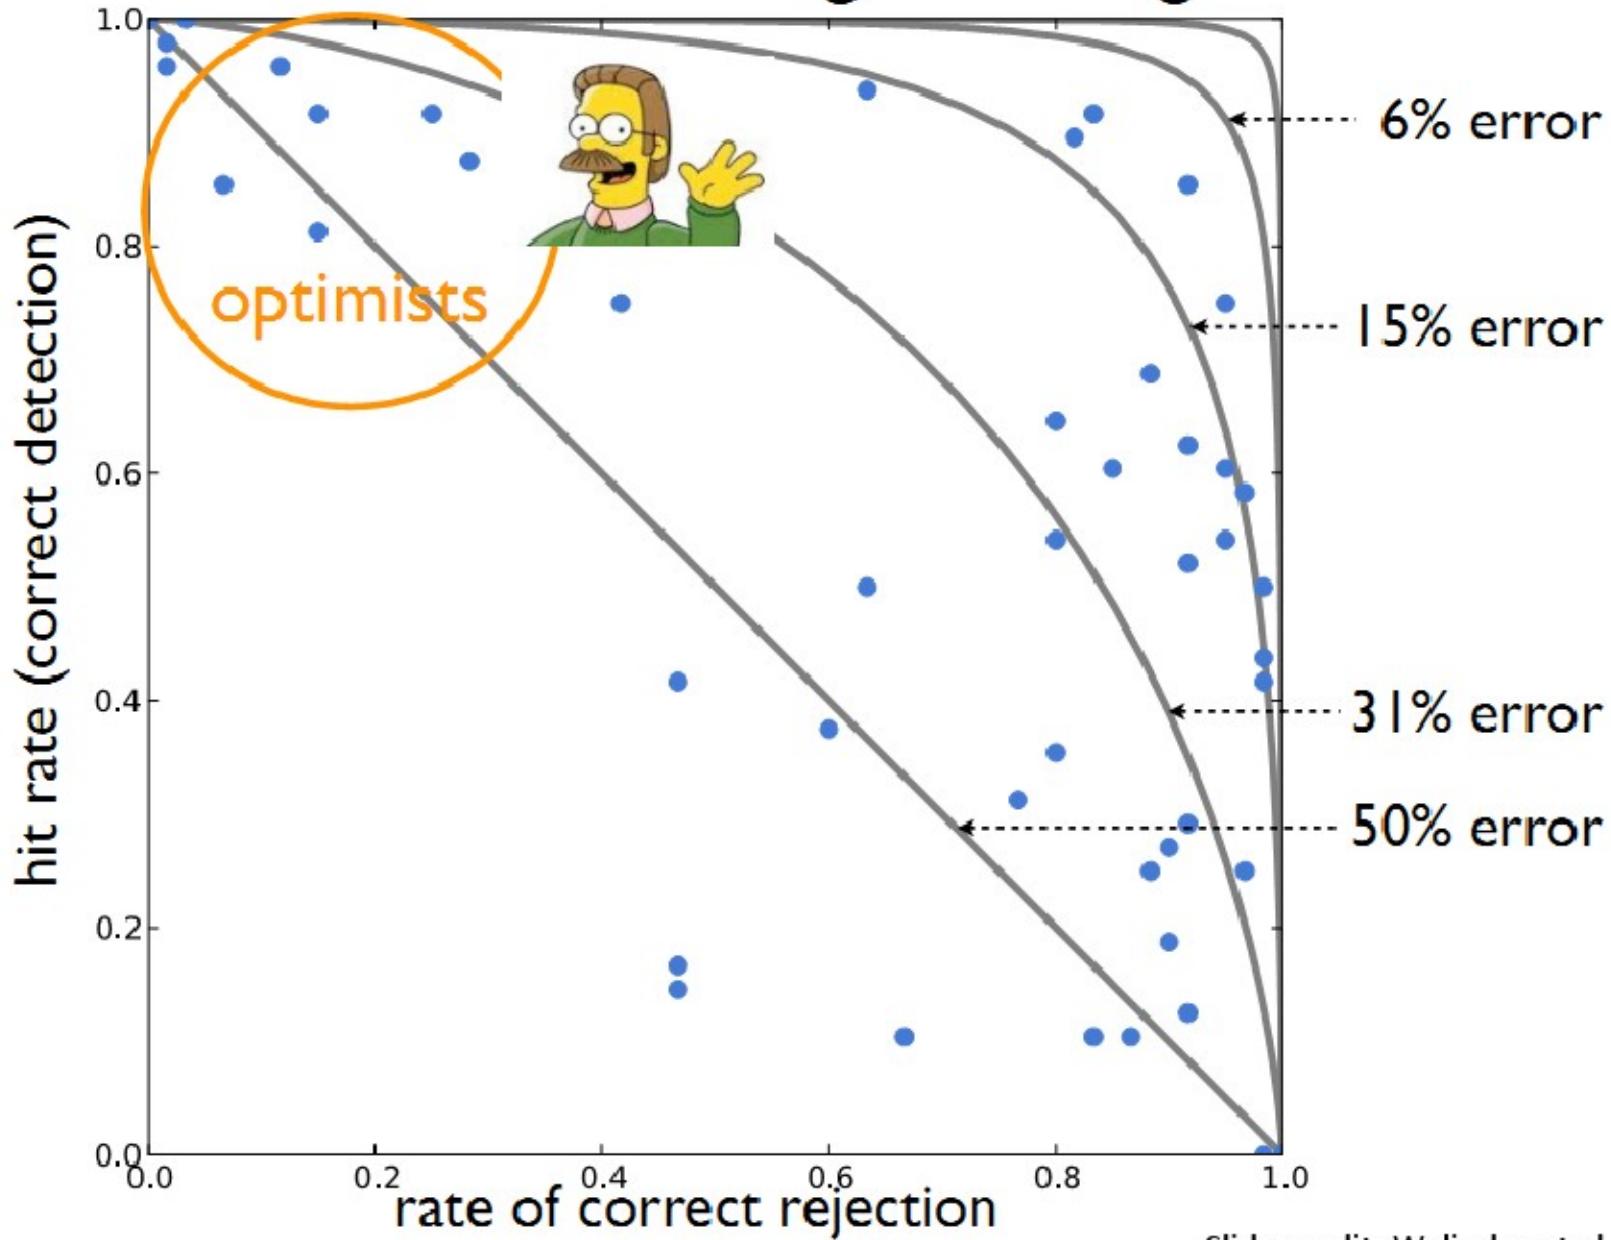
# Task: Find the Indigo Bunting



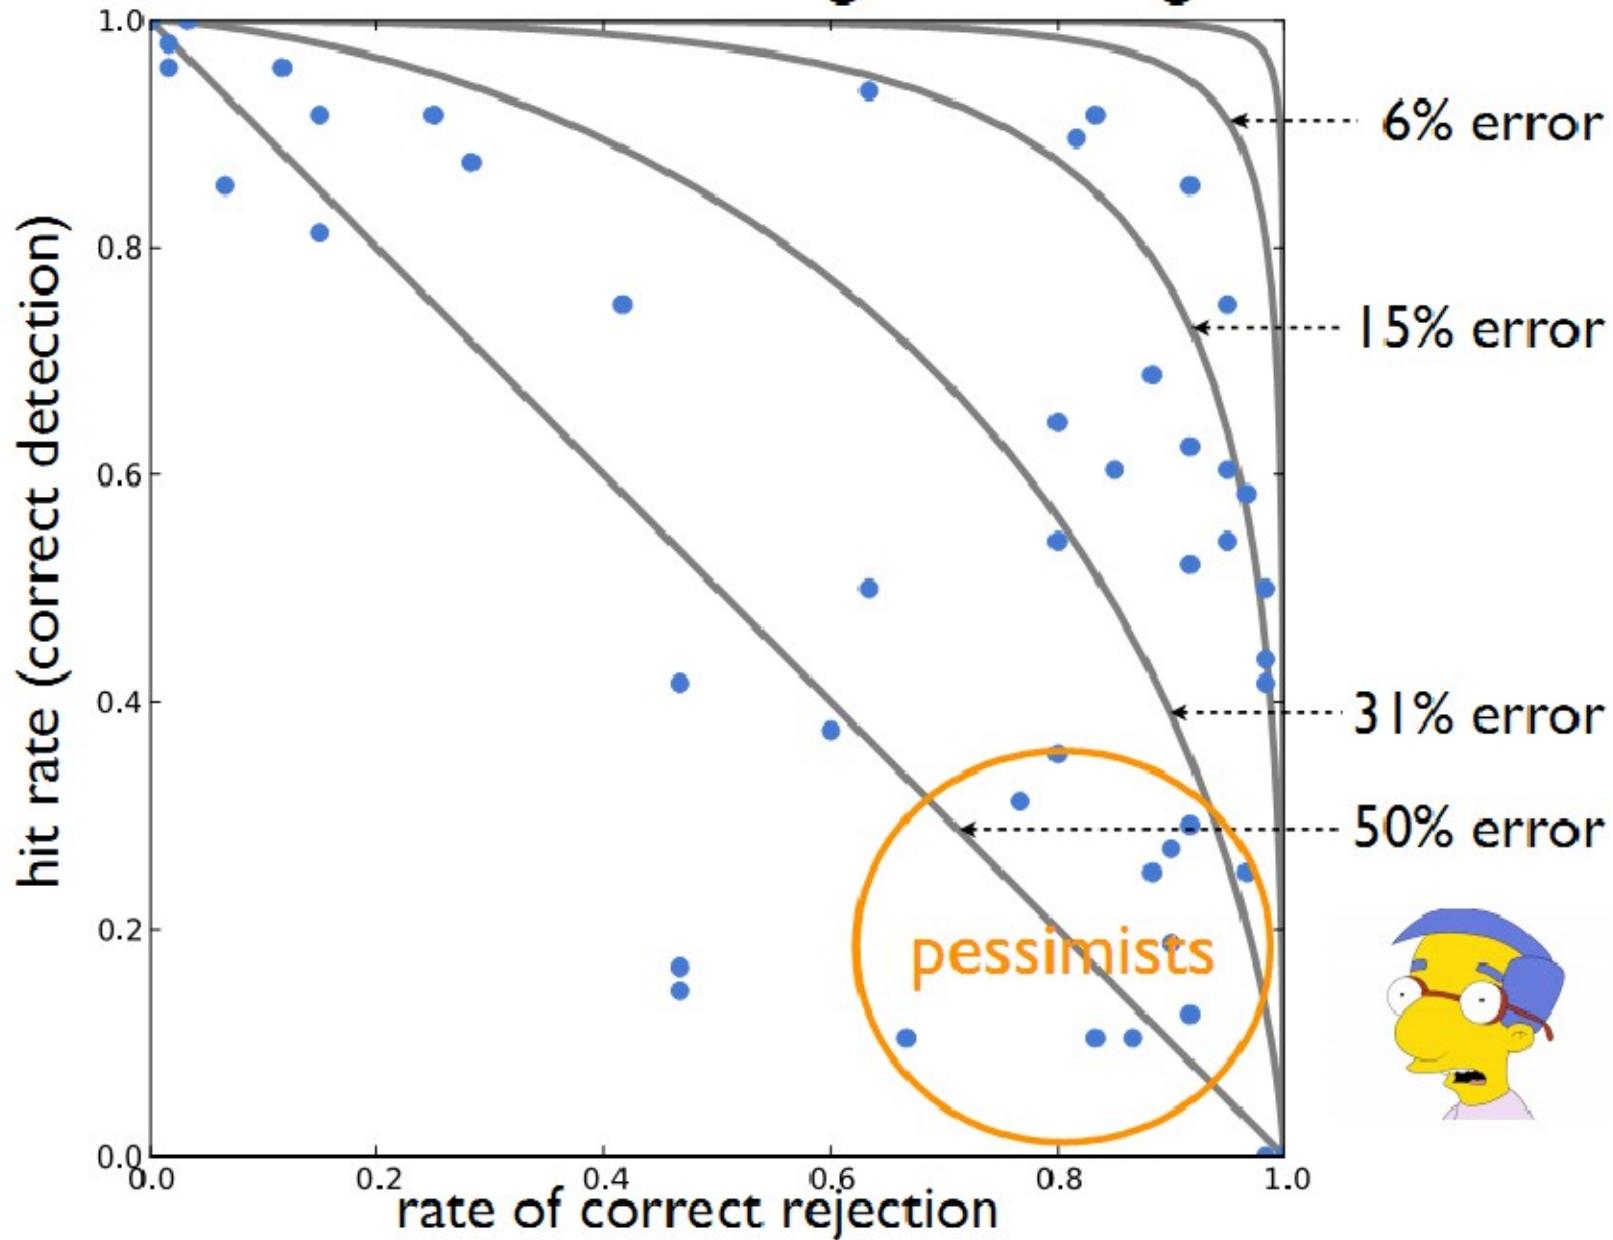
# Task: Find the Indigo Bunting



# Task: Find the Indigo Bunting



# Task: Find the Indigo Bunting



# Task: Find the Indigo Bunting

