


Training Neural Networks


Learning the weight matrices W

Gradient descent




General approach

Pick random starting point.




General approach

Compute gradient at point (analytically or by finite differences)



General approach


Move along parameter space in direction of negative gradient



γ = amount to move
= *learning rate*

General approach


Move along parameter space in direction of negative gradient.



γ = amount to move
= *learning rate*

General approach

Stop when we don't move any more.




$$a_{stop}: \\ a_{n-1} - \gamma \nabla f(a_{n-1}) = 0$$

Gradient descent



1847!!

- Optimizer for functions.
- Guaranteed to find optimum for convex functions.
 - Non-convex = find *local* optimum.
 - Most vision problems aren't convex.
- Works for multi-variate functions.
 - Need to compute matrix of *partial derivatives* ("*Jacobian*")



Convex functions

If a function is convex downward, the midpoint B of each chord (A_1, A_2) lies above the corresponding point A_0 of the graph of the function or coincides with this point.



Why would I use this over Least Squares?

If my function is convex,
why can't I just use linear least squares?

$$A\mathbf{x} - \mathbf{b} = 0$$

$$F(\mathbf{x}) = \|A\mathbf{x} - \mathbf{b}\|^2.$$

$$\nabla F(\mathbf{x}) = 2A^T (A\mathbf{x} - \mathbf{b}).$$

Analytic solution = normal equations $\mathbf{x} = (A^T A)^{-1} A^T \mathbf{b}$

You can, yes.

Why would I use this over Least Squares?

But now imagine that you have 1,000,000 data points.

Matrices are ridiculously large. You will need to invert a (1 000 000 x 1 000 000) matrix (have fun).

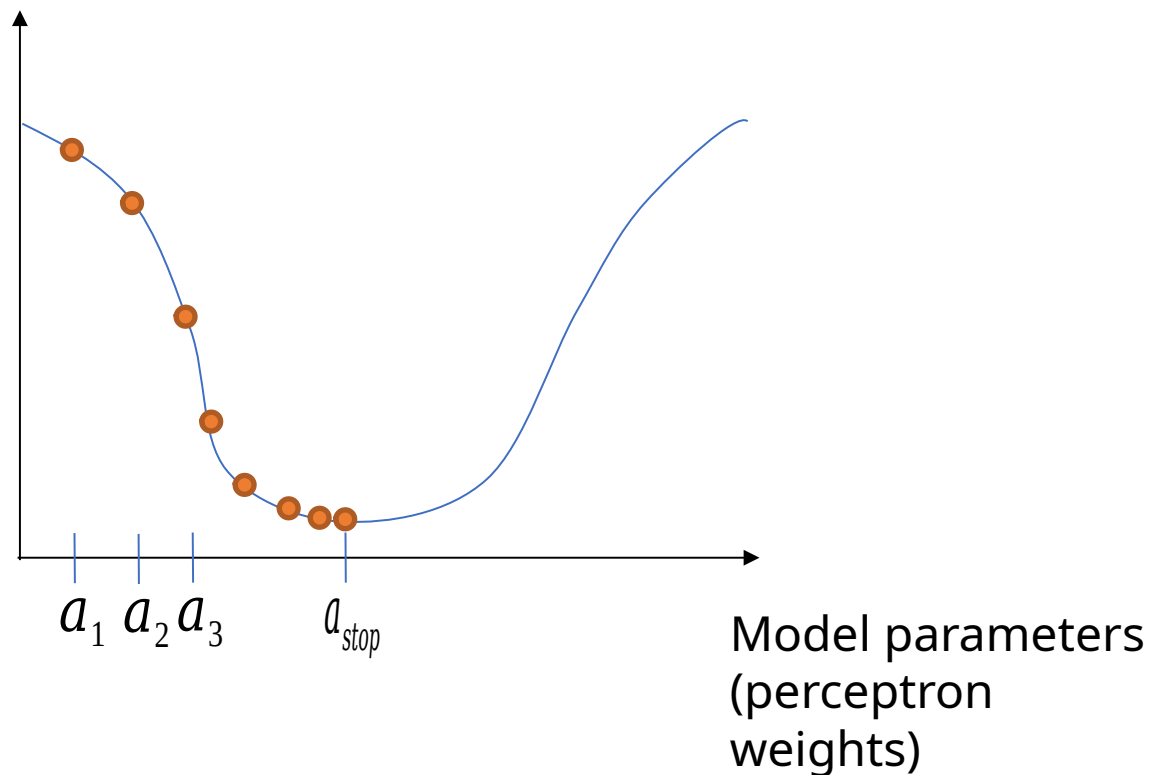
Even for convex functions, gradient descent allows me to iteratively solve the solution without requiring very large matrices.

Train NN with Gradient Descent

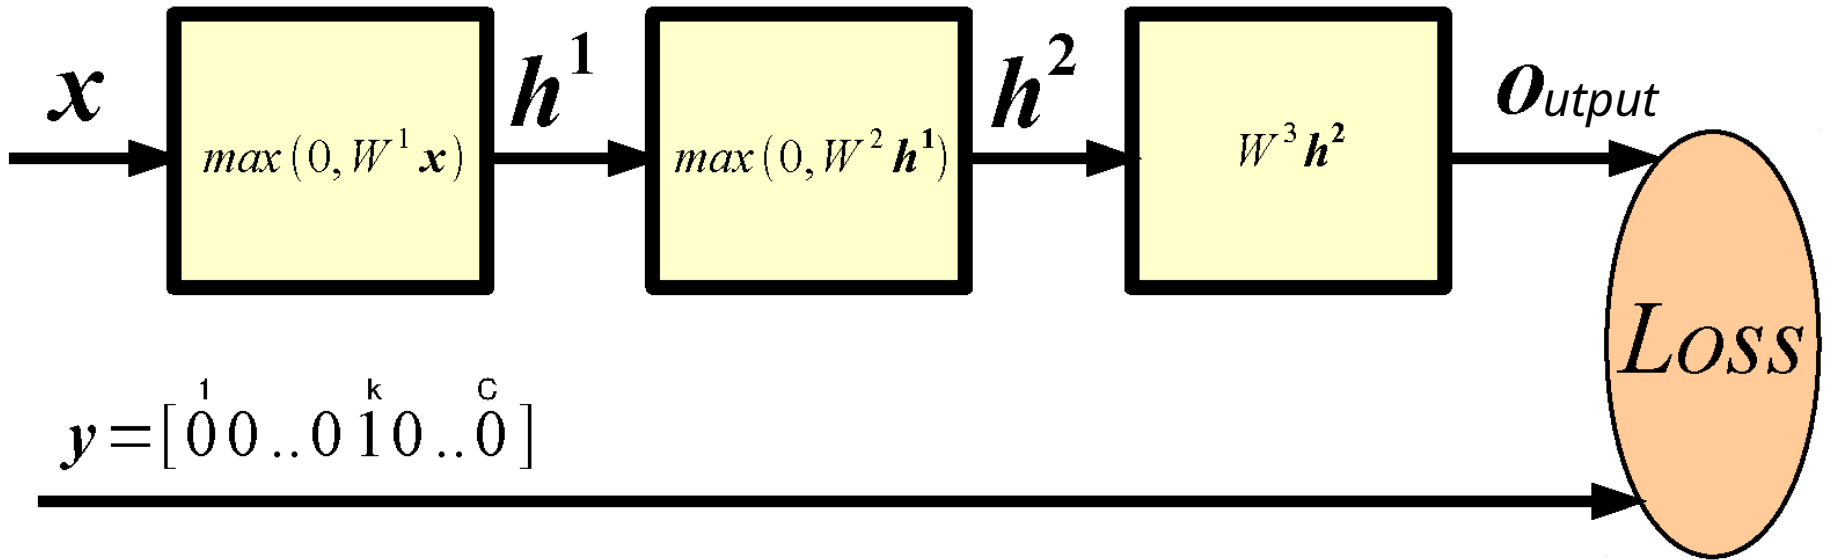
- $x^i, y^i = n$ training examples
- $f(\mathbf{x})$ = feed forward neural network
- $L(\mathbf{x}, y; \boldsymbol{\theta})$ = some *loss function*
- *Loss function* measures how ‘good’ our network is at classifying the training examples wrt. the parameters of the model (the perceptron weights).

Train NN with Gradient Descent

Loss function
(Evaluate NN
on training
data)



How Good is a Network?



What is an appropriate loss?

- Define some output threshold on detection
- Classification: compare training class to output class
- Zero-one loss (per class)

$$\begin{array}{l} y = \text{true label} \\ \hat{y} = \text{predicted label} \end{array} \quad L(\hat{y}, y) = I(\hat{y} \neq y),$$

- Is it good?
 - Nope – it's a step function.
 - I need to compute the *gradient* of the loss.
 - This loss is not differentiable, and 'flips' easily.

Classification as probability

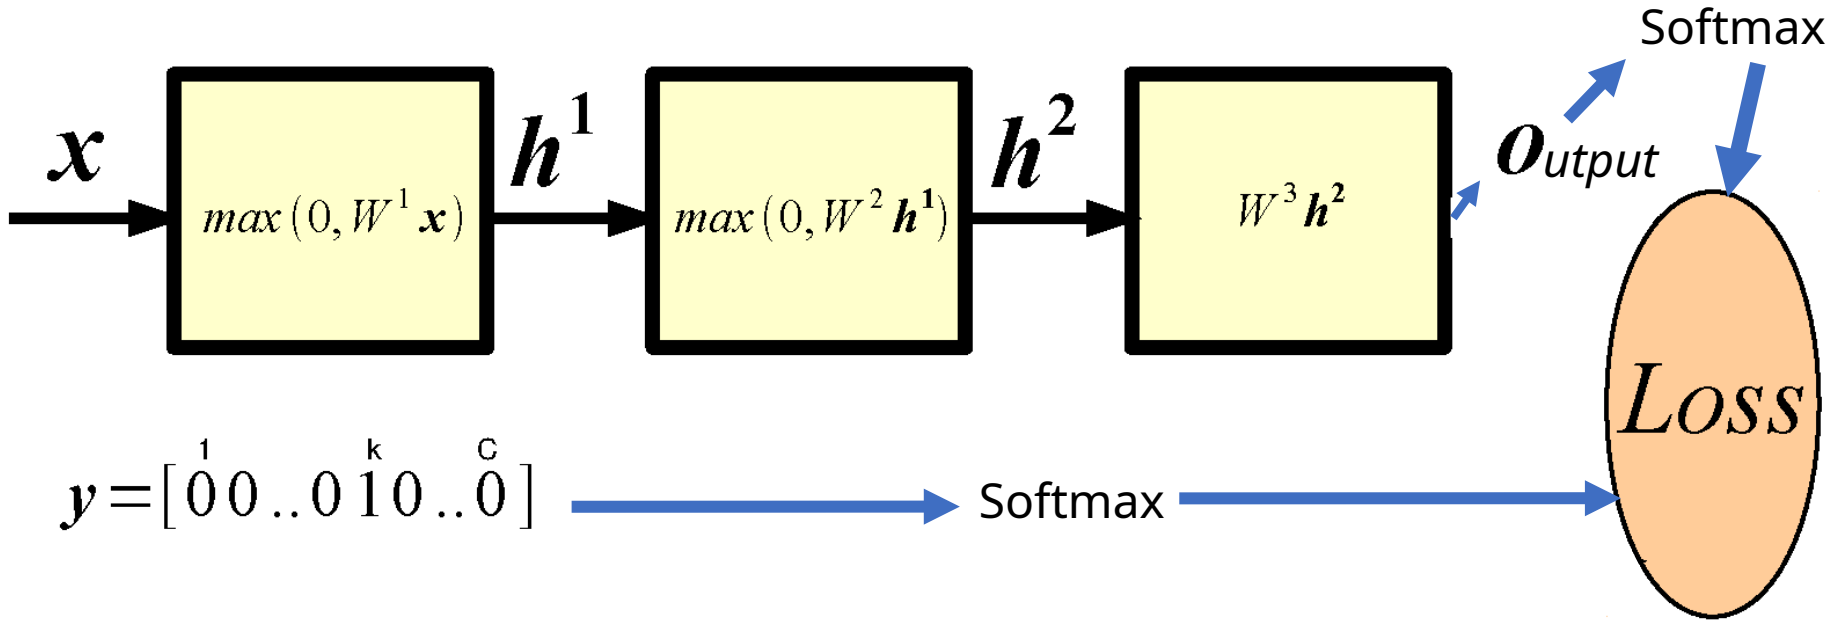
Special function on last layer - '*Softmax*':

"Squashes" a C -dimensional vector \mathbf{o} of arbitrary real values to a C -dimensional vector $\sigma(\mathbf{o})$ of real values in the range $(0, 1)$ that add up to 1.

Turns the output into a probability distribution on classes.

$$p(c_k = 1 | \mathbf{x}) = \frac{e^{o_k}}{\sum_{j=1}^C e^{o_j}}$$

How Good is a Network?



Probability of class k given input (softmax):

$$p(c_k = 1 | \mathbf{x}) = \frac{e^{o_k}}{\sum_{j=1}^C e^{o_j}}$$

Example: your model is classifying 20 classes, the final output vector will have a size 20x1. Softmax transforms the outputs to positive numbers with their sum = 1. Each row now corresponds to the probability of its associated class.


Cross-entropy loss function

- Negative log-likelihood

y - binary indicator (0 or 1) if class label j is the correct classification for observation

$$L(\mathbf{x}, y; \boldsymbol{\theta}) = - \sum_j y_j \log p(c_j | \mathbf{x})$$

- Minimizing this is similar to minimizing KL-divergence on predicted and target probability distributions
- Is it a good loss?
 - Differentiable
 - Cost decreases as probability increases



Training

Learning consists of minimizing the loss (plus some regularization term) w.r.t. parameters over the whole training set.

$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta}} \sum_{n=1}^P L(\mathbf{x}^n, y^n; \boldsymbol{\theta})$$

The diagram includes five text labels with blue arrows pointing to specific parts of the equation:

- training set**: Points to the upper limit P of the summation.
- model weights**: Points to the parameter vector $\boldsymbol{\theta}$ inside the loss function.
- label associated with the training data point n**: Points to the output variable y^n .
- input vector associated with the training data point n**: Points to the input variable \mathbf{x}^n .
- sum over all training data**: Points to the summation symbol \sum .

Training

Learning consists of minimizing the loss (plus some regularization term) w.r.t. parameters over the whole training set.

$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta}} \sum_{n=1}^P L(\mathbf{x}^n, y^n; \boldsymbol{\theta})$$

Question: How to minimize a complicated function of the parameters?

Answer: Chain rule, a.k.a. **Backpropagation**! That is the procedure to compute gradients of the loss w.r.t. parameters in a multi-layer neural network.

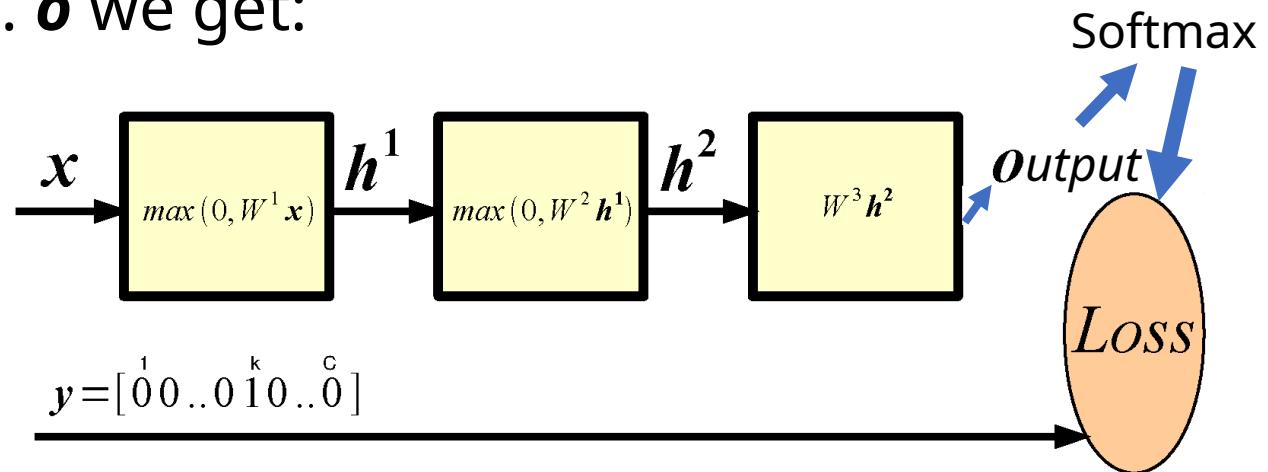
Derivative w.r.t. Input of Softmax

$$p(c_k=1|\mathbf{x}) = \frac{e^{o_k}}{\sum_j e^{o_j}}$$

$$L(\mathbf{x}, \mathbf{y}; \boldsymbol{\theta}) = -\sum_j y_j \log p(c_j|\mathbf{x}) \quad \mathbf{y} = [\overset{1}{0} \overset{2}{0} \dots \overset{k}{1} \overset{C}{0} \dots \overset{C}{0}]$$

By substituting the first formula in the second, and taking the derivative wr.t. \mathbf{o} we get:

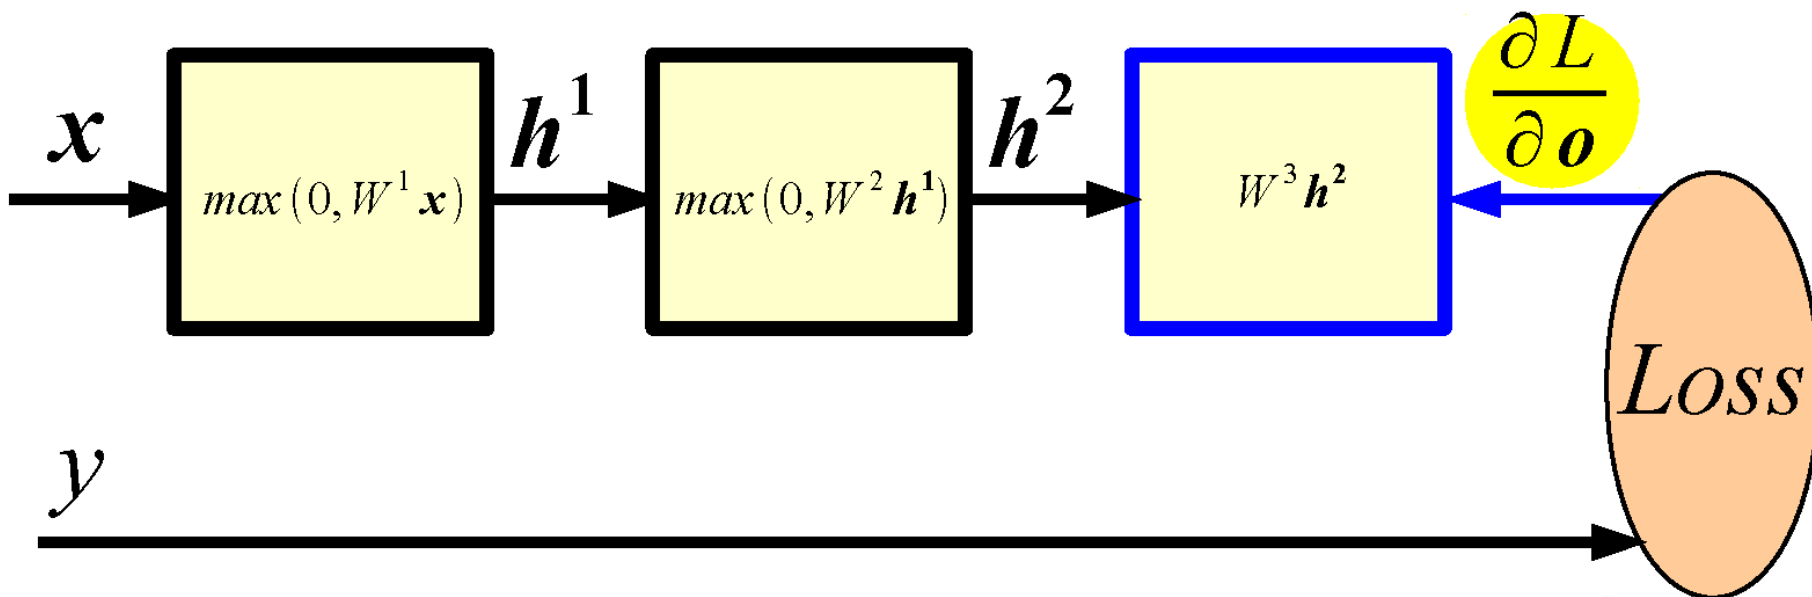
$$\frac{\partial L}{\partial \mathbf{o}} = p(c|\mathbf{x}) - \mathbf{y}$$



Proof:

$$\begin{aligned}\frac{\partial \xi}{\partial z_i} &= - \sum_{j=1}^C \frac{\partial t_j \log(y_j)}{\partial z_i} \\&= - \sum_{j=1}^C t_j \frac{\partial \log(y_j)}{\partial z_i} = - \sum_{j=1}^C t_j \frac{1}{y_j} \frac{\partial y_j}{\partial z_i} \\&= - \frac{t_i}{y_i} \frac{\partial y_i}{\partial z_i} - \sum_{j \neq i}^C \frac{t_j}{y_j} \frac{\partial y_j}{\partial z_i} = - \frac{t_i}{y_i} y_i (1 - y_i) - \sum_{j \neq i}^C \frac{t_j}{y_j} (-y_j y_i) \\&= -t_i + t_i y_i + \sum_{j \neq i}^C t_j y_i = -t_i + \sum_{j=1}^C t_j y_i = -t_i + y_i \sum_{j=1}^C t_j \\&= y_i - t_i\end{aligned}$$

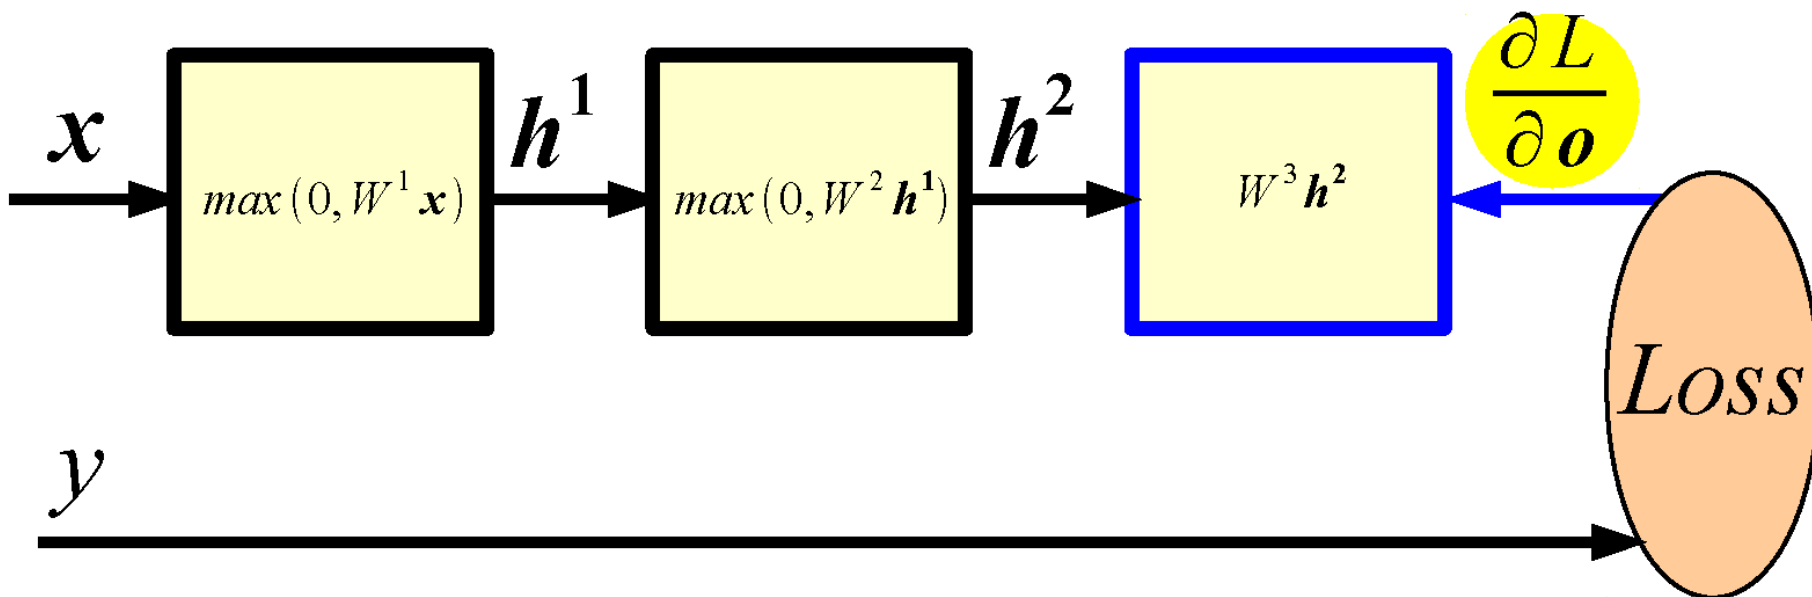
Backward Propagation



Given $\partial L / \partial \mathbf{o}$ and assuming we can easily compute the Jacobian of each module, we have:

$$\frac{\partial L}{\partial W^3} = \frac{\partial L}{\partial \mathbf{o}} \frac{\partial \mathbf{o}}{\partial W^3}$$

Backward Propagation

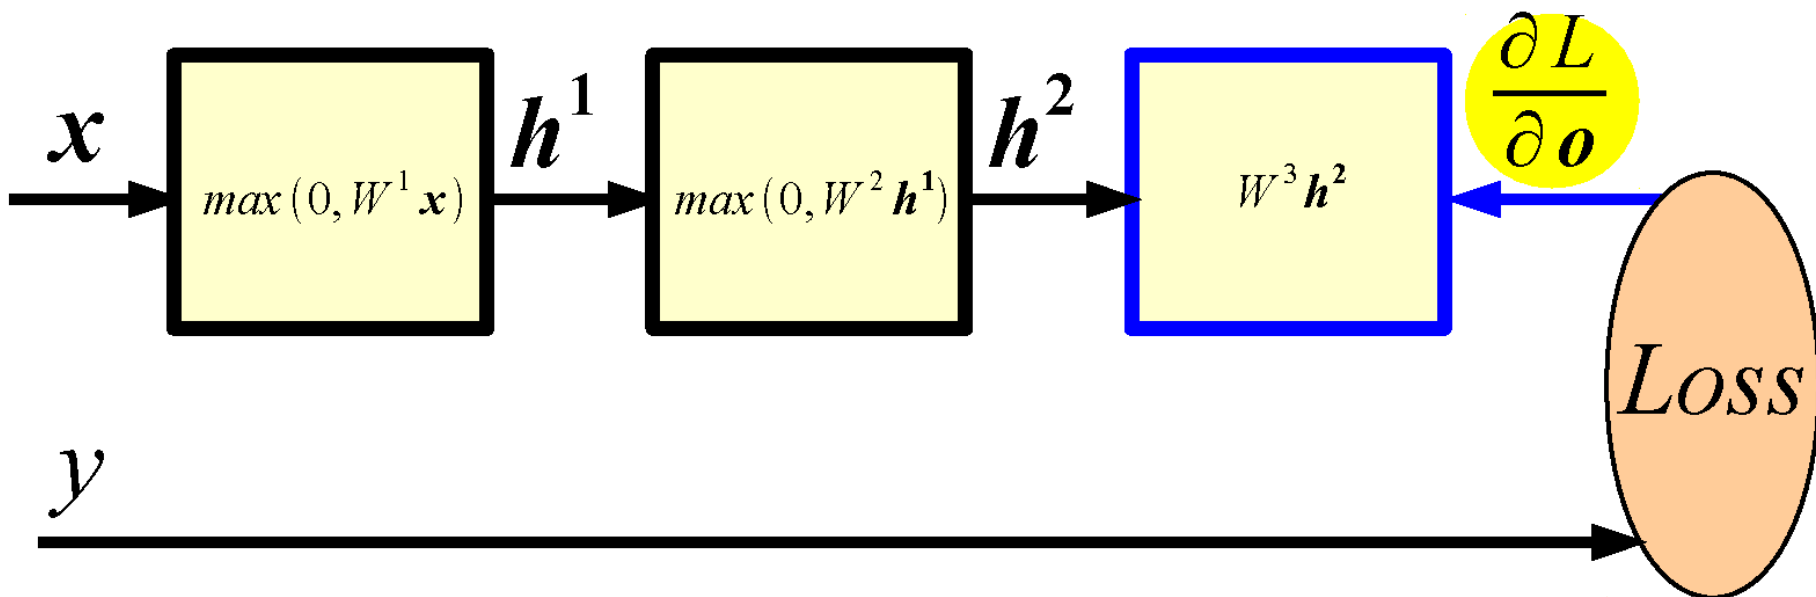


Given $\partial L / \partial o$ and assuming we can easily compute the Jacobian of each module, we have:

$$\frac{\partial L}{\partial W^3} = \frac{\partial L}{\partial o} \frac{\partial o}{\partial W^3}$$

$$\frac{\partial L}{\partial h^2} = \frac{\partial L}{\partial o} \frac{\partial o}{\partial h^2}$$

Backward Propagation



Given $\partial L / \partial o$ and assuming we can easily compute the Jacobian of each module, we have:

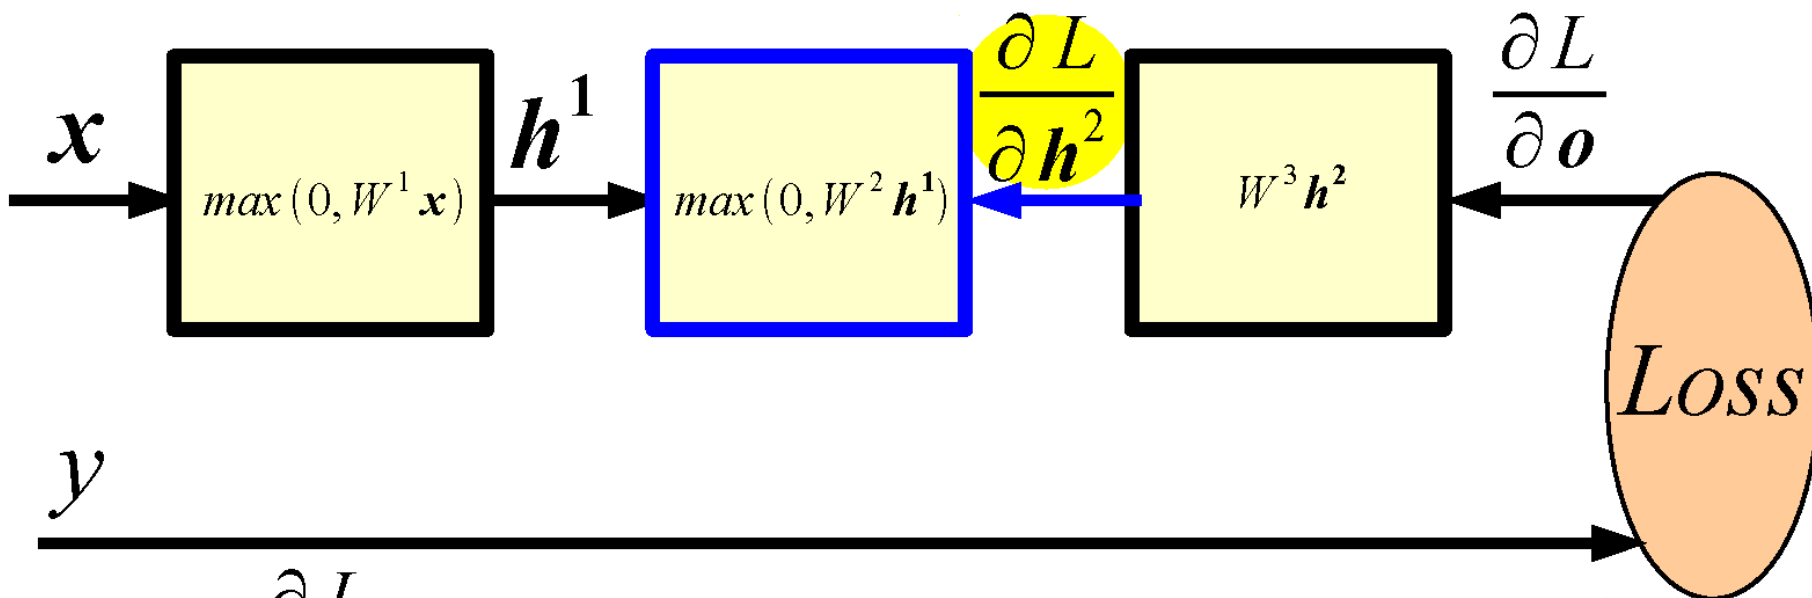
$$\frac{\partial L}{\partial W^3} = \frac{\partial L}{\partial o} \frac{\partial o}{\partial W^3}$$

$$\frac{\partial L}{\partial h^2} = \frac{\partial L}{\partial o} \frac{\partial o}{\partial h^2}$$

$$\frac{\partial L}{\partial W^3} = (p(c|x) - y) h^{2T}$$

$$\frac{\partial L}{\partial h^2} = W^{3T} (p(c|x) - y)$$


Backward Propagation



Given $\frac{\partial L}{\partial h^2}$ we can compute now:

$$\frac{\partial L}{\partial W^2} = \frac{\partial L}{\partial h^2} \frac{\partial h^2}{\partial W^2} \qquad \frac{\partial L}{\partial h^1} = \frac{\partial L}{\partial h^2} \frac{\partial h^2}{\partial h^1}$$

Backward Propagation



Given $\frac{\partial L}{\partial \mathbf{h}^1}$ we can compute now:

$$\frac{\partial L}{\partial W^1} = \frac{\partial L}{\partial \mathbf{h}^1} \frac{\partial \mathbf{h}^1}{\partial W^1}$$

Toy Code (Matlab): Neural Net Trainer

% F-PROP

```
for i = 1 : nr_layers - 1
    [h{i} jac{i}] = nonlinearity(W{i} * h{i-1} + b{i});
end
h{nr_layers-1} = W{nr_layers-1} * h{nr_layers-2} + b{nr_layers-1};
prediction = softmax(h{1-1});
```

% CROSS ENTROPY LOSS

```
loss = - sum(sum(log(prediction) .* target)) / batch_size;
```

% B-PROP

```
dh{1-1} = prediction - target;
for i = nr_layers - 1 : -1 : 1
    Wgrad{i} = dh{i} * h{i-1}';
    bgrad{i} = sum(dh{i}, 2);
    dh{i-1} = (W{i}' * dh{i}) .* jac{i-1};
end
```

% UPDATE

```
for i = 1 : nr_layers - 1
    W{i} = W{i} - (lr / batch_size) * Wgrad{i};
    b{i} = b{i} - (lr / batch_size) * bgrad{i};
end
```

Stochastic Gradient Descent

- Dataset can be too large to strictly apply gradient descent.
- Instead, randomly sample a data point, perform gradient descent per point, and iterate.
 - True gradient is approximated only
 - Picking a subset of points: "*mini-batch*"

Pseudo code:

Pick starting W and learning rate γ

While not at minimum:


- Shuffle training set
- For each data point $i=1\dots n$ (*maybe as mini-batch*)
 - *Gradient descent*

} "Epoch"


Stochastic Gradient Descent

Loss will not
always decrease
(locally) as
training data
point is random.


Still converges
over time.



Gradient descent oscillations



Gradient descent oscillations



Slow to
converge to
the (local)
optimum

Momentum


- Adjust the gradient by a weighted sum of the previous amount plus the current amount.

- Without momentum: $\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \gamma \frac{\partial L}{\partial \boldsymbol{\theta}}$

- With momentum (new α parameter):

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \gamma \left(\alpha \left[\frac{\partial L}{\partial \boldsymbol{\theta}} \right]_{t-1} + \left[\frac{\partial L}{\partial \boldsymbol{\theta}} \right]_t \right)$$



Lowering the learning rate = smaller steps in SGD



-Less 'ping pong'

-Takes longer to get to the optimum

Low learning rate may get stuck in flat regions in energy landscape




Large learning rate can inadvertently increase rather than decrease the training error (can jump over the sought minima).

have fun tuning :)

Problem of fitting


- Too many parameters = overfitting
- Not enough parameters = underfitting
- More data = less chance to overfit
- How do we know what is required?

Data fitting problem




Which is better?

Which is better *a priori*?



9th order polynomial



1st order polynomial

Both can describe the data...

- ...but one is simpler.
- Occam's razor (c. 1287-1347):
"Among competing hypotheses, the one with the fewest assumptions should be selected"

For us:

Large weights cause large changes in behaviour in response to small changes in the input.

Simpler models (or smaller changes) are more robust to noise => generalizes better.

Regularization

- Idea:
Penalize the use of parameters to prefer small weights.
- Attempt to guide solution to *not overfit*
- But still give freedom with many parameters
- add a cost to having high weights
 λ = regularization parameter

$$C = C_0 + \lambda \sum_W W^2$$

Regularization

- Idea: add a cost to having high weights
- λ = regularization parameter

$$C = C_0 + \lambda \sum_W W^2$$

$$C = \underbrace{-\frac{1}{n} \sum_j y_j \log p(c_j | \mathbf{x})}_{\text{Normal cross-entropy loss}} + \underbrace{\lambda \sum_W W^2}_{\text{Regularization term}}$$

averaged over
number of
training data


Normal cross-entropy loss

Regularization term

Regularization: Dropout

- Our networks typically start with random weights.
- Every time we train = slightly different outcome.


- Why random weights?
- If weights are all equal, response across filters will be equivalent.
 - Network doesn't train.




Regularization

- Our networks typically start with random weights.
- Every time we train = slightly different outcome.
- Why not train 5 different networks with random starts and vote on their outcome?
 - Works fine!
 - Helps generalization because error is averaged.
 - BUT you need to run 5 different networks ! (slow)

Regularization: Dropout



Regularization: Dropout



At each mini-batch:

- Randomly select a subset of neurons.
- Ignore them.

On test: half weights outgoing to compensate for training on half neurons.

Effect:

- Neurons become less dependent on output of connected neurons.
- Forces network to learn more robust features that are useful to more subsets of neurons.
- Like averaging over many different trained networks with different random initializations.
- Except cheaper to train.


Many forms of 'regularization'

- Adding more data is a kind of regularization
- Pooling is a kind of regularization
- Data augmentation is a kind of regularization

...I thought we were going to treat machine learning like a black box? I like black boxes.

Deep learning is:

- a black box
- *also a black art.*



Many approaches and hyperparameters:

Activation functions, learning rate, mini-batch size, momentum...

Often these need tweaking, and you need to know what they do to change them intelligently.